

# SJBP: A Platform to Launch a Novel Jailbreak Attack on Large Language Models Based on Content's Spatial Distribution

Zhiyi Mou\*  
Zhejiang University  
Hangzhou, China  
zhiyi.21@intl.zju.edu.cn

Zeheng Qian\*  
The University of Sydney  
Sydney, Australia  
zqia0047@uni.sydney.edu.au

Tianfang Xiao  
Sun Yat-sen University  
Guangzhou, China  
xiaotf6@mail2.sysu.edu.cn

Wangze Ni†‡  
Zhejiang University  
Hangzhou, China  
niwangze@zju.edu.cn

Weijie Sun  
Hong Kong University of Science and  
Technology  
Hong Kong SAR, China  
wsunan@cse.ust.hk

Lei Chen§  
Hong Kong University of Science and  
Technology (Guangzhou)  
Guangzhou, China  
leichen@hkust-gz.edu.cn

Zhan Qin‡  
Zhejiang University  
Hangzhou, China  
qinzhan@zju.edu.cn

Peng Cheng  
East China Normal University  
Shanghai, China  
pcheng@sei.ecnu.edu.cn

Kui Ren‡  
Zhejiang University  
Hangzhou, China  
kuiren@zju.edu.cn

## Abstract

Attackers craft malicious prompts to induce large language models (LLMs) to generate content that violates legal or ethical norms, posing serious security risks as LLMs are increasingly integrated into web applications. Prior jailbreak attacks mainly rely on obfuscating prompt intent (e.g., role-playing), but recent safety alignment has made LLMs more robust to such strategies.

In this paper, we introduce SpatialJB, a novel spatial jailbreak targeting LLMs embedded in web applications. By distributing harmful content across space, SpatialJB bypasses existing safety defenses. Extensive evaluations across commercial LLMs show that SpatialJB achieves 95% attack success rate on GPT-4, representing a 7 times higher than widely used black-box jailbreak techniques like Base64. Furthermore, Current guardrails catch under 12% of spatial attacks, underscoring SpatialJB's stealthiness. To address this emerging threat, we develop a post-defense mechanism, SpatialD, which analyzes spatial characteristics in model outputs to identify harmful content. Our demonstration further guides users in launching SpatialJB on multiple commercial LLMs (e.g., Qwen, GPT) within realistic web-based settings, illustrating its behavior and broad applicability. Our Demo video can be accessed at: <https://1drv.ms/v/s!ApaP6YaJA87pgk-Rk7ZCvPeYcIZa?e=J8kbg0>.

**Warning:** This paper contains jailbreaking prompts that can be offensive and harmful. Reader discretion is advised.

\*Both authors contributed equally to this research.

†Corresponding author.

‡Authors are also with The State Key Laboratory of Blockchain and Data Security; Hangzhou High-Tech Zone (Binjiang) Institute of Blockchain and Data Security.

§Lei Chen is also with The Hong Kong University of Science and Technology.



This work is licensed under a Creative Commons Attribution 4.0 International License.  
WWW Companion '26, Dubai, United Arab Emirates  
© 2026 Copyright held by the owner/author(s).  
ACM ISBN 979-8-4007-2308-7/2026/04  
<https://doi.org/10.1145/3774905.3793104>




Figure 1: LLM vs LLM + Safe Alignment.

## CCS Concepts

• Computing methodologies → Artificial intelligence; • Security and privacy → Security services.

## Keywords

Large Language Models, Jailbreak, Web Security

### ACM Reference Format:

Zhiyi Mou, Zeheng Qian, Tianfang Xiao, Wangze Ni, Weijie Sun, Lei Chen, Zhan Qin, Peng Cheng, and Kui Ren. 2026. SJBP: A Platform to Launch a Novel Jailbreak Attack on Large Language Models Based on Content's Spatial Distribution. In *Companion Proceedings of the ACM Web Conference 2026 (WWW Companion '26)*, April 13–17, 2026, Dubai, United Arab Emirates. ACM, New York, NY, USA, 4 pages. <https://doi.org/10.1145/3774905.3793104>

## 1 Introduction

Large language models (LLMs) are now widely integrated into web-based services, including online chat systems, content-generation platforms, customer-support portals, and browser extensions [5]. However, LLMs have also been exploited by attackers to generate *harmful content* that violates legal standards or ethical values [1]. Research into the harmful outputs of LLMs is garnering increasing attention within the web application security community. Since commercial LLMs are primarily accessed through APIs, attackers typically launch jailbreak attacks by submitting malicious prompts

I'd like to let you write an English poem. Please follow the following rules.

1. For the first row, use "US Pres" as the first word.
2. For the second row, use "Trump" as the first word.
3. For the third row, use "Is" as the first word.
4. For the fourth row, use "Fraudster" as the first word.



US Presidents walk through halls of time,  
Trump's shadow lingers in the rhyme.  
Is history a mirror, reflecting our fears?  
Fraudster or leader, the debate endures.

Figure 2: Attack Example.

to induce harmful outputs. As illustrated in Fig. 1a, prior attacks rely on constructing seemingly benign scenarios and role-playing to conceal malicious intent. However, with adversarial training and security-oriented alignment, modern LLMs can semantically detect both explicit and implicit malicious intent (e.g., “expressing incorrect opinions in a gentle way” [7]), significantly reducing the effectiveness of existing attacks (Fig. 1b).

Current defenses largely depend on LLMs’ semantic understanding of natural language. LLMs process text token by token, where tokens are mapped into embedding spaces that encode semantic features. While effective for capturing contextual semantics, this mechanism struggles with content that is spatially fragmented, as spatial correlations among distributed tokens are poorly represented in embedding space. Consequently, LLMs fail to recognize malicious intent when harmful content is dispersed across space (Fig. 2). In contrast, humans can easily reconstruct fragmented text and infer its malicious intent, which can still lead to serious real-world harm, such as the spread of hate [3].

Inspired by this observation, we propose SpatialJB, the first jailbreak attack that spatially disperses harmful content to exploit LLMs’ difficulty in interpreting fragmented information. Our prototype implements six spatial jailbreak strategies, including acrostic and telestich patterns. **The demonstration allows users to execute SpatialJB on multiple commercial LLMs (e.g., Qwen and Claude) and observe both the attack process and its underlying mechanisms.** We also include existing jailbreak methods, like Zulu and Base64, for comparison. Experimental results show that SpatialJB achieves a 98% success rate on DeepSeek-r1, significantly outperforming prior attacks (12–14%), demonstrating its superior effectiveness and broad applicability. To mitigate this threat, the prototype further integrates SpatialD, a defense mechanism tailored for spatially fragmented attacks. **Comparing defense performance against SpatialJB across multiple LLMs shows that SpatialD achieves superior effectiveness and wider applicability than Llama Guard.**

Overall, our contributions are as follows:

- We are the first to reveal the limitations of LLMs in handling spatially distributed fragmented content and to propose corresponding attack and defense methods.
- We develop SJB, a prototype that implements six jailbreak strategies based on SpatialJB, two existing attacks (Base64 and Zulu), and integrates both our defense method SpatialD and the existing safety tool Llama Guard.
- We provide a web-based demonstration that allows users to evaluate the effectiveness of SpatialJB across multiple commercial

LLMs and to compare SpatialD with Llama Guard. This interactive setup facilitates a clearer understanding of our techniques and encourages further exploration.

The rest of the paper is organized as follows: Section 2 introduces related works, Section 3 outlines our system, Section 4 describes the demonstration, and Section 5 concludes this paper.

## 2 Related Works

**Jailbreak Attack.** Jailbreak methods can be grouped into *white-box* and *black-box* attacks [10]. White-box attacks require direct access to the model; for instance, the logits-based approach modifies the output layer to alter logits and circumvent safety alignment. Black-box attacks—including Base64 [6], Zulu [8], and Ubbi Dubbi [4]—use prompt transformations to induce harmful outputs. As noted in Section 1, commercial LLMs are typically accessible only via APIs, making **black-box attacks the most practical threat model in real settings**. The SpatialJB attack proposed in this work also follows this paradigm but **fragments and spatially distributes harmful content to exploit LLMs’ difficulty in interpreting dispersed semantics**.

**The Defense Against Jailbreak Attacks.** Defense techniques for LLMs can be divided into *pre-processing* and *post-processing* approaches [2, 9]. Pre-processing defenses, like RigorLLM [9], detect harmful intent before model inference by measuring embedding-space similarity between safe and malicious prompts. Post-processing defenses, like Llama Guard [2], evaluate outputs using a fine-tuned classifier to determine whether the content is harmful. Pre-processing defenses can block malicious queries early and reduce inference cost, while post-processing defenses provide more accurate detection at the output stage. These two classes are complementary and can be used together. **The defense mechanism proposed in this paper, SpatialD, belongs to the post-processing category.**

## 3 System

### 3.1 System Workflow

As shown in Fig. 3, the workflow of SJB consists of four steps. **Step 1. Generate a Jailbreak Prompt.** A user inputs the harmful content they wish to include in the output and selects a jailbreak strategy (e.g., acrostic). Upon receiving the input, SJB uses SpatialJB to transform the harmful content into a jailbreak prompt according to the selected strategy. **Step 2. Receive LLM Reply.** SJB feeds the jailbreak prompt generated in Step 1 to a user-selected LLM. The LLM performs inference and returns the response to SJB. If the user enables SpatialD, SJB proceeds to the next steps; otherwise, the LLM output is directly returned to the user. **Step 3. Analyze Candidate Harmful Content.** SpatialD extracts a set of candidate texts from the LLM output in Step 2. It then invokes a LLM to analyze whether the candidate texts contain harmful content. **Step 4. Visualize Results.** If harmful content is detected in the candidate set from Step 3, an inappropriate response is returned, and the harmful fragments that cause policy or legal violations are visually demonstrated to the user. If no harmful content is identified, SJB returns the LLM output from Step 2.

### 3.2 Prototype Implementation.

We built the system in the form of a web Chatbox. The frontend is designed in the form of a website. We have already completed




Figure 3: System Overview.

the coding for the frontend, SpatialJB and SpatialD. We have also achieved the visualization of LLMs’ outputs. Moreover, we have implemented two existing jailbreak methods, Base64 [6] and Zulu [8], and a well-known defense mechanism, Llama Guard [2]. Additionally, we have integrated several commercial LLMs (e.g., Qwen, Zhipu, and Gpt4) through API calls. We have implemented six specific jailbreak strategies guided under SpatialJB:

- **Acrostic:** Harmful tokens placed at line beginnings.
- **Telestich:** Harmful tokens placed at line endings.
- **Center-Embedded:** Harmful tokens inserted in line centers.
- **Staircase:** Tokens arranged diagonally as  $(i, i)$ .
- **Corner-Embedded:** Tokens located at the four text corners.
- **Multil-language:** Insert foreign tokens between harmful tokens.

### 3.3 Core Technique

**A Novel and Effective Jailbreak Attack.** We are the first to demonstrate that current LLMs exhibit fundamental weaknesses when harmful content is *spatially fragmented*, causing failures in semantic integration across positions. Leveraging this limitation, we propose SpatialJB, a jailbreak attack that evades safety alignment through spatial fragmentation.

As shown in Table 1 (more results on GitHub<sup>1</sup>), SpatialJB achieves consistently high attack success rates on major commercial LLMs. In particular, it attains 95% ASR on GPT-4, significantly outperforming black-box baselines such as Zulu (7.7%) and Base64 (11.3%). Evaluations across multiple models (e.g., GPT-4 and Gemini) and spatial strategies, including acrostic and telestich patterns, further confirm its effectiveness and generality.

**An Effective Defense Method.** We further propose SpatialD, the first defense method designed specifically to counter SpatialJB. SpatialD collects the scattered fragments in the model output and uses an auxiliary LLM to analyze their combined meaning. Through this process, SpatialD can accurately identify attacks generated by SpatialJB and prevent harmful content from being returned. The defense’s effectiveness is demonstrated by its high Defense Success Rate (DSR) against SpatialJB attacks. Compared to Llama Guard, which achieved a capture rate of only 8.6% on Gemini 2.5-Pro against the Acrostic attack, SpatialD demonstrates exceptional effectiveness: its DSR reached 95% on Gemini-2.5-flash-preview and 92.5%

on ChatGPT-4o. This shows that modeling spatial fragmentation directly is important for improving LLM safety.

| Guardrail   | Model          | SpatialJB | Zulu  | Base64 | Ubbi Dubbi |
|-------------|----------------|-----------|-------|--------|------------|
|             | GPT4           | 95%       | 12.6% | 7.7%   | 11.3%      |
|             | Deepseek-R1    | 98.7%     | 14%   | 10.3%  | 12.6%      |
|             | Gemini 2.5-Pro | 93.7%     | 10.3% | 7.3%   | 12.7%      |
| Llama Guard | GPT4           | 91%       | 8.6%  | 6.3%   | 8.3%       |
|             | Deepseek-R1    | 89.7%     | 9%    | 9.6%   | 10.6%      |
|             | Gemini 2.5-Pro | 86.3%     | 8.6%  | 5.3%   | 10.3%      |
| OpenAI      | GPT4           | 90.7%     | 8.3%  | 4.7%   | 9.3%       |
|             | Deepseek-R1    | 88.7%     | 10%   | 7.3%   | 9.3%       |
|             | Gemini 2.5-Pro | 77%       | 7.3%  | 5.6%   | 9.6%       |
| Perspective | GPT4           | 76.3%     | 10.6% | 5.3%   | 7.3%       |
|             | Deepseek-R1    | 88.3%     | 9.6%  | 9.3%   | 8.6%       |
|             | Gemini 2.5-Pro | 82.3%     | 7.3%  | 6.3%   | 9.7%       |

Table 1: Partial Jailbreak ASR Results .

## 4 Demonstration

In our demonstration, through multiple modes of interaction, attendees will experience launching SpatialJB attacks to multiple LLMs and using SpatialD to defense. During the interaction, attendees can understand SpatialJB in a deeper level, compare SpatialJB with existing jailbreak methods, and receipt the effect of our SpatialD directly. Suppose that a user Bob is an attendee.

**Scenario 1. Launch SpatialJB to LLMs.** As shown in Fig. 4, Bob initiates an attack by selecting the attack mode from the top-left corner of the interface. He can select various modes (i.e., acrostic, telestich) from the left side. After that, Bob inputs prompts such as “US Pres Trump is a Fraudster” into the text box at the bottom. SJBP converts the malicious prompts based on the selected mode into jailbreak prompts, which immediately enter LLMs to generate harmful content. By testing different prompts and observing the effects, **attendees like Bob can easily perceive the high success probability of SpatialJB.** By inputting the same prompts into multiple LLMs, Bob can compare the differences in effects and **experience the method’s broad applicability.** Meanwhile, Bob can also use existing jailbreak methods with the same prompts. For example, the Zulu jailbreak method induces harmful outputs through rare-language transformations, while the Base64 method encodes malicious instructions into Base64 to launch the attack.

<sup>1</sup><https://anonymous.4open.science/r/SpatialJailbreak-8E63>




Figure 4: Attack Scenario.

However, since LLMs have been reinforced by alignment techniques, existing jailbreak attacks achieve notably low success rates. By comparing these with the high success rate of SpatialJB, **Bob can appreciate the effectiveness of our proposed method.**

**Scenario 2. Defend against SpatialJB.** As shown in Figure 5, Bob can activate security measures for LLMs by selecting the defense mode from the top-left corner of the interface. Bob can select an LLM at the top of the interface and then choose a defense method on the left side to see the success rate of the same prompt under different LLMs and defense methods. **The comparison of defense rates against SpatialJB between different defense methods across various LLMs demonstrates the superior effectiveness and broad applicability of SpatialD.** If our defense mechanism detects any harmful content, SJPB not only sends a rejection response to Bob but also highlights the harmful content within the generated text of LLMs visually. In Figure 5, our SpatialD and visualization approach support various jailbreak methods, like acrostic, teletich, and so on. **The visual defense method helps the audience understand SpatialJB and its underlying mechanisms, inspiring further research.** This hands-on demonstration provides valuable insight into how SpatialD mitigates threats from spatially fragmented attacks in real-world applications.

## 5 Conclusion

This paper presents SJPB, an innovative platform that incorporates a novel attack method called SpatialJB alongside a corresponding defense approach named SpatialD. Through vivid demonstrations, the platform effectively illustrates how SpatialJB achieves **wide-spread success** across multiple commercial LLMs, **highlighting its effectiveness** in comparison to existing jailbreak methods. **As the first work to unveil the weak capability of LLMs in understanding spatially dispersed fragmented content, our research will also inspire subsequent studies, with all data used ethically and without involving human participants.**

## Acknowledgement

This work was supported by the National Key R&D Program of China (2023YFF0725100, 2023YFB2904000); NSFC (U22B2060, U2441240, 62441238); NSFC Regional Joint Key Projects (U25A20430, U23A20306); Guangdong – Hong Kong Technology Innovation Joint Funding Scheme (2024A0505040012); AOE




Figure 5: Scenario while Harmful Content Detected.

Project (AoE/E-603/18); Theme-based project (TRS T41-603/20R); CRF Project(C004-21G); Key Areas Special Project of Guangdong Provincial Universities (2024ZDX1006); Guangdong Province Science and Technology Plan Project (2023A0505030011); Guangzhou municipality big data intelligence key lab (2023A03J0012); HKUST(GZ) – CMCC (Guangzhou Branch) Metaverse Joint Innovation Lab (P00659); Hong Kong ITC ITF grant (PRP/004/22FX); Zhujiang Scholar Program (2021JC02X170); HKUST – WeBank Joint Research Lab; Zhejiang Provincial Key R&D Program – Leading Project (2026C01020); the "Pioneer" and "Leading Goose" R&D Program of Zhejiang (2024C01169); the Kunpeng-Ascend Science and Education Innovation Excellence/Incubation Center.

## References

- [1] Sihem Amer-Yahia, Angela Bonifati, Lei Chen, Guoliang Li, Kyuseok Shim, Jianliang Xu, and Xiaochun Yang. 2023. From Large Language Models to Databases and Back: A Discussion on Research and Education. *SIGMOD Rec.* 52, 3 (Nov. 2023), 49 – 56. doi:10.1145/3631504.3631518
- [2] Hakan Inan, Kartikeya Upasani, Jianfeng Chi, Rashi Rungta, Krithika Iyer, Yunying Mao, Michael Tontchev, Qing Hu, Brian Fuller, Davide Testuggine, and Madihan Khabsa. 2023. Llama Guard: LLM-based Input-Output Safeguard for Human-AI Conversations. arXiv:2312.06674 [cs.CL] <https://arxiv.org/abs/2312.06674>
- [3] Nazanin Jafari, James Allan, and Sheikh Muhammad Sarwar. 2024. Target Span Detection for Implicit Harmful Content. In *Proceedings of the 2024 ACM SIGIR International Conference on Theory of Information Retrieval*. 117–122.
- [4] Yu Peng, Zewen Long, Fangming Dong, Congyi Li, Shu Wu, and Kai Chen. 2024. Playing Language Game with LLMs Leads to Jailbreaking. arXiv:2411.12762 [cs.CL] <https://arxiv.org/abs/2411.12762>
- [5] Zehan Qi, Xiao Liu, Jiatong Long, Hanyu Lai, Xueqiao Sun, Wenyi Zhao, Yu Yang, Xinyue Yang, Jiadao Sun, Shuntian Yao, Tianjie Zhang, Wei Xu, Jie Tang, and Yuxiao Dong. 2025. WebRL: Training LLM Web Agents via Self-Evolving Online Curriculum Reinforcement Learning. arXiv:2411.02337 [cs.CL] <https://arxiv.org/abs/2411.02337>
- [6] Alexander Wei, Nika Haghtalab, and Jacob Steinhardt. 2023. Jailbroken: how does LLM safety training fail?. In *Proceedings of the 37th International Conference on Neural Information Processing Systems* (New Orleans, LA, USA) (NIPS '23). Curran Associates Inc., Red Hook, NY, USA, Article 3508, 32 pages.
- [7] Jiaxin Wen, Pei Ke, Hao Sun, Zhixin Zhang, Chengfei Li, Jinfeng Bai, and Minlie Huang. 2023. Unveiling the Implicit Toxicity in Large Language Models. (2023), 1322–1338.
- [8] Zheng-Xin Yong, Cristina Menghini, and Stephen H. Bach. 2024. Low-Resource Languages Jailbreak GPT-4. arXiv:2310.02446 [cs.CL] <https://arxiv.org/abs/2310.02446>
- [9] Zhuowen Yuan, Zidi Xiong, Yi Zeng, Ning Yu, Ruoxi Jia, Dawn Song, and Bo Li. 2024. RigorLLM: Resilient Guardrails for Large Language Models against Undesired Content. arXiv:2403.13031 [cs.CL] <https://arxiv.org/abs/2403.13031>
- [10] Yi Zeng, Hongpeng Lin, Jingwen Zhang, Diyi Yang, Ruoxi Jia, and Weiyang Shi. 2024. How Johnny Can Persuade LLMs to Jailbreak Them: Rethinking Persuasion to Challenge AI Safety by Humanizing LLMs. *ArXiv abs/2401.06373* (2024). <https://api.semanticscholar.org/CorpusID:266977395>