

Marketing Hosting: From Fixed to Endogenous Budgets

Bingzhe Wang*
Gaoling School of Artificial
Intelligence, Renmin University of
China
Beijing, China
wbz2022@ruc.edu.cn

Tianyu Wang
Alibaba Group
Beijing, China
yves.wty@alibaba-inc.com

Qi Qi†‡§
Gaoling School of Artificial
Intelligence, Renmin University of
China
Beijing, China
qi.qi@ruc.edu.cn

Xiaoxuan Deng
Alibaba Group
Beijing, China
dengxiaoxuan.dxx@alibaba-inc.com

Zhilin Zhang
Alibaba Group
Beijing, China
zhangzhilin.pt@alibaba-inc.com

Chuan Yu
Alibaba Group
Beijing, China
yuchuan.yc@alibaba-inc.com

Abstract

The canonical model for multi-channel marketing, Bandits with Knapsacks (BwK), optimizes cumulative rewards subject to resource constraints but typically assumes a fixed, exogenous budget. This assumption is tenuous in real-world systems requiring performance-adaptive investment, where pre-committing to a budget is challenging and suboptimal. We introduce *Marketing Hosting*, a paradigm modeling the budget as an endogenous, performance-dependent variable. This yields a new problem class, Bandits with Endogenous Knapsacks (BwEK), characterized by a challenging feedback loop coupling rewards with constraints. We develop a specialized primal-dual algorithm to manage this coupling. For settings with hard, per-round constraints, we design a novel risk-aware algorithm that mitigates the path-dependent risk of ruin, providing the first high-probability safety guarantee for such problems. Finally, we solve the strategic bi-level problem of learning the optimal reinvestment rate. We validate our theoretical results through extensive simulations, real-world data experiments, and a live A/B test.

CCS Concepts

• Information systems → Computational advertising.

Keywords

Online Marketing, Multi-channel Budget Allocation

ACM Reference Format:

Bingzhe Wang, Tianyu Wang, Qi Qi, Xiaoxuan Deng, Zhilin Zhang, and Chuan Yu. 2026. Marketing Hosting: From Fixed to Endogenous Budgets. In *Proceedings of the ACM Web Conference 2026 (WWW '26)*, April 13–17, 2026, Dubai, United Arab Emirates. ACM, New York, NY, USA, 12 pages. <https://doi.org/10.1145/3774904.3792519>

*Work is done during the internship at Alibaba Group

†Corresponding Author

‡Also with, Beijing Key Laboratory of Research on Large Models and Intelligent Governance

§Also with, Engineering Research Center of Next-Generation Intelligent Search and Recommendation, MOE



This work is licensed under a Creative Commons Attribution 4.0 International License. *WWW '26, Dubai, United Arab Emirates*

© 2026 Copyright held by the owner/author(s).

ACM ISBN 979-8-4007-2307-0/2026/04

<https://doi.org/10.1145/3774904.3792519>

Dubai, United Arab Emirates. ACM, New York, NY, USA, 12 pages. <https://doi.org/10.1145/3774904.3792519>

1 Introduction

In the modern digital economy, effective marketing is paramount to a merchant's success [16, 17, 22, 25]. Consider a merchant on a large e-commerce platform offering a diverse suite of marketing channels, such as sponsored search, social media advertising, recommendations, and promotional coupons. By investing a budget across these channels, the merchant can significantly boost revenue, which in turn funds further marketing, creating a virtuous cycle of growth. This paper aims to assist merchants in optimizing decisions within this complex environment [8, 10–12, 15].

However, the high dimensionality of channels and performance uncertainty make it difficult for merchants to devise optimal strategies independently. This challenge has fostered *Performance Marketing* or *Marketing Automation*, where the platform's experts manage the merchant's multi-channel allocation to maximize returns.

Academic research has historically been dominated by a fixed-budget paradigm. Here, the merchant commits to a fixed, upfront budget, and the platform's task is modeled as a combinatorial multi-armed bandit (CMAB) for round-by-round allocation [7, 14, 18, 19, 23, 29, 31] or a bandit with knapsacks (BwK) for long-term budget management [2, 3, 6, 9]. While theoretically elegant, these models assume an exogenously fixed budget. In practice, determining the optimal upfront budget is non-trivial: overly conservative budgets miss market opportunities, while aggressive ones risk inefficient spending. This reveals a critical gap between fixed-budget theoretical models and the strategic needs of merchants requiring adaptive investment levels.

To bridge this gap, we introduce *Marketing Hosting*, a paradigm where the budget is not fixed but hosted and scaled based on performance. In marketing, budgets are inherently endogenous: revenue funds future marketing. Our problem originates from a real "Marketing Hosting" service where the platform acts as the authorized decision-maker. For instance, if a merchant sells an item for \$200 with a 10% reinvestment rate, the platform retains \$20 as the marketing budget. The platform allocates this budget across channels (e.g., setting advertising bids or coupon values) to drive further sales. Increased sales generate more budget, creating a virtuous feedback loop that traditional fixed-budget models cannot capture.

The reinvestment rate (e.g., 10%) is set by the merchant—varying per merchant rather than a theoretical assumption. Crucially, the platform manages the risk of ruin: if costs exceed the generated budget (e.g., cost \$60 vs. budget \$40 from \$400 revenue), the merchant will not cover the deficit. To prevent overdrafts (e.g., \$60 > \$40), a strict per-round constraint ensures cumulative cost remains less than or equal to the cumulative budget at all times. In this setting, effort levels represent interventions like placing a \$5 bid or \$15 coupon; higher efforts cost more but increase conversion probability.

This endogenous framework introduces a critical challenge: risk management. The marketing budget effectively functions as a commission on revenue. The platform must expend resources to execute campaigns *before* revenue—and thus the budget—is realized. If a campaign performs poorly, costs may exceed the generated budget, causing a net loss. In high-stakes environments with hard constraints, a single violation could terminate the process. Therefore, the platform must rigorously manage this spend-then-earn dynamic to mitigate the risk of ruin.

Furthermore, even with safe budget management, the system relies on a predetermined reinvestment parameter (rate ι). This parameter embodies a trade-off: a high rate drives gross revenue but yield minimal net profit (revenue minus marketing costs), while a low rate ensures safety but may fail to capture market potential. Since the optimal rate depends on unknown market dynamics, this motivates a higher-level strategic problem: learning the optimal reinvestment parameter online to maximize long-term *net profit*.

This paper addresses this hierarchy of problems with the following contributions:

(1) We formalize the problem of sequential allocation where the budget is **endogenously generated** via past rewards, termed Bandits with Endogenous Knapsacks (BwEK). We develop a specialized primal-dual algorithm to manage the feedback loop between rewards and constraints.

(2) For the stringent setting with hard, per-round constraints, we design a novel **risk-aware algorithm** (RAPD) to mitigate the path-dependent risk of ruin. By integrating pessimism into safety checks while maintaining optimism in action selection, we provide, to our knowledge, the first high-probability safety guarantee for bandits with aggregate, path-dependent endogenous constraints.

(3) We solve the strategic bi-level problem of **learning the optimal reinvestment rate**. Formulated as a hierarchical bandit problem, our algorithm (HBL) provably converges to a near-optimal policy, bridging operational allocation with financial strategy.

1.1 Related Work

Our work sits at the confluence of online optimization for marketing, bandits with knapsacks, and combinatorial bandits.

Multi-channel marketing has been approached via various algorithmic perspectives. Classical methods often use convex optimization assuming known response curves [17], mixed-integer programming [27], or multiple-choice knapsack problems (MCKP) to handle discrete allocations [1, 24, 26], occasionally accounting for channel complementarities [26]. Research on joint bid/budget optimization in sponsored search [20, 21, 27] typically assumes

static environments, neglecting the exploration-exploitation trade-off. Recent reinforcement learning approaches [8, 16, 25] address dynamics but often lack formal regret guarantees.

The Bandits with Knapsacks (BwK) model [6] is the canonical framework for learning under long-term constraints. Literature includes robust, distribution-free regret bounds of $O(\sqrt{T})$ [6] and distribution-dependent $O(\log T)$ bounds [9]. Advancements include generalizations to concave rewards and convex knapsacks (BwCR) [3], and extensions incorporating contextual information [2, 13], non-stationary distributions [28], and adversarial settings [5].

To handle simultaneous multi-channel selection, we draw on the Combinatorial Multi-Armed Bandit (CMAB) framework [31], where agents select a *super-arm* to maximize cumulative reward [7]. This has been extended to Contextual Combinatorial Bandits (CCB) [11, 14, 18, 23, 29] and applied to influence maximization [4], recommender systems [18, 19], and wireless resource management [30]. Our SC-CBwK formulation integrates long-term budget constraints into the contextual combinatorial setting. While prior work addresses multi-channel allocation [1, 4, 24, 26] or long-term budgets [13, 28], this paper is first to synthesize these elements and generalize them to endogenous, performance-dependent constraints.

2 Preliminaries and Notation

This section establishes the mathematical foundation for sequential multi-channel resource allocation. We begin by defining the online learning environment, encompassing actions, observations, and the information structure. Next, we formalize statistical estimation, the offline benchmark, and the regret metric used to quantify performance. Finally, we state the core assumptions enabling our rigorous theoretical analysis.

A decision-maker operates over a finite horizon of T rounds, indexed by $\tau \in \{1, \dots, T\} = \mathcal{T}$. At each round τ , the decision-maker allocates resources across a set of N channels, denoted by the index set $\mathcal{N} = \{1, \dots, N\}$.

Actions: The decision for each channel $n \in \mathcal{N}$ is a base arm $x_{n,\tau}$ chosen from a discrete set of effort levels $\mathcal{X}$. The collective action at round τ is a super-arm $\mathbf{x}_\tau = (x_{1,\tau}, \dots, x_{N,\tau}) \in \mathcal{S}$, where $\mathcal{S} \subseteq \mathcal{X}^N$ is the set of feasible super-arms defined by combinatorial constraints.

Context: At the beginning of each round τ , a stochastic context vector $\phi_\tau \in \mathbb{R}^d$ is drawn from an unknown distribution $\mathcal{D}_\phi$.

Observations: Upon playing super-arm $\mathbf{x}_\tau$, the agent observes a reward $r_{n,\tau}$ and a cost $c_{n,\tau}$ for each channel n . The underlying reward and cost distributions, conditioned on the context and action, are unknown. Let $\mu_{r,n,x}(\phi_\tau) = \mathbb{E}[r_{n,\tau} | \mathbf{x}_\tau, \phi_\tau]$ and $\mu_{c,n,x}(\phi_\tau) = \mathbb{E}[c_{n,\tau} | \mathbf{x}_\tau, \phi_\tau]$ denote the true, unknown expected reward and cost functions, respectively.

Aggregates: We denote the total reward, cost, and (if applicable) budget generated in round τ as $r^\tau = \sum_{n \in \mathcal{N}} r_{n,\tau}$, $c^\tau = \sum_{n \in \mathcal{N}} c_{n,\tau}$, and $b^\tau = \sum_{n \in \mathcal{N}} b_{n,\tau}$. Their cumulative counterparts up to and including round τ are $R^\tau = \sum_{t=1}^\tau r^t$, $C^\tau = \sum_{t=1}^\tau c^t$, and $B^\tau = \sum_{t=1}^\tau b^t$.

A key component of online learning algorithms is the estimation of unknown quantities from data. Let $N_{n,x}(\tau)$ be the number of times the base arm x_n has been played up to round $\tau - 1$. The empirical mean reward is $\bar{r}_{n,x}(\tau) = \frac{1}{N_{n,x}(\tau)} \sum_{t=1}^{\tau-1} r_{n,t} \cdot \mathbb{I}(x_{n,t} = x)$. The empirical mean cost $\bar{c}_{n,x}(\tau)$ is defined analogously. To manage the

exploration-exploitation trade-off, we employ confidence bounds around these means.

Definition 2.1 (Confidence Bounds). For any base arm x_n at round τ , given a risk parameter $\delta' \in (0, 1)$, we define the confidence radius as: $\rho_{n,x}(\tau) = \sqrt{\frac{\ln(1/\delta')}{2N_{n,x}(\tau)}}$. The Lower Confidence Bound (LCB) and Upper Confidence Bound (UCB) for the true expected reward $\mu_{r,n,x}$ are defined respectively as:

$$\check{r}_{n,x}(\tau) = \bar{r}_{n,x}(\tau) - \rho_{n,x}(\tau), \hat{r}_{n,x}(\tau) = \bar{r}_{n,x}(\tau) + \rho_{n,x}(\tau).$$

The bounds $\check{c}_{n,x}(\tau)$ and $\hat{c}_{n,x}(\tau)$ for the true expected cost $\mu_{c,n,x}$ are defined identically.

To evaluate an online algorithm, we compare its performance to a clairvoyant oracle with full knowledge of the problem parameters. The performance of this oracle is captured by a linear programming (LP) relaxation, which represents an idealized, time-averaged optimal strategy. In the fixed-budget problem with a single global budget B , let ξ_s be the expected number of times a super-arm $s \in \mathcal{S}$ is played over the horizon T . The optimal offline performance, OPT_{LP} , is the value of the following linear program:

$$\begin{aligned} \max_{\xi} \quad & \sum_{s \in \mathcal{S}} \xi_s \left(\sum_{n \in \mathcal{N}} \mu_{r,n,s_n} \right) \\ \text{s.t.} \quad & \sum_{s \in \mathcal{S}} \xi_s \left(\sum_{n \in \mathcal{N}} \mu_{c,n,s_n} \right) \leq B, \\ & \sum_{s \in \mathcal{S}} \xi_s \leq T, \\ & \xi_s \geq 0, \quad \forall s \in \mathcal{S}. \end{aligned} \quad (P)$$

Definition 2.2 (Regret). The performance of an online policy π is measured by its expected cumulative regret: $\text{Regret}(T) = OPT_{LP} - \mathbb{E}_{\pi}[R^T]$. An algorithm is considered efficient if its regret is sublinear in T , i.e., $\text{Regret}(T) = o(T)$.

LEMMA 2.3 (THE OFFLINE BENCHMARK). *The value of the LP relaxation, OPT_{LP} , is a valid upper bound on the expected total reward of any adaptive online policy for the fixed-budget problem.*

A complete proof is provided in Appendix A.1. Our analysis also relies on a foundational lemma that bounds the number of times a suboptimal arm can be played before its empirical mean is statistically distinguishable from that of a better arm.

LEMMA 2.4 (SAMPLE COMPLEXITY OF SUBOPTIMAL BASE ARMS). *Consider a setting where the per-round objective is to maximize the sum of expected rewards, $\sum_n \mu_{r,n,x_n}$. Let x_n be a base arm that is sub-optimal with respect to another arm x'_n for the same channel, such that $\Delta = \mu_{r,n,x'} - \mu_{r,n,x} > 0$. Any UCB-based algorithm that selects super-arms by maximizing $\sum_n \hat{r}_{n,\tau}(x_n)$ will play the suboptimal base arm x_n an expected number of times bounded by $\mathbb{E}[N_{n,x}(T)] = \tilde{O}(1/\Delta^2)$, where $\tilde{O}(\cdot)$ hides polylogarithmic factors in T .*

A rigorous proof is detailed in Appendix A.2. Finally, our framework relies on standard technical assumptions.

Assumption 2.5 (Offline Oracle). We assume access to a $(1 - \epsilon)$ -approximation oracle for the single-round deterministic version of the combinatorial optimization problem.

Assumption 2.6 (Bounded Outcomes). All rewards $r_{n,\tau}$ and costs $c_{n,\tau}$ are bounded in $[0, 1]$.

Assumption 2.7 (Null Arm). We assume the existence of a "null action" $x_0 \in \mathcal{X}$ for each channel, which yields zero reward and zero cost: $\mu_{r,n,x_0} > \mu_{c,n,x_0} = 0$ for all $n \in \mathcal{N}$. This ensures the allocation problem is always feasible.

3 A Hierarchy of Budget Allocation Problems

In this section, we systematically analyze a hierarchy of sequential allocation problems, each adding a layer of realism and complexity critical for practical applications. We begin with a foundational fixed-budget model to establish core learning techniques for combinatorial actions and multiple constraints. We then advance to a dynamic endogenous budget model where constraints are stochastically coupled with performance, necessitating a sophisticated primal-dual approach. Finally, we address the most challenging scenario: managing the per-round risk of ruin under hard, path-dependent endogenous constraints. This structured progression builds the theoretical machinery required to solve our primary problem, demonstrating why simpler models are ultimately insufficient.

3.1 Foundation: The Fixed-Budget Problem

We begin our analysis with the foundational problem of allocating resources under a predetermined, fixed budget for the entire time horizon. This scenario models situations with pre-committed spending and serves as a crucial theoretical foundation for the more complex settings that follow. We formalize this as the Stochastic Contextual Combinatorial Bandit with Knapsacks (SC-CBwK).

Definition 3.1 (The SC-CBwK Problem). Let the total budget over the horizon T be a fixed constant $B > 0$. The Stochastic Contextual Combinatorial Bandit with Knapsacks problem is to find a policy π that maximizes the expected total reward over a horizon T , subject to a single global knapsack constraint on the total expected cost:

$$\begin{aligned} \max_{\pi} \quad & \mathbb{E}_{\pi}[R^T] = \mathbb{E}_{\pi} \left[\sum_{\tau \in \mathcal{T}} r^{\tau} \right] \\ \text{s.t.} \quad & \mathbb{E}_{\pi}[C^T] = \mathbb{E}_{\pi} \left[\sum_{\tau \in \mathcal{T}} c^{\tau} \right] \leq B. \end{aligned}$$

The Challenge. This problem presents a significant challenge due to the confluence of three factors: (1) the combinatorial action space $\mathcal{S}$, which makes brute-force exploration intractable; (2) the contextual information ϕ_{τ} , which requires the policy to be adaptive; (3) the long-term temporal coupling introduced by the single budget constraint over the horizon T . Canonical CMAB algorithms are unsuited to handle long-term constraints, as they typically focus on per-round optimization. Conversely, standard BwK algorithms are not designed for combinatorial and contextual action spaces. A naive application of existing methods would either ignore the budget constraint, leading to infeasible solutions, or fail to effectively explore the combinatorial action space, resulting in high regret.

Algorithmic Insight: A Primal-Dual Approach. To address this, we propose an algorithm, SCB-K, based on the primal-dual methodology, a powerful technique for constrained online optimization. The core insight is to transform the constrained optimization

Algorithm 1 SCB-K Algorithm

```

1: Initialize: History  $\mathcal{H}_{n,x} \leftarrow \emptyset$  for all  $(n, x)$ . Dual variable  $\lambda_1 \leftarrow \lambda_0$ . Learning rate  $\eta > 0$ .
2: for  $\tau = 1, 2, \dots, T$  do
3:   Observe context  $\phi_\tau$  and compute  $\hat{r}, \bar{r}, \bar{c}, \check{c}$  by  $\mathcal{H}_{n,x}$ .
4:    $\mathbf{x}_\tau \leftarrow \arg \max_{x \in \mathcal{S}} \{ \sum_n \hat{r}_{n,\tau}(x_n) - \lambda_\tau \sum_n \check{c}_{n,\tau}(x_n) \}$ .
5:   Play  $\mathbf{x}_\tau$ , observe rewards  $r_{n,\tau}$  and costs  $c_{n,\tau}$  for each  $n$ .
6:   For each  $n \in \mathcal{N}$ , update history  $\mathcal{H}_{n,x_{n,\tau}}$ .
7:   Update dual variable:  $\lambda_{\tau+1} \leftarrow \lambda_\tau \cdot \exp(\eta(c^\tau - B/T))$ .
8: end for

```

problem into a sequence of unconstrained ones. We introduce a single non-negative dual variable, λ_τ , which represents the dynamic *shadow price* of the budget at round τ . Instead of directly enforcing the budget constraint, the algorithm seeks to maximize an optimistic estimate of the Lagrangian objective at each round:

$$\mathcal{L}_\tau^{\text{CBwK}}(x, \lambda_\tau) = \sum_{n \in \mathcal{N}} r_{n,\tau}(x_n) - \lambda_\tau \sum_{n \in \mathcal{N}} c_{n,\tau}(x_n).$$

The dual variable λ_τ is updated using a multiplicative rule based on the observed cost relative to an amortized budget, effectively learning the correct price for the budget over time. This approach effectively decouples the immediate goal of reward maximization from the long-term necessity of satisfying the budget constraint.

This algorithm is provably efficient, achieving sublinear regret against the optimal offline policy.

THEOREM 3.2 (REGRET FOR SCB-K). *The SCB-K algorithm achieves an expected cumulative regret of $\text{Regret}(T) = \tilde{O}(\sqrt{N|X|T})$ for the SC-CBwK problem, where $\tilde{O}(\cdot)$ hides polylogarithmic factors in T, N , and $|X|$.*

The proof, detailed in Appendix B.1, relies on a standard sample complexity analysis for UCB-based algorithms coupled with a potential function argument for the dual variable. While foundational, the fixed-budget model's assumption of a known a priori budget is restrictive in dynamic environments where investment should ideally scale with success.

3.2 Intermediate Step: The Endogenous Budget

We now advance to a more dynamic model where the budget is not fixed but is endogenously determined as a fraction of the revenue generated by the algorithm's actions. This captures the realistic scenario where successful marketing efforts fund further investment.

Definition 3.3 (The BwEK Problem). Given a reinvestment parameter $\iota \in (0, 1)$, the budget generated in round τ is $b^\tau = \iota \cdot r^\tau$. The Bandits with Endogenous Knapsacks (BwEK) problem is to find a policy π that maximizes the total expected reward, subject to satisfying the endogenous budget constraint in expectation:

$$\begin{aligned} \max_{\pi} \quad & \mathbb{E}_{\pi} [R^T] \\ \text{s.t.} \quad & \mathbb{E}_{\pi} [C^T] \leq \mathbb{E}_{\pi} [B^T] = \mathbb{E}_{\pi} \left[\sum_{\tau \in \mathcal{T}} \iota \cdot r^\tau \right]. \end{aligned}$$

The Endogenous Challenge. The core difficulty of the BwEK problem lies in the feedback loop where the budget constraint is

Algorithm 2 SCB-EK Algorithm

```

1: Initialize: History  $\mathcal{H}_{n,x} \leftarrow \emptyset$ . Dual variable  $\lambda_1 \leftarrow \lambda_0$ . Learning rate  $\eta > 0$ .
2: for  $\tau = 1, 2, \dots, T$  do
3:   Observe context  $\phi_\tau$  and compute  $\hat{r}, \check{r}, \check{c}$  by  $\mathcal{H}_{n,x}$ .
4:   Select super-arm  $\mathbf{x}_\tau \leftarrow \arg \max_{x \in \mathcal{S}} \mathcal{L}_\tau^{\text{BwEK}}(x, \lambda_\tau)$ .
5:   Play super-arm  $\mathbf{x}_\tau$ , observe total reward  $r^\tau$  and costs  $c^\tau$ .
6:   For each  $n \in \mathcal{N}$ , update history  $\mathcal{H}_{n,x_{n,\tau}}$ .
7:    $b^\tau \leftarrow \iota \cdot r^\tau$ .
8:    $\lambda_{\tau+1} \leftarrow \lambda_\tau \cdot \exp(\eta(c^\tau - b^\tau))$ .
9: end for

```

coupled with the algorithm's performance. The capacity of the knapsack, $\sum b^\tau$, is no longer a fixed quantity but a random variable whose distribution is directly influenced by the policy. This endogenous relationship invalidates the core assumptions of standard BwK algorithms. An optimistic policy based on UCBs for rewards might overestimate performance, leading to an overestimated budget of the available budget. This, in turn, could justify selecting a higher-cost (and presumably higher-reward) action, creating a dangerous self-reinforcing cycle that can lead to severe budget violations.

Algorithmic Insight: A Conservative Primal-Dual Approach.

To break this cycle, we develop a specialized Primal-Dual algorithm, SCB-EK. The key insight is to be optimistic about the primal objective (reward) but pessimistic about the endogenous component of the constraint (budget generation). This is achieved by modifying the Lagrangian to use a UCB for the direct reward term but an LCB for the reward term that contributes to the budget. The optimistic Lagrangian for round τ becomes:

$$\mathcal{L}_\tau^{\text{BwEK}}(x, \lambda_\tau) = \sum_{n \in \mathcal{N}} \hat{r}_{n,\tau}(x_n) - \lambda_\tau \left(\sum_{n \in \mathcal{N}} \check{c}_{n,\tau}(x_n) - \iota \sum_{n \in \mathcal{N}} \check{r}_{n,\tau}(x_n) \right).$$

This conservative estimation of the generated budget prevents the algorithm from over-leveraging its anticipated success, thereby stabilizing the learning process. The dual variable is updated based on the realized net consumption $c^\tau - b^\tau$.

THEOREM 3.4 (REGRET FOR SCB-EK). *The SCB-EK algorithm achieves a sublinear expected cumulative regret for the BwEK problem of $\text{Regret}(T) = O(\sqrt{T})$.*

The proof is detailed in Appendix B.2. The SCB-EK algorithm effectively manages long-term resources in expectation. However, it permits temporary budget deficits, a luxury unavailable in mission-critical scenarios where a single round of overspending can lead to termination. This motivates our final and primary problem.

3.3 Main Challenge: Per-Round Constraints

We now address the most stringent and practical scenario, where budget constraints are hard, cumulative, and must be respected in every single round. This elevates the problem from managing expected resources to actively managing the per-round risk of ruin.

Definition 3.5 (The Risk of Ruin Problem). The process is governed by a stopping time, $\mathcal{T}_{\text{stop}}$, defined as the first round in which the cumulative budget is violated:

$$\mathcal{T}_{\text{stop}} = \min\{\tau \in \mathcal{T} \mid C^\tau > B^\tau\}.$$

Algorithm 3 Risk-Aware Primal-Dual Algorithm (RAPD)

```

1: Initialize: History  $\mathcal{H}_{n,x} \leftarrow \emptyset$ . Risk schedule  $\{\kappa_\tau\}_{\tau \in \mathcal{T}}$ . Dual
   variable  $\lambda_1 \leftarrow \lambda_0$ . Learning rate  $\eta > 0$ .
2: for  $\tau = 1, 2, \dots, T$  do
3:   Observe context  $\phi_\tau$  and compute  $\hat{r}, \check{r}, \hat{c}, \check{c}$  by  $\mathcal{H}_{n,x}$ .
4:   Let  $C^{\tau-1} = \sum_{t=1}^{\tau-1} c^t$  and  $B^{\tau-1} = \sum_{t=1}^{\tau-1} b^t$ .
5:   Stage 1: Identify Safe Action Set
6:    $S_{\text{safe}}(\tau) \leftarrow \emptyset$ .
7:   for each super-arm  $x \in \mathcal{S}$  do
8:      $\hat{c}^\tau(x) \leftarrow \sum_{n \in \mathcal{N}} \hat{c}_{n,\tau}(x_n)$ .
9:      $\check{b}^\tau(x) \leftarrow \iota \cdot \sum_{n \in \mathcal{N}} \check{r}_{n,\tau}(x_n)$ .
10:    if  $C^{\tau-1} + \hat{c}^\tau(x) \leq (1 - \kappa_\tau)(B^{\tau-1} + \check{b}^\tau(x))$  then
11:       $S_{\text{safe}}(\tau) \leftarrow S_{\text{safe}}(\tau) \cup \{x\}$ .
12:    end if
13:  end for
14:  Stage 2: Optimistic Lagrangian Selection
15:  if  $S_{\text{safe}}(\tau)$  is empty then
16:    Play null super-arm  $\mathbf{x}_\tau \leftarrow (x_0, \dots, x_0)$ .
17:  else
18:     $\mathbf{x}_\tau \leftarrow \arg \max_{x \in S_{\text{safe}}(\tau)} \left\{ \sum_n [\hat{r}_{n,\tau}(x_n) - \lambda_\tau(\check{c}_{n,\tau}(x_n) - \iota \check{r}_{n,\tau}(x_n))] \right\}$ .
19:  end if
20:  Play  $\mathbf{x}_\tau$ , observe  $r^\tau$  and  $c^\tau$ , and update histories  $\mathcal{H}_{n,x_{n,\tau}}$ .
21:  Stage 3: Update Dual Variable
22:   $b^\tau \leftarrow \iota \cdot r^\tau$ .
23:   $\lambda_{\tau+1} \leftarrow \lambda_\tau \cdot \exp(\eta(c^\tau - b^\tau))$ .
24: end for

```

If no violation occurs, $\mathcal{T}_{\text{stop}} = T$. The objective is to find a policy π that maximizes the expected total reward accumulated before this stopping time: $\max_\pi \mathbb{E}_\pi [R^{\mathcal{T}_{\text{stop}}}]$.

The Path-Dependent Challenge. This hard, path-dependent constraint fundamentally alters the problem. Algorithms designed for expected-value constraints, like SCB-EK, are rendered inadequate, as a single period of significant over-consumption can trigger $\mathcal{T}_{\text{stop}}$ and terminate the process irreversibly. The algorithm must now be prudently optimistic, controlling the tail behavior of both cost and budget deviations at every step to ensure safety with high probability.

Algorithmic Insight: Optimism in a Safeguarded Lagrangian.

To address this, we introduce the principle of *Optimism in a Safe-guarded Set*. Our Risk-Aware Primal-Dual (RAPD) algorithm first identifies a subset of actions that are *certifiably safe* with high probability with a time-dependent risk tolerance $\kappa_\tau \in (0, 1)$, and only then selects the most promising action from this restricted set using a primal-dual mechanism. The safety check is constructed with extreme conservatism: it verifies if the UCB of cumulative cost will remain below the LCB of the cumulative generated budget, even after taking the current action. This double pessimism in the constraint evaluation ensures a high-probability safety guarantee.

This design decouples risk management from opportunity seeking, ensuring that exploration only occurs within a space of actions that are highly unlikely to cause ruin.

THEOREM 3.6 (PERFORMANCE AND SAFETY OF RAPD). *For a decreasing risk schedule $\kappa_\tau = O(1/\tau^2)$ and learning rate $\eta = O(1/\sqrt{T})$,*

the RAPD algorithm ensures that the probability of premature termination is bounded, $\mathbb{P}(\mathcal{T}_{\text{stop}} < T) = O(1/T)$, and achieves an expected cumulative regret of $\text{Regret}(T) = \tilde{O}(\sqrt{N|X|T})$.

This result provides the first high-probability safety guarantee alongside a sublinear regret bound for a bandit problem with aggregate, endogenous, path-dependent constraints. The proof, found in Appendix B.3, requires a careful analysis combining martingale concentration inequalities for the budget surplus with the regret analysis of the primal-dual learning component.

4 Optimizing the Reinvestment Parameter

In the preceding sections, we developed algorithms to manage endogenous budgets governed by a fixed, predetermined reinvestment parameter, ι . This parameter represents a critical strategic decision: what fraction of the generated revenue should be reinvested into ongoing marketing efforts? An overly aggressive ι may drive high gross revenue but yield minimal net profit, while a conservative ι may be profitable but fail to capture market potential. The optimal value of ι is unknown a priori and depends on the complex, uncertain dynamics of the market environment.

This section elevates our analysis to address this critical strategic question. We move beyond simply operating with a given ι to actively learning the optimal ι online. This transforms the problem into a challenging bi-level optimization, where the outer loop seeks to optimize a strategic parameter whose performance can only be evaluated by running a separate, inner-loop online learning process.

4.1 The Bi-Level Optimization Challenge

The decision-maker's ultimate goal is to maximize the cumulative net revenue (total reward minus total cost), which is the portion of the gross revenue not reinvested as budget. For a given reinvestment parameter $\iota \in (0, 1)$, let $\pi^*(\iota)$ be an optimal online policy for the corresponding endogenous budget allocation problem, and let $\mathbb{E}[r_{\pi^*(\iota)}^\tau]$ be the expected total reward it generates at round τ . The strategic optimization problem is to find the parameter that maximizes the total expected net revenue over the horizon T .

Definition 4.1 (Optimal Reinvestment Parameter). The optimal reinvestment parameter $\iota^* \in (0, 1)$ is the solution to the following bi-level optimization problem:

$$\begin{aligned}
 & \text{maximize}_{\iota \in (0,1)} \quad J(\iota) := \mathbb{E} \left[\sum_{\tau \in \mathcal{T}} (1 - \iota) r_{\pi^*(\iota)}^\tau \right] \\
 & \text{subject to} \quad \pi^*(\iota) = \arg \max_{\pi} \mathbb{E}_\pi [R^T] \\
 & \quad \quad \quad \text{for the BwEK problem with parameter } \iota.
 \end{aligned}$$

The Black-Box Optimization Hurdle. This bi-level structure poses a formidable challenge that lies outside the scope of traditional bandit algorithms. The outer objective function, $J(\iota)$, is a classic black-box function:

No Closed-Form Expression: The value of $J(\iota)$ depends on the complex interplay between the policy $\pi^*(\iota)$ and the stochastic environment. There is no analytical expression for it.

Expensive, Noisy Evaluations: To obtain even a single, noisy sample of $J(\iota)$, one must execute the entire inner-loop online algorithm (e.g., SCB-EK) for a significant number of rounds.

Unavailable Gradients: The function $J(\iota)$ is not necessarily convex, and its gradients are unavailable, rendering standard continuous optimization techniques inapplicable.

Solving this problem requires a method that can intelligently explore the parameter space of ι using only noisy, expensive function evaluations, which is precisely the domain of bandit optimization.

4.2 A Hierarchical Bandit Learning Approach

We propose a novel framework, Hierarchical Bandit Learning (HBL), to solve this bi-level problem. The core insight is to reframe the outer optimization loop as a *meta-bandit* problem, where each potential value of the reinvestment parameter is treated as a *meta-arm*. An online bandit algorithm can then be used to efficiently learn the best meta-arm.

Algorithmic Design Rationale. The key challenge is managing the exploration-exploitation tradeoff at two levels simultaneously: the meta-level (choosing ι) and the base-level (choosing allocations $\mathbf{x}_t$ for a given ι). Our hierarchical approach addresses this by modularizing the problem. A meta-bandit algorithm makes high-level strategic decisions about ι , and for each choice, it deploys a base-level bandit algorithm as a subroutine to execute the operational task. The HBL algorithm is structured as follows:

Parameter Discretization. We first discretize the continuous parameter space $(0, 1)$ into a finite set of M representative values, $\mathcal{I} = \{\iota_1, \dots, \iota_M\}$. Each $\iota_j \in \mathcal{I}$ is now a meta-arm.

Episodic Learning. The total horizon T is divided into K episodes, each of length L (so $T = K \cdot L$). In each episode $k \in \{1, \dots, K\}$, the meta-bandit selects one meta-arm ι_{j_k} to play.

Meta-Arm Selection (UCB1). The meta-bandit employs a UCB1 strategy to select arms. At the start of episode k , it chooses the meta-arm ι_{j_k} that maximizes an upper confidence bound on the expected net revenue: $j_k = \arg \max_{j \in \{1, \dots, M\}} \left(\bar{J}_j(k-1) + \sqrt{\frac{2 \ln k}{N_j(k-1)}} \right)$, where $\bar{J}_j(k-1)$ is the average net revenue observed from meta-arm j in previous episodes, and $N_j(k-1)$ is the number of times it has been played.

Inner-Loop Execution. Playing the chosen meta-arm ι_{j_k} involves running a base-level online algorithm (e.g., the SCB-EK algorithm) for L rounds with the reinvestment parameter set to ι_{j_k} .

Observe Meta-Reward. At the end of the L rounds, the total net revenue for the episode, $\sum_{t=1}^L (1 - \iota_{j_k}) r^t$, is observed. This value serves as a noisy sample of the *reward* for the chosen meta-arm. The statistics $\bar{J}_{j_k}$ and N_{j_k} are then updated.

This design effectively addresses the black-box optimization challenge by transforming it into a well-understood multi-armed bandit problem, allowing for provably efficient exploration of the parameter space. The full pseudocode is provided in Algorithm 4.

4.3 Convergence Analysis

We provide a theoretical guarantee that HBL efficiently converges to a near-optimal reinvestment parameter. The performance is measured by the meta-regret, which is the cumulative loss in net revenue compared to an oracle that knew the optimal parameter $\iota^* \in \mathcal{I}$.

THEOREM 4.2 (META-REGRET OF HBL). *Assume the base-level algorithm has an expected regret bound of $\text{Reg}_{base}(L) = \tilde{O}(\sqrt{L})$. The total expected meta-regret of the HBL algorithm after T rounds is*

Algorithm 4 Hierarchical Bandit Learning (HBL)

```

1: Initialize: Discretized parameter set  $\mathcal{I} = \{\iota_1, \dots, \iota_M\}$ . Episode length  $L$ .
2: For each  $j \in \{1, \dots, M\}$ , initialize observed net revenue sum  $\mathcal{J}_j \leftarrow 0$ , number of plays  $N_j \leftarrow 0$ .
3: for episode  $k = 1, 2, \dots, T/L$  do
4:   Meta-Bandit Level: Select a parameter
5:   Select  $j_k \leftarrow \arg \max_{j \in \{1, \dots, M\}} \left( \frac{\mathcal{J}_j}{N_j} + \sqrt{\frac{2 \ln k}{N_j}} \right)$  (if  $N_j = 0$ , play arm  $j$ ).
6:   Let chosen parameter be  $\iota_{j_k}$ .
7:   Base-Bandit Level: Execute for one episode
8:   Initialize a base-level online algorithm  $\pi_{base}$  (e.g., SCB-EK).
9:    $J_{episode} \leftarrow 0$ .
10:  for  $t = 1, 2, \dots, L$  do
11:    Using  $\pi_{base}$  with  $\iota_{j_k}$ , select action  $\mathbf{x}_t$ .
12:    Observe total reward  $r^t$ .
13:     $J_{episode} \leftarrow J_{episode} + (1 - \iota_{j_k}) r^t$ .
14:    Update  $\pi_{base}$  with observations.
15:  end for
16:  Meta-Bandit Level: Update statistics
17:   $N_{j_k} \leftarrow N_{j_k} + 1$ .
18:   $\mathcal{J}_{j_k} \leftarrow \mathcal{J}_{j_k} + J_{episode}$ .
19: end for
```

bounded by:

$$\text{Meta-Regret}(T) = \tilde{O}(\sqrt{MTL}) + \tilde{O}(T/\sqrt{L})$$

where $M = |\mathcal{I}|$, L is the episode length, and $\tilde{O}(\cdot)$ hides polylogarithmic factors. For an optimal choice of $L \approx \sqrt{T/M}$, this yields a total regret of $\tilde{O}(M^{1/4} T^{3/4})$.

The proof, detailed in Appendix C, decomposes the total regret into the regret from the meta-bandit's exploration and the cumulative regret from the inner-loop base-level algorithm. This theorem demonstrates that the HBL framework provides a principled and provably efficient solution to the bi-level optimization problem, converging to a near-optimal strategic policy.

5 Empirical Evaluation

To validate our theoretical guarantees and assess the practical effectiveness of the proposed algorithms, we conduct a comprehensive empirical study. We first evaluate our framework in a controlled, synthetic environment to precisely measure performance against clairvoyant oracles. We then conduct a detailed ablation study on our main risk-aware algorithm to analyze its core components. Finally, we validate our findings using a simulation environment derived from real-world e-commerce data and report the results of a live online experiment.

5.1 Experimental Setup

Simulation Environment. We construct a detailed synthetic environment that captures the essential dynamics of multi-channel resource allocation while maintaining controlled conditions for rigorous evaluation. (1) **Environment Architecture:** The system comprises $N = 10$ distinct marketing channels, each with

Table 1: Performance comparison in the simulation studies.

Config	Algorithm	Reward	Cost	Budget Utilization	ROI	Violation Rate (%)	Net Profit	Regret
Endogenous	SCB-EK	108676	21785	1.002	4.989	48.16	86891	1556
	SCB-K	108445	21721	0.999	4.993	-	86724	1061
	OPT-EK	110232	21947	0.995	5.023	51.94	88285	0
	OPT-FB	109506	21730	0.999	5.039	-	87776	0
	PPO	108925	21754	0.998	5.007	49.93	87171	581
Per-Round	RAPD	91704	16303	0.889	5.625	0	75401	7164
	SCB-K	97845	18267	0.996	5.356	-	79578	907
	OPT-PR	98868	18704	0.946	5.286	0	80164	0
	OPT-FB	98752	18359	0.999	5.379	-	80393	0
	PPO	95038	17429	0.917	5.453	14.2	77609	3830
Reinvestment	HBL	106113	20999	0.995	5.053	41.09	85113	3172
	OPT-HBL-EK	110232	21947	0.995	5.023	51.94	88285	0

unique performance characteristics to model real-world heterogeneity. Experiments are conducted over a long time horizon of $T = 10,000$ rounds to capture long-term dynamics. At each round, a 10-dimensional contextual vector is observed. (2) **Action Space:** Each channel possesses a discrete action space of $|X| = 20$ effort levels. A critical combinatorial constraint is imposed: the sum of effort levels across all channels must not exceed a global limit, i.e., $\sum_{n \in \mathcal{N}} x_n \leq 1$, simulating realistic resource allocation scenarios. A *null* super-arm (zero allocation to all channels) is available, yielding zero cost and a baseline reward. (3) **Reward and Cost Functions:** To create a non-trivial optimization landscape, we employ concave reward and convex cost functions. Each channel n is assigned unique efficiency multipliers (e.g., 0.8, 1.0, 1.2) and cost factors (e.g., 2.5, 3.75, 5.0) to model differential performance and diminishing returns on investment. (4) **Stochasticity:** In each round τ , the observed rewards $r_{n,\tau}$ and costs $c_{n,\tau}$ incorporate multiplicative noise drawn from a normal distribution with a controlled standard deviation ($\sigma = 0.1$), simulating real-world uncertainty. All offline experimental results are averaged over 20 independent runs to ensure statistical robustness.

Evaluation Metrics. The following metrics are used, where π denotes the policy being evaluated and π^* denotes the corresponding offline optimal policy. (1) **Reward:** The total reward $R_\pi^T = \sum_{\tau \in \mathcal{T}} r^\tau$ accumulated by the policy, reflecting gross performance. (2) **Cost:** The total cost $C_\pi^T = \sum_{\tau \in \mathcal{T}} c^\tau$ incurred by the policy, reflecting the total investment. (3) **Budget Utilization:** The ratio of total cost to total generated budget, C_π^T/B_π^T , indicating fiscal efficiency. A value near 1.0 suggests the policy effectively uses its available resources. (4) **Return on Investment (ROI):** The ratio of total reward to total cost, R_π^T/C_π^T . (5) **Violation Rate (%):** For endogenous budget models, this is the fraction of rounds in which the cumulative cost exceeds the cumulative budget, i.e., the empirical probability of the event $C^\tau > B^\tau$. This metric is not applicable to fixed-budget models. (6) **Net Profit:** The total reward less the total cost, $R_\pi^T - C_\pi^T$. (7) **Regret:** The performance gap relative to the oracle, defined as $R_{\pi^*}^T - R_\pi^T$. For the bi-level optimization problem, regret is measured on net profit.

5.2 Performance Analysis

In Table 1, OPT-EK and OPT-PR are oracles for the endogenous budget problems with expected and per-round constraints, respectively. OPT-FB is the fixed-budget oracle. OPT-HBL-EK is the oracle for the bi-level optimization problem. To provide a comprehensive benchmark, we include the Proximal Policy Optimization (PPO) algorithm, trained to convergence in the simulation environment described in Section 5.1. This comparison is not entirely fair because PPO learns its policy within a known environment and can be trained repeatedly. In contrast, our proposed bandit algorithms learn their policies simultaneously with the environment, eliminating the need for repeated training. We include PPO as a baseline to demonstrate the best performance achievable by state-of-the-art reinforcement learning algorithms in a known environment.

Endogenous Budgets with Endogenous Constraints. Our first experiment evaluates performance in the BwEK setting (Section 3.2), where the budget is endogenously generated and the constraint must hold only in expectation. We set the reinvestment rate to a representative $\iota = 0.2$. We compare our specialized algorithm, SCB-EK, against two key benchmarks: the clairvoyant offline optimal policy for this setting (OPT-EK) and a standard fixed-budget algorithm, SCB-K. To ensure a fair comparison, the fixed-budget algorithm and its corresponding oracle (OPT-FB) are allocated a total budget equal to the budget endogenously generated by our SCB-EK algorithm. The results, presented in the *Endogenous* section of Table 1, are illuminating. SCB-EK significantly outperforms SCB-K in cumulative reward and net profit, demonstrating the inadequacy of fixed-budget models when the underlying budget process is dynamically coupled with performance. The fixed-budget algorithm, blind to the performance-feedback loop, operates more conservatively and fails to capitalize on opportunities as effectively. Notably, SCB-EK achieves its superior performance by strategically overspending in nearly half the rounds (48.16%). This behavior, while permissible in an expected-value setting, underscores the distinct challenge of this problem and highlights why such an algorithm would be entirely unsuitable for settings with hard constraints.

Endogenous Budgets with Per-Round Hard Constraints. Next, we evaluate performance in the more stringent *risk of ruin*

Table 2: Ablation study on the RAPD safety check components: (Cost, Budget, δ_τ). The full RAPD uses $(\hat{c}, \check{b}, \delta_\tau > 0)$.

Config. (cost, budget, δ)	Reward	Cost	Total Rounds
$(\hat{c}, \check{b}, \delta > 0)$	91704	16304	10000
$(\hat{c}, \check{b}, \delta > 0)$	18.89	3.49	3
$(\hat{c}, \check{b}, \delta = 0)$	5.29	1.12	1

setting (Section 3.3), where the process terminates upon the first violation of the cumulative budget constraint. Here, the core challenge shifts from long-term resource management to per-round risk mitigation. We compare our risk-aware algorithm, RAPD, against the fixed-budget SCB-K and their respective oracles (OPT-PR and OPT-FB). The fixed-budget models are given a total budget equal to that realized by RAPD. Table 1 (*Per-Round* section) demonstrates the critical importance of our risk-aware design. RAPD successfully navigates the environment, achieving high reward and profit while maintaining a perfect safety record (0% violation rate), empirically confirming our theoretical guarantees. An algorithm like SCB-EK, with its high violation rate, would have terminated almost immediately in this setting. This result clearly demonstrates that managing the risk of ruin requires the specialized, conservative safety mechanism embedded in RAPD. Simply applying an algorithm designed for expected-value constraints is insufficient and leads to catastrophic failure.

Learning the Optimal Reinvestment Strategy. Finally, we evaluate HBL on the bi-level optimization problem of learning the optimal reinvestment parameter ι (Section 4). This addresses the highest-level strategic question: what fraction of revenue should be reinvested to maximize long-term net profit? We discretized the parameter space $\iota \in [0, 1]$ into 11 equally spaced points. By running the OPT-EK oracle for each point, we identified the true optimal parameter as $\iota = 0.2$. As shown in Table 1 (*Reinvestment* section), HBL effectively solves the bi-level problem. It achieves a low meta-regret, converging to a net profit close to that of the clairvoyant oracle that knew the optimal ι from the start. The experiment confirmed that HBL successfully identified $\iota = 0.2$ as the best parameter. This demonstrates that our hierarchical bandit-on-bandit approach can efficiently navigate the complex, black-box objective landscape to find a near-optimal strategic policy.

5.3 Ablation Study of the RAPD Algorithm

To analyze the contributions of the core components of our most novel algorithm, RAPD, we conduct a targeted ablation study. The key insight of RAPD is its safety check: an action is deemed safe only if an upper confidence bound on its cumulative cost, $C^{\tau-1} + \hat{c}^\tau(x)$, remains below a risk-adjusted lower confidence bound on its cumulative budget, $(1 - \delta_\tau)(B^{\tau-1} + \check{b}^\tau(x))$. We investigate how performance changes when the pessimistic LCB on the budget ($\check{b}$) is replaced with a neutral empirical mean ($\bar{b}$) or an optimistic UCB ($\hat{b}$), and when the explicit risk-tolerance parameter (δ_τ) is removed ($\delta_\tau = 0$). Performance is measured by reward, cost, and the number of rounds completed before a budget violation forces termination.

The results, presented in Table 2, are unequivocal. The full RAPD configuration is the only one to survive the entire 10,000-round horizon. Any deviation from its conservative design leads to catastrophic failure. Specifically:

Removing Pessimism on Budget: Replacing the LCB of the budget ($\check{b}$) with the empirical mean ($\bar{b}$) causes the algorithm to overestimate its future resources. This leads to overly aggressive actions and a budget violation within just 3 rounds.

Removing Risk Tolerance: Removing the safety margin δ_τ (i.e., setting $\delta = 0$), even while using the LCB for the budget, results in immediate termination after a single round.

This study empirically validates the central thesis of our risk-aware design: the *double pessimism* of using a UCB for cost and an LCB for the budget, combined with an explicit safety margin, are critical and non-negotiable components for achieving both safety and efficiency under hard, path-dependent endogenous constraints.

5.4 Validation on Real-World and Live Data

To bridge the gap between simulation and practice, we conducted two final validation steps. First, we constructed a high-fidelity simulation environment using a 360-day dataset from a large e-commerce platform, containing performance metrics for two marketing channels (paid search and promotional coupons) under various investment levels. Using this data to learn the environment dynamics, we re-ran our experiments. The results closely mirrored our synthetic findings, with RAPD again demonstrating a superior ability to maximize rewards while adhering to per-round constraints, confirming the robustness of our conclusions.

Second, we deployed our framework in a live A/B experiment within a production environment. A small fraction (5%) of user traffic was directed to a test group where budgets were managed by the RAPD algorithm, while the control group used the platform's existing uplift-based allocation strategy. Over a two-week period, the RAPD-managed portfolio achieved a statistically significant increase in revenue and significantly reduced the incidence of daily budget over-expenditure compared to the control group. These results from a live setting provide strong evidence for the practical value and real-world impact of our proposed framework.

6 Conclusion

This paper introduced the *Marketing Hosting*, a hierarchical set of problems for sequential budget allocation that generalizes existing bandit models by incorporating endogenous, performance-dependent constraints. We provided novel formulations and algorithms for each setting with rigorous theoretical guarantees.

Acknowledgments

This work was supported by National Natural Science Foundation of China (No.62472428), Public Computing Cloud, Renmin University of China, the fund for building world-class universities (disciplines) of Renmin University of China and Alibaba Group through the Alibaba Research Intern Program.

References

- [1] Vahideh Sadat Abedi. 2017. Allocation of advertising budget between multiple channels to support sales in multiple markets. *Journal of the Operational Research Society* 68, 2 (2017), 134–146.
- [2] Shipra Agrawal and Nikhil Devanur. 2016. Linear contextual bandits with knapsacks. *Advances in neural information processing systems* 29 (2016).
- [3] Shipra Agrawal and Nikhil R Devanur. 2014. Bandits with concave rewards and convex knapsacks. In *Proceedings of the fifteenth ACM conference on Economics and computation*. 989–1006.
- [4] Noga Alon, Iftah Gamzu, and Moshe Tennenholtz. 2012. Optimizing budget allocation among channels and influencers. In *Proceedings of the 21st international conference on World Wide Web*. 381–388.
- [5] Peter Auer, Nicolo Cesa-Bianchi, Yoav Freund, and Robert E Schapire. 2002. The nonstochastic multiarmed bandit problem. *SIAM journal on computing* 32, 1 (2002), 48–77.
- [6] Ashwinkumar Badanidiyuru, Robert Kleinberg, and Aleksandr Slivkins. 2018. Bandits with knapsacks. *Journal of the ACM (JACM)* 65, 3 (2018), 1–55.
- [7] Lixing Chen, Jie Xu, and Zhuo Lu. 2018. Contextual combinatorial multi-armed bandits with volatile arms and submodular reward. *Advances in Neural Information Processing Systems* 31 (2018).
- [8] Zhijian Duan, Yusen Huo, Tianyu Wang, Zhilin Zhang, Yeshu Li, Chuan Yu, Jian Xu, Bo Zheng, and Xiaotie Deng. 2025. An Adaptable Budget Planner for Enhancing Budget-Constrained Auto-Bidding in Online Advertising. *arXiv preprint arXiv:2502.05187* (2025).
- [9] Arthur Flajolet and Patrick Jaillet. 2015. Logarithmic regret bounds for bandits with knapsacks. *arXiv preprint arXiv:1510.01800* (2015).
- [10] Briti Gangopadhyay, Zhao Wang, Alberto Silvio Chiappa, and Shingo Takamatsu. 2025. Adaptive Budget Optimization for Multichannel Advertising Using Combinatorial Bandits. *arXiv preprint arXiv:2502.02920* (2025).
- [11] Lin Ge, Yang Xu, Jianing Chu, David Cramer, Fuhong Li, Kelly Paulson, and Rui Song. 2025. Multi-task combinatorial bandits for budget allocation. In *Proceedings of the 31st ACM SIGKDD Conference on Knowledge Discovery and Data Mining V. 1*. 2247–2258.
- [12] Sahin Cem Geyik, Abhishek Saxena, and Ali Dasdan. 2014. Multi-touch attribution based budget allocation in online advertising. In *Proceedings of the eighth international workshop on data mining for online advertising*. 1–9.
- [13] Hengquan Guo and Xin Liu. 2025. On stochastic contextual bandits with knapsacks in small budget regime. In *The Thirteenth International Conference on Learning Representations*.
- [14] Benjamin Han and Jared Gabor. 2020. Contextual bandits for advertising budget allocation. *Proceedings of the ADKDD* 17 (2020), 2.
- [15] Randall Lewis and Jeffrey Wong. 2022. Incrementality bidding and attribution. *arXiv preprint arXiv:2208.12809* (2022).
- [16] Mengfan Li, Jian Zhang, Roohallah Alizadehsani, and Pawel Plawiak. 2024. A multi-channel advertising budget allocation using reinforcement learning and an improved differential evolution algorithm. *IEEE Access* (2024).
- [17] Pengcheng Li, Ammar Hawbani, et al. 2018. An efficient budget allocation algorithm for multi-channel advertising. In *2018 24th International Conference on Pattern Recognition (ICPR)*. IEEE, 886–891.
- [18] Shuai Li, Baoxiang Wang, Shengyu Zhang, and Wei Chen. 2016. Contextual combinatorial cascading bandits. In *International conference on machine learning*. PMLR, 1245–1253.
- [19] Xutong Liu, Jinhang Zuo, Siwei Wang, John CS Lui, Mohammad Hajiesmaili, Adam Wierman, and Wei Chen. 2023. Contextual combinatorial bandits with probabilistically triggered arms. In *International Conference on Machine Learning*. PMLR, 22559–22593.
- [20] Alessandro Nuara, Francesco Trovo, Nicola Gatti, and Marcello Restelli. 2018. A combinatorial-bandit algorithm for the online joint bid/budget optimization of pay-per-click advertising campaigns. In *Proceedings of the AAAI Conference on Artificial Intelligence*, Vol. 32.
- [21] Alessandro Nuara, Francesco Trovo, Nicola Gatti, and Marcello Restelli. 2022. Online joint bid/daily budget optimization of internet advertising campaigns. *Artificial Intelligence* 305 (2022), 103663.
- [22] Guangyuan Shen, Shenjie Sun, Dehong Gao, Shaolei Li, Libin Yang, Yongping Shi, and Wei Ning. 2023. Cross-channel Budget Coordination for Online Advertising System. *arXiv preprint arXiv:2305.06883* (2023).
- [23] Tatsuhiko Shimizu, Koichi Tanaka, Ren Kishimoto, Haruka Kiyohara, Masahiro Nomura, and Yuta Saito. 2024. Effective Off-Policy Evaluation and Learning in Contextual Combinatorial Bandits. In *Proceedings of the 18th ACM Conference on Recommender Systems*. 733–741.
- [24] Stamatis Thomaidou, Kyriakos Liakopoulos, and Michalis Vazirgiannis. 2014. Toward an integrated framework for automated development and optimization of online advertising campaigns. *Intelligent Data Analysis* 18, 6 (2014), 1199–1227.
- [25] Bingkun Wang and Pourya Zareehemat. 2025. Multi-channel advertising budget allocation: A novel method using Q-learning and mutual learning-based artificial bee colony. *Expert Systems With Applications* 271 (2025), 126649.
- [26] Haifeng Zhang and Yevgeniy Vorobeychik. 2017. Multi-channel marketing with budget complementarities. In *International Conference on Autonomous Agents and Multiagent Systems*.
- [27] Weinan Zhang, Ying Zhang, Bin Gao, Yong Yu, Xiaojie Yuan, and Tie-Yan Liu. 2012. Joint optimization of bid and budget allocation in sponsored search. In *Proceedings of the 18th ACM SIGKDD international conference on Knowledge discovery and data mining*. 1177–1185.
- [28] Xilin Zhang and Wang Chi Cheung. 2024. Piecewise-stationary bandits with knapsacks. *Advances in Neural Information Processing Systems* 37 (2024), 56001–56040.
- [29] Xiaojin Zhang, Shuai Li, Weiwen Liu, and Shengyu Zhang. 2019. Contextual combinatorial conservative bandits. *arXiv preprint arXiv:1911.11337* (2019).
- [30] Xiaoxi Zhang, Jinhang Zuo, Zhe Huang, Zhi Zhou, Xu Chen, and Carlee Joe-Wong. 2023. Learning with side information: Elastic multi-resource control for the open RAN. *IEEE Journal on Selected Areas in Communications* 42, 2 (2023), 295–309.
- [31] Jinhang Zuo and Carlee Joe-Wong. 2021. Combinatorial multi-armed bandits for resource allocation. In *2021 55th Annual Conference on Information Sciences and Systems (CISS)*. IEEE, 1–4.

A Omitted Details for Section 2

A.1 Proof of Lemma 2.3

PROOF. The proof employs a martingale-based argument, a standard technique for establishing optimality bounds in online settings with long-term constraints. We begin by formulating the dual of the linear program defined in problem (P). Let $\lambda \geq 0$ be the dual variable associated with the budget constraint B , and let $\alpha \geq 0$ be the dual variable for the time horizon constraint T .

Definition A.1 (The Dual LP). The dual of the LP benchmark is:

$$\begin{aligned} & \text{minimize}_{\lambda, \alpha} \quad B\lambda + T\alpha \\ & \text{subject to} \quad \lambda \left(\sum_{n \in N} \mu_{c,n,s_n} \right) + \alpha \geq \sum_{n \in N} \mu_{r,n,s_n}, \quad \forall s \in S, \\ & \quad \quad \quad \lambda \geq 0, \quad \alpha \geq 0. \end{aligned} \tag{D}$$

Let (λ^*, α^*) be an optimal solution to the dual program. By strong duality for linear programs, the optimal value of the primal (P) is equal to the optimal value of this dual program: $OPT_{LP} = B\lambda^* + T\alpha^*$. We can interpret λ^* as an optimal "shadow price" for consuming the budget and α^* as the optimal "shadow price" of time.

Let π be any feasible adaptive online policy. Let $\mathcal{F}_{\tau-1}$ be the sigma-algebra representing the history of contexts, actions, and observations up to the beginning of round τ . Let T_{stop} be the random stopping time when the policy halts (either because a budget is exhausted or T is reached). The total expected reward of this policy is $\mathbb{E}_{\pi}[R^{T_{stop}}]$.

We define a stochastic process $\{Z_{\tau}\}_{\tau=0}^{T_{stop}}$. Let the initial value be the optimal value of the dual objective: $Z_0 = B\lambda^* + T\alpha^*$. For any $\tau \geq 1$, we define Z_{τ} as the sum of the cumulative reward obtained up to round $\tau - 1$ and the value of the remaining resources (budget and time) at the start of round τ , priced at their optimal dual values:

$$Z_{\tau} = R^{\tau-1} + (B - C^{\tau-1})\lambda^* + (T - (\tau - 1))\alpha^*.$$

We now demonstrate that $\{Z_{\tau}\}$ is a supermartingale with respect to the filtration $\{\mathcal{F}_{\tau-1}\}$, i.e., $\mathbb{E}[Z_{\tau+1}|\mathcal{F}_{\tau-1}] \leq Z_{\tau}$. Consider the conditional expectation of the one-step change:

$$\begin{aligned} \mathbb{E}[Z_{\tau+1} - Z_{\tau}|\mathcal{F}_{\tau-1}] &= \mathbb{E}[r_{\tau} - c^{\tau}\lambda^* - \alpha^*|\mathcal{F}_{\tau-1}] \\ &= \mathbb{E}\left[\sum_{n=1}^N r_{n,\tau} - \lambda^* \sum_{n=1}^N c_{n,\tau} - \alpha^* \middle| \mathcal{F}_{\tau-1}\right]. \end{aligned}$$

Let $\mathbf{x}_\tau$ be the super-arm chosen by policy π at round τ . The choice of $\mathbf{x}_\tau$ is $\mathcal{F}_{\tau-1}$ -measurable. By linearity of expectation and the definition of the expected rewards and costs:

$$\mathbb{E}[Z_{\tau+1} - Z_\tau | \mathcal{F}_{\tau-1}] = \sum_{n=1}^N \mu_{r,n,\mathbf{x}_{n,\tau}} - \lambda^* \sum_{n=1}^N \mu_{c,n,\mathbf{x}_{n,\tau}} - \alpha^*.$$

The dual feasibility constraint (D) states that for any super-arm $\mathbf{s} \in \mathcal{S}$, we have $\sum_n \mu_{r,n,\mathbf{s}_n} - \lambda^* \sum_n \mu_{c,n,\mathbf{s}_n} - \alpha^* \leq 0$. Since $\mathbf{x}_\tau \in \mathcal{S}$, it follows that: $\mathbb{E}[Z_{\tau+1} - Z_\tau | \mathcal{F}_{\tau-1}] \leq 0$.

Thus, $\{Z_\tau\}$ is a supermartingale. By Doob's Optional Stopping Theorem, for any stopping time $T_{stop} \leq T$, we have $\mathbb{E}[Z_{T_{stop}}] \leq Z_0$.

At the stopping time T_{stop} , the policy terminates. The value of the remaining budget and time is non-negative.

$$Z_{T_{stop}} = R^{T_{stop}-1} + (B - C^{T_{stop}-1})\lambda^* + (T - (T_{stop} - 1))\alpha^* \geq R^{T_{stop}-1}.$$

Taking expectations over the policy and environment randomness:

$$\mathbb{E}[R^{T_{stop}-1}] \leq \mathbb{E}[Z_{T_{stop}}] \leq Z_0 = OPT_{LP}.$$

Since the reward in the final round T_{stop} is not counted in the regret definition for BwK, the total reward is $R^{T_{stop}-1}$. This holds for any policy π , including the optimal one. Therefore, OPT_{LP} is a valid upper bound on the total expected reward of any online policy. $\square$

A.2 Proof of Lemma 2.4

PROOF. Let $\mathbf{x}_n$ be a suboptimal base arm and $\mathbf{x}'_n$ be a better base arm for the same channel n , such that $\Delta = \mu_{r,n,\mathbf{x}'} - \mu_{r,n,\mathbf{x}} > 0$.

An algorithm based on the principle of optimism in the face of uncertainty will select a super-arm $\mathbf{x}_\tau$ that maximizes an optimistic objective. For the specific case in the lemma statement, this objective is $\sum_{j=1}^N \hat{r}_{j,\tau}(\mathbf{x}_j)$.

Suppose at round τ , the algorithm selects a super-arm $\mathbf{x}_\tau$ that includes the suboptimal base arm $\mathbf{x}_{n,\tau} = \mathbf{x}$. Let us construct an alternative super-arm $\mathbf{x}'_\tau$ which is identical to $\mathbf{x}_\tau$ except that the action for channel n is replaced by the better base arm $\mathbf{x}'$, i.e., $\mathbf{x}'_{n,\tau} = \mathbf{x}'$ and $\mathbf{x}'_{j,\tau} = \mathbf{x}_{j,\tau}$ for all $j \neq n$. Since the algorithm chose $\mathbf{x}_\tau$, its optimistic value must be at least as high as that of $\mathbf{x}'_\tau$:

$$\sum_{j=1}^N \hat{r}_{j,\tau}(\mathbf{x}_{j,\tau}) \geq \sum_{j=1}^N \hat{r}_{j,\tau}(\mathbf{x}'_{j,\tau}).$$

Substituting the components of the super-arms, we get:

$$\hat{r}_{n,\tau}(\mathbf{x}) + \sum_{j \neq n} \hat{r}_{j,\tau}(\mathbf{x}_{j,\tau}) \geq \hat{r}_{n,\tau}(\mathbf{x}') + \sum_{j \neq n} \hat{r}_{j,\tau}(\mathbf{x}_{j,\tau}).$$

This simplifies to the condition that the suboptimal base arm must have a higher UCB than the better one: $\hat{r}_{n,\tau}(\mathbf{x}) \geq \hat{r}_{n,\tau}(\mathbf{x}')$.

Let us define a high-probability "good event" $\mathcal{E}$ where all confidence bounds hold for all arms, channels, and rounds, i.e., $|\bar{r}_{j,y}(\tau) - \mu_{r,j,y}| \leq \rho_{j,y}(\tau)$ for all (j, y, τ) . This event holds with probability at least $1 - O(T^{-c})$ for some $c > 1$ by a standard union bound argument over Hoeffding's inequality. On this event, we have:

$$\hat{r}_{n,\tau}(\mathbf{x}) \leq \mu_{r,n,\mathbf{x}} + 2\rho_{n,\mathbf{x}}(\tau), \quad \hat{r}_{n,\tau}(\mathbf{x}') \geq \mu_{r,n,\mathbf{x}'}.$$

Combining these with the selection condition $\hat{r}_{n,\tau}(\mathbf{x}) \geq \hat{r}_{n,\tau}(\mathbf{x}')$ gives: $\mu_{r,n,\mathbf{x}} + 2\rho_{n,\mathbf{x}}(\tau) \geq \mu_{r,n,\mathbf{x}'}$.

Rearranging this inequality, we have: $2\rho_{n,\mathbf{x}}(\tau) \geq \mu_{r,n,\mathbf{x}'} - \mu_{r,n,\mathbf{x}} = \Delta$.

Now, we substitute the definition of the confidence radius $\rho_{n,\mathbf{x}}(\tau) = \sqrt{\frac{\ln(1/\delta')}{2N_{n,\mathbf{x}}(\tau)}} : 2\sqrt{\frac{\ln(1/\delta')}{2N_{n,\mathbf{x}}(\tau)}} \geq \Delta$.

Squaring both sides and rearranging for $N_{n,\mathbf{x}}(\tau)$ yields:

$$\frac{4\ln(1/\delta')}{2N_{n,\mathbf{x}}(\tau)} \geq \Delta^2 \implies N_{n,\mathbf{x}}(\tau) \leq \frac{2\ln(1/\delta')}{\Delta^2}.$$

Let $C_0 = \frac{2\ln(1/\delta')}{\Delta^2}$. This shows that, on the good event $\mathcal{E}$, the sub-optimal arm $\mathbf{x}_n$ can be chosen at most C_0 times for this reason. The total number of plays $N_{n,\mathbf{x}}(T)$ can be decomposed into an initial exploration phase (e.g., playing each arm once) and the subsequent choices. The number of optimistic plays is bounded by C_0 . Thus, the total number of plays is bounded by $1 + C_0 = O(\ln(T)/\Delta^2)$ when we set $\delta' = 1/T^k$. Taking the expectation over the (low-probability) bad event adds a small constant term. Therefore, the expected number of pulls is $\mathbb{E}[N_{n,\mathbf{x}}(T)] = \tilde{O}(1/\Delta^2)$. $\square$

B Omitted Details for Section 3

B.1 Proof of Theorem 3.2

PROOF. We condition the analysis on a high-probability "good event" $\mathcal{E}$, where all confidence bounds hold for all arms and all rounds, i.e., $|\bar{r}_{n,\mathbf{x}}(\tau) - \mu_{r,n,\mathbf{x}}| \leq \rho_{n,\mathbf{x}}(\tau)$ and similarly for costs. By a union bound over all arms and time steps, $\mathbb{P}(\mathcal{E}) \geq 1 - O(T^{-2})$.

The regret is decomposed as follows:

$$\begin{aligned} OPT_{LP} - \mathbb{E}[R^T] &= \mathbb{E} \left[\sum_{\tau=1}^T \left(\frac{OPT_{LP}}{T} - r^\tau \right) \right] \\ &\leq \mathbb{E} \left[\sum_{\tau=1}^T (\hat{r}^\tau(\mathbf{x}_\tau) - r^\tau) \right] + \mathbb{E} \left[\sum_{\tau=1}^T \left(\frac{OPT_{LP}}{T} - \hat{r}^\tau(\mathbf{x}_\tau) \right) \right] \end{aligned}$$

The first term, $\sum \mathbb{E}[\hat{r}^\tau(\mathbf{x}_\tau) - r^\tau]$, is the sum of estimation errors. It can be bounded by the sum of confidence radii, which, by a standard sample complexity argument, is $\tilde{O}(\sqrt{N|X|T})$.

For the second term, we use a potential function argument on the dual variable. Let $\Phi_\tau = \frac{1}{2\eta}(\lambda_\tau - \lambda^*)^2$.

$$\begin{aligned} \mathbb{E}[\Phi_{\tau+1} - \Phi_\tau | \mathcal{F}_{\tau-1}] &\leq \mathbb{E}[(\lambda_\tau - \lambda^*)(c^\tau - B/T) + \frac{\eta}{2}(c^\tau - B/T)^2] \\ &\leq \mathbb{E}[\lambda_\tau(c^\tau - B/T)] - \mathbb{E}[\lambda^*(c^\tau - B/T)] + O(\eta). \end{aligned}$$

The algorithm's choice of $\mathbf{x}_\tau$ maximizes the optimistic Lagrangian. On event $\mathcal{E}$, $\hat{r}^\tau(\mathbf{x}_\tau) - \lambda_\tau \bar{c}^\tau(\mathbf{x}_\tau) \geq \hat{r}^\tau(\mathbf{x}^*) - \lambda_\tau \bar{c}^\tau(\mathbf{x}^*)$, where $\mathbf{x}^*$ is the optimal offline solution. Rearranging and using the properties of confidence bounds, we get an upper bound on $\lambda_\tau c^\tau$. By summing over τ , telescoping the potential function, and combining the terms, we show that the second regret term is also controlled, leading to the final bound $Regret(T) = \tilde{O}(\sqrt{N|X|T})$. The contribution from the bad event $\neg\mathcal{E}$ is at most $T \cdot \mathbb{P}(\neg\mathcal{E}) = O(1/T)$, which is negligible. $\square$

B.2 Proof of Theorem 3.4

PROOF. The analysis extends the primal-dual framework to the endogenous setting. The key difficulty is that the offline benchmark is no longer a simple LP, and the budget's dependence on rewards must be carefully handled.

The benchmark is the solution to the following problem:

$$\begin{aligned} \text{maximize}_{\xi} \quad & R^* = T \sum_{n,x} \xi_{n,x} \mu_{r,n,x} \\ \text{s.t.} \quad & T \sum_{n,x} \xi_{n,x} \mu_{c,n,x} \leq \iota R^*, \\ & \sum_{n,x} \xi_{n,x} \leq 1, \quad \xi_{n,x} \geq 0. \end{aligned}$$

Let OPT_{EK} be the optimal value.

The proof structure follows that of Theorem 3.2, but with a modified Lagrangian and potential function analysis. We again condition on the good event $\mathcal{E}$. The algorithm maximizes the optimistic Lagrangian:

$$\mathcal{L}_\tau(x, \lambda_\tau) = \sum_n \hat{r}_{n,\tau}(x_n) - \lambda_\tau \left(\sum_n \check{c}_{n,\tau}(x_n) - \iota \sum_n \check{r}_{n,\tau}(x_n) \right).$$

The regret decomposition now involves the difference between the algorithm's reward and the optimal reward OPT_{EK}/T . The crucial step is to show that the choice of LCB for the budget generation term $\check{r}$ provides sufficient pessimism. On the good event $\mathcal{E}$, we have:

$$\sum_n \mu_{r,n,x_{n,\tau}} \leq \sum_n \hat{r}_{n,\tau}(x_{n,\tau})$$

and for the optimal policy x^* ,

$$\begin{aligned} & \sum_n \mu_{r,n,x_n^*} - \lambda^* \left(\sum_n \mu_{c,n,x_n^*} - \iota \sum_n \mu_{r,n,x_n^*} \right) \\ & \leq \sum_n \hat{r}_{n,\tau}(x_n^*) - \lambda_\tau \left(\sum_n \check{c}_{n,\tau}(x_n^*) - \iota \sum_n \check{r}_{n,\tau}(x_n^*) \right) \end{aligned}$$

where λ^* is the optimal dual price for the endogenous problem. By the algorithm's choice of x_τ , we have:

$$\begin{aligned} & \sum_n \hat{r}_{n,\tau}(x_\tau) - \lambda_\tau \left(\sum_n \check{c}_{n,\tau}(x_\tau) - \iota \sum_n \check{r}_{n,\tau}(x_\tau) \right) \\ & \geq \sum_n \hat{r}_{n,\tau}(x^*) - \lambda_\tau \left(\sum_n \check{c}_{n,\tau}(x^*) - \iota \sum_n \check{r}_{n,\tau}(x^*) \right) \end{aligned}$$

Combining these inequalities and summing over τ allows us to bound the total reward in terms of OPT_{EK} , the sum of confidence radii (estimation error), and a term related to the dual updates. The dual update uses the realized net consumption $c^\tau - b^\tau$. The potential function analysis on $\Phi_\tau = (\lambda_\tau - \lambda^*)^2$ is more complex but follows a similar logic, showing that the cumulative expected constraint violation $\sum \mathbb{E}[c^\tau - b^\tau]$ is controlled. This again leads to an overall regret of $O(\sqrt{T})$. $\square$

B.3 Proof of Theorem 3.6

LEMMA B.1 (SAFETY OF THE OPTIMAL POLICY). *Conditioned on the good event $\mathcal{E}$, the optimal super-arm for the unconstrained problem, x^* , is an element of the safe set $\mathcal{S}_{\text{safe}}(\tau)$ for all $\tau \in \mathcal{T}$ with high probability.*

PROOF. Let $M_\tau = B^\tau - C^\tau$ be the budget surplus at the end of round τ . We need to show that for the optimal policy x^* , the safety condition holds:

$$C^{\tau-1} + \hat{c}^\tau(x^*) \leq (1 - \kappa_\tau)(B^{\tau-1} + \check{b}^\tau(x^*)).$$

This is equivalent to:

$$\hat{c}^\tau(x^*) - \check{b}^\tau(x^*) + \kappa_\tau(B^{\tau-1} + \check{b}^\tau(x^*)) \leq M_{\tau-1}.$$

On the good event $\mathcal{E}$, we have $\hat{c}^\tau(x^*) \geq \mu_c(x^*)$ and $\check{b}^\tau(x^*) \leq \mu_b(x^*)$. The term $\mu_c(x^*) - \mu_b(x^*) \leq 0$ since the optimal policy must be feasible in expectation. The difference $\hat{c}^\tau(x^*) - \check{b}^\tau(x^*)$ is bounded by the sum of confidence radii. The surplus $M_{\tau-1}$ is a sum of random variables $b^t - c^t$. It can be shown that M_τ is a submartingale under the RAPD policy. Using Azuma-Hoeffding inequality for submartingales, we can show that $M_{\tau-1}$ is highly concentrated around its (non-negative) expectation and is unlikely to be substantially negative. The safety margin $\kappa_\tau = O(1/\tau^2)$ provides a buffer that decays over time. A careful combination of these concentration bounds shows that the inequality holds with high probability for any given τ . A union bound over all τ completes the argument that x^* is always in the safe set with high probability. $\square$

PROOF OF THEOREM 3.6. The proof consists of two main parts: a safety analysis to bound the probability of ruin, and a regret analysis conditioned on survival.

Part 1: Safety Analysis. The goal is to prove $\mathbb{P}(\mathcal{T}_{\text{stop}} < T) = O(1/T)$.

Since $\mathcal{S}_{\text{safe}}(\tau)$ contains the optimal policy x^* (and thus is non-empty) with high probability, the algorithm will not be forced to play the null action. A ruinous action that violates the hard constraint can only be taken if the confidence bounds fail in a specific, unfavorable way (i.e., if event $\mathcal{E}$ does not hold). The probability of this is bounded by $\mathbb{P}(\neg\mathcal{E}) = O(T^{-2})$. Taking a union bound over the T rounds, the probability of ruin at any point is $\sum_{\tau=1}^T O(\tau^{-2}) = O(1)$, but a more careful analysis on the sequential nature of the event yields $\mathbb{P}(\mathcal{T}_{\text{stop}} < T) \leq T \cdot \mathbb{P}(\neg\mathcal{E}) = O(1/T)$.

Part 2: Regret Analysis. Conditioned on survival for all T rounds, the regret analysis mirrors that of Theorem 3.4. The key difference is that the optimization at each step is over the restricted set $\mathcal{S}_{\text{safe}}(\tau)$. The regret can be decomposed into two main parts:

$$\begin{aligned} \text{Regret}(T) = & \underbrace{\mathbb{E} \left[\sum_{\tau=1}^T (r(x^*) - r(x_\tau)) \cdot \mathbb{I}(x^* \in \mathcal{S}_{\text{safe}}(\tau)) \right]}_{\text{Regret from learning}} \\ & + \underbrace{\mathbb{E} \left[\sum_{\tau=1}^T r(x^*) \cdot \mathbb{I}(x^* \notin \mathcal{S}_{\text{safe}}(\tau)) \right]}_{\text{Regret from safety restriction}}. \end{aligned}$$

The safety analysis shows that the second term is small, as the event $x^* \notin \mathcal{S}_{\text{safe}}(\tau)$ is rare. The first term is bounded using the same primal-dual potential function argument as in the proof of Theorem 3.4. The logic holds because the analysis only requires that the optimal solution x^* is a feasible candidate for the argmax at each step, which is guaranteed with high probability by the safety analysis. Combining these components yields the final regret bound of $\tilde{O}(\sqrt{N|X|T})$. $\square$

C Proof of Theorem 4.2

Definition C.1 (Meta-Regret). Let $J(\iota)$ be the true expected net revenue per round for a fixed parameter ι . The oracle's total expected net revenue is $T \cdot J(\iota^*)$, where $\iota^* = \arg \max_{\iota \in I} J(\iota)$. The

meta-regret is defined as:

$$\text{Meta-Regret}(T) = TJ(i^*) - \mathbb{E} \left[\sum_{k=1}^K \sum_{\tau=(k-1)L+1}^{kL} (1 - \iota_{j_k}) r^\tau \right]$$

where $K = T/L$ is the number of episodes and ι_{j_k} is the parameter chosen in episode k .

LEMMA C.2 (SAMPLE COMPLEXITY FOR META-BANDIT). *Let $\Delta_i = J(i^*) - J(i)$ be the suboptimality gap for a meta-arm i . The number of times a suboptimal meta-arm i is pulled, $N_i(K)$, is bounded in expectation by $\mathbb{E}[N_i(K)] = O(\ln(K)/\Delta_i^2)$.*

PROOF. This is the standard sample complexity result for the UCB1 algorithm. A suboptimal arm i is played at episode k only if its UCB is higher than the true mean of the optimal arm i^* . This can happen only if the confidence interval for the suboptimal arm is wide enough to mask its inferiority. For the UCB of a Δ_i -suboptimal arm to exceed the true mean of an optimal arm, its confidence radius must be at least $\Delta_i/2$. This implies $\sqrt{2 \ln(k)/N_i(k-1)} \geq \Delta_i/2$, which leads to $N_i(k-1) \leq 8 \ln(k)/\Delta_i^2$. Summing over all episodes gives the result. $\square$

PROOF OF THEOREM 4.2. We analyze the total expected meta-regret, which compares the performance of HBL-Opt to an oracle that uses the best parameter $i^* \in \mathcal{I}$ for all T rounds.

Let $J_{\text{episode}}(k)$ be the total net revenue realized in episode k . The meta-regret can be rewritten and decomposed:

$$\begin{aligned} \text{Meta-Regret}(T) &= TJ(i^*) - \sum_{k=1}^K \mathbb{E}[J_{\text{episode}}(k)] \\ &= L \sum_{k=1}^K \left(J(i^*) - \frac{\mathbb{E}[J_{\text{episode}}(k)]}{L} \right) \\ &= L \sum_{k=1}^K \mathbb{E}[J(i^*) - J(\iota_{j_k})] + L \sum_{k=1}^K \mathbb{E} \left[J(\iota_{j_k}) - \frac{J_{\text{episode}}(k)}{L} \right] \\ &= \underbrace{L \cdot \mathbb{E}[\text{Regret}_{\text{meta}}(K)]}_{\text{Term A: Meta-Bandit Regret}} + \underbrace{L \sum_{k=1}^K \mathbb{E}[\text{Bias}_k]}_{\text{Term B: Cumulative Bias}} \end{aligned}$$

The first term is the expected regret of the meta-bandit (UCB1) scaled by the episode length L . The second term, Term B, is the cumulative bias between the true mean net revenue and the observed average net revenue in an episode.

Bounding Term B: The Estimation Bias. Term B arises because the base-level algorithm does not achieve the optimal reward within an episode; it incurs its own regret. For any episode where parameter i is used, let R_{base}^τ be the reward achieved by the base algorithm at time τ within that episode, and let $R_{\pi^*}^\tau(i)$ be the reward of an optimal policy for that i . The total expected gross reward in the episode is:

$$\mathbb{E} \left[\sum_{\tau=1}^L R_{\text{base}}^\tau \right] = L \cdot \mathbb{E}[R_{\pi^*}(i)] - \text{Regret}_{\text{base}}(L)$$

The expected total net revenue in the episode is:

$$\begin{aligned} \mathbb{E}[J_{\text{episode}}] &= (1 - \iota) \mathbb{E} \left[\sum_{\tau=1}^L R_{\text{base}}^\tau \right] \\ &= (1 - \iota) (L \cdot \mathbb{E}[R_{\pi^*}(i)] - \text{Regret}_{\text{base}}(L)) \\ &= L \cdot J(i) - (1 - \iota) \text{Regret}_{\text{base}}(L) \end{aligned}$$

The per-episode bias is thus $\text{Bias}_k = \mathbb{E}[J(\iota_{j_k}) - J_{\text{episode}}(k)/L] = \frac{(1 - \iota_{j_k}) \text{Regret}_{\text{base}}(L)}{L}$. Since $1 - \iota < 1$ and we assume $\text{Regret}_{\text{base}}(L) = \tilde{O}(\sqrt{L})$:

$$\begin{aligned} |\text{Term B}| &= \left| L \sum_{k=1}^K \mathbb{E} \left[\frac{(1 - \iota_{j_k}) \text{Regret}_{\text{base}}(L)}{L} \right] \right| \\ &\leq K \cdot \text{Regret}_{\text{base}}(L) = K \cdot \tilde{O}(\sqrt{L}) \end{aligned}$$

Substituting $K = T/L$, we get $|\text{Term B}| = \tilde{O}(T/\sqrt{L})$.

Bounding Term A: The Meta-Bandit Regret. The first term is $L \cdot \mathbb{E}[\text{Regret}_{\text{meta}}(K)]$. The meta-bandit is a standard UCB1 algorithm running for $K = T/L$ episodes on M arms.

The standard regret bound for UCB1 over K episodes with M arms, for rewards bounded in $[0, L]$ (since episode rewards are sums over L rounds), is:

$$\mathbb{E}[\text{Regret}_{\text{meta}}(K)] = \tilde{O}(L\sqrt{MK})$$

Scaling this by the episode length L is incorrect; the regret is the sum of gaps, which are already scaled. The regret of the meta-bandit is:

$$\mathbb{E}[\text{Regret}_{\text{meta}}(K)] = \sum_{i \in \mathcal{I}, i \neq i^*} \Delta_i \mathbb{E}[N_i(K)] = \tilde{O}(\sqrt{MK})$$

The total regret from Term A is this meta-regret, but each unit of meta-regret corresponds to one episode of length L . So, the total loss in net revenue is:

$$L \cdot \mathbb{E}[\text{Regret}_{\text{meta}}(K)] = \tilde{O}(L\sqrt{MK}) = \tilde{O}(L\sqrt{M(T/L)}) = \tilde{O}(\sqrt{MTL})$$

Combining and Optimizing the Bounds. The total meta-regret is the sum of the scaled meta-regret and the cumulative bias:

$$\text{Meta-Regret}(T) = \tilde{O}(\sqrt{MTL}) + \tilde{O}(T/\sqrt{L})$$

To find the optimal episode length L that minimizes this bound, we balance the two terms by setting:

$$\sqrt{MTL} \approx \frac{T}{\sqrt{L}} \implies \sqrt{MTL} \approx T \implies L \approx \frac{T}{\sqrt{MT}} = \sqrt{\frac{T}{M}}$$

Substituting this optimal choice of L back into the total regret expression gives:

$$\begin{aligned} \text{Meta-Regret}(T) &= \tilde{O} \left(\sqrt{MT\sqrt{T/M}} \right) = \tilde{O} \left(\sqrt{M^{1/2}T^{3/2}} \right) \\ &= \tilde{O}(M^{1/4}T^{3/4}) \end{aligned}$$

This concludes the proof of Theorem 4.2. $\square$