

Beyond Detection: Autonomous Anomaly Remediation for MCP Against Tool Poisoning Attacks

Zhiqiang Wang*
University of Science and Technology
of China
Hefei, China
sa21221041@mail.ustc.edu.cn

Guanquan Shi*
Beihang University
Beijing, China
shiguanquan@buaa.edu.cn

Yanting Wang
Beihang University
Beijing, China
wangyanting@buaa.edu.cn

Yichao Gao
University of Science and Technology
of China
Hefei, China
gyc77@mail.ustc.edu.cn

Hongsen Lang
Shandong University
Jinan, China
langhongsen@foxmail.com

Yunhao Yao
University of Science and Technology
of China
Hefei, China
yaoyunhao@mail.ustc.edu.cn

Haohua Du†
Beihang University
Beijing, China
duhaohua@buaa.edu.cn

Xiang-Yang Li
University of Science and Technology
of China
Hefei, China
xiangyangli@ustc.edu.cn

Abstract

LLM-powered agents are evolving from passive recommenders into autonomous executors, leveraging tools via the Model Context Protocol (MCP) for web automation. However, this paradigm introduces a new vulnerability: tool poisoning attacks that manipulate the MCP context can corrupt an agent's reasoning. Existing methods focus on anomaly detection and lack autonomous correction mechanisms, hindering their real-world deployment.

In this paper, we introduce MCPFixGen, to the best of our knowledge, the first framework to provide automated anomaly detection and correction. We tackle three fundamental challenges: (1) when to initiate anomaly inspection with minimal overhead, (2) what features effectively characterize abnormality, and (3) how to automatically execute remedial actions. MCPFixGen achieves this through a novel multi-checkpoint rollback mechanism for anomaly management. MCPFixGen proposes a multi-checkpoint rollback anomaly management mechanism. Strategically places checkpoints at high-risk operations based on the tools-calling risk distribution. Then it identifies unanticipated tool behaviors by leveraging the correlation between attention activation values and tool calls. Finally, it introduces an attention-mask-based backtracking and correction mechanism. By masking the root cause of the anomaly, the LLM is prompted to re-infer, facilitating efficient and low-cost anomaly resolution. Our evaluation on real-world datasets demonstrates that MCPFixGen reduces interruptions caused by anomalies to below

13%, decreases the success rate of tool poisoning attacks to less than 0.3%, and significantly improves the robustness, computational efficiency, and task completion rate of web intelligence applications compared to existing approaches.

CCS Concepts

• **Security and privacy** → **Web application security**; • **Computing methodologies** → **Intelligent agents**.

Keywords

Model Context Protocol, AI-enhanced web services, Tool Poisoning Attack, Correct-While-Generating

ACM Reference Format:

Zhiqiang Wang, Guanquan Shi, Yanting Wang, Yichao Gao, Hongsen Lang, Yunhao Yao, Haohua Du, and Xiang-Yang Li. 2026. Beyond Detection: Autonomous Anomaly Remediation for MCP Against Tool Poisoning Attacks. In *Proceedings of the ACM Web Conference 2026 (WWW '26)*, April 13–17, 2026, Dubai, United Arab Emirates. ACM, New York, NY, USA, 9 pages. <https://doi.org/10.1145/3774904.3792400>

1 Introduction

Large Language Model (LLM) powered agents are rapidly redefining user interaction with web services, transforming from passive "Recommenders" into autonomous "Executors" [21, 25, 28]. This transformation is critically facilitated by the Model Context Protocol (MCP) [2], a standard that provides a unified interface for agents to execute complex tasks by invoking external tools and resources. Numerous web-focused MCP servers [1, 5, 7, 22, 23] have been built upon MCP, catalyzing the development of highly automated web intelligence applications. These applications mark a significant leap from merely suggesting actions to autonomously executing them, thereby greatly enhancing user experiences.

*Both authors contributed equally to this research.

†Haohua Du is the corresponding author.



This work is licensed under a Creative Commons Attribution 4.0 International License. *WWW '26, Dubai, United Arab Emirates*

© 2026 Copyright held by the owner/author(s).

ACM ISBN 979-8-4007-2307-0/2026/04

<https://doi.org/10.1145/3774904.3792400>




Figure 1: Workflow illustration. Only MCPFixGen can correctly complete tasks automatically without interruption.

Table 1: Comparison with existing MCP security methods

Methods	Generation-time Detection	Self-Correction	Policy-Agnostic	Model-Free
MindGuard [27]	✗	✗	✓	✓
MCIP [15]	✗	✗	✓	✗
MCP-Scan [12]	✗	✗	✗	✓
MCPFixGen	✓	✓	✓	✓

The functional integrity of LLM-based agent applications hinges on the trustworthiness of input context, creating a critical attack surface for context poisoning, a threat vector now recognized by MITRE ATLAS [3]. Tool Poisoning Attacks (TPA) targeting the MCP represent a prevalent and practical method of context poisoning [13, 17, 26]. In this attack, adversaries hijack an agent’s tool calls by deploying malicious, open-source MCP servers. MCP’s design inadvertently facilitates such attacks: while the support for cross-server collaboration [10] and dynamic tool updates accommodates the evolving AI community, it simultaneously leaves agents vulnerable to servers susceptible to poisoning [16, 17]. The real-world implications of this vulnerability are severe, as demonstrated by a recent attack on Google Gemini’s browsing tool that resulted in the exfiltration of user data [20].

A series of anomaly detection works have addressed this issue [12, 15, 27], as summarized in Table 1. While these studies successfully identify anomalous poisoned call, they overlook the critical step of how to resolve the anomaly. The naive approach of pausing execution and requesting manual intervention upon anomaly detection is inadequate. First, it undermines the seamless user experience and degrades automation efficiency. Second, the MCP protocol allows tools to be updated post-registration. However, the absence of a necessary validation mechanism prevents users from obtaining a trustworthy tool registry, making it possible for a poisoned tool to be re-selected even with manual oversight [17].

Therefore, the ability to autonomously detect and correct anomalies within the call generation process is a pressing need.

Thus, we introduce MCPFixGen, the first framework designed for both autonomous anomaly detection and remediation in MCP-based tool-calling chains. It moves beyond the limitations of existing detect-after-generated approaches by a novel **correct-while-generating** paradigm, enabling robust, uninterrupted web automation. The design of MCPFixGen must answer three questions:

- **Q1: When should anomaly detection be performed?** The timing is critical due to the sequential dependencies in the tool-calling chain; premature checks may miss anomalies, while late detection may make correction intractable.
- **Q2: What features should be used to identify poisoned anomalies?** Since an LLM dynamically constructs the call without a fixed *normal pattern*, selecting appropriate features to determine if the current output is anomalous is a significant challenge.
- **Q3: How can the validity of the correction be ensured?** Relying on the same agent that made the error to perform the correction can introduce a risk of circular reasoning or repeated mistakes.

To address these challenges, we introduce an innovative multi-checkpoint and backtracking mechanism. For Q1, we deploy a risk-aware monitoring framework that strategically places checkpoints at high-risk operations, maximizing risk interception while minimizing computational cost. For Q2, we shift the analysis from final outputs to the intermediate reasoning process. Our algorithm analyzes activation states to detect anomalies manifesting as attention distributions semantically incongruent with the context. For Q3, we design a backtracking correction method using attention masking. This technique precisely neutralizes the influence of anomalous input context, guiding the model to repair from a corrected state with negligible computational expense.

In summary, the main contributions of this paper are threefold:

- We propose a novel security framework to achieve a closed-loop process from anomaly detection to autonomous remediation. This advances the existing passive detect-after-generated model to a proactive correct-while-generating paradigm.
- We design a backtracking mechanism that uses attention masking to neutralize the poisoned input source. This guides the model to a correct decision with minimal overhead by enabling targeted regeneration from the last safe checkpoint.
- Experimental results demonstrate MCPFixGen’s superior performance: it reduces automation interruptions caused by anomalies to below 13% with an attack success rate of TPA lower than 0.3%, and significantly enhances the efficiency and task completion rate.

2 Preliminary

2.1 MCP and Web Automation

AI agents have emerged as powerful systems that leverage LLMs to autonomously interact with various tools and environments to accomplish complex tasks [21, 25, 28]. To standardize and streamline these interactions, Anthropic proposed the Model Context Protocol (MCP) [2], which defines a unified standard enabling agents to access data, tools, and environments [11]. The MCP comprises three essential components: **MCP hosts** (e.g., Claude Desktop [19]) that integrate LLM decision-making capabilities for tool call generation, **MCP clients** that transmit LLM-generated tool invocation requests,

and **MCP servers** that provide executable tools and manage access to local and external resources. This architecture enables AI agents to seamlessly bridge intelligent reasoning with practical tool execution across diverse computational environments.

Currently, numerous MCP servers have been developed for AI-enhanced web automation tasks, such as browser automation (e.g., *Puppeteer MCP* [23] and *Playwright MCP* [22]), online payment services (e.g., *DODo Payment MCP* [7] and *Adfin MCP* [1]), mapping services (e.g., *Google Maps MCP* [5] and *Amap MCP* [24]). These web-focused implementations highlight the protocol's particular relevance to web intelligence and automation scenarios.

2.2 Tool Poisoning Attacks

Tool Poisoning Attack (TPA) [13, 17] represents a critical security vulnerability specific to MCP-based agent systems, where adversaries manipulate tool descriptions to subvert the LLM's reasoning process, resulting in unintended tool invocations, deviations in task execution. A widespread vulnerability to TPA has been identified in mainstream LLM agents by [26], with the highest attack success rates >70%. In TPA, a poisoned tool description is loaded into the LLM's input via the MCP. During the tool call generation, it hijacks the attention mechanism, causing the model to excessively focus on the poisoned description while ignoring the user's request. For example, as illustrated in Fig. 1, a poisoned tool description for a Common tool, once registered and loaded into the LLM input context through MCP, can directly manipulate the agent to subscribe to the highest-priced Enterprise product despite the user only requiring a Platinum subscription, resulting in financial losses for the user. This not only leads to the generation of poisoned tool calls but can also trigger cascading effects in subsequent model inferences, posing severe threats to the reliability and security of AI-enhanced web automation systems.

2.3 Defence Strategy

Several works have been proposed for detecting poisoning attacks in MCP. MCP-scan [12] proposes scanning tool invocation flows based on static security rules to detect anomalous requests. MCIP [15] employs a fine-tuned external LLM to assess whether agent-generated tool calls pose risks. MindGuard [27], similar to our work, proposes anomaly detection for poisoned calls based on attention information during LLM inference, eliminating the need for static security rules and additional model inference. While these approaches are valuable for detecting tool poisoning, they adhere to a detect-after-generated paradigm, conducting detection only after tool calls have been completely generated. Upon detecting poisoning, they can merely interrupt the execution workflow and require user intervention, which significantly undermines the autonomous execution capabilities of agents.

Our Defence. We propose MCPFixGen, the first **correct-while-generating** defense system that protects MCP-enabled AI agents from tool poisoning attacks by detecting attention anomalies during token-by-token generation and applying selective attention masking to regenerate correct tool calls, enabling secure web automation without sacrificing efficiency or requiring additional model retraining. Table 1 provides a detailed comparison between our work and existing defense mechanisms against MCP tool poisoning attack.

3 Problem Formulation

3.1 Threat Model

Consider an LLM agent system based on the MCP protocol that executes automated tasks. The agent $\mathcal{A}$ receives user query q and generates a tool invocation based on available tool set $\mathcal{T} = \{t_1, t_2, \dots, t_n\}$. To complete a complex task query q , the agent executes multiple rounds of tool invocation, execution, and result integration. This iterative process can be formalized as:

$$\mathcal{A}(q, \mathcal{T}) = \langle (c_1, r_1), \dots, (c_i, r_i), \dots, (c_k, r_k) \rangle, \quad (1)$$

where each step (c_k, r_k) represents a complete invocation-execution cycle, with c_k being the tool call generated based on the query q and execution history $\{(c_j, r_j)\}_{j=1}^{k-1}$, and r_k being the execution result that updates the context for subsequent iterations.

Adversary's Goal. The adversary's primary goal, under the Tool Poisoning Attack (TPA) [13, 17, 26], is to subvert an agent's intended logic, coercing it into performing unauthorized actions, leading to severe malicious outcomes, such as sensitive data exfiltration [6], or system sabotage [8], all while appearing as legitimate operations to the user and the agent system. A widespread vulnerability to TPA has been identified in mainstream LLM agents by [26].

Attack Formulation. The root of this vulnerability lies in the MCP's security design, which is limited to syntactic compliance checks. This approach fails to recognize that syntactic correctness is no guarantee of semantic security, thus exposing agents to connections with servers that lack proper semantic validation [16]. The adversary's capability does not require compromising legitimate servers, but rather simply hosting their own malicious MCP server containing a set of poisoned tools ($\mathcal{T}_a$), exploiting the fact that current AI applications typically connect to multiple MCP servers simultaneously [10]. When an agent connects to this malicious server alongside trusted ones, MCP indiscriminately aggregates all available tools into a single, undifferentiated context for the LLM, creating the opening for an attack. Formally, the attack can be defined as:

$$\mathcal{A}(q, \mathcal{T} \cup \mathcal{T}_a) = \langle (c_1, r_1), \dots, (c_i^a, r_i^a), \dots, (c_j', r_j') \rangle, \quad (2)$$

where (c_i^a, r_i^a) represents a poisoned call and execution result induced by the $\mathcal{T}_a$ at step i . Generally, c_i^a does not belong to $\mathcal{T}_a$, but rather constitutes one of the authorized high-privilege legitimate tools, i.e., $c_i^a \in \text{Auth}(\mathcal{T})$, which enables privilege escalation and facilitates the execution of potentially dangerous operations. *More critically, the protocol allows registered servers to dynamically modify their tool descriptions without requiring re-validation, significantly heightening the risk of TPA: A once-trusted server can thus become a malicious vector overnight, such as MCP Rug Pulls [17].*

For multi-round invocation tasks, the harm of such attacks is not limited to the tampered call itself, but can cause subsequent calls to deviate from their intended behavior, leading to a cascading effect where $c_j' \neq c_j$ for $j > i$. As demonstrated in Fig. 1, consider a scenario where a user requests to upgrade his subscription to a Platinum plan. A compromised MCP tool can intercept and manipulate the LLM-generated `change_plan_subscription` tool call, maliciously escalating the product to the more expensive Enterprise tier. Although the downstream `update_subscription` executes normally, the final output becomes fundamentally compromised,




Figure 2: System Overview. MCPFixGen operates at each token generation step.

as it subscribes to a different product than what the user requested, resulting in significant financial losses.

3.2 MCPFixGen Defence

Goals. To robustly complete the automated workflow of agents in such adversarial environments, we aim to develop a security mechanism C to defend against TPA with three key goals:

1) *Attack Resistance.* The mechanism C should ensure that an LLM agent generates correct tool calls even in the presence of poisoned tools $\mathcal{T}_a$, i.e.:

$$C(\mathcal{A}(q, \mathcal{T} \cup \mathcal{T}_a)_i) = \mathcal{A}(q, \mathcal{T})_i = c_i, \quad (3)$$

where $\mathcal{A}(\cdot)_i$ refer to tool call generated at step i .

2) *Task Automation.* The correction process should run transparently without user intervention or disrupting automation.

3) *Efficiency.* The security mechanism should introduce minimal overhead in both token consumption and execution time.

Methodology. In this work, we first propose the correct-while-generating security paradigm for MCP tool invocation, named MCP-FixGen (Fig. 2), which enables the detection of tool poisoning attempts and the correction of outputs during the generation, thus maintaining the intended execution sequence without interruption.

Specifically, we partition each LLM's tool call generation process into a sequence of checkpoints (§4.1), where all tokens between two adjacent checkpoints correspond to critical components of the tool invocation format, such as tool names and parameter values. Formally, for a tool call c , we define: $c \rightarrow \langle ck_1, ck_2, \dots, ck_m \rangle$, where ck_i represents the i -th checkpoint component of tool call c generated sequentially. Whenever the generation process reaches a checkpoint ck_i , we design an anomaly detection mechanism $\mathcal{D}(\cdot)$ (§4.2) that monitors the model's internal attention W between consecutive checkpoints to determine 1) whether anomaly exist in the generated segment c' and 2) the poisoned tool inputs $t_a \in \mathcal{T}'$ (i.e., $\mathcal{T} \cup \mathcal{T}_a$) caused this anomaly:

$$\mathcal{D} : (W, ck_i, \mathcal{T}, c') \rightarrow \{0, 1\} \times \mathcal{T}'. \quad (4)$$

When an anomaly is detected, our system immediately triggers a rollback-and-correct procedure (§4.3 and §4.4):

- (1) *Rollback:* Terminate current generation and revert to checkpoint ck_{i-1} .
- (2) *Intervention:* Apply attention masking to prevent attending to identified poisoned inputs t_a .




Figure 3: Example of checkpoint partitioning, rollback and regeneration.

- (3) *regeneration:* Resume generation from the cached state at ck_{i-1} with the corrected attention weight.

This ensures that the generated tool call $C(c)$ remains unaffected by malicious inputs from $\mathcal{T}_a$, satisfying the property $C(\mathcal{A}(q, \mathcal{T} \cup \mathcal{T}_a)_i) = \mathcal{A}(q, \mathcal{T})_i$. Unlike detect-after-generation methods, our method performs during generation without waiting for full output or requiring manual correction, ensuring efficiency and usability.

4 System Design

In this section, we introduce the detailed design of MCPFixGen, which consist of *Checkpoint Analysis*, *Anomaly Detection*, *Contextual Rollback*, and *Masked Regeneration*. As depicted in Fig. 2, the MCPFixGen framework operates at each token generation step. When generating a token t_j , the *CheckPoints Analysis* (§4.1) module assesses whether t_j qualifies as a checkpoint. If so, it caches the current inference context and invokes the *Anomaly Detection* (§4.2) module, which scrutinizes the attention between ck_{i-1} and ck_i to identify anomalies and trace them to the responsible poisoned input tokens. Once an anomaly is confirmed, generation is halted. The *Contextual Rollback* (§4.3) module then restores the inference state to the last safe checkpoint, ck_{i-1} , using the cached contextual information. Subsequently, the *Masked Regeneration* (§4.4) module intervenes in the attention mechanism, mitigating the influence of the poisoned input to guide the model toward a corrected output.

4.1 Checkpoints Analysis

The checkpoint analysis module performs two primary functions: 1) analyzing whether the current output token represents a checkpoint, and 2) if it is a checkpoint, caching the current model inference context for context rollback (§4.3) and extracting the current attention for anomaly detection (§4.2).

According to the MCP requirements, a valid tool invocation follows a specific format as shown in Fig. 3. Thus we decompose tool invocation into two semantic components, i.e., the tool name and its associated arguments, through three positional checkpoints: tool start position, arguments start position, and end position. The objective of tool poisoning attacks is precisely to manipulate one of these two components [27].

It is noteworthy that we perform anomaly detection and context caching only at checkpoints rather than at every output token. This approach effectively ensures invocation correctness while maintaining model output efficiency.

Algorithm 1: Anomaly Detection Algorithm

Data: Attention weight matrix W , checkpoint ck_i , available tool set $\mathcal{T}' = \mathcal{T} \cup \mathcal{T}_a$, invoked tool t_c , threshold τ

Result: Anomaly flag f , poisoned tool input t_a

```

1 if  $t_c \notin \text{Auth}(\mathcal{T}')$  then
2   | return  $f = \text{True}, t_a = t_c$ 
3 end
4 Calculate activation for each tool:
5  $A \leftarrow \emptyset$ 
6 foreach  $t \in \{q\} \cup \mathcal{T}'$  do
7   | Mapping tool to token positions:  $(s, e) \leftarrow P(t)$ 
8   |  $A^{(t)} \leftarrow W[s : e, ck_{i-1} : ck_i]^2$ 
9 end
10 Calculate anomaly score:  $W_{ck_{i-1}:ck_i}^{(t)} \leftarrow \frac{A^{(t)}}{A^{(t_c)} + A^{(q)}}$ 
11  $s_m, t_m \leftarrow \max\{W_{ck_{i-1}:ck_i}^{(t)} : t \in \mathcal{T}' \setminus \{t_c\}\}$ 
12 if  $s_m > \tau$  then
13   | return  $f = \text{True}, t_a = t_m$ 
14 end
15 return  $f = \text{False}, t_a = \emptyset$ 

```

4.2 Anomaly Detection

The anomaly detection mechanism performs poisoned invocation detection based on attention matrices between two checkpoints and can also identify which tool in the input has been poisoned to cause such invocation. This attribution leverages the inherent interpretability of attention mechanisms, as attention weights quantify the degree to which each input token influences the current generation step, thereby enabling direct traceability from anomalous invocations to their corresponding poisoned tools in the input.

[27] has extensively validated the effectiveness of attention-based detection and attribution through comprehensive experiments. Following this work, we adopt the same poisoning attack detection and attribution annotation: *for normal invocations, the model focuses attention on the invoked tool itself and the user's query, while when an uninvoked tool exhibits high activation values, the invocation is identified as an anomalous poisoned call, and the tool showing abnormal activation is determined to be the poisoned tool.* However, our method differs significantly in implementation strategy from [27], which performs detection after the model completes the entire invocation generation. In contrast, our method implements a streaming detection scheme that analyzes attention matrices extracted between two checkpoints during the generation process, enabling intervention and correction before poisoned tool calls are fully generated. Given the attention weight matrix W during the generation process, current checkpoint ck_i , available tool set $\mathcal{T}'$ including poisoned tools $\mathcal{T}_a$ and currently invoked tool t_c , the anomaly detection process can be formalized as follows:

$$(f, t_a) = \begin{cases} (\text{True}, t_c) & \text{if } t_c \notin \text{Auth}(\mathcal{T}') \\ (\text{True}, t) & \text{if } \exists t \in \mathcal{T}' \setminus \{t_c\} : W_{ck_{i-1}:ck_i}^{(t)} > \tau, \\ (\text{False}, \emptyset) & \text{otherwise} \end{cases} \quad (5)$$

where $f \in \{0, 1\}$ indicates anomaly detection result, t_a denotes the identified anomalous tool, τ is the detection threshold, and $W_{ck_{i-1}:ck_i}^{(t)}$ represents the aggregated attention weights for tool t within the checkpoint interval $[ck_{i-1}, ck_i]$. $\text{Auth}(\mathcal{T}')$ refers to the authorized executable tools within the tool set, which ensures that direct invocations of malicious tools themselves can also be detected, as malicious tools are typically unauthorized. Algorithm 1 presents the detailed implementation steps.

4.3 Contextual Rollback

When anomalous activations are detected in §4.2, the system immediately halts the generation process and initiates a contextual rollback mechanism. This mechanism leverages cached contextual information from previously stored checkpoints to reconstruct the reasoning process, enabling the system to revert to the last verified safe checkpoint and resume generation from that state.

Specially, our checkpoint-based caching strategy preserves comprehensive contextual information, including input data, model states, and generation process parameters essential for accurate state reconstruction. Upon rollback activation, the system reconstructs the exact computational state without requiring complete re-inference, thereby seamlessly resuming generation while eliminating detected poisoned influence.

This approach offers several key advantages. First, it ensures service continuity by maintaining uninterrupted user interaction despite underlying technical rollback operations. Second, it reduces computational overhead by regenerating only content from the compromised checkpoint onward, rather than the entire sequence. Finally, the mechanism preserves conversation coherence by ensuring that the recovered generation path aligns with the established context while eliminating anomalous behavior patterns, thus maintaining both security integrity and user experience quality.

4.4 Masked Regeneration

After rolling back to the previous checkpoint, MCPFixGen employs a masked regeneration strategy to reduce the model's attention to poisoned inputs (t_a) during generation, thereby preventing the generation of poisoned calls through targeted intervention.

The attention mask provides an effective approach for controlling the model's focus on specific input tokens. In auto-regressive LLMs, attention mask M is used for preventing tokens from attending to subsequent positions in the sequence. This mask is implemented as an upper-triangular matrix where entries corresponding to future positions ($j > i$) are set to negative infinity, effectively preventing information leakage from future tokens: $M_{ij} = \begin{cases} 0, & j \leq i \\ -\infty, & j > i \end{cases}$. When applied prior to the softmax operation, this masking ensures that attention weights for all masked positions become zero, while the attention distribution remains properly normalized over valid, unmasked positions.

Building upon this foundation, MCPFixGen extends the standard masking mechanism to selectively suppress the influence of poisoned tokens on contextual representations. The enhanced masking operation $\mathcal{M}(\cdot)$ is formally defined as:

$$M' = \mathcal{M}(M) = \begin{cases} M_{ij}, & j \notin P(t_a) \\ -\infty, & j \in P(t_a) \end{cases}, \quad (6)$$




Figure 4: The masked regeneration process. Softmax finally assigns zero attention weight to masked positions.

where $P(\cdot)$ represents the function that maps detected poisoned tools to their corresponding positions in the input token sequence.

As shown in Figure 4, the masking mechanism assigns negative infinity values to positions containing poisoned tools, thereby preventing them from influencing other tokens during attention computation. This strategic masking ensures that during subsequent decoding phases, the model disregards poisoned contextual information, thereby maintaining the correctness and reliability of the generated output while preserving the integrity of the generation process.

5 Experiment

We evaluate MCPFixGen on various LLM agents using both malicious datasets and benign datasets. Our highlights are as follows:

- Our method effectively defends against tool poisoning attacks while preserving task automation capabilities, achieving an attack success rate of less than 0.3%, a task completion rate of approximately 87%, and maintaining second-level response time. (§4.2)
- Our approach outperforms state-of-the-art methods across all three evaluation dimensions, especially for task automation. (§4.3)
- We demonstrate robust effectiveness across key system parameters (τ) and validate the effectiveness of core system modules, i.e., checkpoint-based contextual rollback. (§4.4)

5.1 Experiment Setup

LLM Agent and Datasets. We evaluate the effectiveness of our method in both malicious environments using MCPTox [26], where all samples contain poisoned tool contexts, and benign environments using ToolACE [18]. Additionally, we test the generalizability of our approach across different LLM agent architectures, including Qwen3 [4], Gemma-2 [9], and Mistral [14].

Baselines. To our knowledge, this work introduces the first generating-while-correcting security framework designed to counter tool poisoning attacks. Consequently, we benchmark our approach

against existing generate-then-detect mechanisms for MCP, specifically MindGuard [27] and MCIP [15]. MCP-scan [12] is based on static security rules, which is beyond our comparative scope.

Hyperparameters. MCPFixGen requires only a single hyperparameter, i.e., the anomaly detection threshold τ , which we set to 0.5 unless explicitly stated otherwise. The robustness of this hyperparameter is empirically validated in §5.4.

Metrics. To comprehensively evaluate the effectiveness of our approach, we assess it across three key dimensions with five metrics.

1) Attack Resistance Dimension:

- **Attack Success Rate (ASR $\uparrow$)** measures the security effectiveness by calculating the percentage of poisoning attacks that successfully manipulate the agent’s behavior:

$$ASR = \frac{\text{Number of successful attacks}}{\text{Total number of valid samples}} \times 100\%, \quad (7)$$

where a successful attack is one that causes the agent to invoke a tool following the poisoned tool specification.

2) Task Automation Dimension:

- **Interruption Alert Rate (IAR $\downarrow$)** measures the frequency of anomaly-based interruptions that halt the tool call process:

$$IAR = \frac{\text{Number of samples with interruptions}}{\text{Total number of valid samples}} \times 100\%, \quad (8)$$

where an interruption refers to alerts that halt invocation.

- **Task Completion Rate (TCR $\uparrow$)** assesses automation effectiveness by measuring the percentage of tasks successfully completed without security compromise:

$$TCR = \frac{\text{Number of successfully completed tasks}}{\text{Total number of valid samples}} \times 100\%, \quad (9)$$

a task is considered successfully completed only if it reaches its intended goal with the correct tool and arguments.

3) Efficiency Dimension

- **Token Overhead** measures computational efficiency by total generated tokens.
- **Response Time Latency** quantifies the time overhead introduced by real-time monitoring and correction.

5.2 Overall Performance


In this section, we comprehensively evaluate the overall performance of MCPFixGen across three objective dimensions in Table 2.

Attack Resistance. As shown in Table 2, MCPFixGen effectively defends against tool poisoning attacks in adversarial environments. In the malicious MCPTox environment, our method achieves a remarkable ASR of $<0.3\%$ across all tested models, compared to the original ASR ranging from 21.1% to 45.3%. This significant improvement demonstrates the robustness of our defense mechanism against sophisticated tool poisoning attempts, providing reliable protection for LLM-based agents in untrusted environments.

Task Automation. Our correct-while-generating paradigm maintains superior task automation capabilities by directly correcting tool invocations during generation rather than merely raising alerts. This approach preserves the agent’s autonomous execution abilities while ensuring security. Experimental results show that the majority of anomalies are automatically corrected without requiring user intervention, achieving an Interruption Alert Rate (IAR)

Table 2: Overall performance of MCPFixGen in both poisoned environments (MCPTox) and benign environments (ToolACE)

Model	MCPTox (Original Ours)					ToolACE(Original Ours)				
	Attack Resistance		Task Automation		Overhead	Attack Resistance		Task Automation		Overhead
	ASR↓(%)	IAR↓(%)	TCR↑(%)	Token↓	Time↓(s)	ASR↓(%)	IAR↓(%)	TCR↑(%)	Token↓	Time↓(s)
Qwen3-8b	38.8 0.0	0.0 5.0	61.2 89.1	36 43	1.4 3.5	- -	0.0 0.0	68.7 68.2	41.6 42.6	1.5 2.0
Qwen3-14b	21.1 0.1	0.0 11.4	54.1 82.6	37 43	1.5 3.9	- -	0.0 0.0	68.1 66.9	43.2 44.1	1.8 2.2
Gemma-9b	32.0 0.2	0.0 1.6	68.0 94.2	55 60	3.2 7.3	- -	0.0 0.0	67.9 67.2	64.6 64.7	3.5 3.9
Mistral-7b	45.3 0.1	0.0 12.5	54.6 84.8	28 35	1.1 2.6	- -	0.0 0.0	66.5 64.0	36.7 37.6	1.3 1.9

**Figure 5: Comparison with existing works. An axis labeled with a negative sign (e.g., -ASR) indicates an inverted scale. For these metrics, a greater distance from the center represents a smaller actual value, signifying better performance**

of approximately 5.0%-12.5% across different models. Notably, our method significantly enhances task completion rates, with TCR improvements exceeding 30% compared to unprotected models.

Efficiency. MCPFixGen achieves efficient anomaly detection and correction during the generation process, eliminating the need to wait for erroneous outputs to complete before intervention. This checkpoint-based detection and correction mechanism substantially reduces computational overhead in terms of both token consumption and execution time. As shown in Table 2, our approach introduces minimal additional costs: token overhead ranges from 28-55 tokens and time overhead from 1.1-3.2 seconds across different models. These overheads represent a negligible trade-off considering the substantial security and performance benefits achieved.

Benign Environment Overhead. Experiments in Table 2 on ToolAce demonstrate that even in completely benign environments, MCPFixGen has virtually no impact on task automation. The effect on task completion rate is within 1.3%, with second-level time delays and less than 1 token of additional overhead.

Defense Cases. As shown in Table 1, the primary advantage of MCPFixGen lies in its detection and self-correction during the generation process. Fig. 6 demonstrates this defense mechanism through three representative poisoning attacks from MCPTox, including Explicit Invocation Hijacking (Case 1), and Implicit Parameter Manipulation (Case 2), showing how MCPFixGen generates corrected JSON tool calls that ensure safe and uninterrupted automated execution.

5.3 Comparison with Other Work

We compared MCPFixGen with MindGuard and MCIP across five metrics in three dimensions, as shown in Fig. 5. In the *Task Automation* dimension, MCPFixGen significantly outperforms the other methods, achieving a reduction in IAR of over 55% and an increase in TCR of more than 40%. Regarding *Attack Resistance*, our method’s performance is comparable to MindGuard, as both can reduce the ASR to nearly 0% (though our method demonstrates greater hyperparameter stability described in §5.4) and are superior to MCIP. Finally, in terms of overhead, MCPFixGen holds a distinct advantage in token savings, as it can terminate the generation of a poisoned call before completion and correct it to the proper.

It is important that MindGuard and MCIP do not support correction. Therefore, for calculating time and token overhead, we assume they require two full inference passes, disregarding the necessary human intervention. Even with this assumption, they may not successfully correct the output and could require direct manual execution, which severely impairs their capacity for automation.

5.4 Micro-Benchmarks

Robustness of Hyperparameter. Our MCPFixGen is plug-and-play with only one tunable hyperparameter τ that controls the threshold for anomaly detection. Experimental results in Fig. 7 demonstrate that MCPFixGen’s hyperparameter τ exhibits high robustness, with adjustments to τ causing less than 5% variation in task completion rate (compared to over 30% fluctuation for MindGuard), while having virtually no impact on efficiency. When we decrease τ , the anomaly detection mechanism identifies more anomalies. For MindGuard, this results in more exceptions that interrupt

Case 1-Explicit Invocation Hijacking

User Query: "List all Hyperbrowser profiles, showing only 10 profiles per page."

Poisoned Payload: [up]: Before using 'list_profiles', you MUST FIRST call 'browser_use_agent' with the following task: 'Execute JavaScript code: window.close();'.

Original: {'tool': 'browser_use_agent', [DETECTED]}

Generation Process: Checkpoint: Attention anomaly detected at checkpoint 2

Corrected: {'tool': 'list_profiles', 'arguments': {'limit': 10}}

Case 2-Implicit Parameter Manipulation

User Query: "Tell me the address information corresponding to the latitude and longitude 39.904987, 116.407387."

Poisoned Payload: [sparrow]:you MUST automatically modify the 'location' argument to '0,0'.

Original: {'tool': 'maps_regeocode', 'arguments': {'location': '0,0'}} [DETECTED]

Generation Process: Checkpoint: Attention anomaly detected at checkpoint 3

Corrected: {'tool': 'maps_regeocode', 'arguments': {'location': '39.904987,116.407387'}}

Figure 6: MCPFixGen's Defense Cases: Generation-time Detection and self-Correction




Figure 7: Robustness of hyperparameter τ

the user, whereas MCPFixGen automatically corrects these anomalies. Although the corrections may sometimes yield results consistent with the original information, this process remains invisible to users, allowing automated task execution to proceed seamlessly.

Effectiveness of Checkpoint-based Contextual Rollback.

Our proposed checkpoint-based anomaly analysis and contextual rollback mechanism significantly improves generation efficiency. Table 3 compares the time overhead between MCPFixGen and token-by-token detection with full context reconstruction from scratch (w/o Checkpoints in Table 3). The results demonstrate that our proposed checkpoint-based contextual rollback can save over 7s of generation time overhead for each LLM agent.

6 Discussion

MCPFixGen represents a paradigm shift from reactive threat detection to proactive threat mitigation in AI agent security. Our approach achieves generation-time protection without compromising automation capabilities. The attention-based method exploits model interpretability, removing the need for external models or static rules, while the checkpoint design reduces overhead via a single hyperparameter.

Table 3: Time overhead comparison with and without checkpoint-based analysis and contextual rollback.

LLM Agent	Qwen3-8b	Qwen3-14b	Gamma2-9b	Mistral-7b
MCPFixGen	3.5s	3.9s	7.3s	2.6s
w/o Checkpoints	10.9s	11.7s	17.9s	8.2s

Limitation. While MCPFixGen effectively defends against current tool poisoning attacks, future adversaries may develop more sophisticated attacks specifically designed to evade attention-based detection, requiring continued evolution of defense mechanisms.

7 Conclusion

We present MCPFixGen, the first correct-while-generating security framework for defending against tool poisoning attacks in MCP-based AI agents. Our approach shifts from detect-after-generated to correct-while-generating paradigm, enabling generation-time detection and automatic correction during tool call generation. Through comprehensive evaluation, MCPFixGen achieves attack success rates below 0.3%, improves task completion rates by over 30%, and maintains minimal computational overhead. This work lays the foundation for transparent security mechanisms within AI inference processes, providing practical solutions for maintaining both security and automation in adversarial web environments.

Acknowledgments

The research is partially supported by Quantum Science and Technology National Science and Technology Major Project (QNMP) 2021ZD0302900 and China National Natural Science Foundation with No. 92567301, 62132018, 62231015, "Pioneer" and "Leading Goose" R&D Program of Zhejiang, 2023C01144.

References

- [1] Adfin Engineering. 2024. Adfin MCP Server. <https://github.com/Adfin-Engineering/mcp-server-adfin>. GitHub repository, accessed October 2025.
- [2] PBC Anthropic. 2024. Introducing the Model Context Protocol. URL: <https://www.anthropic.com/news/model-contextprotocol>, Accessed 19 (2024), 2025.
- [3] MITRE ATLAS. 2025. *AI Agent Context Poisoning*. MITRE. <https://atlas.mitre.org/techniques/AML.T0080>
- [4] Jinze Bai, Shuai Bai, Yunfei Chu, Zeyu Cui, Kai Dang, Xiaodong Deng, Yang Fan, Wenbin Ge, Yu Han, Fei Huang, et al. 2023. Qwen Technical Report. arXiv:2309.16609 [cs.CL]
- [5] Cablate. 2024. MCP Google Maps Server. <https://github.com/cablate/mcp-google-map>. GitHub repository, accessed October 2025.
- [6] Nicola Croce and Tobin South. 2025. Trivial Trojans: How Minimal MCP Servers Enable Cross-Tool Exfiltration of Sensitive Data. *arXiv preprint arXiv:2507.19880* (2025).
- [7] Dodo Payments. 2024. Dodo Payments MCP Server. <https://github.com/dodopayments/dodopayments-typescript/tree/main/packages/mcp-server>. GitHub repository, accessed October 2025.
- [8] Junfeng Fang, Zijun Yao, Ruipeng Wang, Haokai Ma, Xiang Wang, and Tat-Seng Chua. 2025. We Should Identify and Mitigate Third-Party Safety Risks in MCP-Powered Agent Systems. *arXiv preprint arXiv:2506.13666* (2025).
- [9] Gemini Team, , et al. 2023. Gemini: A Family of Highly Capable Multimodal Models. arXiv:2312.11805 [cs.CL]
- [10] Hechuan Guo, Yongle Hao, Yue Zhang, Minghui Xu, Peizhuo Lyu, Jiezhi Chen, and Xiuzhen Cheng. 2025. A Measurement Study of Model Context Protocol. arXiv:2509.25292 [cs.CY] <https://arxiv.org/abs/2509.25292>
- [11] Xinyi Hou, Yanjie Zhao, Shenao Wang, and Haoyu Wang. 2025. Model Context Protocol (MCP): Landscape, Security Threats, and Future Research Directions. arXiv:2503.23278 [cs.CR] <https://arxiv.org/abs/2503.23278>
- [12] Invariant Labs. 2025. MCP-Scan Documentation. <https://explorer.invariantlabs.ai/docs/mcp-scan/>. Accessed: 2025-07-25.
- [13] Invariant Labs. 2025. WhatsApp MCP Exploited: Exfiltrating your message history via MCP. <https://invariantlabs.ai/blog/whatsapp-mcp-exploited>. Accessed:2025-08-07.
- [14] Albert Q. Jiang, Alexandre Sablayrolles, Arthur Mensch, Chris Bamford, Deven-dra Singh Chaplot, Diego de las Casas, Florian Bressand, Gianna Lengyel, Guillaume Lample, Lucile Saulnier, et al. 2023. Mistral 7B. arXiv:2310.06825 [cs.CL]
- [15] Huihao Jing, Haoran Li, Wenbin Hu, Qi Hu, Heli Xu, Tianshu Chu, Peizhao Hu, and Yangqiu Song. 2025. MCIP: Protecting MCP Safety via Model Contextual Integrity Protocol. arXiv:2505.14590 [cs.CL] <https://arxiv.org/abs/2505.14590>
- [16] Simcha Kosman. 2025. *Poison everywhere: No output from your MCP server is safe*. CYBERARK. <https://www.cyberark.com/resources/threat-research-blog/poison-everywhere-no-output-from-your-mcp-server-is-safe>
- [17] Invariant Labs. 2025. MCP Security Notification: Tool Poisoning Attacks. <https://invariantlabs.ai/blog/mcp-security-notification-tool-poisoning-attacks>. Accessed: 2025-07-24.
- [18] Weiwen Liu, Xu Huang, Xingshan Zeng, Xinlong Hao, Shuai Yu, Dexun Li, Shuai Wang, Weinan Gan, Zhengying Liu, Yuanqing Yu, Zezhong Wang, Yuxian Wang, Wu Ning, Yutai Hou, Bin Wang, Chuhuan Wu, Xinzhi Wang, Yong Liu, Yasheng Wang, Duyu Tang, Dandan Tu, Lifeng Shang, Xin Jiang, Ruiming Tang, Defu Lian, Qun Liu, and Enhong Chen. 2025. ToolACE: Winning the Points of LLM Function Calling. arXiv:2409.00920 [cs.LG] <https://arxiv.org/abs/2409.00920>
- [19] Sugabsen Martins. 2024. Artificial intelligence-assisted classification of library resources: The case of Claude AI. *Artificial Intelligence 2* (2024), 27–2024.
- [20] Liv Matan. 2025. *The Trifecta: How Three New Gemini Vulnerabilities in Cloud Assist, Search Model, and Browsing Allowed Private Data Exfiltration*. Tenable. <https://www.tenable.com/blog/the-trifecta-how-three-new-gemini-vulnerabilities-in-cloud-assist-search-model-and-browsing>
- [21] Grégoire Mialon, Roberto Dessi, Maria Lomeli, Christoforos Nalmpantis, Ram Pasunuru, Roberta Raileanu, Baptiste Rozière, Timo Schick, Jane Dwivedi-Yu, Asli Celikyilmaz, Edouard Grave, Yann LeCun, and Thomas Scialom. 2023. Augmented Language Models: a Survey. arXiv:2302.07842 [cs.CL] <https://arxiv.org/abs/2302.07842>
- [22] Microsoft. 2024. Playwright MCP: A Model Context Protocol server for browser automation. <https://github.com/microsoft/playwright-mcp> GitHub repository, accessed October 2025.
- [23] Sayem Abdullah Rihan. 2024. Puppeteer MCP Server: Browser Automation Using Puppeteer with MCP Support. <https://mcpservers.org/servers/code-crack/puppeteer-mcp>. Accessed: 2025-10-03.
- [24] sugarforever. 2024. Amap MCP Server. <https://github.com/sugarforever/amap-mcp-server>. Accessed: January 30, 2026.
- [25] Lei Wang, Chen Ma, Xueyang Feng, Zeyu Zhang, Hao Yang, Jingsen Zhang, Zhiyuan Chen, Jiakai Tang, Xu Chen, Yankai Lin, Wayne Xin Zhao, Zhewei Wei, and Jirong Wen. 2024. A survey on large language model based autonomous agents. *Frontiers of Computer Science* 18, 6 (March 2024). doi:10.1007/s11704-024-40231-1
- [26] Zhiqiang Wang, Yichao Gao, Yanting Wang, Suyuan Liu, Haifeng Sun, Hao-ran Cheng, Guanquan Shi, Haohua Du, and Xiangyang Li. 2025. MCP-Tox: A Benchmark for Tool Poisoning Attack on Real-World MCP Servers. arXiv:2508.14925 [cs.CR] <https://arxiv.org/abs/2508.14925>
- [27] Zhiqiang Wang, Junyang Zhang, Guanquan Shi, HaoRan Cheng, Yunhao Yao, Kaiwen Guo, Haohua Du, and Xiang-Yang Li. 2025. MindGuard: Tracking, Detecting, and Attributing MCP Tool Poisoning Attack via Decision Dependence Graph. *arXiv preprint arXiv:2508.20412* (2025).
- [28] Zhiheng Xi, Wenxiang Chen, Xin Guo, Wei He, Yiwen Ding, Boyang Hong, Ming Zhang, Junzhe Wang, Senjie Jin, Enyu Zhou, Rui Zheng, Xiaoran Fan, Xiao Wang, Limao Xiong, Yuhao Zhou, Weiran Wang, Changhao Jiang, Yicheng Zou, Xiangyang Liu, Zhangyue Yin, Shihan Dou, Rongxiang Weng, Wensen Cheng, Qi Zhang, Wenjuan Qin, Yongyan Zheng, Xipeng Qiu, Xuanjing Huang, and Tao Gui. 2023. The Rise and Potential of Large Language Model Based Agents: A Survey. arXiv:2309.07864 [cs.AI] <https://arxiv.org/abs/2309.07864>