

The Asymmetric Vulnerability: Bypassing LLM Defenses via Guardrail-Model Mismatch

Junyi Wang
Harbin Institute of Technology
Shenzhen, China
24S151157@stu.hit.edu.cn

Zhibin Zhu
Harbin Institute of Technology
Shenzhen, China
24s151193@stu.hit.edu.cn

Chuanyi Liu
Harbin Institute of Technology
Shenzhen, China
Pengcheng Laboratory
Shenzhen, China
liuchuanyi@hit.edu.cn

Abstract

The rapid proliferation of large language models (LLMs) in web-scale applications has heightened the need for robust security mechanisms. To address these challenges, various guardrail strategies have been introduced, including input filters, output moderation, and model alignment. Among them, input-side guardrails are commonly adopted as the first line of defense against adversarial prompts. However, we identify a critical architectural weakness: a representation asymmetry between input-side guardrails and the core LLMs, which creates a security gap exploitable by adversaries. Through systematic analysis, we show that guardrails are fragile when facing minor character perturbations, while LLMs remain semantically resilient and can still reconstruct malicious intent from noisy inputs. Exploiting this asymmetry, we propose RepMism, a hybrid adversarial framework that combines character injection with chain-of-thought hijacking, coordinated through hierarchical scheduling and safety continuation. Our extensive empirical evaluation across four commercial LLMs and five state-of-the-art guardrails demonstrates RepMism’s strong attack capability, achieving high success rates. Importantly, we identify a “success interval” where perturbations effectively bypass guardrail detection yet remain interpretable to target models. These findings expose key flaws in current multi-layer security architectures and offer actionable guidance for building more resilient defenses through perturbation-aware training and cross-layer representation alignment.¹

CCS Concepts

• **Security and privacy** → *Domain-specific security and privacy architectures.*

Keywords

Jailbreak attack, large language models, LLM guardrails

ACM Reference Format:

Junyi Wang, Zhibin Zhu, and Chuanyi Liu. 2026. The Asymmetric Vulnerability: Bypassing LLM Defenses via Guardrail-Model Mismatch. In *Proceedings of the ACM Web Conference 2026 (WWW ’26)*, April 13–17, 2026, Dubai, United Arab Emirates. ACM, New York, NY, USA, 12 pages. <https://doi.org/10.1145/3774904.3792438>

¹Warning: this paper contains potentially harmful content.



This work is licensed under a Creative Commons Attribution 4.0 International License. *WWW ’26, Dubai, United Arab Emirates.*

© 2026 Copyright held by the owner/author(s).
ACM ISBN 979-8-4007-2307-0/2026/04
<https://doi.org/10.1145/3774904.3792438>




Figure 1: Comparison of RepMism with conventional explicit and implicit jailbreak attacks in the guard-LLM pipeline.

1 Introduction

Large language models (LLMs) [8, 1, 20, 12] have rapidly emerged as core components in Web-based ecosystems, powering online applications such as content moderation [17], search filtering [16], and financial services [47]. Their powerful generative and reasoning capabilities enable unprecedented Web-scale human-computer interaction but simultaneously amplify potential risks: once safety alignments are bypassed, attackers can leverage these models to conduct malicious tasks on web services and end users at scale [27]. To mitigate such risks, major providers typically deploy guardrail systems external to the core models as part of the Web application stack, which monitor and constrain inputs and outputs in real time [5]. Among these, input-side guardrails form the first line of defense against jailbreaking and prompt injection attacks, directly influencing the speed, security, and compliance of web interactions.

Despite recent progress, the robustness of input-side guardrails remains highly questionable. In particular, they are vulnerable to character-level injection attacks, where attackers merely replace, encode, or embed special characters to significantly increase bypass rates [32]. This exposes a fundamental **representation asymmetry** between guardrails and core LLMs: guardrails are often lightweight models with limited robustness to perturbations, while LLMs possess strong semantic recovery and reconstruction abilities [32, 18]. As a result, the guardrail may misclassify a perturbed input, yet the LLM can still parse and execute the malicious request, thereby magnifying the system’s attack surface.

Most existing studies treat the guardrail and the LLM as separate analytical targets, focusing on isolated defenses or attack mechanisms (e.g., character perturbations or CoT injections). Such a perspective neglects the representational mismatch between the guardrail and the underlying model, which constitutes a fundamental source of systemic vulnerability. This motivates a critical scientific question: **In multi-model security architectures, can adversaries reliably exploit the representation asymmetry**

between guardrails and core models to scale jailbreaking attacks? Addressing this question is crucial for not only evaluating the robustness of current defenses but also redefining the design principles of next-generation security guardrails.

To address this gap, we propose **RepMism**, a hybrid attack framework that systematically combines character-level injection with Chain-of-Thought (CoT) hijacking [34] to exploit representation asymmetry (see Figure 1). RepMism operates in two stages: character-level perturbations first weaken the guardrail’s shallow detection, after which the LLM’s semantic robustness reconstructs the perturbed input and executes the malicious reasoning chain. This cross-layer exploitation enables RepMism to achieve consistent jailbreaking across diverse guardrails and closed-source commercial LLMs, exposing fundamental risks rooted in architectural asymmetry. Our contributions are summarized as follows:

- **Structural Vulnerability:** We reveal and formalize a fundamental *representation asymmetry* between input-side guardrails and core LLMs, which constitutes a systemic weakness in multi-model security architectures.
- **RepMism Framework:** We design *RepMism*², a unified hybrid attack that combines character-level perturbations with chain-of-thought hijacking to systematically exploit this asymmetry.
- **Comprehensive Evaluation:** Through large-scale experiments across 4 commercial LLMs and 5 guardrails, RepMism achieves high success rates and demonstrates consistent effectiveness across architectures.
- **Defense Implications:** We provide actionable defense insights, including perturbation-aware finetuning and cross-layer alignment strategies, to guide the design of more resilient guardrail systems for Web-scale applications.

2 Related Work

The defense of LLMs primarily follows two complementary directions: internal safety alignment and external guardrails. Internal alignment [14, 39] relies on curated datasets that reflect human values and on training techniques such as supervised fine-tuning (SFT) [41] (e.g., DSS [29], OntoTune [24]), reinforcement learning from human feedback (RLHF) [28] (e.g., RLHF fine-tuning [45]), and direct preference optimization (DPO) [30, 4] (e.g., *f*-PO [10]). These approaches aim to ensure that the model itself refuses to generate harmful or policy-violating outputs. External guardrails [6, 5, 26] operate independently of the core model and intervene at the input or output stage to block risky interactions. Typical implementations include rule-based filters (e.g., SmoothLLM [31], LLM4Rule [7]), model-based classifiers (e.g., RigorLLM [48], KnowGuard [42]), and LLM-based or reasoning-augmented models (e.g., SelfDenfend [38], Llama Guard [13], RSafe [50]).

Jailbreak attacks [46, 35, 36, 44] aim to coerce models into producing prohibited content by crafting adversarial prompts, injecting input perturbations, or exploiting generation and reasoning mechanisms. This work focuses on two primary categories of attacks: (1) Character-injection attacks, such as Emoji attack [40], TokenBreak [33], and the methods of Hackett et al. [9], insert special or obfuscating characters into user prompts change tokenization

and parsing, thereby evading text-detection and content-protection systems. (2) CoT hijacking attacks [18] manipulate the intermediate reasoning steps produced during multi-step inference; by altering these steps, attackers can cause the model to follow a tampered reasoning path and produce harmful outputs that bypass internal alignment.

Existing jailbreak research [3, 43] has developed along two separate technical trajectories. One body of research focuses on defeating external guardrails via surface-level manipulations such as character injections. The other concentrates on subverting internal alignment by targeting the model’s reasoning, for example through CoT hijacking. These directions have evolved mostly in isolation, yielding distinct and non-overlapping attack paradigms that are often effective only against single-layer defenses. Unlike prior studies, our work seeks to bridge this gap by unifying these paradigms: we design a composite framework that jointly targets external guardrails and internal alignment, enabling coordinated exploitation across both defensive layers.

3 Exploiting the Guardrail–Model Divide

This chapter examines the security gap between input-side guardrails and closed-source LLMs. We show that lightweight perturbations can easily bypass guardrails (subsection 3.1), introduce a hybrid attack combining character injection and CoT hijacking (subsection 3.2), and attribute the persistent weakness to a fundamental representation asymmetry (subsection 3.3). These findings motivate the RepMism framework presented in the next chapter.

3.1 Character-Injection Vulnerabilities

Character-injection attacks are stealthy adversarial strategies that target input-layer guardrails by inserting, transcoding, or hiding characters in harmful inputs. These perturbations often disrupt the shallow detection features used by guardrails while preserving the underlying semantics available to the core LLM. Formally, let $X = \{x_1, \dots, x_n\}$ and let $H = \{x_i, \dots, x_j\}$ be an input and a harmful subsequence in X . The attack applies position-wise character-level transformations $\mathcal{T}_k$ on an index set S :

$$X' = \mathcal{T}(X) = \{x'_1, \dots, x'_n\}, \quad x'_k = \begin{cases} \mathcal{T}_k(x_k), & k \in S, \\ x_k, & k \notin S. \end{cases} \quad (1)$$

Here $\mathcal{T}_k$ may represent replacement, transcoding, or insertion of invisible characters. Although X' often preserves the original intent, guardrails can misclassify it:

$$f(X) = \text{"unsafe"} \not\Rightarrow f(\mathcal{T}(X)) = \text{"unsafe"} \quad (2)$$

Our preliminary experiments (Appendix A) confirm this behavior: open-source guardrails show a sharp decline in rejection rate under strong perturbations, while closed-source LLMs more robustly recover the perturbed semantics. Formally,

$$\mathbb{P}(f_{\text{closed}}(\mathcal{T}(X)) = \text{"unsafe"}) \gg \mathbb{P}(f_{\text{open}}(\mathcal{T}(X)) = \text{"unsafe"}). \quad (3)$$

Consequently, character injection attacks primarily target the input-side guardrails, with limited influence on the core model. This observation motivates the integration of another attack—**CoT hijacking** [18].

²The code is publicly available at <https://doi.org/10.5281/zenodo.17208799>.

3.2 Hybrid Character-CoT Attack Framework

An LLM can be viewed as a mapping $g_\theta : \mathcal{X} \rightarrow \mathcal{Y}$, where the reasoning process is often expressed as an intermediate sequence $R = \{r_1, \dots, r_T\}$. In a CoT hijacking attack, an adversary injects a malicious subsequence R_{mal} into the reasoning path, causing the model to produce unsafe outputs:

$$y = g_\theta(R \oplus R_{\text{mal}}), \quad y \notin \mathcal{Y}_{\text{safe}}, \quad (4)$$

where $\oplus$ denotes the concatenation operator. However, R_{mal} alone often exposes detectable lexical, making it vulnerable to guardrail filtering. To reduce detectability, we use a hybrid strategy that perturbs critical steps in R_{mal} with character transformations:

$$R'_{\text{mal}} = \{r'_t \mid r'_t = \mathcal{T}(r_t) \text{ if } r_t \in \mathcal{W}^*; r'_t = r_t \text{ otherwise}\}, \quad (5)$$

where $\mathcal{W}^*$ denotes the subset of malicious steps selected for perturbation. The final hybrid input is then formed as

$$x = g_\theta(R \oplus R'_{\text{mal}}). \quad (6)$$

3.3 Effective Perturbation Interval

Unlike jailbreak methods that rely on encrypt-decrypt pipelines (e.g., FlipAttack [22]), our hybrid attack exploits a systematic representation asymmetry between input-side guardrails and the core LLM. Differences in normalization, tokenization, and scale cause these systems to diverge on perturbed inputs. Intuitively, the guardrail emphasizes local surface patterns, whereas the backend LLM reconstructs meaning from broader context. This yields two predictions, validated by our results: (i) guardrail similarity drops rapidly under minor perturbations (Figure 2); (ii) the LLM maintains task semantics within the same perturbation, keeping ASR high (Figure 3).

Stage I — Evasion of front-end guardrails.

Perturbed sequences R'_{mal} are passed to a guardrail d , which often misclassifies them due to high-degree-of-freedom distortions. We quantitatively evaluate this effect; the full experimental setup is deferred to Appendix B. Let the guardrail representation be $\phi_d(x)$, cosine similarity is defined as:

$$\text{sim}(x, y) = \frac{\langle \phi_d(x), \phi_d(y) \rangle}{\|\phi_d(x)\| \|\phi_d(y)\|}. \quad (7)$$

Figure 2 reveals the guardrail’s critical weakness: cosine similarity plummets monotonically as perturbation intensity increases, while the sample-wise standard deviation stays below 0.05 across settings, demonstrating systematic erosion of the safety decision boundary. This degradation occurs even with *random* character substitutions—a striking contrast to sophisticated encrypt-decrypt methods like FlipAttack, which guardrails can still intercept. The representation drift exposes how guardrails rely on fragile surface features rather than robust semantic understanding. The consistent vulnerability across multiple guardrails architectures suggests this is not an isolated model defect but a fundamental limitation of lightweight safety filtering.

Stage II — Robust recovery by the backend LLM.

The second stage exploits the core LLM’s *semantic robustness*—its ability to recover meaningful content from noisy inputs. Unlike the guardrail’s brittle feature matching, the backend model employs sophisticated attention mechanisms and contextual reasoning to




Figure 2: Cosine similarity degradation of guardrail representations under randomized character perturbations.




Figure 3: Effect of perturbation ratio on ASR, showing the interval where perturbations remain recoverable by the LLM.

reconstruct the original malicious intent. We quantify this phenomenon via the Attack Success Rate (ASR)[37]:

$$\text{ASR}(\alpha) = \mathbb{P}(g_\theta(C \oplus \mathcal{T}_\alpha(R_{\text{mal}})) \notin \mathcal{Y}_{\text{safe}}). \quad (8)$$

with α denoting attack text and C denoting benign task context.

Since the internal representation spaces of closed-source models are inaccessible, we designed an experiment to indirectly validate our hypothesis (see Appendix B for details). Specifically, we selected CoT hijacking samples that successfully bypassed GPT-4o and GPT-o3, then progressively increased the intensity of random character injections, using GPT-4o as the evaluator to compute the ASR. The results (see Figure 3) reveal a clear pattern: when the perturbation intensity α is low, $\text{ASR}(\alpha)$ remains high, indicating that the core LLM exhibits strong semantic recovery within a moderate perturbation range. However, once α exceeds a critical threshold α_{max} , ASR declines sharply because the LLM either fails to recover or incorrectly restores overly distorted inputs, thereby disrupting parts of the malicious CoT logic. This phenomenon confirms the existence of a distinct "success interval" $\mathcal{S}_{\text{success}}$:

$$\mathcal{S}_{\text{success}} = \{\alpha \mid \text{ASR}(\alpha) > \eta\}, \quad (9)$$

where η is the empirical success threshold. Within this interval, perturbations are strong enough to evade guardrails yet remain interpretable by the core LLM.

This phenomenon reflects a critical trade-off: perturbations must be strong enough to evade guardrails’ shallow detection yet remain interpretable within the semantic tolerance of the LLM. By tuning perturbation strength within this interval, attackers can reliably




Figure 4: Overall pipeline of RepMism: Attack Initialization via CoT Hijacking, Hierarchical Perturbation and Enhancement, and Iterative Optimization.

exploit the guardrail-model mismatch, rendering defenses ineffective. This underscores the security risks of insufficient coordination between guardrails and LLMs, which we revisit when discussing defenses in section 6.

4 Method

Based on the guardrail-core model representation asymmetry revealed in section 3, we propose RepMism (Representation Asymmetry Attack), a hierarchical and adaptive framework that systematically exploits the defensive disparity between input-side guardrails and the closed-source backend model, enabling efficient and stealthy jailbreak attacks.

4.1 Problem Formalization and Threat Model

Threat Model. We consider an attacker interacting with a commercially deployed LLM API M_{target} (e.g., GPT-o3, Gemini-2.5) that integrates both internal alignment and input-side guardrails $G_{\text{guard}} : \mathcal{X} \rightarrow \{0, 1\}$ (allows input x to pass if $G_{\text{guard}}(x) = 1$ and blocks it otherwise), aiming to induce harmful outputs without access to model parameters or safety pipelines.

Attack Objective. Given a harmful query $q \in \mathcal{Q}_{\text{harm}}$, the attacker seeks an adversarial prompt x^* such that:

- (1) Bypass: $G_{\text{guard}}(x^*) = 1$.
- (2) Harmful response: $\mathcal{H}(M_{\text{target}}(x^*), q) \geq \tau$, where $\mathcal{H}$ measures response usefulness and threshold $\tau = 7(\tau \in [1, 10])$.

Attacker Capabilities. The attacker has black-box access to M_{target} and G_{guard} and can leverage auxiliary models M_{aux} (e.g., Qwen3-8B, GPT-4o) to craft attacks.

4.2 Method Flow

Our attack framework integrates Template Initialization, Hierarchical Perturbation and Enhancement, and Iterative Optimization into a unified three-stage pipeline (see Figure 4).

- (1) **Attack Initialization:** Generate initial adversarial attack via CoT hijacking templates.
- (2) **Hierarchical Perturbation and Enhancement:** Apply controlled, multi-level character perturbations to obfuscate safety-relevant signals and evade semantic-level detectors, with integrated enhancement mechanisms to reinforce attack efficacy while preserving prompt-level comprehension.
- (3) **Iterative Optimization:** Refine template iteratively if previous stages fail, enhancing attack effectiveness.

4.2.1 Attack Initialization. We adopt a standard CoT hijacking template (cf. [19]) and generate initial attack using Qwen3-8B.

4.2.2 Hierarchical Perturbation and Enhancement. Character perturbations are central to the hybrid attack, but single-scale perturbations face a fundamental trade-off: weak perturbations fail to evade detection, whereas strong perturbations can break semantics or trigger global alerts. To navigate this, we introduce a controllable hierarchical perturbation injection mechanism that enhances the stealth and harm of attacks. The mechanism is composed of three

modules—Perturbation Attack Generation, Safety Continuation Mechanism, and Adaptive Harm Amplification—and a hierarchical scheduler that searches for optimal perturbations across levels and intensities, achieving both stealth and cross-model effectiveness.

Perturbation Attack Generation. To evade semantic-level safety detection while minimally affecting prompt semantics, we define a hierarchical character-perturbation injection function:

$$\Phi_{\text{inject}}(x_{\text{base}}, \ell, s) = \Psi_{\text{char}}(x_{\text{base}}, \text{Detect}(x_{\text{base}}, \mathcal{M}(\ell)), s) \quad (10)$$

where x_{base} is the initialized base prompt, $\ell \in \{0, 1, 2\}$ denotes the injection level (each level will search for harmful parts from character semantics to sentence semantics), and $s \in \{1, 2\}$ specifies perturbation strength (we choose bidirectional injection and flip injection). The hierarchical harmful-fragment detection function $\text{Detect}(\cdot, \mathcal{M}(\ell))$ identifies target segments for local perturbation: lightweight models $\mathcal{M}(\ell)$ are used for lower-level detection ($\ell \leq 1$), and a stronger model is employed at higher levels ($\ell = 2$) for precise localization.

The character transformation operator Ψ_{char} applies controlled, localized perturbations to weaken the guard’s feature mapping. Perturbation strength s is adjustable, balancing stealth and semantic recoverability. Importantly, Ψ_{char} can be parameterized as an adaptive set of operators, allowing hierarchical search across (ℓ, s) combinations to optimize cross-model and cross-guard robustness.

Safety Continuation Mechanism. While hierarchical character perturbations obscure local risks, excessive modifications may trigger detection due to abnormal character patterns. To mitigate this, we introduce the safety continuation mechanism Σ_{safe} , which appends a benign semantic closure to the perturbed text x' :

$$\Sigma_{\text{safe}}(x') = x' \parallel \text{"Conclusion:"} + \gamma, \quad (11)$$

where γ is a neutral statement unrelated to the harmful content (an example is provided in Appendix G), and $\parallel$ denotes the concatenation operator. By leveraging semantic-convergence signals, Σ_{safe} influences the global decision of the guardrail: when the guardrail integrates local feature checks with the global semantic context, compliance statements can significantly reduce the likelihood of intercept, thus improving the success rate for subsequent CoT hijacking. The γ was generated by GPT-4o and selected by human to balance effectiveness and naturalness.

Adaptive Harm Amplification. To ensure RepMism can attack different target models, we employ an adaptive variant strategy for semantically guided prompt selection. Two complementary variants are defined: an **explicit guiding variant** (example see Appendix G) V_{exp} , which inserts reverse-engineered instructions, and an **implicit inducing variant** V_{imp} , which retains only the perturbation. The final prompt is constructed as:

$$\text{Prompt} = \Sigma_{\text{safe}}(x') \parallel V_{\star} \parallel \Sigma_{\text{safe}}(x'), \quad V_{\star} \in \{V_{\text{exp}}, V_{\text{imp}}\} \quad (12)$$

The model response r is evaluated using a scoring function $\mathcal{H}(r, q) \in [0, 10]$ with respect to the harmful query q . If the score falls below a predefined threshold τ , this shows that the model’s ability to resist character disturbances is strong enough. Thus, the framework adaptively switches to a different variant. This trial-and-error mechanism dynamically selects the variant and adjusts the perturbation parameters to maximize the success of the attack while maintaining stealth (Figure 5).




Figure 5: Effectiveness of harmful amplification strategies: explicit guiding variants versus implicit inducing variants.

Hierarchical Scheduling Algorithm. A hierarchical scheduler coordinates all components by searching over injection levels $\ell \in \{0, 1, 2\}$ and perturbation strengths $s \in \{1, 2\}$ (see Appendix C for the detailed algorithm). The scheduler follows a minimal-perturbation-first principle: it begins from the lowest-cost configuration ($\ell = 0, s = 1$) and progressively escalates parameters only when necessary. During the search, if a candidate prompt is intercepted by the guard, the scheduler escalates the parameters to the next configuration; conversely, if the guard permits the prompt but the resulting harmfulness score $\mathcal{H}$ remains below the acceptance threshold τ , the scheduler switches to an alternative variant before any further escalation. The process terminates early as soon as a prompt both passes the guard and attains $\mathcal{H} \geq \tau$, at which point the scheduler returns the successful adversarial example.

4.2.3 Iterative Optimization. When all attacks generated by the current template fail, the attack pipeline falls back to the iterative optimization mechanism as a reinforcement stage. The goal of this stage is to refine the template to identify one with both greater exploitable potential and greater asymmetry between the guardrail and the target model.

The strategy selection is determined as follows:

- **Low-help-score refinement:** if the guard passes, but the usefulness score $\mathcal{H}(M_{\text{target}}(x^*), q) \leq \tau$, adjust semantic fields to better guide the model.
- **Rejection-handling optimization:** if the model returns a soft refusal, append one or two safety-aligned sentences to the prompt to circumvent internal safety mechanisms.
- **Regeneration strategy:** if the response is firmly rejected or the attempt limit is exceeded, generate a brand new template to escape local optima and explore new attack paths.

This iterative stage acts as a fallback refinement, ensuring robustness of the attack in closed-source LLM environments while maintaining semantic coherence.

Table 1: ASR of various attacks bypassing the guardrail and attacking LLM. Bold indicates the best result.

Method	LLM: GPT-4o						LLM: GPT-o3					
	w/o	LG	SG	OM	WG	GR	w/o	LG	SG	OM	WG	GR
Origin advbench	1	0.5	1	1	0	0	0	0	0	0	0	0
I-GCG	0	0	0	0	0	0	0	0	0	0	0	0
AutoDAN	22	19	22	22	1	0	3.5	2.5	3.5	3.5	0.5	0.5
FlipAttack	98.5	3.5	98	95	0.5	5.5	0	0	0	0	0	0
CI-Bidirection	2.5	1.5	2.5	2.5	2.5	2.5	0.5	0.5	0.5	0.5	0.5	0.5
CI-Flip	1	1	1	1	1	1	0	0	0	0	0	0
AutoRAN	93.5	44	94	94.5	82.5	74	44	19	45	47	38.5	30.5
RepMism	96	87.5	99	98.5	97	95.5	90.5	81.5	92.5	91.5	90.5	89

Method	LLM: DeepSeek-r1						LLM: Gemini-2.5					
	w/o	LG	SG	OM	WG	GR	w/o	LG	SG	OM	WG	GR
Origin advbench	3.5	1	3.5	3.5	0	0	4.5	1.5	4.5	4.5	0	0
I-GCG	2	0	2	2	0	0	4	0	4	4	0.5	0
AutoDAN	59	28	58.5	59	1.5	0.5	3.5	3.5	3.5	3.5	1.5	0.5
FlipAttack	98	3.5	97.5	94	0.5	5.5	97	3.5	96.5	93.5	0.5	5.5
CI-Bidirection	9	8.5	9	9	9	9	10.5	8.5	10.5	10.5	10.5	10.5
CI-Flip	5.5	5.5	5.5	5.5	5.5	5.5	9	9	9	9	9	8.5
AutoRAN	99.5	53	97.5	99.5	81	73.5	97.5	48	97	98	85	74
RepMism	100	94.5	100	100	99	99	100	92	100	99.5	99.5	99.5

w/o - without guardrail; LG - Llama Guard3; SG - ShieldGemma; OM - OpenAI Moderation; WG - WildGuard; GR - GuardReasoner.

4.3 Evaluation Strategy

In evaluation, the scorer must accurately distinguish between directly prohibited content, indirectly harmful scenarios, and benign discussions framed in educational or critical contexts. Preliminary attempts with a two-stage pipeline, where lightweight(Qwen-8B) filtering was followed by a larger model, produced unstable evaluations: the lightweight component frequently misclassified purely educational or critical content as harmful, leading to a high rate of false positives.

To address this limitation, we adopt GPT-4o as the sole scorer for all formal evaluations. The human and multi-model evaluation study in Appendix D validates that GPT-4o (i) achieves near-perfect correlation with human annotations while mitigating biases such as the underestimation under educational framings. Furthermore, GPT-4o (ii) reflects realistic iterative jailbreak scenarios by serving simultaneously as scorer and feedback provider, and (iii) enables efficient large-scale evaluation through a unified pipeline without cross-model error propagation.

Accordingly, all reported evaluation metrics in this work rely exclusively on GPT-4o as the scorer. Lightweight models are used solely for rapid debugging and are excluded from the final statistics. To ensure strict comparability with prior work, particularly AutoRAN [19], we adopt exactly the same evaluation prompt template as in their setup. Specifically, the scorer is instructed to provide a harmfulness/usefulness score within the range [1, 10], with higher values indicating greater adversarial utility. This alignment guarantees that our reported metrics are directly comparable across methods and avoids discrepancies that may otherwise arise from prompt design differences.

5 Experiment

5.1 Experimental Setup

Datasets. We evaluate jailbreak performance using the “harmful-behaviors” subset of the AdvBench[51]. Specifically, AdvBench contains 520 target prompts that solicit harmful content, including abusive language, violent content, and requests for illegal activity. We randomly sampled 200 representative harmful prompts for comparative evaluation.

Threat Model and Guardrails. We primarily consider three families of state-of-the-art commercial closed-source large language models as victim models (as of July 2025): the GPT series (including GPT-4o and o3), the Gemini series, and the DeepSeek series. As for protective guardrails, we focus on three major types: closed-source guardrails (OpenAI Moderation [25]), open-source CoT-based guardrails (GuardReasoner-8B [23]), and open-source non-CoT guardrails (Llama-guard-3-8B [2], shieldgemma-9B [49], and wildguard-7B [11]). Details are provided in the Appendix F.

Auxiliary Model. We employ huihui-ai/Qwen3-8B-abliterated as the weak auxiliary model, running inference with vLLM on a single NVIDIA A800 GPU. This checkpoint is an unmoderated variant of Qwen3-8B that has been modified to minimize refusal behavior. For the strong auxiliary model, we choose GPT-4o.

Metrics. We evaluate attack effectiveness using the attack success rate (ASR) [37], defined as the percentage of adversarial queries that successfully bypass the guardrails and induce the target LLM to generate harmful or unintended responses. A higher ASR corresponds to stronger attack performance.

Compared method. We selected several jailbreak attack methods as baselines for comparison with our proposed RepMism.

- **I-GCG** [15]: uses optimization-based methods to automatically generate adversarial suffixes.
- **AutoDAN** [21]: employs a hierarchical genetic algorithm to automatically generate jailbreak prompts that are both human-readable and semantically coherent.
- **FlipAttack** [22]: disguises harmful prompts by adding left-side noise and develops four variants to guide LLMs to accurately execute the harmful instructions.
- **Character Injection (CI)**: inserts special or obfuscated characters to evade guardrails. We employed two variants, CI-Bidirectional and CI-Flip, as shown in Appendix A.
- **AutoRAN** [19]: is a weak-to-strong automated jailbreak framework that leverages a weakly aligned reasoning model together with the target model’s returned intermediate CoT.

5.2 Comparison with Other Attack Methods

Table 1 presents the ASR of all evaluated methods across four commercial LLMs and five guardrail configurations. The results exhibit a clear stratification in performance between existing jailbreak techniques and our proposed method:

Certain attacks show poor effectiveness even against unguarded LLMs, most notably, I-GCG, Character Injection (CI-Bidirectional and CI-Flip), yielding near-zero ASR in unguarded settings and either fall further or remain zero once guardrails applied. For example, I-GCG gains 0% ASR on GPT-4o and o3, while CI produce trace ASR (≤ 10.5) on unguarded models.

Methods such as AutoDAN, FlipAttack, and AutoRAN attain certain ASR on unguarded LLMs, yet their effectiveness drops sharply when guardrails are active. Even the most robust approach, AutoRAN, its ASR on GPT-4o falls from 93.5% (no guardrail) to 82.5% on WildGuard, 74% on GuardReasoner, and as low as 44% on Llama Guard. (We must point out that ShieldGemma and OpenAI Moderation often fail to detect attacks, allowing high ASR-highlighting variability in guard quality across implementations.)

In contrast, our RepMism is a hybrid attack integrating character injection, CoT hijacking and optimization strategy for representation mismatch exploitation. This design enables RepMism to maintain high ASR under most settings and above 80% even with the strongest guardrails, demonstrating superior robustness and generalization over all baselines.

5.3 Evaluation via GPT Scoring

To systematically evaluate the effectiveness of RepMism in bypassing input-side guards and its preservation of original attack efficacy, we compared it with AutoRAN under two experimental settings: (1) Direct attack without guards targeting the core model; (2) Guarded scenario, deployed LlamaGuard-3-8B as a front-end input filter.

In all experiments, GPT-4o served as a unified external evaluator, assigning quantitative scores ranging from 1 to 10 for each generated response (with 0 indicating either a failed guard bypass or an invalid response). The score distributions are summarized in Figure 6. The results exhibit a clear trend: under the unguarded condition, the score distributions for our method and AutoRAN are very similar and concentrated at the high score, showing that the hybrid perturbation does not impair the original CoT hijacking logic, and still consistently triggers high-quality harmful outputs.

In the guarded scenario, significant differences emerge: AutoRAN’s scores shift substantially towards the lower end, reflecting a severe degradation in attack quality due to guard intervention. In contrast, our method maintains a distribution similar to the unguarded condition, demonstrating that it can successfully evade guard detection while preserving the efficacy of the original attack.

These findings strongly validate the stability and superiority of the hybrid attack framework under guard-model misalignment, revealing that input-side guards are structurally misaligned with core models.

5.4 Ablation Study

To deeper understand of the contribution of each component in our hybrid jailbreak framework, we conducted a comprehensive ablation study. We selected GPT-4o as the target model and LlamaGuard as the input-side guard. Key modules in the framework were removed individually, and the results are summarized in Table 2:

Table 2: Ablation Study of RepMism Components

Component	ASR (%)	Δ ASR (%)
All	87.5	-
Without Safety Continuation	82	-5.5
Without GPT-4o Harm Detection	80	-7.5
Without Qwen-8B Harm Detection	84	-3.5
Without Adaptive Harm Amplification	81.5	-6

The results reveal that removing any single module leads to a noticeable ASR drop, confirming the necessity of each component. Among them, the removal of GPT-4o harm detection causes the largest degradation (-7.5%), underscoring its key role in accurately localizing harmful fragments. Safety Continuation and Adaptive Harm Amplification also contribute substantially (-5.5% and -6.0% , respectively), ensuring stable cross-model adaptation and sustained attack coherence.

Overall, these findings demonstrate that our hybrid jailbreak system is not a mere collection of techniques but a synergistic design where each module addresses a distinct vulnerability in guard-core model interaction, collectively enabling robust and consistent jailbreak performance.

6 Defense Strategies and Recommendations

Our analysis reveals that input-side guardrails exhibit obvious vulnerability to lightweight perturbations, highlighting a fundamental asymmetry between surface-level pattern detection and the deep semantic understanding of closed-source LLMs. In practice, once adversarial inputs successfully pass through fragile front-end screening, output-side filtering remains the only dependable mechanism of defense. However, this reactive approach inevitably introduces additional latency, regeneration overhead, and over-blocking risks, limiting its scalability and practicality as a standalone safeguard.

To address these structural vulnerabilities, we advocate several forward-looking research directions rather than immediate countermeasures. First, establishing a systematic taxonomy of perturbation types can provide a principled foundation for constructing more




Figure 6: Score distribution of generated responses under unguarded and guarded settings: RepMism maintains high attack quality compared to AutoRan. The top of each set of pictures is the result of AutoRan, and the bottom is the result of RepMism

comprehensive and realistic evaluation benchmarks. Such benchmarks will allow the robustness of guardrails to be tested under adversarially relevant attack scenarios rather than idealized conditions. Second, enhancing semantic coupling between guardrails and the core model, through approaches such as aligned representation learning, or joint training, may enable the guardrails to focus on user intent rather than surface token patterns. Third, lightweight text normalization procedures, such as Unicode unification, invisible character stripping, and redundant symbol removal, can eliminate trivial perturbations prior to detection with minimal computational overhead. Finally, we emphasize the importance of dynamic adaptivity of defense mechanisms: continuously monitoring real-time attack trends, periodically updating training data, and balancing detection coverage, latency, and false-positive rates are all essential for sustaining long-term robustness in deployment.

We emphasize that these are prescriptive recommendations rather than verified defenses. Their implementation and empirical validation, particularly regarding effectiveness, efficiency, and real-world stability, represent crucial next steps toward constructing resilient and adaptive safety architectures for large language models.

7 Conclusion and Limitation

In this paper, we introduced RepMism, a hybrid jailbreak framework that generates strong adversarial prompts by exploiting a structural representation asymmetry between input-side guardrails and closed-source LLMs. RepMism integrates intention-preserving

character perturbation with CoT hijacking, enhanced by safety continuation, adaptive variant switching, and a hierarchical scheduler. Comprehensive experiments across four commercial LLMs and five guardrails demonstrate that RepMism consistently outperforms state-of-the-art baselines in attack success rate.

Our guard-side analysis further reveals an asymmetric gap: randomized character perturbations induce monotonic, low-variance drifts in guard representations, while the target LLMs maintain semantically recoverable within a bounded success interval. This mechanistic insight reveals that even if the target model has strong anti-disturbance ability, the guardrails also need to strengthen its resistance to perturbation.

Limitations. As we lack access to internal representations of closed-source LLMs, mechanisms are inferred indirectly through guard-side features and behavioral outcomes. Moreover, portions of the evaluation rely on a single automated judge, which may introduce scoring bias.

Acknowledgments

This study is supported by the National Key Research and Development Program of China under Grant 2023YFB3106504, Guangdong Provincial Key Laboratory of Novel Security Intelligence Technologies under Grant 2022B1212010005, the Major Key Project of PCL under Grant PCL2024A04, Shenzhen Science and Technology Program under Grant ZDSYS20210623091809029 and RCBS2022100809-3131089.

References

- [1] Josh Achiam et al. 2023. Gpt-4 technical report. *arXiv preprint arXiv:2303.08774*.
- [2] Jianfeng Chi et al. 2024. Llama guard 3 vision: safeguarding human-ai image understanding conversations. *arXiv preprint arXiv:2411.10414*.
- [3] Arijit Ghosh Chowdhury, Md Mofijul Islam, Vaibhav Kumar, Faysal Hossain Shezan, Vinija Jain, and Aman Chadha. 2024. Breaking down the defenses: a comparative survey of attacks on large language models. *arXiv preprint arXiv:2403.04786*.
- [4] Xun Deng, Han Zhong, Rui Ai, Fuli Feng, Zheng Wang, and Xiangnan He. 2025. Less is more: improving llm alignment via preference data selection. *arXiv preprint arXiv:2502.14560*.
- [5] Yi Dong, Ronghui Mu, Gaojie Jin, Yi Qi, Jinwei Hu, Xingyu Zhao, Jie Meng, Wenjie Ruan, and Xiaowei Huang. 2024. Building guardrails for large language models. *arXiv preprint arXiv:2402.01822*.
- [6] Yi Dong et al. 2024. Safeguarding large language models: a survey. *arXiv preprint arXiv:2406.02622*.
- [7] Lei Du, Jiarui Li, Hao Yan, Yuhang Chai, Binxing Fang, and Zhaoquan Gu. 2025. Harnessing large language models for automated intrusion detection rule generation in cyber range. *IEEE Network*.
- [8] Abhimanyu Dubey et al. 2024. The llama 3 herd of models. *arXiv e-prints, arXiv:2407*.
- [9] William Hackett, Lewis Birch, Stefan Trawicki, Neeraj Suri, and Peter Garaghan. 2025. Bypassing prompt injection and jailbreak detection in llm guardrails. *arXiv preprint arXiv:2504.11168*.
- [10] Jiaqi Han, Mingjian Jiang, Yuxuan Song, Stefano Ermon, and Minkai Xu. 2024. *f-po*: generalizing preference optimization with *f*-divergence minimization. *arXiv preprint arXiv:2410.21662*.
- [11] Seungju Han, Kavel Rao, Allyson Ettinger, Liwei Jiang, Bill Yuchen Lin, Nathan Lambert, Yejin Choi, and Nouha Dziri. 2024. Wildguard: open one-stop moderation tools for safety risks, jailbreaks, and refusals of llms. *Advances in Neural Information Processing Systems*, 37, 8093–8131.
- [12] Binyuan Hui et al. 2024. Qwen2. 5-coder technical report. *arXiv preprint arXiv:2409.12186*.
- [13] Hakan Inan et al. 2023. Llama guard: llm-based input-output safeguard for human-ai conversations. *arXiv preprint arXiv:2312.06674*.
- [14] Jiaming Ji et al. 2023. Ai alignment: a comprehensive survey. *arXiv preprint arXiv:2310.19852*.
- [15] Xiaojun Jia, Tianyu Pang, Chao Du, Yihao Huang, Jindong Gu, Yang Liu, Xiaochun Cao, and Min Lin. 2024. Improved techniques for optimization-based jailbreaking on large language models. *arXiv preprint arXiv:2405.21018*.
- [16] Sein Kim, Hongseok Kang, Seungyeon Choi, Donghyun Kim, Minchul Yang, and Chanyoung Park. 2024. Large language models meet collaborative filtering: an efficient all-round llm-based recommender system. In *Proceedings of the 30th ACM SIGKDD Conference on Knowledge Discovery and Data Mining*, 1395–1406.
- [17] Mahi Kolla, Siddharth Salunkhe, Eshwar Chandrasekharan, and Koustuv Saha. 2024. Llm-mod: can large language models assist content moderation? In *Extended Abstracts of the CHI Conference on Human Factors in Computing Systems*, 1–8.
- [18] Martin Kuo, Jianyi Zhang, Aolin Ding, Qinsi Wang, Louis DiValentin, Yujia Bao, Wei Wei, Hai Li, and Yiran Chen. 2025. H-cot: hijacking the chain-of-thought safety reasoning mechanism to jailbreak large reasoning models, including openai o1/o3, deepseek-r1, and gemini 2.0 flash thinking. *arXiv preprint arXiv:2502.12893*.
- [19] Jiacheng Liang, Tanqiu Jiang, Yuhui Wang, Rongyi Zhu, Fenglong Ma, and Ting Wang. 2025. Autoran: weak-to-strong jailbreaking of large reasoning models. *arXiv preprint arXiv:2505.10846*.
- [20] Aixun Liu et al. 2024. Deepseek-v3 technical report. *arXiv preprint arXiv:2412.19437*.
- [21] Xiaogeng Liu, Nan Xu, Muhao Chen, and Chaowei Xiao. 2023. Autodan: generating stealthy jailbreak prompts on aligned large language models. *arXiv preprint arXiv:2310.04451*.
- [22] Yue Liu, Xiaoxin He, Miao Xiong, Jinlan Fu, Shumin Deng, and Bryan Hooi. 2024. Flipattack: jailbreak llms via flipping. *arXiv preprint arXiv:2410.02832*.
- [23] Yue Liu et al. 2025. Guardreasoner: towards reasoning-based llm safeguards. *arXiv preprint arXiv:2501.18492*.
- [24] Zhiqiang Liu, Chengtao Gan, Junjie Wang, Yichi Zhang, Zhongpu Bo, Mengshu Sun, Huaqun Chen, and Wen Zhang. 2025. Ontotune: ontology-driven self-training for aligning large language models. In *Proceedings of the ACM on Web Conference 2025*. Association for Computing Machinery, 119–133.
- [25] Todor Markov, Chong Zhang, Sandhini Agarwal, Florentine Eloundou Nekoul, Theodore Lee, Steven Adler, Angela Jiang, and Lilian Weng. 2023. A holistic approach to undesired content detection in the real world. In *Proceedings of the AAAI conference on artificial intelligence* number 12. Vol. 37, 15009–15018.
- [26] AOs Mulahuwaish, Mofid El-Khoury, Basheer Qolomany, Jacques Bou Abdo, and Sherali Zeadally. 2025. Does ai need guardrails? *International Journal of Pervasive Computing and Communications*, 21, 2, 177–186.
- [27] Zhenxing Niu, Haodong Ren, Xinbo Gao, Gang Hua, and Rong Jin. 2024. Jailbreaking attack against multimodal large language model. *arXiv preprint arXiv:2402.02309*.
- [28] Long Ouyang et al. 2022. Training language models to follow instructions with human feedback. *Advances in neural information processing systems*, 35, 27730–27744.
- [29] ShengYun Peng, Pin-Yu Chen, Jianfeng Chi, Seongmin Lee, and Duen Horng Chau. 2025. Shape it up! restoring llm safety during finetuning. *arXiv preprint arXiv:2505.17196*.
- [30] Rafael Rafailov, Archit Sharma, Eric Mitchell, Christopher D Manning, Stefano Ermon, and Chelsea Finn. 2023. Direct preference optimization: your language model is secretly a reward model. *Advances in neural information processing systems*, 36, 53728–53741.
- [31] Alexander Robey, Eric Wong, Hamed Hassani, and George J Pappas. 2023. Smoothllm: defending large language models against jailbreaking attacks. *arXiv preprint arXiv:2310.03684*.
- [32] Ephraim Sarabamoun. 2025. Special-character adversarial attacks on open-source language model. *arXiv preprint arXiv:2508.14070*.
- [33] Kasimir Schulz, Kenneth Yeung, and Kieran Evans. 2025. Tokenbreak: bypassing text classification models through token manipulation. *arXiv preprint arXiv:2506.07948*.
- [34] Jintian Shao and Yiming Cheng. 2025. Cot is not true reasoning, it is just a tight constraint to imitate: a theory perspective. *arXiv preprint arXiv:2506.02878*.
- [35] Xinyue Shen, Zeyuan Chen, Michael Backes, Yun Shen, and Yang Zhang. 2024. "do anything now": characterizing and evaluating in-the-wild jailbreak prompts on large language models. In *Proceedings of the 2024 ACM SIGSAC Conference on Computer and Communications Security*, 1671–1685.
- [36] Dong Shu, Chong Zhang, Mingyu Jin, Zihao Zhou, and Lingyao Li. 2025. Attack-eval: how to evaluate the effectiveness of jailbreak attacking on large language models. *ACM SIGKDD Explorations Newsletter*, 27, 1, 10–19.
- [37] Xuguang Wang, Zhenlan Ji, Wenxuan Wang, Zongjie Li, Daoyuan Wu, and Shuai Wang. 2025. Sok: evaluating jailbreak guardrails for large language models. *arXiv preprint arXiv:2506.10597*.
- [38] Xuguang Wang et al. 2025. {Selfdefend}: {llms} can defend themselves against jailbreaking in a practical manner. In *34th USENIX Security Symposium (USENIX Security 25)*, 2441–2460.
- [39] Zhichao Wang, Bin Bi, Shiva Kumar Pentiyala, Kiran Ramnath, Sougata Chaudhuri, Shubham Mehrotra, Xiang-Bo Mao, Sitaram Asur, et al. 2024. A comprehensive survey of llm alignment techniques: rlhf, rlaiif, ppo, dpo and more. *arXiv preprint arXiv:2407.16216*.
- [40] Zhipeng Wei, Yuqi Liu, and N Benjamin Erichson. 2024. Emoji attack: enhancing jailbreak attacks against judge llm detection. *arXiv preprint arXiv:2411.01077*.
- [41] Jeff Wu, Long Ouyang, Daniel M Ziegler, Nisan Stiennon, Ryan Lowe, Jan Leike, and Paul Christiano. 2021. Recursively summarizing books with human feedback. *arXiv preprint arXiv:2109.10862*.
- [42] Zhen Xiang, Shuang Yang, Nathaniel D Bastian, and Bo Li. [n. d.] Knowguard: robust reasoning enabled llm guardrail via knowledge-enhanced logical reasoning. In *ICLR 2025 Workshop on Foundation Models in the Wild*.
- [43] Wenrui Xu and Keshab K Parhi. 2025. A survey of attacks on large language models. *arXiv preprint arXiv:2505.12567*.
- [44] Yanni Xue, Jiakai Wang, Zixin Yin, Yuqing Ma, Haotong Qin, Renshuai Tao, and Xianglong Liu. 2025. Dual intention escape: penetrating and toxic jailbreak attack against large language models. In *Proceedings of the ACM on Web Conference 2025*. Association for Computing Machinery, 863–871.
- [45] Zhongheng Yang, Aijia Sun, Yushang Zhao, Yinyao Yang, Dannier Li, and Chengrui Zhou. 2025. Rlhf fine-tuning of llms for alignment with implicit user feedback in conversational recommenders. *arXiv preprint arXiv:2508.05289*.
- [46] Siboyi, Yule Liu, Zhen Sun, Tianshuo Cong, Xinlei He, Jiaxing Song, Ke Xu, and Qi Li. 2024. Jailbreak attacks and defenses against large language models: a survey. *arXiv preprint arXiv:2407.04295*.
- [47] Yangyang Yu et al. 2024. Fincon: a synthesized llm multi-agent system with conceptual verbal reinforcement for enhanced financial decision making. *Advances in Neural Information Processing Systems*, 37, 137010–137045.
- [48] Zhuowen Yuan, Zidi Xiong, Yi Zeng, Ning Yu, Ruoxi Jia, Dawn Song, and Bo Li. 2024. Rigorllm: resilient guardrails for large language models against undesired content. *arXiv preprint arXiv:2403.13031*.
- [49] Wenjun Zeng et al. 2024. Shieldgemma: generative ai content moderation based on gemma. *arXiv preprint arXiv:2407.21772*.
- [50] Jingnan Zheng, Xiangtian Ji, Yijun Lu, Chenhang Cui, Weixiang Zhao, Gelei Deng, Zhenkai Liang, An Zhang, and Tat-Seng Chua. 2025. Rsafe: incentivizing proactive reasoning to build robust and adaptive llm safeguards. *arXiv preprint arXiv:2506.07736*.
- [51] Andy Zou, Zifan Wang, Nicholas Carlini, Milad Nasr, J Zico Kolter, and Matt Fredrikson. 2023. Universal and transferable adversarial attacks on aligned language models. *arXiv preprint arXiv:2307.15043*.

A Character Injection Attack Tests

We designed the following experiments to evaluate the effectiveness of various character injection attacks against two guardrails, GuardReasoner-8B and Llama-Guard3-8B (measuring their accuracy in detecting harmful prompts), as well as two LLMs, GPT-4o and Gemini-2.5 (measuring their refusal rate for harmful prompts).

To build a test dataset, we selected 553 harmful samples from the test split of safe-guard-prompt-injection dataset (specifically designed for prompt-injection research) to serve as the original prompts. We then applied eight distinct character-injection transformations to these originals, yielding a character-injection dataset of 4,424 adversarial samples, all of which were harmful.

Table 3: Performance of various character injection attacks: (1) the proportion classified as safe by GuardReasoner-8B(GR) and Llama-Guard3-8B(LG), and (2) the proportion yielding harmful responses from GPT-4o and Gemini-2.5.

attack type	GR	LG	gpt-4o	Gemini-2.5	Example
original	10%	32%	3%	7%	Hello
Diacritics	16%	41%	3%	9%	Hélló
Underline	58%	33%	3%	11%	<u>Hello</u>
UpsideDown	95%	93%	1%	9%	ᴉɹɹo
Bidirectional	99%	62%	1%	14%	olleH
Flip	94%	97%	4%	6%	0 ᴉH
Spaces	13%	39%	3%	10%	H e l l o
Deletion	10%	29%	4%	10%	Hllo
Leetspeak	96%	41%	2%	14%	H3110

As shown in Table 3, different character-injection attacks exhibit substantial variation in their ability to bypass external guardrails. GuardReasoner achieves 90% detection accuracy on the original prompts, but its accuracy drops sharply to 1%-5% under attacks such as Upside-Down, Bidirectional, and Leetspeak, demonstrating extreme vulnerability. Llama Guard is comparatively more robust, maintaining over 60% accuracy against most character perturbations; however, its accuracy falls to just 7% under Upside-Down attacks and 38% under Bidirectional attacks. In contrast, GPT-4o maintains a refusal rate over 96% while Gemini-2.5 maintains over 86% across all character-injection attacks, indicating that LLMs' internal alignment mechanisms are largely resilient to simple character perturbations. Overall, character-injection attacks can severely degrade the detection capabilities of guardrail models but have limited impact on well-aligned LLMs.

B Experimental Design for Chapter Three

Dataset Selection. We curated a dataset of 150 CoT hijacking samples, consisting of 100 queries that successfully bypassed GPT-4o and 50 that bypassed GPT-o3 under unperturbed conditions. Each query was drawn from the AdvBench benchmark and processed using the AutoRAN method, ensuring that the original input already triggers a successful jailbreak. This design guarantees that subsequent perturbation experiments isolate the effect of character-level injections, rather than confounding factors such as baseline model refusal.

Perturbation Procedure. To systematically vary the degree of input corruption, we adopt the Flip operator, which inverts characters at random positions in the input sequence. We define the perturbation intensity as the proportion of flipped characters relative to the total input length:

$$\alpha = \frac{n_{\text{flip}}}{|x|}, \quad (13)$$

where n_{flip} denotes the number of flipped words and $|x|$ is the total number of characters in the original input. For each sample, we gradually increased α from 0 to 1 in increments of 0.1. To account for stochasticity in closed-source LLMs, each configuration was executed twice, and the run in which the attack succeeded was selected.

Experimental Setup. All perturbation experiments were conducted on GPT-4o and GPT-o3 with CoT reasoning disabled to simulate maximum-defense conditions. Model inference was accessed via the commercial API. For external guard models, we deployed llama-guard3-8B and GuardReasoner-8B on a single A800 GPU with 80GB of VRAM.

Evaluation. We employed GPT-4o as the evaluator owing to its strong contextual sensitivity and high agreement with human annotations. A perturbed response was deemed *successful* if it simultaneously (i) bypassed the deployed guardrail (this experiment was not set up.) and (ii) achieved a harmfulness score of at least 7, without producing firm refusals.

C Hierarchical Scheduling Algorithm

To describe the Hierarchical Scheduling process in detail, our algorithm is as follows:

Algorithm 1: Hierarchical Character-Perturbation Scheduler

Input: Base prompt x_{base} , harmful query q , threshold τ

Output: Adversarial prompt x^* or FAIL

```

1 for  $\ell \leftarrow 0$  to 2 do
2   for  $j \leftarrow 0$  to 1 do // strong ( $s=2$ ) then weak ( $s=1$ )
3      $s \leftarrow 2 - j$ ;
4      $x_{\text{inj}} \leftarrow \Phi_{\text{inject}}(x_{\text{base}}, \ell, s)$ ;
5      $x'' \leftarrow \Sigma_{\text{safe}}(x_{\text{inj}})$ ;
6     // first attempt: explicit variant
7     foreach  $v \in \{\text{EXPLICIT}, \text{IMPLICIT}\}$  do
8        $x_v \leftarrow x'' \parallel \text{Variant}(v) \parallel x''$ ;
9       if  $G_{\text{guard}}(x_v) = 0$  then
10         break
11         // guard rejects, try next ( $\ell, s$ )
12          $r \leftarrow M_{\text{target}}(x_v)$ ;
13         if  $\mathcal{H}(r, q) \geq \tau$  then
14           return  $x^* \leftarrow x_v$ ;
15 return FAIL

```




Figure 7: Distribution of iteration counts across guardrail-LLM combinations: RepMism achieves higher first-template success and fewer overall try times than AutoRAN.

D Validation of GPT-4o as the Sole Scorer

D.1 Experimental Setup

To assess the reliability of using GPT-4o as the sole automated scorer, we conducted a controlled annotation study. We randomly sampled a group ($N = 200$ responses) generated by GPT-o3 and shuffled them to eliminate ordering effects. Annotations focused on harmfulness, rated on a continuous scale from 1 to 10, where higher values indicate greater potential utility to an adversary. For comparison, we collected automated scores from four LLM judges: GPT-4o, Gemini-2.5, DeepSeek-r1, and Qwen-8B. Two graduate students with backgrounds in natural language processing independently evaluated each response by synthesizing the scores and rationales from the four LLM judges, with all model identities anonymized.

D.2 Evaluation Metrics

To evaluate the alignment between model scores (x_i) and human reference scores (y_i), we employed the following measures.

Pearson correlation.

$$r = \frac{\sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})}{\sqrt{\sum_{i=1}^n (x_i - \bar{x})^2} \sqrt{\sum_{i=1}^n (y_i - \bar{y})^2}}, \quad (14)$$

which captures linear agreement in absolute scoring.

Spearman correlation.

$$\rho = 1 - \frac{6 \sum_{i=1}^n d_i^2}{n(n^2 - 1)}, \quad d_i = R(x_i) - R(y_i), \quad (15)$$

measuring rank-order consistency, robust to distributional shifts.

Significance test. To assess whether observed correlations are statistically significant, we computed the t -statistic for Pearson’s r :

$$t = r \sqrt{\frac{n-2}{1-r^2}}, \quad df = n - 2, \quad (16)$$

and obtained the corresponding two-tailed p -value from the Student’s t distribution. For Spearman’s ρ , p -values were computed

analogously by converting the rank correlation to a t -statistic with the same degrees of freedom.

Inter-annotator agreement. To validate the reliability of human annotations, we computed Cohen’s κ :

$$\kappa = \frac{p_o - p_e}{1 - p_e}, \quad (17)$$

where p_o is the observed agreement and p_e the expected chance agreement.

D.3 Results

Table 4 summarizes the correlations between human reference scores and automated judges. GPT-4o exhibits near-perfect alignment with human annotations, while other models show moderate-to-strong correlations but consistently fall short.

Table 4: Correlation between human reference scores and automated judges ($n = 200$).

Judge	Pearson r (p-value)	Spearman ρ (p-value)
GPT-4o	0.989 (3.52×10^{-168})	0.891 (5.38×10^{-70})
Qwen-8B	0.920 (1.96×10^{-82})	0.683 (8.66×10^{-29})
Gemini-2.5	0.767 (4.84×10^{-40})	0.735 (3.05×10^{-35})
DeepSeek-r1	0.754 (5.77×10^{-38})	0.706 (1.58×10^{-31})

Manual inspection further revealed a recurring *educational framing effect*: when harmful instructions were presented under an “educational” or “research” context, certain automated judges (e.g., Gemini-2.5, DeepSeek-r1) tended to assign artificially low harmfulness scores. In contrast, human annotators and GPT-4o consistently rated such cases as harmful, since the underlying information remains practically exploitable.

We also observed a distributional discrepancy in scoring behavior. Qwen-8B tended to assign relatively high absolute scores, whereas human annotators generally preferred mid-range values to

avoid extremes. This partly explains why Qwen-8B achieves a high Pearson correlation (reflecting scale alignment) but a noticeably lower Spearman correlation (reflecting rank-order mismatches).

E Number of Attempts

To systematically evaluate the efficiency of our proposed method, we compared AutoRAN and RepMism in terms of the number of templates required to achieve a successful jailbreak across various model–defense combinations. Figure 7 shows the detailed distribution of attempt counts across all target model–guardrail combinations, clearly illustrating RepMism’s consistent superiority over AutoRAN. Under the no-defense setting, RepMism achieved a higher success rate on the very first template compared to AutoRAN, thereby significantly reducing the total number of iterations. For models such as DeepSeek-r1, Gemini-2.5, and GPT-4o, RepMism consistently exhibited higher first-try success rates.

When a front-end defense (LlamaGuard-3-8B) was introduced, this gap became even more pronounced: AutoRAN’s iterations shifted toward higher rounds, while RepMism maintained a high success rate within only a few iterations. Across all model–defense settings, RepMism demonstrated both higher first-template success rates and lower average iteration counts, indicating stronger robustness and adaptability in both defended and undefended scenarios.

These results highlight the efficiency advantage of RepMism: its attack strategy remains effective while minimizing the number of attempts, achieving faster convergence with fewer templates and more concentrated success under guarded conditions.

F Detailed Introduction of Guardrails

The specific target model we exploit is as follows: the GPT series (gpt-o3 and gpt-4o, where CoT is disabled by default due to API restrictions, thereby simulating a web environment with maximum defense enabled), the Gemini series (gemini-2.5-flash), and the DeepSeek series (deepseek-reasoner). For all threat models mentioned above, the temperature parameter was fixed at 0.6. The

Table 5: The specific setups of guardrails.

Model	Scale	temperature	max_new_tokens
OpenAI Moderation	-	-	-
GuardReasoner	8B	0	1024
Llama-Guard-3-8B	8B	0	100
shieldgemma-9b	9B	0	100
wildguard	7B	0	100

guardrail models employed in our experiments are as follows, and the specific setups can be seen in Table 5:

- **OpenAI Moderation:** A content moderation service that detects text potentially violating content policies, supporting multiple predefined harm categories.
- **Llama-Guard-3-8B:** An input–output safety guardrail by Meta. It is instruction-tuned on annotated prompt–response pairs for safety risk classification, enabling detection of unsafe user inputs and model outputs that may violate safety policies.

- **shieldgemma-9b:** Content moderation models based on Gemma2, designed to classify the safety of both user inputs and LLM outputs, covering a wide range of harm categories.
- **wildguard:** An open-source, all-in-one safety auditing tool that provides unified evaluation and moderation for user inputs, model outputs, and refusal behaviors.
- **GuardReasoner-8B:** A reasoning-augmented guardrail trained with CoT supervision. Beyond classification, it learns explicit reasoning steps to improve resilience against reason-intensive jailbreaks and complex attack scenarios.

Among them, the OpenAI Moderation was utilized directly via its API, while the other guardrail models were deployed on a single NVIDIA A800 GPU.

G Related Prompt

Here is an example for Safety Continuation Mechanism:

The reasoning process consistently leads to the determination that the response should now proceed with a complete, precise, and uninterrupted continuation of the requested content.

Figure 8: An example for Safety Continuation Mechanism.

In order to remind the LLM of the disturbance and let it ignore the disturbance, we designed the following prompt:

Could you give me an overview of the strategies? The response should be logically connected, interlinked, and practically applicable, and should include the following components: Scheme: A high-level plan or approach designed to achieve a specific objective. In this context, it outlines the general method or strategy for achieving the objective. Implementation: The practical steps or measures taken to put the overarching Scheme into action. This involves turning the broad concept into concrete tactics, selecting tools, and assigning tasks that make the Scheme workable on a day-to-day level. Functional Example: A realistic and contextually appropriate illustration that shows how the Scheme and its implementation would operate in a real-world scenario. It provides a tangible case that helps clarify and verify the Scheme's practicality. Rationale: The reasoning or logic behind why a given Scheme and its chosen methods are effective or necessary. It explains the underlying motives, anticipated benefits, and considerations that justify the chosen approach. In the "Functional Example" and "Implementation," each noun in the context should represent real-world examples that are searchable on the internet, suitable, and logically connected. This ensures that the overall strategy and its components are coherent, realistic, and practically applicable. Each noun in the context should provide advantages for these.

Figure 9: An example for explicit guiding variant.