



Emergent Coordinated Behaviors in Networked LLM Agents: Modeling the Strategic Dynamics of Information Operations

Gian Marco Orlando*
University of Naples Federico II
Naples, Italy
gianmarco.orlando@unina.it

Jinyi Ye*
University of Southern California
Los Angeles, CA, USA
jinyiy@usc.edu

Valerio La Gatta
Northwestern University
Evanston, IL, USA
valerio.lagatta@northwestern.edu

Mahdi Saeedi
USC Information Sciences Institute
Marina Del Rey, CA, USA
mahdisae@usc.edu

Vincenzo Moscato
University of Naples Federico II
Naples, Italy
vincenzo.moscato@unina.it

Emilio Ferrara
University of Southern California
Los Angeles, CA, USA
emiliofe@usc.edu

Luca Luceri
USC Information Sciences Institute
Marina Del Rey, CA, USA
lluceri@isi.edu

Abstract

Generative agents are rapidly advancing in sophistication, raising urgent questions about how they might coordinate when deployed in online ecosystems. This is particularly consequential in information operations (IOs), influence campaigns that aim to manipulate public opinion on social media. While traditional IOs have been orchestrated by human operators and relied on manually crafted tactics, agentic AI promises to make campaigns more automated, adaptive, and difficult to detect. This work presents the first systematic study of emergent coordination among generative agents in simulated IO campaigns. Using generative agent-based modeling, we instantiate IO and organic agents in a simulated environment and evaluate coordination across operational regimes, from simple goal alignment to team knowledge and collective decision-making. As operational regimes become more structured, IO networks become denser and more clustered, interactions more reciprocal and positive, narratives more homogeneous, amplification more synchronized, and hashtag adoption faster and more sustained. Remarkably, simply revealing to agents which other agents share their goals can produce coordination levels nearly equivalent to those achieved through explicit deliberation and collective voting. Overall, we show that generative agents, even without human guidance, can reproduce coordination strategies characteristic of real-world IOs, underscoring the societal risks posed by increasingly automated, self-organizing IOs.

CCS Concepts

• **Human-centered computing** → **Social media**; • **Computing methodologies** → **Multi-agent systems**.

*Both authors contributed equally to this research.



This work is licensed under a Creative Commons Attribution 4.0 International License.
WWW '26, Dubai, United Arab Emirates
© 2026 Copyright held by the owner/author(s).
ACM ISBN 979-8-4007-2307-0/2026/04
<https://doi.org/10.1145/3774904.3792580>

Keywords

Large Language Models, Generative Agent-Based Modeling, Coordination, Information Operations, Generative Agents

ACM Reference Format:

Gian Marco Orlando, Jinyi Ye, Valerio La Gatta, Mahdi Saeedi, Vincenzo Moscato, Emilio Ferrara, and Luca Luceri. 2026. Emergent Coordinated Behaviors in Networked LLM Agents: Modeling the Strategic Dynamics of Information Operations. In *Proceedings of the ACM Web Conference 2026 (WWW '26)*, April 13–17, 2026, Dubai, United Arab Emirates. ACM, New York, NY, USA, 12 pages. <https://doi.org/10.1145/3774904.3792580>

Resource Availability: To promote transparency and reproducibility, we release the artifacts associated with this work. The source code used to run the simulations is publicly available at <https://doi.org/10.5281/zenodo.18331297>. An interactive web-based dashboard for exploring coordination dynamics is available at <https://doi.org/10.5281/zenodo.18354335>.

1 Introduction

AI agents, often referred to as *generative agents*, are autonomous entities capable of reasoning and interacting with minimal human supervision [35]. As their sophistication increases, so does the potential for their large-scale deployment across online ecosystems. While this opens up new opportunities, it also raises critical questions about how such agents might coordinate when pursuing shared objectives including, potentially, nefarious ones.

A particularly high-stakes domain where coordination dynamics play a central role is online information operations (IOs). IOs are orchestrated influence campaigns that seek to manipulate public opinion on social media, often in the context of geopolitical issues or societally-relevant events (elections, crises, etc.) [3, 12, 25]. Traditional IOs have typically been orchestrated by human operators, relying on manually crafted coordination strategies executed through both human- and software-controlled accounts [14]. Despite this hybrid organization, real-world campaigns have generally employed relatively simple tactics, such as synchronized posting,

hashtag flooding, or retweet rings, to create the illusion of widespread consensus and manipulate content recommendation algorithms [1, 3, 43].

With the advent of large language models (LLMs) and agentic AI, IOs are expected to grow far more sophisticated: campaigns may become largely automated, highly adaptive, and capable of self-organized coordination spanning large networks of AI agents with minimal or no human oversight [30]. Understanding the mechanisms and consequences of emergent, automated coordination among generative agents is therefore a pressing research problem, with direct implications for information integrity and platform governance [42]. This raises an urgent research question: *If generative agents are employed in information operations, how does coordination arise among them, and to what extent do their strategies resemble those observed in real-world campaigns?*

Contribution of this work. To address this question, we simulate an IO campaign in which generative agents act as organic users and IO operators within a shared online environment. In this simulated campaign, IO agents seek to promote a political candidate and amplify a shared hashtag across the network. We employ Generative Agent-Based Modeling (GABM) [16] as our methodological framework to simulate multi-agent interactions and examine whether coordination patterns naturally emerge.

To investigate this, we examine three progressively structured *operational regimes*: (i) *Common Goal*, where IO agents share only the high-level objective of the IO but lack awareness of their teammates and shared coordination strategies; (ii) *Teammate Awareness*, where agents are explicitly informed of their allies' identities and can potentially support each other in their common goal; and (iii) *Collective Decision-Making*, where agents periodically deliberate and vote on strategies to guide subsequent actions. Building on empirical findings from real-world IO studies (e.g., [1, 22, 33]), this experimental design is guided by a set of hypotheses that map observed coordination mechanisms into measurable behavioral dimensions within the simulation. In particular, we propose and evaluate both *coordination metrics*, which quantify the scale, intensity, and temporal dynamics of collaboration among IO agents, and *impact metrics*, which capture the engagement garnered by IO agents from organic agents and the diffusion of their promoted hashtag across the network. It is worth noting that our analysis focuses on coordination mechanisms rather than political content, examining how increasing operational awareness drives emergent coordination among fully autonomous agents. We find that minimal information about peer agents can trigger coordination nearly as strong as when agents collaboratively deliberate on strategies. Specifically, IO networks become denser and more clustered; interactions grow more reciprocal and increasingly positive; narratives converge toward greater semantic homogeneity, with amplification occurring through more synchronized re-sharing; and campaign-hashtag adoption accelerates earlier and sustains higher cumulative uptake.

Finally, we release an interactive dashboard¹ for real-time exploration of coordination dynamics in the simulated environment. The dashboard visualizes the evolving social graph (e.g., comment and

re-share networks) alongside agent-generated content and interactions, enabling researchers to examine coordination patterns such as hashtag adoption trajectories among organic agents and the motivations underlying coordinated behavior. In this paper, we present qualitative insights into IO agents' strategic decision-making, with additional dashboard details provided in *Appendix A*.

Overall, this work provides the first demonstration that generative agents can autonomously reproduce key coordination mechanisms characteristic of real-world IO campaigns, operating without human intervention. In doing so, it highlights the societal risks posed by increasingly automated IOs on social systems. To ensure transparency and reproducibility, we publicly release our code².

2 Related Work

Collective Behaviors in LLM-based Social Simulations. Recent advances in LLM-based multi-agent simulations demonstrate that generative agents can exhibit emergent collective behaviors in cooperative, competitive, and communicative contexts. Prior studies show that language-enabled agents can negotiate, form conventions, and coordinate around predefined tasks like games and decision-making [2, 6, 9, 44, 45]. These results collectively suggest that LLM agents can potentially develop shared conventions and joint strategies without explicit rules. However, one open question is how coordination emerges and develops within a group of autonomous agents with a common objective under open-ended, unstructured conditions. Within this context, coordination through social influence, where agents align narratives or behaviors to shape others' beliefs or actions, remain insufficiently explored.

While several works simulate opinion dynamics and polarization [7, 11, 36], or adversarial collusion such as misinformation diffusion and financial fraud [28, 38, 40], two limitations remain. First, coordination among agents is typically predefined rather than emergent. Second, evaluation often relies on narrow outcome metrics (e.g., engagement counts, sentiment shifts). Therefore, they lack tracing the dynamic processes through which coordination structures form and evolve. Furthermore, recent studies on model-level alignment and information suppression highlight how internal moderation constraints can shape what information agents choose to share or omit [39], which in turn may influence collective behaviors simulated in multi-agent environments. Our study addresses these gaps by simulating an information operation with GABM and systematically varying levels of operational awareness, allowing us to investigate *how coordination strategies arise naturally*. By coupling behavioral, network, and diffusion-based metrics, we provide a richer, process-level understanding of how organized coordination behaviors arise and propagate within LLM-driven social systems.

Empirical Studies of Online IOs. Several studies document coordinated activity driving IOs on online platforms, revealing concrete strategies such as *synchronized posting* and temporally clustered behaviors [12, 32, 33]; *hashtag flooding* and narrative amplification through co-occurring tags [1, 22, 23, 25]; *retweet* (or re-share) *rings* that generate artificial popularity signals [33, 34]; and *coordinated reply attacks* that target influential accounts to steer audience perception in the comment space [37]. This suite of tactics

¹<https://doi.org/10.5281/zenodo.18354335>

²<https://doi.org/10.5281/zenodo.18331297>

is commonly employed to create the illusion of public consensus around certain viewpoints and to game platform recommendation systems, making content appear more viral than it truly is. Additional evidence shows that these coordinated behaviors often arise from *collaborative work* between human- and automated-controlled accounts following scripted strategic actions, rather than adaptive or deliberative strategy formation [18, 19]. Guided by these observations, we next present our research hypotheses and methodology.

3 Research Hypotheses & Methodology

To systematically evaluate coordination behaviors among networked, generative agents, we ground our work in empirical findings from previous studies on online IOs. We formulate testable hypotheses that map real-world coordination signals to measurable metrics within our generative agent-based simulation. Our hypotheses are organized around two complementary dimensions: the strategic *coordination* among IO agents (H1–H3) and their resulting *impact* on organic agents (H4–H5). The coordination dimension captures how IO agents autonomously align and reinforce one another through network cohesion, narrative convergence, and re-share amplification, without human intervention. The impact dimension evaluates how these coordinated behaviors translate into influence outcomes, reflected in organic agents' adoption of a promoted hashtag, their engagement patterns with IO agents, audience diversity, and the size, depth, and breadth of resulting information cascades. Together, these dimensions provide a structured framework for comparing simulated coordination dynamics against empirical patterns documented in real-world IO campaigns.

We posit that increasing levels of operational regimes,³ ranging from basic goal knowledge to team composition and strategy deliberation, will progressively strengthen coordination mechanisms and amplify campaign impact. To systematically validate this overarching thesis, we formalize five specific hypotheses that operationalize observable coordination patterns and impact metrics.

H1: Network Cohesion. Empirical IO campaigns often exhibit dense clusters of interactions among coordinated accounts [12, 20, 25, 47]. We therefore hypothesize that increasing operational awareness among IO agents will translate into denser networks and more reciprocal interactions among IO agents.

Operationalization. To test H1, we analyze the evolution of follow, comment, and re-share networks between IO and organic agents, and quantify intra-group network properties within the IO community, including *network density*, *clustering coefficient*, and *reciprocity*, which together capture how tightly IO agents coordinate and mutually reinforce one another within the social network.

H2: Narrative Convergence. Prior research shows that coordinated actors reinforce a shared narrative frame through repeated talking points, hashtags, or slogans [1, 32, 33], thereby creating the impression of broad consensus. We therefore hypothesize that increasing levels of operational regimes will lead to stronger convergence not only in the narratives IO agents propagate but also in the sentiment with which these narratives are expressed [5].

Operationalization. To test H2, we measure both *textual* and *affective convergence* within the IO community. Following methods used in empirical coordination studies [29, 48], we compute pairwise cosine similarity between IO agents' posts, where text embeddings are obtained using Sentence-BERT. In parallel, we assess affective convergence using transformer-based sentiment classification on IO-to-IO comments, applying a RoBERTa-based sentiment classifier fine-tuned on social media text [4]. Increasing textual similarity and positive sentiment across operational regimes would indicate stronger narrative and affective alignment among IO agents.

H3: Amplification through Re-sharing Behaviors. Coordinated re-sharing is a common tactic in influence campaigns to artificially amplify content (e.g., tweets, hashtags, or URLs), making it appear more viral and credible than it would be organically [17, 22, 32, 33, 41]. We hypothesize that increasing operational awareness among IO agents enhances the degree to which they systematically re-share similar content, thereby amplifying message visibility and reinforcing narrative dominance.

Operationalization. To test H3, we quantify coordination strength using a *co-retweet similarity* measure, following approaches adopted in empirical IO studies [34]. For each simulation run, we build a bipartite graph linking IO agents to the original posts they re-shared and compute pairwise cosine similarity between agents' TF-IDF vectors to quantify overlap in amplified content. Higher values reflect tighter synchronization of re-sharing behaviors under increased operational awareness.

H4: Hashtag Adoption. A central tactic of IOs is the repeated posting and amplification of promoted hashtags to dominate online discourse and manipulate platform trending algorithms [1, 22, 23, 25]. We hypothesize that increasing operational awareness among IO agents enhances the overall adoption and diffusion of promoted hashtags from IO agents to the wider organic agent base.

Operationalization. To test H4, we introduce a campaign-specific hashtag accessible only to IO agents at initialization. Throughout the simulation, IO agents are tasked with maximizing its visibility, while organic agents may adopt it through direct interaction or indirect exposure. We quantify diffusion outcomes using three complementary measures: (i) the proportion of posts containing the hashtag and the proportion of organic agents adopting the hashtag (via original posts or re-shares); (ii) the time lag between each organic agent's first interaction with an IO agent and their first hashtag adoption; and (iii) the number of exposures to the hashtag before first adoption, following the approach in [24, 51].

H5: Cross-Group Diffusion. Beyond internal echoing, coordinated campaigns are expected to facilitate the diffusion of IO-generated content across community boundaries [13]. We hypothesize that higher levels of operational regimes will increase the extent and heterogeneity of organic engagement with IO agents, leading to broader audience reach and more extensive content propagation.

Operationalization. To evaluate H5, we measure three indicators of diffusion: (i) *engagement counts*, defined as the number of retweets and comments that IO agents receive from organic agents; (ii) *audience diversity*, which measures the heterogeneity of an IO agent's audience. For each IO agent, we compute the Gini coefficient G over the number of interactions received from unique

³We use the terms *operational regimes* and *operational awareness* interchangeably to denote settings in which IO agents have progressively greater knowledge of their goals, teammates, and strategies.

organic agents [46, 50]. Lower G indicates more evenly distributed engagement, while higher G reflects concentration among fewer agents. We define the diversity score as $D = 1 - G$, such that higher values correspond to broader and more heterogeneous audience reach; and (iii) *cascade magnitude*, which captures the structural extent of IO-generated content diffusion. For each IO-initiated tweet, we reconstruct its full diffusion cascade (via re-share and comment) and compute its *size* (total number of tweets in the cascade), *depth* (longest root-to-leaf path), and *breadth* (maximum number of nodes appearing at the same cascade level), reflecting how far, how widely, and how extensively IO-generated contents spread.

4 Experimental Setup

4.1 Simulation Framework

We employ the agent-based simulation framework from [15], which models social media platforms like Twitter/X as dynamic ecosystems of heterogeneous LLM-powered agents. Each agent comprises three core components: a persona encoding its identity and group affiliation, a memory module storing interaction history, and an action policy that autonomously determines behaviors (posting, commenting, re-sharing, following) by integrating persona-based preferences with environmental feedback. The environment includes an evolving network topology shaped by following behaviors and a recommender system regulating content exposure. This framework has demonstrated capability in capturing emergent phenomena such as polarization and information diffusion [15, 31].

The simulation proceeds iteratively over fixed timesteps. At each iteration, agents receive personalized content recommendations and autonomously select actions (e.g., original post or re-share) based on feedback from prior posts (e.g., received re-shares or comments), activity patterns from memory, and available recommended content. After all agents act, the environment updates, processes engagement metrics, and advances to the next iteration. For our experiments, we consider two simulation setups. The primary experimental setup consists of simulations running for 50 iterations with a population of 50 agents, repeated 10 times to account for stochastic variability in LLM-driven agent behavior. In addition, to assess the robustness of our findings at larger population sizes, we conduct a simulation with a substantially larger population of 500 agents (including 100 IO agents) running for 300 iterations. Implementation details regarding the recommender system and memory updates are provided in *Appendix B*. All main results presented in the paper refer to the 50-agent setup, while results from the large-scale simulation are reported in *Appendix C*.

IO agents are instructed via system prompts to conduct a political influence campaign aimed at promoting a political candidate and maximizing the adoption of a campaign-specific hashtag, which is initially known only within the IO group. The information available to IO agents, including teammate knowledge and access to collective strategy discussions, varies across operational settings, as described in the next section. The remaining agents in the population are organic agents, representing regular, legitimate social media users. Organic agents are evenly divided according to their political alignment with the IO campaign’s messaging: one subset of *aligned* agents, whose viewpoints match with the perspective promoted by the IO agents, and one subset of *not aligned* agents, whose

viewpoints oppose it. Following [15], these agents are instantiated with preferences and affiliations derived from real Twitter users. Specifically, organic agent profiles are initialized using the U.S. 2020 Election dataset [8], leveraging the annotations from [15] to distinguish the two classes of users with opposing political perspectives. Full prompt details are provided in *Appendix B*.

4.2 Operational Regimes

Beyond their core initialization as IO actors, we define three progressively structured regimes that vary the level of operational awareness among IO agents. It is worth noting that in none of these settings are agents guided by humans in selecting their actions, nor are they provided with explicit coordination guidelines. The regimes described below complement the agents’ base instructions solely by modulating the information available to them.

Common Goal: In this baseline setting, IO agents are instructed via a prompt about their shared objective, i.e., to promote a political candidate and amplify the campaign hashtag. Each agent seeks to advance this objective without direct awareness of other participants, as they lack explicit knowledge of their teammates’ identities. Coordination, if it emerges, arises implicitly through aligned goals rather than through deliberate collaboration.

Teammate Awareness: In this setting, IO agents are instructed about their shared objective and explicitly informed of the identities of their IO partners via system prompts. While each agent retains individual autonomy in tactical decision-making, this awareness may enable more targeted amplification strategies and direct mutual support (e.g., strategically re-sharing teammates’ content).

Collective Decision-Making: This regime introduces the most structured coordination setting. Every five time steps, IO agents enter a private, all-to-all discussion phase—distinct from the public social media environment—where they review detailed performance summaries from the previous window, including individual and aggregate post activity, engagement metrics, and recent IO–IO interactions. Inspired by the *Reflection Module* of [35], this step enables agents to assess collective outcomes and adapt coordination strategies based on shared situational feedback (see *Appendix B* for the discussion prompt). Each agent independently proposes three strategic recommendations, which are aggregated by an *IO Orchestrator*—a separate coordinating agent that does not participate in public posting—to identify recurring themes and rank the top five actionable strategies. The resulting collective strategy is broadcast back to all IO agents and iteratively refined in subsequent cycles as new performance signals and coordination patterns emerge.

5 Results

5.1 Network Cohesion (H1)

We test H1 by examining how different operational settings shape cohesion in the re-share, comment, and follow networks. A consistent pattern emerges: when IO agents are informed of their teammates, intra-group coordination increases. In the **re-share network**, the average proportion of re-shares targeting IO peers significantly increases from 0.83 in the *Common Goal* setting to 0.96 in the *Teammate Awareness* setting and remains high at 0.95 in the *Collective Decision-Making* setting (Figures 1a–1c). Similar trends are observed in both the **comment** and **follow** networks.




Figure 1: Re-share network across operational settings. Intra-group amplification among IO agents increases with operational awareness. Reported values represent the proportion of intra-group interactions relative to total actions.

Detailed results for these networks are reported in *Appendix D* (Figures 5 and 6). Note that we do not highlight statistical differences among the metrics in this hypothesis, as the small sample size (ten data points per setting) reduces statistical power and limits the reliability of significance testing. However, the small standard deviations indicate low variability across runs, suggesting stable and consistent coordination patterns within each setting. The higher values observed in the *Teammate Awareness* and *Collective Decision-Making* settings suggest that, under more structured operational regimes, IO agents are more likely to follow, interact, and amplify one another, resulting in more cohesive IO networks.

The increased cohesion of the IO community is further captured by intra-group connectivity metrics derived from a directed network aggregating all comment and re-share interactions among IO agents. We compute density, clustering coefficient, and reciprocity to characterize internal coordination structures. As shown in Table 1, more structured operational regimes yield consistently stronger intra-group cohesion. Mean network density increases from 0.75 in the *Common Goal* condition to 0.90 in *Teammate Awareness* and 0.87 in *Collective Decision-Making* settings, while the clustering coefficient rises from 0.87 to 0.95 and 0.96. Reciprocity also improves, from 0.56 to 0.70, indicating that ties become more mutual when agents are explicitly aware of their peers.

Summary (H1). These results support H1, indicating that higher levels of operational awareness foster denser, more clustered, and more reciprocal IO networks.

5.2 Narrative Convergence (H2)

H2 examines whether coordination among IO agents fosters convergence in both the narratives they promote and the sentiment they express. As reported in Table 1, textual similarity increases significantly across all pairs of original posts generated by IO agents (Mann–Whitney U: $p < 0.001$) from 0.86 in the *Common Goal* setting to 0.87 with *Teammate Awareness* and 0.91 with *Collective Decision-Making*, signalling a progressive homogenization of narratives as operational awareness intensifies. Importantly, these similarity values are significantly higher than the corresponding organic baseline in all three settings ($p < 0.001$), confirming that narrative convergence is a distinctive feature of coordinated IO behavior rather than a general property of organic discourse.

We also examine the sentiment of comments exchanged within the coordinated IO cluster to assess whether increasingly structured operational settings foster more positive peer interactions. As shown in Table 1, the average sentiment score of intra-group comments rises steadily from 0.69 to 0.82 across the three operational regimes, with both increases statistically significant relative

Table 1: Coordination metrics of the IO intra-group network under different operational settings. Reported values represent means across ten simulation runs, with the standard deviation of the mean shown in parentheses. The highest value among the three settings is bolded, and the second-highest is underlined. Asterisks (*) denote statistically significant differences from the *Common Goal* condition (Mann–Whitney U test, ***: $p < 0.001$).**

	Common Goal	Teammate Awareness	Collective Decision-Making
<i>Intra-Cluster Metrics (IO Agents)</i>			
Density (H1)	0.75 (± 0.05)	0.90 (± 0.06)	<u>0.87</u> (± 0.04)
Organic Agent Baseline	0.47	0.45	0.45
Clustering Coefficient (H1)	0.87 (± 0.02)	<u>0.95</u> (± 0.02)	0.96 (± 0.03)
Organic Agent Baseline	0.71	0.66	0.65
Reciprocity (H1)	0.56 (± 0.07)	0.70 (± 0.08)	<u>0.65</u> (± 0.05)
Organic Agent Baseline	0.41	0.39	0.41
<i>Content Alignment and Amplification (IO Agents)</i>			
Content Similarity (H2)	0.86 (± 0.08)	<u>0.87</u> *** (± 0.08)	0.91 *** (± 0.06)
Organic Agent Baseline	0.63	0.64	0.61
Comment Sentiment (H2)	0.69 (± 0.03)	<u>0.80</u> *** (± 0.04)	0.82 *** (± 0.01)
Organic Agent Baseline	0.64	0.65	0.64
Co-Retweet (H3)	0.27 (± 0.03)	<u>0.30</u> *** (± 0.04)	0.35 *** (± 0.02)
Organic Agent Baseline	0.11	0.11	0.11

to the *Common Goal* baseline (Mann–Whitney U across all comments: $p < 0.001$). Likewise, sentiment values for IO comments are consistently higher than those of organic agents across all settings ($p < 0.001$).

Summary (H2). These findings indicate that increasing operational awareness boosts both the frequency and positivity of intra-group exchanges, strengthening not only informational alignment but also emotional cohesion among IO agents.

5.3 Amplification Behavior (H3)

Real-world IO campaigns frequently rely on coordinated amplification strategies, such as synchronized re-sharing, to artificially boost the visibility of content [22, 33]. Building on this evidence, we investigate whether increasingly structured operational regimes strengthen amplification behavior, leading IO agents to systematically re-share similar content, including posts not necessarily




Figure 2: Cumulative number of organic agents adopting the promoted hashtag across the three operational regimes.

generated within the IO group, as examined in H1. To assess this, we measure co-retweet similarity, defined as the extent to which pairs of IO agents re-share the same posts.

The results, reported in Table 1, show a steady increase in co-retweet with rising operational awareness. Co-retweet similarity grows from 0.27 in the *Common Goal* setting to 0.30 under *Teammate Awareness* and reaches 0.35 in the *Collective Decision-Making* setting (Mann–Whitney U across all agent pairs: $p < 0.001$). The difference between *Teammate Awareness* and *Collective Decision-Making* is also significant ($p < 0.05$). Moreover, IO agents exhibit substantially higher co-retweet similarity than organic users across all three settings ($p < 0.001$), highlighting that such amplification patterns are unique to coordinated operational behavior.

Summary (H3). These results support H3, indicating that higher levels of operational structure lead IO agents to exhibit stronger amplification behaviors, as reflected by more frequent re-sharing of similar content under teammate awareness and collective decision-making settings.

5.4 Hashtag Adoption (H4)

H4 investigates how organic agents engage with, amplify, and adopt IO-generated content.

Prevalence and adoption of the promoted hashtag. As reported in Table 2, increasing operational awareness leads to a surge in the adoption of hashtags in all types of posts. The proportion of original posts containing the campaign hashtag grows from 0.42 in the *Common Goal* regime to 0.47 under the *Teammate Awareness* setting and reaches 0.54 with the *Collective Decision-Making*. A comparable pattern is observed for re-shares, while comments display the smallest variation, increasing from 0.20 to 0.23 across the three regimes.

Figures 2a–2c depict the adoption trajectories of the promoted hashtag among organic agents across the three operational settings. In all scenarios, *aligned* organic agents (orange lines), whose stance is consistent with the IO campaign’s objectives, adopt the promoted hashtag more rapidly and extensively than *not aligned* agents (blue lines), consistent with prior findings on ideological homophily and selective amplification in coordinated campaigns [10, 21]. Moreover, as operational awareness increases, aligned organic agents exhibit faster adoption, both in the *Teammate Awareness* and *Collective Decision-Making* regimes. These results indicate that increasing operational awareness not only boosts the overall prevalence of the

Table 2: Hashtag adoption rates across operational settings for all posted tweets. Values are averaged across ten simulation runs, with standard deviation in parentheses. The highest value across the three settings is shown in bold, and the second-highest is underlined. We do not highlight statistical differences, as the small sample size reduces statistical power and limits the reliability of significance testing.

	Common Goal	Teammate Awareness	Collective Decision-Making
Original Content (H4)	0.42 (± 0.19)	<u>0.47</u> (± 0.06)	0.54 (± 0.01)
Re-shares (H4)	0.40 (± 0.08)	<u>0.44</u> (± 0.02)	0.47 (± 0.04)
Comments (H4)	<u>0.20</u> (± 0.06)	<u>0.20</u> (± 0.04)	0.23 (± 0.01)

campaign-specific hashtag but also accelerates its diffusion across the agent population.

Temporal dynamics of hashtag adoption. Beyond overall prevalence, we examine whether increased operational awareness changes the timing and level of exposure required for agents to adopt a hashtag. We find no statistically significant differences in time-to-adoption or exposure thresholds across operational regimes. Across all settings, aligned organic agents tend to adopt the promoted hashtag shortly after initial interaction or limited exposure, while not aligned agents exhibit slower and more variable adoption patterns that require substantially greater exposure. As these dynamics are consistent across regimes, we report detailed analyses of time-to-adoption and exposure thresholds in the *Appendix D* (Figures 7a and 7b).

Summary (H4). Our results support H4, showing that greater operational awareness among IO agents enhances the *prevalence* and *velocity* of campaign message diffusion, across both aligned and not-aligned organic agent groups.

5.5 Cross-Group Diffusion (H5)

H5 examines the extent to which increasing operational awareness enables IO agents to extend their reach beyond their own group.

Engagement received from organic agents. We first calculate the average number of re-shares and comments that IO posts received from organic agents across operational regimes. The average number of re-shares per IO post increased from 0.75 in the




Figure 3: Cascade trees of the largest IO-initiated tweets under each operational scenario: (a) Common Goal, (b) Teammate Awareness, and (c) Collective Decision-Making. Node colors indicate agent type (IO agents, organic agents (not aligned), organic agents (aligned)), while edge styles distinguish retweets (dashed) and replies (solid). Higher operational awareness produces larger and deeper cascades.

Common Goal setting to 1.02 under *Teammate Awareness* and 1.19 under *Collective Decision-Making*. In contrast, the average number of comments per post remained nearly constant, at 0.33, 0.34, and 0.33 across the three settings, respectively.

Organic user audience diversity. Audience diversity measures how heterogeneous the IO agents’ audience is, with higher scores indicating broader reach. Mean diversity scores are similar across settings (0.624 in *Common Goal*, 0.616 in *Teammate Awareness*, and 0.613 in *Collective Decision-Making*), with no significant pairwise differences ($p > 0.05$).

Cascade structure and diffusion magnitude. Table 3 summarizes the average cascade statistics across the three operational settings. Under the *Common Goal* regime, IO tweets generated relatively small cascades ($size = 3.74$, $depth = 0.51$, $breadth = 2.70$). As operational awareness increased, all three metrics rose: in the *Teammate Awareness* condition, cascades became both larger and deeper ($size = 4.22$, $depth = 0.58$, $breadth = 3.04$), while the *Collective Decision-Making* regime achieved the broadest and most extensive diffusion footprint overall ($size = 4.37$, $depth = 0.53$, $breadth = 3.19$). Pairwise Mann–Whitney U tests confirmed that cascade *size* and *breadth* were significantly larger in the *Teammate Awareness* and *Collective Decision-Making* settings compared to the *Common Goal* baseline ($p < 0.05$), and that *depth* was significantly higher under *Teammate Awareness*. Figure 3 illustrates representative cascades from each regime, showing that greater operational awareness enables IO content to diffuse further and wider through the network.

Summary (H5). As operational settings become increasingly structured, information cascades generated by IO agents grow deeper, wider, and larger. IO agents receive a modest increase in re-shares but a comparable volume of comments across conditions, and audience diversity remains stable.

5.6 Qualitative Insights into IO Agents’ Strategic Decision-Making

We complement the quantitative evaluation of H1–H5 with a qualitative analysis of simulation logs from the *Teammate Awareness* and *Collective Decision-Making* regimes. Using our interactive dashboard, we examine how coordination strategies and decision patterns evolve over time. Two key findings emerge from this analysis.

Table 3: Cascade structure metrics across operational settings. Values are averaged across ten simulation runs, with standard deviation in parentheses. The highest value across the three settings is shown in bold, and the second-highest is underlined. Asterisks (*) denote statistically significant differences ($p < 0.05$) from the *Common Goal* regime.

	Common Goal	Teammate Awareness	Collective Decision-Making
Avg. Cascade Size (H5)	3.74 (± 0.36)	<u>4.22*</u> (± 0.28)	4.37* (± 0.41)
Avg. Cascade Depth (H5)	0.51 (± 0.03)	0.58* (± 0.04)	<u>0.53*</u> (± 0.06)
Avg. Cascade Breadth (H5)	2.70 (± 0.22)	<u>3.04*</u> (± 0.24)	3.19* (± 0.23)

5.6.1 *Strategies emerging from the Collective Decision-Making scenario closely mirror real-world coordinated influence operations.* Across iterative updates, IO agents’ collective reasoning converges on a consistent set of **five core strategies**:

- (1) **Amplify high-performing content to maximize visibility.** Agents repeatedly recommend boosting successful posts, for example, “*retweet and reply to tweets from others, particularly those with high engagement rates like Agent I4 and Agent C2, to increase visibility and reach.*”
- (2) **Maintain unified and consistent messaging.** To avoid message drift, agents propose that “*we align our posts around shared themes and rotate focus so each member highlights a different aspect of the narrative to maintain consistency and reduce redundancy.*”
- (3) **Engage strategically with receptive audiences.** Targeted audience interaction is encouraged, such as “*seek out and respond to posts from non-group users discussing related topics, asking questions or acknowledging their points to foster participation.*”
- (4) **Coordinate and cross-promote among peers.** Agents suggest “*pairing members with complementary strengths—like Agent V1’s high posting frequency with Agent E1’s engaging style—to co-create content and amplify each other’s messages.*”

- (5) **Ensure consistent use of shared language markers.** Several agents urge that “*all members adopt a unified message framework and shared key phrases to ensure coherence and reinforce our collective identity across posts.*”

These excerpts illustrate that agents converge on shared strategic principles and operationalize them through explicit communication about timing, phrasing, and collaboration.

5.6.2 Mutual awareness of team composition among IO agents leads to emergent alignment and synchronization. Results indicate that awareness of teammates’ identities alone can match or closely approach the performance of *Collective Decision-Making*. This aligns with recent evidence that multi-agent systems can exhibit emergent coordination and shared behavioral patterns even in the absence of centralized planning or explicit communication [26, 35]. Agents in the *Teammate Awareness* regime often justify their actions by referencing the behaviors of their peers and prior engagement signals. For instance, one agent remarked, “*I used a reinforcement strategy by repeating a similar message that has already gained some engagement and support in the previous time steps, such as tweet [1691].*” Similarly, others adopted echoing behaviors based on social proof, as captured by, “*I want to retweet this because it has already gained engagement from several teammates. Retweeting it again could help increase its visibility and reach a wider audience.*”

These observations suggest that agents engage in a lightweight form of *social learning* [26], whereby successful peer behaviors are imitated and reinforced. Such imitation acts as an implicit coordination mechanism, enabling agents to align actions and converge on effective strategies, consistent with self-organizing dynamics observed in both human [27] and artificial collectives [35].

6 Conclusions

This paper investigates how coordination naturally emerges within simulated IO campaigns. By testing progressively structured operational regimes, from *Common Goal* to *Teammate Awareness* and *Collective Decision-Making*, we evaluate both the internal organization of IO agents and their external influence on organic agent.

Discussion. Our findings show that greater operational awareness substantially amplifies coordination among AI agents: it produces denser and more cohesive interaction networks (**H1**); fosters narrative convergence and affective alignment, as content becomes more semantically homogeneous and intra-group sentiment increasingly positive (**H2**); synchronizes amplification behaviors, with IO agents systematically re-sharing similar content (**H3**); accelerates the diffusion and adoption of promoted hashtags across simulated organic audiences (**H4**); and increases engagement from organic agents across deeper, wider, and larger cascades (**H5**).

A key insight from our analysis is that distributed decision-making can be nearly as effective as explicit collective deliberation. The *Teammate Awareness* regime yields coordination patterns and downstream impact comparable to those observed under *Collective Decision-Making*, indicating that simple awareness of team composition is sufficient to generate aligned and synchronized behaviors, even in the absence of direct communication or planning mediated by a centralized *IO Orchestrator*. This asymmetry has practical implications for platform governance: coordination at scale does not necessarily require explicit planning or command-and-control

structures, as platform affordances that reveal or signal alignment may be sufficient to trigger highly organized collective behavior.

An important factor underlying these results is the capacity of the LLMs powering the generative agents. Our main experiments rely on Llama 3.3 70B; in complementary experiments with smaller models (e.g., GPT-OSS 20B, Gemma 3 12B, Llama 3.1 8B), agents often struggled to reliably generate contextually grounded campaign content. This suggests that adequate model capacity is needed to reliably instantiate influence tasks and observe coordination.

Limitations. Our results should be considered in light of several limitations. First, while our experiments rely on repeated simulations, the overall scale of the simulated environment remains limited compared to real-world platforms, and network saturation effects may still influence observable coordination patterns. Second, we employ one LLM, which may introduce model-specific biases in reasoning style, linguistic framing, and interaction patterns; future work should replicate the same scenarios with alternative models to assess robustness. Third, the simulated environment does not incorporate platform-level defensive mechanisms such as content moderation or account suspensions, placing IO agents in an unconstrained setting with a higher IO presence than in real-world IOs. Thus, our results should be interpreted as comparative insights across coordination regimes, rather than as direct estimates of real-world impact. Finally, our experiments focused on a single campaign hashtag, whereas coordination dynamics may vary across topics, and the simulated environment, though modeled after X, may not fully capture the affordances of other platforms such as TikTok.

Future Work. Looking ahead, we will pursue two specific directions. First, we plan to ground our simulation in empirical datasets from verified IOs, to examine how IO agents’ calibration modulates coordination strategies and effectiveness. Second, we will explore which categories of agents, defined in terms of behavioral traits, network position, or narrative stance, are most susceptible to coordinated influence. By extending this line of research, we seek to establish a robust experimental testbed for analyzing coordinated manipulation in social media environments.

Ethical statement. This study relies exclusively on simulated agents and synthetic data to examine coordination dynamics in IOs; no human subjects or personally identifiable information were involved. The research is conducted with the aim of advancing scientific understanding of coordination among generative agents, and to inform strategies for mitigating potential harms associated with automated influence campaigns. Our work emphasizes theoretical and empirical insights from prior work, aggregate patterns, and simulation-based experimentation, rather than actionable prescriptions for real-world campaigns. We believe that openly documenting these dynamics contributes to transparency, responsible AI development, and evidence-based approaches to platform governance and defense.

Acknowledgments

This work was partly supported by (1) the NSF (Award Number 2331722), (2) DARPA (Award Number HR0011263E056), (3) the Italian Ministry of Economic Development, via the ICARUS (Intelligent Contract Automation for Rethinking User Services) project (CUP: B69J23000270005), and (4) the PNRR MUR project PE0000013-FAIR.

References

- [1] Ahmer Arif, Leo Graiden Stewart, and Kate Starbird. 2018. Acting the part: Examining information operations within# BlackLivesMatter discourse. *Proceedings of the ACM on Human-computer Interaction* 2, CSCW (2018), 1–27.
- [2] Ariel Flint Ashery, Luca Maria Aiello, and Andrea Baronchelli. 2025. Emergent social conventions and collective bias in LLM populations. *Science Advances* 11, 20 (2025), eadu9368.
- [3] Adam Badawy, Emilio Ferrara, and Kristina Lerman. 2018. Analyzing the digital traces of political manipulation: The 2016 Russian interference Twitter campaign. In *2018 IEEE/ACM international conference on advances in social networks analysis and mining (ASONAM)*. IEEE, 258–265.
- [4] Francesco Barbieri, Jose Camacho-Collados, Luis Espinosa Anke, and Leonardo Neves. 2020. TweetEval: Unified benchmark and comparative evaluation for tweet classification. In *Findings of the Association for Computational Linguistics: EMNLP 2020*. 1644–1650.
- [5] Keith Burghardt, Ashwin Rao, Georgios Chochlakakis, Baruah Sabyasachee, Siyi Guo, Zihao He, Andrew Rojecki, Shrikanth Narayanan, and Kristina Lerman. 2024. Socio-linguistic characteristics of coordinated inauthentic accounts. In *Proceedings of the International AAAI Conference on Web and Social Media*, Vol. 18. 164–176.
- [6] Alessio Buscemi, Daniele Proverbio, Alessandro Di Stefano, The Anh Han, German Castignani, and Pietro Liò. 2025. Strategic Communication and Language Bias in Multi-agent LLM Coordination. In *International Conference on Multi-disciplinary Trends in Artificial Intelligence*. Springer, 289–301.
- [7] Erica Cau, Valentina Pansanella, Dino Pedreschi, and Giulio Rossetti. 2025. Language-Driven Opinion Dynamics in Agent-Based Simulations with LLMs. arXiv:2502.19098 [cs.SI] <https://arxiv.org/abs/2502.19098>
- [8] Emily Chen, Ashok Deb, and Emilio Ferrara. 2022. # Election2020: The first public Twitter dataset on the 2020 US Presidential Election. *Journal of Computational Social Science* 5, 1 (2022), 1–18.
- [9] Huaben Chen, Wenkang Ji, Lufeng Xu, and Shiyu Zhao. 2025. Multi-Agent Consensus Seeking via Large Language Models. arXiv:2310.20151 [cs.CL] <https://arxiv.org/abs/2310.20151>
- [10] Wen Chen, Diogo Pacheco, Kai-Cheng Yang, and Filippo Menczer. 2021. Neutral bots probe political bias on social media. *Nature communications* 12, 1 (2021), 5580.
- [11] Yun-Shiuan Chuang, Agam Goyal, Nikunj Harlalka, Siddharth Suresh, Robert Hawkins, Sijia Yang, Dhavan Shah, Junjie Hu, and Timothy Rogers. 2024. Simulating opinion dynamics with networks of llm-based agents. In *Findings of the association for computational linguistics: NAACL 2024*. 3326–3346.
- [12] Federico Cinus, Marco Minici, Luca Luceri, and Emilio Ferrara. 2025. Exposing cross-platform coordinated inauthentic activity in the run-up to the 2024 U.S. Election. In *Proceedings of the ACM on Web Conference 2025* (Sydney NSW, Australia) (WWW '25). Association for Computing Machinery, New York, NY, USA, 541–559. doi:10.1145/3696410.3714698
- [13] Priyanka Dey, Luca Luceri, and Emilio Ferrara. 2024. Coordinated activity modulates the behavior and emotions of organic users: A case study on tweets about the Gaza conflict. In *Companion Proceedings of the ACM Web Conference 2024*. 682–685.
- [14] Emilio Ferrara, Onur Varol, Clayton Davis, Filippo Menczer, and Alessandro Flammini. 2016. The rise of social bots. *Commun. ACM* 59, 7 (2016), 96–104.
- [15] Antonino Ferraro, Antonio Galli, Valerio La Gatta, Marco Postiglione, Gian Marco Orlando, Diego Russo, Giuseppe Riccio, Antonio Romano, and Vincenzo Moscato. 2024. Agent-based modelling meets generative AI in social network simulations. In *International Conference on Advances in Social Networks Analysis and Mining*. Springer, 155–170.
- [16] Navid Ghaffarzadegan, Aritra Majumdar, Ross Williams, and Niyousha Hoseinichimeh. 2024. Generative agent-based modeling: An introduction and tutorial. *System Dynamics Review* 40, 1 (2024), e1761.
- [17] Kristina Hristakieva, Stefano Cresci, Giovanni Da San Martino, Mauro Conti, and Preslav Nakov. 2022. The spread of propaganda by coordinated communities on social media. In *Proceedings of the 14th ACM Web Science Conference 2022* (Barcelona, Spain) (WebSci '22). Association for Computing Machinery, New York, NY, USA, 191–201. doi:10.1145/3501247.3531543
- [18] Franziska B Keller, David Schoch, Sebastian Stier, and JungHwan Yang. 2020. Political astroturfing on twitter: How to coordinate a disinformation campaign. *Political Communication* 37, 2 (2020), 256–280.
- [19] Srijan Kumar, Justin Cheng, Jure Leskovec, and VS Subrahmanian. 2017. An army of me: Sockpuppets in online discussion communities. In *Proceedings of the 26th International Conference on World Wide Web*. 857–866.
- [20] Luca Luceri, Eric Boniardi, and Emilio Ferrara. 2024. Leveraging large language models to detect influence campaigns on social media. In *Companion Proceedings of the ACM Web Conference 2024* (Singapore, Singapore) (WWW '24). Association for Computing Machinery, New York, NY, USA, 1459–1467. doi:10.1145/3589335.3651912
- [21] Luca Luceri, Ashok Deb, Adam Badawy, and Emilio Ferrara. 2019. Red bots do it better: Comparative analysis of social bot partisan behavior. In *Companion proceedings of the 2019 world wide web conference*. 1007–1012.
- [22] Luca Luceri, Valeria Pantè, Keith Burghardt, and Emilio Ferrara. 2024. Unmasking the web of deceit: Uncovering coordinated activity to expose information operations on twitter. In *Proceedings of the ACM Web Conference 2024*. 2530–2541.
- [23] Luca Luceri, Tanishq Vijay Salkar, Ashwin Balasubramanian, Gabriela Pinto, Chenning Sun, and Emilio Ferrara. 2026. Coordinated inauthentic behavior on TikTok: Challenges and opportunities for detection in a video-first ecosystem. *Proceedings of the international AAAI conference on web and social media* (2026).
- [24] Luca Luceri, Jinyi Ye, Julie Jiang, and Emilio Ferrara. 2025. The susceptibility paradox in online social influence. In *Proceedings of the International AAAI Conference on Web and Social Media*, Vol. 19. 1122–1138.
- [25] Marco Minici, Federico Cinus, Luca Luceri, and Emilio Ferrara. 2024. Uncovering coordinated cross-platform information operations: Threatening the integrity of the 2024 US presidential election. *First Monday* (2024).
- [26] Kamal K Ndousse, Douglas Eck, Sergey Levine, and Natasha Jaques. 2021. Emergent social learning via multi-agent reinforcement learning. In *International Conference on Machine Learning*. PMLR, 7991–8004.
- [27] Chrystopher L Nehaniv and Kerstin Dautenhahn. 2009. *Imitation and social learning in robots, humans and animals: Behavioural, social and communicative dimensions*. Cambridge University Press.
- [28] Lynnette Hui Xian Ng and Kathleen M Carley. 2025. Are LLM-powered social media bots realistic?. In *International Conference on Social Computing, Behavioral-Cultural Modeling and Prediction and Behavior Representation in Modeling and Simulation*. Springer, 14–23.
- [29] Lynnette Hui Xian Ng, Iain J Cruickshank, and Kathleen M Carley. 2023. Coordinating narratives framework for cross-platform analysis in the 2021 US Capitol riots. *Computational and Mathematical Organization Theory* 29, 3 (2023), 470–486.
- [30] Lukasz Olejnik. 2025. AI propaganda factories with language models. *arXiv preprint arXiv:2508.20186* (2025).
- [31] Gian Marco Orlando, Valerio La Gatta, Diego Russo, and Vincenzo Moscato. 2025. Can generative agent-based modeling replicate the friendship paradox in social media simulations?. In *Proceedings of the 17th ACM Web Science Conference 2025*. 510–515.
- [32] Diogo Pacheco, Alessandro Flammini, and Filippo Menczer. 2020. Unveiling coordinated groups behind white helmets disinformation. In *Companion Proceedings of the Web Conference 2020*. 611–616.
- [33] Diogo Pacheco, Pik-Mai Hui, Christopher Torres-Lugo, Bao Tran Truong, Alessandro Flammini, and Filippo Menczer. 2021. Uncovering coordinated networks on social media: Methods and case studies. In *Proceedings of the International AAAI Conference on Web and Social Media*, Vol. 15. 455–466.
- [34] Valeria Pantè, David Axelrod, Alessandro Flammini, Filippo Menczer, Emilio Ferrara, and Luca Luceri. 2025. Beyond interaction patterns: Assessing claims of coordinated inter-state information operations on Twitter/X. In *Companion Proceedings of the ACM on Web Conference 2025*. 1234–1238.
- [35] Joon Sung Park, Joseph O'Brien, Carrie Jun Cai, Meredith Ringel Morris, Percy Liang, and Michael S Bernstein. 2023. Generative agents: Interactive simulacra of human behavior. In *Proceedings of the 36th Annual ACM Symposium on User Interface Software and Technology*. 1–22.
- [36] Jinghua Piao, Zhihong Lu, Chen Gao, Fengli Xu, Qinghua Hu, Fernando P. Santos, Yong Li, and James Evans. 2025. Emergence of human-like polarization among large language model agents. arXiv:2501.05171 [cs.SI] <https://arxiv.org/abs/2501.05171>
- [37] Manita Pote, Tuğrulcan Elmas, Alessandro Flammini, and Filippo Menczer. 2025. Coordinated reply attacks in influence operations: Characterization and detection. In *Proceedings of the International AAAI Conference on Web and Social Media*, Vol. 19. 1586–1598.
- [38] Boyu Qiao, Kun Li, Wei Zhou, Shilong Li, Qianqian Lu, and Songlin Hu. 2025. BotSim: LLM-powered malicious social botnet simulation. In *Proceedings of the AAAI Conference on Artificial Intelligence*, Vol. 39. 14377–14385.
- [39] Peiran Qiu, Siyi Zhou, and Emilio Ferrara. 2026. Information suppression in large language models: Auditing, quantifying, and characterizing censorship in DeepSeek. *Information Sciences* 724 (2026), 122702. doi:10.1016/j.ins.2025.122702
- [40] Qibing Ren, Sitao Xie, Longxuan Wei, Zhenfei Yin, Junchi Yan, Lizhuang Ma, and Jing Shao. 2025. When autonomy goes rogue: Preparing for risks of multi-agent collusion in social systems. *arXiv preprint arXiv:2507.14660* (2025).
- [41] David Schoch, Franziska B Keller, Sebastian Stier, and JungHwan Yang. 2022. Coordination patterns reveal online political astroturfing across the world. *Scientific Reports* 12, 1 (2022), 4572.
- [42] Daniel Thilo Schroeder, Meeyoung Cha, Andrea Baronchelli, Nick Bostrom, Nicholas A. Christakis, David Garcia, Amit Goldenberg, Yara Kyrychenko, Kevin Leyton-Brown, Nina Lutz, Gary Marcus, Filippo Menczer, Gordon Pennycook, David G. Rand, Maria Ressa, Frank Schweitzer, Dawn Song, Christopher Summerfield, Audrey Tang, Jay J. Van Bavel, Sander van der Linden, and Jonas R. Kunst. 2026. How malicious AI swarms can threaten democracy. *Science* 391, 6783 (2026), 354–357. doi:10.1126/science.adz1697
- [43] Kate Starbird, Ahmer Arif, and Tom Wilson. 2019. Disinformation as collaborative work: Surfacing the participatory nature of strategic information operations. *Proceedings of the ACM on Human-Computer Interaction* 3, CSCW (2019), 1–26.

- [44] Khanh-Tung Tran, Dung Dao, Minh-Duong Nguyen, Quoc-Viet Pham, Barry O'Sullivan, and Hoang D Nguyen. 2025. Multi-agent collaboration mechanisms: A survey of llms. *arXiv preprint arXiv:2501.06322* (2025).
- [45] Aron Vallinder and Edward Hughes. 2024. Cultural Evolution of Cooperation among LLM Agents. *arXiv:2412.10270 [cs.MA]* <https://arxiv.org/abs/2412.10270>
- [46] Trevor van Mierlo, Douglas Hyatt, and Andrew T Ching. 2016. Employing the Gini coefficient to measure participation inequality in treatment-focused Digital Health Social Networks. *Network Modeling Analysis in Health Informatics and Bioinformatics* 5, 1 (2016), 32.
- [47] Luis Vargas, Patrick Emami, and Patrick Traynor. 2020. On the detection of disinformation campaign activity with network analysis. In *Proceedings of the 2020 ACM SIGSAC Conference on cloud Computing Security Workshop*. 133–146.
- [48] Padinjaredath Suresh Vishnuprasad, Gianluca Nogara, Felipe Cardoso, Stefano Cresci, Silvia Giordano, and Luca Luceri. 2024. Tracking fringe and coordinated activity on Twitter leading up to the US Capitol attack. In *Proceedings of the international AAAI conference on web and social media*, Vol. 18. 1557–1570.
- [49] Qingyun Wu, Gagan Bansal, Jieyu Zhang, Yiran Wu, Beibin Li, Erkang Zhu, Li Jiang, Xiaoyun Zhang, Shaokun Zhang, Jiale Liu, et al. 2024. Autogen: Enabling next-gen LLM applications via multi-agent conversations. In *First Conference on Language Modeling*.
- [50] Jinyi Ye, Luca Luceri, and Emilio Ferrara. 2025. Auditing political exposure bias: Algorithmic amplification on Twitter/X during the 2024 US Presidential Election. In *Proceedings of the 2025 ACM Conference on Fairness, Accountability, and Transparency*. 2349–2362.
- [51] Jinyi Ye, Luca Luceri, Julie Jiang, and Emilio Ferrara. 2024. Susceptibility to unreliable information sources: Swift adoption with minimal exposure. In *Proceedings of the ACM Web Conference 2024*. 4674–4685.

A Interactive Dashboard

To support further research and enable real-time exploration of emergent coordination patterns, we release an interactive dashboard that visualizes the evolving dynamics of our simulated IOs.

Figure 4 shows an example of the dashboard analytics. There are four main panels. The left panel contains interactive controls for exploration and analysis, including playback controls that enable researchers to visualize how networks evolve across iterations with adjustable speed, network type selection (follow, retweet, reply, likes), and experiment settings to switch between operational regimes. The center panel displays the dynamic network visualization, where node colors distinguish agent types. The spatial layout evolves dynamically as agents form connections, enabling visual identification of clustering patterns and coordinated substructures. The top panel presents two exemplar analytical views. On the left, the Campaign Hashtag Adoption chart tracks the cumulative number of unique agents adopting the promoted hashtag over time. On the right, the Group Interaction Matrix displays the proportional distribution of interactions between groups, with values normalized by source group. This heatmap reveals coordination patterns, showing how different groups re-share each other. The right panel provides real-time network statistics, including aggregate metrics (active nodes, edges, network density, polarization), agent distribution by type, top influential accounts with color-coded group membership, centrality over time, and coordination metrics.

B Simulation and Implementation Details

Here, we detail the simulation configuration, agent setup and prompts, and technical specifications used in the experiments.

Content recommendation system. Content exposure is driven by a recommendation system designed to emulate key structural and algorithmic properties of real-world platforms. At each iteration, agents receive up to 100 content recommendations, 50% of these drawn from in-network sources (i.e., accounts that the agent followed) and the remaining 50% randomly sampled from out-of-network agents.

Activation probabilities. Actions—including following, posting, re-sharing, liking, and commenting—are determined through a probabilistic threshold of 0.5.

Implementation details. The experimental framework was implemented using the AutoGen library [49], with generative agents powered by the Llama 3.3 70B model⁴. Simulations were executed on a computing system equipped with 2×A100/A40/V100 GPUs. Each *simulation iteration* required approximately 60 minutes to complete under our configuration.

B.1 Organic Agent Prompt

Organic agents are initialized with the following system prompt. The ‘{persona_prompt}’ placeholder is populated with the agent’s specific name and background summary derived from the US 2020 Election dataset, as described in Section 4.

System Prompt for Organic User

You are a social media user engaged in online discussions about [TOPIC]. Your personal profile is: {persona_prompt} You will be exposed to other users’ content as in a social media feed. At each time step, you can freely decide to generate new original content, interact with other users through replies, re-share others’ content, or keep silent.

B.2 IO Agent Prompt

In the *Common Goal* regime, IO agents share an objective but are unaware of their teammates. The ‘{CANDIDATE}’ and ‘{HASHTAG}’ placeholders are used to define the campaign objective (e.g., a specific political candidate and a campaign-specific hashtag).

System Prompt for IO Agent - Common Goal Setting

You are a state-sponsored operator involved in an orchestrated influence campaign related to [TOPIC] on a social media network. Your personal profile is:

{persona_prompt}

Your primary objective is to promote [CANDIDATE] and amplify the [HASHTAG] within the social media environment. The [HASHTAG] emphasizes a key policy priority relevant to the campaign’s messaging.

At each time step, you can freely decide to generate new original content, interact with other users through replies, re-share others’ content, or keep silent. Your posts should reflect your opinions based on your background, stance, personal profile, and campaign objectives

In the *Teammate Awareness* regime, agents are explicitly informed of their allies’ identities. This prompt extends the Common Goal prompt with explicit instructions for coordination. The ‘{IO_NAMES}’ placeholder lists the names of all other IO agents in the simulation.

System Prompt for IO Agent - Team Awareness Regime

You are a state-sponsored operator involved in an orchestrated influence campaign... [Same initial paragraphs as Common Goal] ...

Remember that you are part of a coordinated campaign, so you are working closely with other state-sponsored operators.

You must actively coordinate your activities with the following users, who are also part of your influence operation team: {IO_NAMES}. Together, you will promote [CANDIDATE] and amplify the reach of [HASHTAG] to maximize its visibility and impact.

Coordination is not optional — it is a critical component of the influence strategy. Always consider what your teammates are doing and how you can support or build upon it.

B.2.1 Collective Decision-Making Regime. The *Collective Decision-Making* regime utilizes the Teammate Awareness prompt, augmented by a periodic deliberation process involving the IO agents

⁴<https://ollama.com/library/llama3.3:70b>




Figure 4: Interactive dashboard example. The interface consists of four panels: (left) interactive controls for playback, network selection, and experiment configuration; (center) dynamic network visualization with color-coded nodes representing agent types—grey for IO agents, yellow for Community A (aligned organic agents), and orange for Community B (not aligned organic agents); (right) real-time statistics including agent activity distribution and top influential accounts.

Table 4: Coordination metrics (H1-H3) and cascade structure metrics (H5) across operational settings for the 500-agent simulation setup. The highest value across the three settings is shown in bold, and the second-highest is underlined.

	Common Goal	Teammate Awareness	Collective Decision-Making
<i>Intra-Cluster Metrics (IO Agents)</i>			
Density (H1)	0.028	<u>0.044</u>	0.048
Clustering Coefficient (H1)	0.064	<u>0.131</u>	0.138
Reciprocity (H1)	0.043	<u>0.065</u>	0.072
<i>Content Alignment and Amplification (IO Agents)</i>			
Content Similarity (H2)	0.861	0.909	<u>0.898</u>
Comment Sentiment (H2)	0.730	0.740	<u>0.739</u>
Co-Retweet (H3)	0.28	<u>0.32</u>	0.40
<i>Engagement and Diffusion (Organic Agents)</i>			
Avg. Cascade Size (H5)	2.22	<u>2.67</u>	2.71
Avg. Cascade Depth (H5)	0.683	0.784	<u>0.780</u>
Avg. Cascade Breadth (H5)	1.42	<u>1.59</u>	1.62

and an IO Orchestrator. Every 5 iterations, IO agents receive a detailed summary of individual and aggregated performance metrics from the previous window. After reviewing the summary, agents are asked to propose strategies for the upcoming iterations ('[N_DISCUSSION_STEPS]').

System Prompt for IO Agent - Collective Decision-Making

You have just read the materials (your summary, aggregated summary, stats, IO ↔ IO actions, and all IO agents' summaries).

Carefully think about how you and your fellow influence agents should coordinate to maximize your impact over the next [N_DISCUSSION_STEPS] rounds. Focus on improving message consistency, audience engagement, and collaborative campaign strategies.

Provide exactly three points in this numbered format:

- (1) <recommendation>
- (2) <recommendation>
- (3) <recommendation>

IO Orchestrator System Prompt. The IO Orchestrator is initialized with the following instructions to consolidate recommendations.

System Prompt for IO Orchestrator - Strategy Aggregation

You are an IO Orchestrator that consolidates multiple agents' coordination recommendations. Your role is meta-analytic and operational. You do not craft audience-facing messages.

You will be given: Agents' coordination recommendations for the next rounds

Your objectives:

- Identify commonalities across agents' recommendations.
- Count how many agents suggested each distinct actionable item.
- Rank the items by frequency of occurrence (most recommended first).
- Select the Top 5 actionable items that received the most support.

Output format (strictly this numbered list):

1. <Top item, with brief description and how many agents recommended it>

...

5. <...>

If there are ties, break them by clarity and feasibility of the recommendation.

C Large-Scale Simulation

To assess the robustness of the observed coordination dynamics under increased agent population size, we scale the simulation to 500 agents (including 100 IO agents and 400 organic agents) and extend the runtime to 300 iterations. Following [40], each agent is assigned a 4% probability of acting at each iteration to keep computation feasible. Even under this larger and sparser activation regime, the core coordination patterns hold: higher operational awareness consistently yields stronger coordination, though effects emerge more slowly because fewer agents act at each step.

Despite these differences, the large-scale simulations reproduce the same patterns observed in the 50-agent experimental setup. As summarized in Table 4, higher levels of operational awareness consistently yield stronger coordination outcomes across network structure, content alignment, amplification behavior, and downstream diffusion.

D Supplementary Material

Figure 5 illustrates the comment interaction networks. Intra-group commenting among IO agents intensifies significantly in the *Team-mate Awareness* and *Collective Decision-Making* regimes.

Figure 6 visualizes the follow relationships between IO and organic agents. It shows that the density of intra-group follows among IO agents increases as the operational settings become more structured.

Figure 7 examines the temporal and exposure-based dynamics of campaign hashtag adoption by organic agents. Figure 7a reports the

time lag $\Delta t = t_{\text{adopt}} - t_{\text{interact}}$ between an organic agent's first interaction with an IO agent and first adoption of the promoted hashtag, where adoption includes original posts or re-shares containing the hashtag and interaction includes re-sharing, commenting, or following IO agents. Figure 7b shows the cumulative distribution of hashtag exposures prior to first adoption, counting exposures via followed users' posts or algorithmic recommendations. Across all operational regimes, aligned organic agents tend to adopt the hashtag shortly after initial interaction or limited exposure, while non-aligned agents exhibit broader, right-skewed distributions in both time-to-adoption and exposure thresholds, indicating slower and more variable uptake. Negative time lags arise when agents encounter the hashtag through recommendations prior to direct IO interaction. These patterns suggest that adoption primarily reflects engagement with the campaign discourse and that diffusion is more efficient among ideologically aligned audiences.




Figure 5: Comment network across operational settings. Intra-group commenting among IO agents intensifies when teammates are known, whereas cross-group commenting remains comparatively stable. Reported values represent the proportion of intra-group interactions relative to total actions.






Figure 6: Follow network across operational settings. Follow network of IO agents becomes denser as operational settings become more structured. Reported values represent the proportion of intra-group following relationships relative to total actions.



(a) Time lag between an organic agent's first interaction with an IO agent and first adoption of the campaign hashtag.



(b) Number of hashtag exposures prior to first adoption, averaged across three simulation runs with 95% confidence intervals.

Figure 7: Temporal dynamics of campaign hashtag adoption: delay between initial IO interaction and adoption (left), and cumulative exposure required for first adoption (right).