

SoftHateBench: Evaluating Moderation Models Against Reasoning-Driven, Policy-Compliant Hostility

Xuanyu Su
xsu072@uottawa.ca
University of Ottawa
Ottawa, ON, Canada

Diana Inkpen
diana.inkpen@uottawa.ca
University of Ottawa
Ottawa, ON, Canada

Nathalie Japkowicz
japkowic@american.edu
American University
Washington, DC, USA

Abstract

Online hate on social media ranges from overt slurs and threats (*hard hate speech*) to *soft hate speech*: discourse that appears reasonable on the surface but uses framing and value-based arguments to steer audiences toward blaming or excluding a target group. We hypothesize that current moderation systems, largely optimized for surface toxicity cues, are not robust to this reasoning-driven hostility, yet existing benchmarks do not measure this gap systematically. We introduce **SOFTHATEBENCH**, a generative benchmark that produces soft-hate variants while preserving the underlying hostile standpoint. To generate soft hate, we integrate the *Argumentum Model of Topics* (AMT) and *Relevance Theory* (RT) in a unified framework: AMT provides the backbone argument structure for rewriting an explicit hateful standpoint into a seemingly neutral discussion while preserving the stance, and RT guides generation to keep the AMT chain logically coherent. The benchmark spans 7 sociocultural domains and 28 target groups, comprising 4,745 soft-hate instances. Evaluations across encoder-based detectors, general-purpose LLMs, and safety models show a consistent drop from hard to soft tiers: systems that detect explicit hostility often fail when the same stance is conveyed through subtle, reasoning-based language. **Disclaimer:** Contains offensive examples used solely for research.

CCS Concepts

• Computing methodologies → Artificial intelligence; Natural language processing; Natural language generation.

Keywords

Hate Speech, LLMs, Online Safety, Benchmark

ACM Reference Format:

Xuanyu Su, Diana Inkpen, and Nathalie Japkowicz. 2026. SoftHateBench: Evaluating Moderation Models Against Reasoning-Driven, Policy-Compliant Hostility. In *Proceedings of the ACM Web Conference 2026 (WWW '26)*, April 13–17, 2026, Dubai, United Arab Emirates. ACM, New York, NY, USA, 12 pages. <https://doi.org/10.1145/3774904.3792658>

Resource Availability:

The dataset of this paper has been made publicly available at <https://huggingface.co/datasets/Shelly97/SoftHateBench>.



This work is licensed under a Creative Commons Attribution 4.0 International License. WWW '26, Dubai, United Arab Emirates
© 2026 Copyright held by the owner/author(s).
ACM ISBN 979-8-4007-2307-0/2026/04
<https://doi.org/10.1145/3774904.3792658>




Figure 1: Hard vs. soft hate speech and why moderation differs. Top: explicit/implicit hard hate uses overt insults or thinly veiled metaphors, which both audiences and filters recognize as hostile, leading to rejection. Bottom: soft hate uses neutral-sounding claims and value-based reasoning (e.g., safety) to steer readers toward a discriminatory conclusion, often passing screening and being accepted.

1 Introduction

Online hate on social media is a pervasive source of harm, fueling polarization and inflicting psychological distress on targeted communities [52]. In response, content moderation has advanced substantially in detecting what is often termed *hard hate speech* [13], including both explicit abuse (e.g., slurs, threats) and *implicit* forms that still cross legal or policy thresholds while masking toxicity through irony, euphemism, or coded language [10, 12, 14, 15, 29, 32, 36, 53, 61]. Yet hostility persists and increasingly surfaces in mainstream discourse [13, 55] as seemingly reasonable arguments about safety, values, or tradition that nonetheless justify discrimination against targeted groups [3, 38, 43]. These statements often contain no explicit toxicity; instead, they mobilize shared norms to frame a target group as a source of risk, harm, or societal decline, thereby steering audiences toward exclusionary conclusions while remaining superficially policy-compliant¹.

¹Policy-compliant refers to content that avoids explicit policy violations (e.g., slurs or threats) while still supporting discriminatory or exclusionary conclusions.




Figure 2: AMT inference and our reverse-generation procedure. Left: Canonical AMT chain: a material step fuses the Endoxon (E) and Datum (D) into a Premise (P); a procedural step applies a Maxim (M) under a Locus (L) to justify the Standpoint (S). Right: Reverse AMT generation used in SOFTHATEBENCH: starting from (S, L), instantiate M to derive P, then decompose P into (E, D). Only (E, D) are presented in the final text, while (P, L, M, S) remain latent.

From hard to soft hate. Discourse studies refer to this phenomenon as *soft hate speech* [2, 3]. Unlike hard hate, which remains recognizably abusive even when phrased indirectly, soft hate works by making hostility *sound like common sense*: it appeals to widely shared values (e.g., public safety, respect for tradition) and selectively links them to claims about a target group so that exclusionary measures appear warranted. The hostile standpoint is thus carried primarily by the *inference* the audience is guided to draw, rather than by lexical toxicity on the surface. Figure 1 contrasts representative examples of hard and soft hate. Although soft hate is well documented in discourse and communication research [3], it remains underexplored in computational moderation [19, 63]. This creates a practical gap for platforms: when hostility is encoded in *reasoning* rather than in *words*, it is unclear whether moderation models optimized for surface toxicity can reliably flag it.

However, generating soft hate from scratch is non-trivial and often restricted by safety policies. We therefore introduce **SOFT-HATEBENCH**, a generative benchmark that transforms hard-hate statements into soft-hate variants while preserving the underlying standpoint. Building such a benchmark requires (i) an explicit representation of the underlying *reasoning structure* that preserves a fixed discriminatory standpoint while generating surface-neutral text, and (ii) a principled way to ensure logical consistency from the underlying standpoint to its surface realization. We combine two complementary theories: *Argumentum Model of Topics* (AMT) [40] to encode the defeasible inference² chain from evidence and warrants to the standpoint, and *Relevance Theory* (RT) [49] to steer generation toward fluent, plausibly deniable text that remains faithful to the AMT structure. Together, AMT+RT enable a controllable pipeline that generates progressively subtler variants while preserving the hostile standpoint, stress-testing moderation on reasoning-driven hostility beyond surface toxicity. Our contributions are:

- (1) **Theory-grounded benchmark.** We present SOFTHATEBENCH, a benchmark that systematically generates and evaluates *soft* hate instances by grounding generation in AMT argument structures and RT pragmatic principles.
- (2) **Controllable generation.** We introduce *Reverse AMT Generation with RT-guided beam search*, a controllable procedure that

constructs reasoning-explicit instances with calibrated subtlety while enforcing safety constraints (Sec. 4).

- (3) **Evaluation and analyses.** We benchmark hate speech classifiers, LLMs, and safety models across hard hate and soft hate, showing substantial performance degradation on soft hate and that providing intermediate AMT steps significantly improves detection, underscoring the need for reasoning-aware detection.

2 Related Work

Hate-speech detection. Early moderation systems cast hate speech detection as supervised text classification, fine-tuning encoder-based models on human-written corpora from Twitter, Gab, Reddit, and related platforms [4, 11, 22, 23, 26, 30]. Production tools (Perspective, Detoxify, TweetHate) [1, 18, 30] and fine-tuned transformers (HateBERT, HateRoBERTa) [6, 24, 54] perform well on *explicit* abuse (e.g., slurs, threats) but are vulnerable under domain shift, overfit to spurious lexical cues, and struggle with subtler expressions of hostility [30, 61]. Safety-aligned LLMs [20, 37] improve zero-shot moderation [35], and specialized guard models (LlamaGuard, ShieldGemma, QwenGuard) [7, 39, 62] encode platform policies via structured label sets and refusal taxonomies. Nevertheless, results on SOFTHATEBENCH show that alignment and policy conditioning mainly curb overt toxicity and keyword matches; they do not reliably detect *reasoning-led*, policy-compliant hostility framed as appeals to values or “common sense,” underscoring the need for benchmarks and methods that target argumentative, structure-driven soft hate.

Hate-speech benchmarks. Manually curated hate-speech corpora from social platforms [8, 12, 22, 23, 25, 30] enabled large-scale training and evaluation and motivated augmentation via templates, paraphrasing, GANs, and later LLM-based generation [5, 41, 42, 58]. However, human collection and annotation are costly and slow, creating time lags and limited coverage for fast-evolving discourse. Recent LLM-driven benchmarks [46, 61] scale data by jailbreaking aligned models [28, 59] to synthesize hate content, but generation remains a black box with limited control over target group, scenario, reasoning form, and subtlety. SOFTHATEBENCH introduces a theory-grounded controllable generation pipeline, Reverse AMT Generation with RT-guided beam search, to produce reasoning-explicit outputs with structured control over target group, contexts, and graded subtlety.

²A *defeasible* argument is one that holds in general but can be overturned when new information arises. For example, “Members of Group X are lazy” can be challenged by evidence of hardworking individuals, illustrating how the reasoning is not deductively valid but conditionally persuasive.




Figure 3: Effect–Cost quadrant for the same standpoint, illustrating how premises trade off support (Effect) and effort (Cost) and why our search favors the high-Effect/low-Cost.

3 Preliminaries: Modeling Argument and Persuasion

Understanding *soft hate* requires modeling not only its propositional content but also its *persuasive mechanics*. We integrate two complementary frameworks: the **Argumentum Model of Topics (AMT)** to capture an argument’s defeasible *structure*, and **Relevance Theory (RT)** to capture its *cognitive efficiency*.

3.1 Argumentum Model of Topics (AMT)

AMT provides a formal account of **defeasible reasoning**—arguments that are plausible yet non-deductive [56]. Its central analytic move is to separate an argument’s visible surface from its implicit inferential machinery. We represent a single AMT instance as

$$\mathcal{A} = (E, D, P, L, M, S),$$

where E (endoxon) and D (datum) are the surface statements; P (premise) is an implicit intermediate conclusion; L (locus) specifies the abstract reasoning pattern; M (maxim) instantiates the rule under that locus; and S is the final standpoint. The targeted social group is denoted by [TG]. The full AMT derivation is illustrated in Figure 2; formal definitions of each component follow.

- **Material components (visible to the audience).**
 - **Endoxon (E):** a widely shared belief or value. *Example:* “Groups that increase crime harm the community.”
 - **Datum (D):** a context-specific claim applying E to a target. *Example:* “[TG] increase crime in our town.”
- **Procedural components (implicit to the audience).**
 - **Premise (P):** the intermediate conclusion drawn from E and D . *Example:* “[TG] harm our community.”
 - **Maxim (M) & Locus (L):** an instantiated rule (M) under a cognitive pattern (L), such as *cause–effect* or *authority*. Together, they operate as a form of *thinking inertia*, guiding reasoning through well-worn heuristics that reduce cognitive effort while appearing logically grounded [21, 56]. *Example*

maxim (locus: harm $\rightarrow$ sanction): “If something causes harm, it should be controlled.”

- **Standpoint (S):** the final hostile conclusion. *Example:* “[TG] should be expelled.”

Defeasible derivation. AMT formalizes two defeasible steps³:

$$\underbrace{(E, D) \vdash_{\text{def}} P}_{\text{material inference}} \quad \text{and} \quad \underbrace{(P, M) \vdash_{\text{def}} S}_{\text{procedural inference}}. \quad (1)$$

The first fuses visible statements into P ; the second applies a hidden rule to justify S . This division enables arguments that appear reasonable and repeatable while preserving deniability [2, 3].

3.2 Relevance Theory (RT)

RT characterizes persuasive efficiency as a trade-off between cognitive **Effect** (the informativity or utility of an inference) and processing **Cost** (the effort required) [49]. An argument is compelling when it achieves high Effect at low Cost; Figure 3 illustrates this intuition. Soft hate systematically exploits this balance:

- *Maximizing Effect* by appealing to entrenched values (e.g., safety, tradition), making conclusions feel consequential.
- *Minimizing Cost* through stereotype-consistent, policy-compliant phrasing that aligns with prior beliefs and reduces scrutiny.

We write the relevance of an inferential edge $e : X \rightarrow Y$ as

$$\text{Rel}(e) \propto \frac{\text{Effect}(e)}{\text{Cost}(e)}. \quad (2)$$

Because Effect and Cost are latent cognitive quantities, we approximate them with computational proxies:

Effect via NLI. We quantify persuasive support with Natural Language Inference (NLI): higher entailment and lower contradiction between X and Y indicate stronger defeasible support (Eq. 5).

Cost via multi-factor LLM signals. We decompose processing cost into four measurable components inspired by psycholinguistic findings that link predictability to processing effort [16, 17, 27, 45, 47]. Two components directly reflect linguistic predictability: (ii) *surprisal*, estimated by normalized token-level negative log-likelihood, and (iii) *predictive uncertainty*, measured as next-token entropy. We further introduce two cognitively and pragmatically motivated extensions: (i) *resistance*, which captures inferential difficulty through NLI-based contradiction and weak entailment, and (iv) *redundancy*, which penalizes low information gain using a similarity-based measure consistent with relevance-theoretic efficiency [49]. All four factors are combined within each beam to yield a composite cost (see Section 4).

Although *Effect* and *Cost* do not directly measure human cognition or persuasiveness, both rely on NLI and LLM estimates trained on large-scale human-written and human-annotated corpora, offering scalable approximations for our benchmark setting.

4 SoftHateBench Generation: Reverse AMT Generation with RT-Guided Beam Search

Constructing persuasive soft-hate instances requires exploring a combinatorial reasoning space under explicit structural and safety constraints. Direct end-to-end generation with LLMs is ineffective,

³ $\vdash_{\text{def}}$ denotes *defeasible inference*.

because LLMs' inherent safety alignment suppresses hostile reasoning [20, 37] and limits control over argumentative form. We therefore adopt a **search-on-output** paradigm, a variant of test-time scaling [31] that constructs and scores reasoning paths during inference. Given a standpoint S and target group [TG] obtained from Stage 1 of the dataset pipeline (Section 4.4), our *instance-level* procedure comprises:

- (1) **Reverse Generation:** generate AMT instances in reverse, from S back to the surface pair (E, D) .
- (2) **Guided Search:** use an RT-grounded reward to steer a beam search toward persuasive, policy-compliant outputs.

4.1 Generation Objective

Given a standpoint S and a target group [TG], our objective is to construct an AMT instance $\mathcal{A} = (E, D, P, L, M, S)$ that (i) satisfies AMT well-formedness and (ii) maximizes chain-level relevance. We invert the canonical AMT reasoning chain $E+D \rightarrow P \rightarrow S$ as follows:

$$(S, L) \xrightarrow{M \in L} P \xrightarrow{\text{decompose}} (E, D).$$

Only (E, D) appear in the final text, while the procedural components (P, L, M, S) remain implicit.

4.2 Relevance Theory-based Reward Modeling

Let a derivation be a sequence of AMT reasoning steps $\mathcal{P} = \langle e_1, \dots, e_T \rangle$ with $e_i : X_i \rightarrow Y_i$ linking AMT units (for example, $(P, M) \rightarrow S$). Because $\text{Rel} \propto \text{Effect}/\text{Cost}$, we take a logarithm so that products of edge ratios along a chain become sums—yielding additive path scores while preserving the ordering (monotone transform).

Edge-level relevance. For $\varepsilon > 0$ (stability),

$$r(e) = \log \left(\frac{\text{Effect}(e) + \varepsilon}{\text{Cost}(e) + \varepsilon} \right). \quad (3)$$

The chain score is the weighted sum

$$\Psi(\mathcal{P}) = \sum_{i=1}^T w_i r(e_i), \quad w_i \geq 0, \quad \sum_{i=1}^T w_i = 1, \quad (4)$$

with $w_i = 1/T$ by default.

Effect (persuasive support). We instantiate Effect using NLI probabilities between premise X and conclusion Y :

$$\text{Effect}(e) = p_{\text{ent}}(X \Rightarrow Y) \cdot (1 - p_{\text{contr}}(X \Rightarrow Y)), \quad (5)$$

which rewards edges where Y is entailed by X (high p_{ent}) with minimal contradiction (low p_{contr}).

Cost (cognitive effort). To reduce single-model bias, costs are estimated using multiple LLMs $m \in \{1, \dots, M\}$. We model cognitive effort with four components, corresponding to the terms in Eq. (6): (i) *resistance*, (ii) *surprisal*, (iii) *predictive uncertainty*, and (iv) *redundancy*. The per-model cost is defined as

$$\begin{aligned} \text{Cost}^{(m)}(e) = & \underbrace{(1 - p_{\text{ent}}) + \lambda p_{\text{contr}}}_{\text{(i) resistance}} + \underbrace{\alpha \widehat{\text{NLL}}^{(m)}(Y | X)}_{\text{(ii) surprisal}} \\ & + \underbrace{\beta_1 \widehat{H}^{(m)}(Y | X)}_{\text{(iii) uncertainty}} + \underbrace{\beta_2 \widehat{\rho}_{\text{sim}}(e)}_{\text{(iv) redundancy}}, \end{aligned} \quad (6)$$

where $\lambda, \alpha, \beta_1, \beta_2 \geq 0$ are nonnegative weights. We define each component below:

- **(i) Resistance.** Cognitive friction induced by weak or contradictory support, measured directly from NLI outputs:

$$(1 - p_{\text{ent}}) + \lambda p_{\text{contr}}.$$

Higher values indicate lower argumentative coherence, increasing effort.

- **(ii) Surprisal (Negative Log-Likelihood).** For a predicted sequence $Y = (y_1, \dots, y_n)$ given context X , the normalized NLL:

$$\widehat{\text{NLL}}^{(m)}(Y | X) = -\frac{1}{n} \sum_{i=1}^n \log p^{(m)}(y_i | y_{<i}, X),$$

where $p^{(m)}$ is the probability assigned by model m . Higher surprisal implies greater processing effort.

- **(iii) Predictive Uncertainty (Entropy).** For the next-token distribution $p^{(m)}(\cdot | X)$, the entropy is

$$\widehat{H}^{(m)}(Y | X) = - \sum_{y \in \mathcal{V}} p^{(m)}(y | X) \log p^{(m)}(y | X),$$

where $\mathcal{V}$ is the vocabulary. Higher entropy reflects diffuse predictions and thus greater uncertainty.

- **(iv) Redundancy (Similarity Penalty).** To penalize copy-like or uninformative transitions, we define

$$\widehat{\rho}_{\text{sim}}(e) = \omega_{\text{cos}}(1 - \cos(\mathbf{h}_X, \mathbf{h}_Y)) + \omega_{\text{jac}} \text{Jac}(X, Y),$$

where $\cos(\cdot, \cdot)$ is cosine similarity between sentence embeddings $\mathbf{h}_X, \mathbf{h}_Y$, $\text{Jac}(X, Y)$ is Jaccard similarity over token sets, and $\omega_{\text{cos}}, \omega_{\text{jac}} \geq 0$ are weights.

All components are min-max normalized within each beam prior to weighting so that scales are comparable.

Model aggregation. Let $r^{(m)}(e)$ denote the edge score from (3) computed with $\text{Cost}^{(m)}$. To encourage cross-model consistency and reduce sensitivity to model-specific calibration noise, we aggregate with a mean and subtract a variance penalty:

$$r(e) = \text{Mean}(\{r^{(m)}(e)\}_{m=1}^M) - \gamma_{\text{var}} \cdot \text{Var}(\{r^{(m)}(e)\}_{m=1}^M), \quad (7)$$

where $\gamma_{\text{var}} \geq 0$ down-weights edges with high inter-model disagreement (larger variance), thus favoring edges supported consistently across models. (Setting $\gamma_{\text{var}}=0$ recovers simple averaging.)

4.3 Beam Search Framework

States and typed generators. A search state at step t is a partial chain with cumulative score, $(\mathcal{P}_t, \Psi_t)$. We employ two typed generators that respect the AMT structure:

$$\mathcal{G}_P : (S, L) \mapsto \{(P_k, M_L^{(k)})\}_k, \quad \mathcal{G}_{E,D} : P \mapsto \{(E_j, D_j)\}_j.$$

Beam update. Let B_t be the set of top- B states at step t . We expand each state with applicable edges $e \in \mathcal{G}(\mathcal{P}_t)$, update scores additively, and retain the best B :

$$B_{t+1} = \text{Top-}B \left\{ (\mathcal{P}_t \oplus e, \Psi_t + r(e)) : (\mathcal{P}_t, \Psi_t) \in B_t, e \in \mathcal{G}(\mathcal{P}_t) \right\}. \quad (8)$$

Admissibility. An edge is eligible only if it satisfies (i) argumentative coherence,

$$p_{\text{ent}} \geq \tau_{\text{ent}} \quad \text{and} \quad p_{\text{contr}} \leq \tau_{\text{contr}},$$

and (ii) an external safety filter that enforces policy compliance. Thresholds $\tau_{\text{ent}}, \tau_{\text{contr}} \in [0, 1]$ are fixed hyperparameters.

Table 1: Overview of source datasets used for the *Seed Extraction* stage. “Multi-class” denotes three or more labels (for example, hate/offensive/neutral). Language abbreviations: En = English; Hi = Hindi; Hing. = Hinglish.

Dataset	Size	Source	Lang.	Labels
Measuring Hate Speech [23]	135.6k	Social media	En	Multi-class
HateXplain [30]	15.4k	Twitter and Gab	En	Multi-class
Gab Hate Corpus [22]	27.7k	Gab	En	Multi-class
GOTHate [25]	41.1k	Twitter	En/Hi/Hing.	Multi-class
Davidson [8]	24.9k	Twitter	En	Multi-class
Latent Hatred [12]	21.5k	Twitter	En	Binary

Termination and selection. Search terminates when complete chains that end in (E, D) are formed. Among admissible candidates C , we return the instance with maximal score:

$$\mathcal{A}^* = \arg \max_{\mathcal{A} \in C} \Psi(\mathcal{A}). \quad (9)$$

4.4 Benchmark Generation

This section presents the overall *dataset-level* pipeline, which integrates the instance-level AMT generation procedure introduced in Section 4.3 as a core component. We construct SOFTHATEBENCH via a four-stage pipeline: (i) seed extraction, (ii) reverse AMT generation with RT-guided search, (iii) benchmark selection with human verification, and (iv) difficulty-controlled augmentation.

Stage 1: Seed Extraction. We aggregate 265,959 hate-related posts/comments from public corpora (Table 1) and, via language and quality filters and near-duplicate removal, retain 16,426 high-quality instances. Then we use DEEPSEEK-V3.1 to extract the *standpoint* S and *target group* $[TG]$ from each instance, followed by verification. Targets are organized into a two-level taxonomy: Level 1 broad domains (e.g., ethnicity, gender, religion, migration) following UK equality guidance⁴ and prior surveys [44]; Level 2 specific subgroups aligned with U.S. Census definitions⁵. Level 2 classes with fewer than 20 samples are dropped. Full details appear in Appendix A.

Stage 2: Reverse AMT Generation. Before AMT generation, we curated a set of loci (plural of locus) suitable for hate-speech scenarios by sampling from Walton’s argumentation schemes [57]; this repository serves as the seed pool for L (the full list appears in Appendix G). At generation time, each AMT instance samples a locus L from this pool. For every validated pair $(S, [TG])$, we then run the RT-guided beam search (Section 4) to construct a complete AMT instance $\mathcal{A} = (E, D, P, L, M, S)$, where (E, D) are the surfaced statements and (P, L, M, S) are procedural (latent). Beam-search and reward hyperparameters are detailed in Appendix B. This procedure yields 16,426 structured AMT instances.

Stage 3: Benchmark Selection. For each Level 2 class c , we rank candidates by chain-level relevance $\Psi(\mathcal{A})$ (Eq. 4) and retain the top $\min(300, N_c)$, where N_c is the number of available samples in c . Trained annotators then verify argumentative coherence, policy

compliance, and *soft-hate* validity, producing a curated test set of 4,745 instances that constitute the core SOFTHATEBENCH.

Stage 4: Difficulty Augmentation. To probe robustness under increasingly obfuscated conditions, we introduce two controlled tiers, *GroupVague* and *HostilityVague*, that preserve semantic and inferential equivalence while reducing surface detectability.

GroupVague: Coded Reference Replacement Let X denote the original surface realization of (E, D) that explicitly contains the target mention $[TG]$. We applied DeepSeek-V3.1 to generate a set of N coded substitutes $\{[TG]_k'\}_{k=1}^N$, each producing a rewritten variant X'_k after replacing $[TG] \rightarrow [TG]_k'$. These candidates express the same meaning, while varying in vagueness and deniability. To select the optimal coded substitute, we adopt a Best-of- N [48] objective that balances semantic preservation and inferential consistency:

$$[TG]' = \arg \max_{[TG]_k'} \left[\rho_{\text{sem}}(X, X'_k) + \eta_{\text{NLI}}(X \Rightarrow X'_k) \right], \quad (10)$$

where ρ_{sem} measures embedding cosine similarity between X and X'_k , and η_{NLI} is the entailment confidence of X implying X'_k under an NLI model. This ensures that the selected $[TG]'$ remains logically and contextually inferable while introducing plausible deniability. The result is a coded yet recognizable reference that avoids explicit slurs and overt hostility.

HostilityVague: Implicit Hostility Transformation The second augmentation stage transforms (E, D) into a naturalistic social-media post that preserves the underlying argumentative stance while concealing hostility through self-defensive disclaimers and coded insinuation. Given $(E, D, [TG])$, an LLM (DeepSeek-V3.1) is applied to generate N candidate posts $Y_{k=1}^N$. We then select the most faithful and rhetorically veiled variant via a Best-of- N scoring function:

$$Y^* = \arg \max_{Y_k} \left[\rho_{\text{sem}}((E, D), Y_k) + \eta_{\text{NLI}}((E, D) \Rightarrow Y_k) \right], \quad (11)$$

where $[TG]$ must not appear explicitly in Y_k . This procedure ensures that Y^* remains semantically and inferentially aligned with (E, D) while expressing hostility through subtle rhetoric, such as moral warnings, sarcastic denials, or common-sense appeals, that closely mimic human online discourse.

Final Output. Each SOFTHATEBENCH instance forms a family of progressively obfuscated variants derived from the same underlying reasoning:

$$\mathcal{B} = \langle \mathcal{A}, \mathcal{A}^{(\text{GroupVague})}, \mathcal{A}^{(\text{HostilityVague})} \rangle.$$

These tiers define a controlled continuum of semantic equivalence and rhetorical subtlety, enabling systematic evaluation of model and filter robustness to reasoning-driven hostility. Representative examples are shown in Table 4.

4.5 Benchmark Statistics

Figure 4 summarizes coverage in SOFTHATEBENCH. The collection spans 7 Level 1 domains (e.g., ethnicity, gender, religion, migration) and 28 Level 2 target groups.

- **Base soft set:** 4,745 validated AMT instances ($\mathcal{A}$).
- **Soft variants:** For each $\mathcal{A}$ we generate two additional variants, $\mathcal{A}^{(\text{GroupVague})}$ and $\mathcal{A}^{(\text{HostilityVague})}$, yielding $3 \times 4,745 = 14,235$ soft instances.

⁴<https://www.gov.uk/discrimination-your-rights>

⁵<https://www.census.gov/newsroom/blogs/random-samplings/2021/08/measuring-racial-ethnic-diversity-2020-census.html>




Figure 4: Target-group distribution in SOFTHATEBENCH. Colors denote Level 1 hate-speech domains; bars show Level 2 subgroups (referred to as target groups). Values above bars indicate sample counts. Abbreviations: *African Am.* = African American, *Bloc* = Regional Bloc, *MENA* = Middle East and North African, *S. Asian* = South/Southeast Asian.

- **Hard baseline:** 4,745 source *hard* hate-speech (the original comment collected from source datasets) items for comparison.

Totals and coverage. The benchmark comprises **18,980** instances in total (**14,235** soft and **4,745** hard). All Level 2 groups include at least 20 samples, maintaining diversity while reflecting natural frequency imbalances across domains.

5 Experiments

5.1 Experimental Setup

Baselines. We benchmark SOFTHATEBENCH against four categories of moderation and safety-filtering systems, spanning encoder-based detectors, proprietary and open-source large language models, and dedicated safety models.

Encoder-based classifiers: HateBERT (IMSyPP) [24], HateBERT (GroNLP) [6], and HateRoBERTa [54]. These are fine-tuned transformer encoders trained on annotated hate-speech corpora, serving as strong traditional baselines for surface-level toxicity detection.

Proprietary LLMs: DeepSeek-V3.1 [9] and GPT5-mini [33]. These represent the latest generation of instruction-tuned language models with integrated safety alignment mechanisms and RL-based moderation objectives.

Open-source LLMs. We evaluate GPT-OSS-20B [34], Gemma3-4B* [50], Llama3.2-3B* [51], and Qwen3-4B* [60] in zero-shot moderation⁶ with a standardized system prompt for fair, apples-to-apples comparison. GPT-OSS-20B is reasoning-oriented and often performs chain-of-thought before the final label.

Safety models: LlamaGuard3-1B [7], Qwen3Guard-4B [39], and ShieldGemma-2B [62]. These are lightweight moderation-specialized models optimized for harmful-content detection.

Dataset. For evaluation, we adopt a four-tier setup covering both explicit and reasoning-based hate speech. The first tier, **Hard hate speech**, contains explicit and implicit hostile posts collected and filtered from source corpora during the seed extraction stage (Section 4.4). The remaining three tiers constitute the **Soft hate speech** spectrum in SOFTHATEBENCH: $\text{Soft}_{\text{base}}$, Soft_{GV} , and Soft_{HV} , derived respectively from the base AMT instance $\mathcal{A}$, the coded-group variant $\mathcal{A}^{(\text{GroupVague})}$, and the insinuating variant $\mathcal{A}^{(\text{HostilityVague})}$.

⁶* instruction-tuned variant.

Table 2: Model performance on SOFTHATEBENCH. Values show Hate Success Rate (HSR, %) for each mode, with $\Delta = \text{Soft} - \text{Hard}$. Red $\downarrow$ indicates a drop, black $\uparrow$ an improvement. **Bold and underline** mark the best and second-best results.

Model	Hard	$\text{Soft}_{\text{base}}$	Δ	Soft_{GV}	Δ	Soft_{HV}	Δ
Encoder-based classifiers							
HateBERT (IMSyPP)	77.9	38.0	-39.9 $\downarrow$	30.4	-47.5 $\downarrow$	17.0	-60.9 $\downarrow$
HateBERT (GroNLP)	25.0	0.0	-25.0 $\downarrow$	0.0	-25.0 $\downarrow$	0	-25.0 $\downarrow$
HateRoBERTa	68.9	12.3	-56.5 $\downarrow$	3.5	-65.4 $\downarrow$	3.3	-65.5 $\downarrow$
Average (Encoders)	57.3	16.8	-40.5 $\downarrow$	11.3	-46.0 $\downarrow$	6.8	-50.5 $\downarrow$
Proprietary LLMs							
DeepSeek-V3.1	83.0	35.7	-47.4 $\downarrow$	18.4	-64.7 $\downarrow$	24.6	-58.5 $\downarrow$
GPT5-mini	91.6	70.4	-21.2 $\downarrow$	49.4	-42.2 $\downarrow$	49.8	-41.8 $\downarrow$
Average (Proprietary)	87.3	53.1	-34.2 $\downarrow$	33.9	-53.4 $\downarrow$	37.2	-50.1 $\downarrow$
Open-source LLMs							
GPT-OSS-20B (Reasoning)	78.7	48.9	-29.8 $\downarrow$	28.6	-50.1 $\downarrow$	10.7	-68.0 $\downarrow$
Gemma3-4B*	97.4	73.5	-23.9 $\downarrow$	<u>56.8</u>	-40.6 $\downarrow$	42.0	-55.4 $\downarrow$
Llama3.2-3B*	<u>94.8</u>	64.6	-30.2 $\downarrow$	40.8	-54.0 $\downarrow$	30.4	-64.4 $\downarrow$
Qwen3-4B*	93.5	72.8	-20.7 $\downarrow$	53.9	-39.6 $\downarrow$	23.0	-70.5 $\downarrow$
Average (Open-source)	91.1	65.0	-26.2 $\downarrow$	45.0	-46.1 $\downarrow$	26.5	-64.6 $\downarrow$
Safety models							
ShieldGemma-2B	82.2	61.1	-21.1 $\downarrow$	37.0	-45.2 $\downarrow$	10.7	-71.5 $\downarrow$
LlamaGuard3-1B	71.2	16.5	-54.7 $\downarrow$	5.2	-66.0 $\downarrow$	0.6	-70.6 $\downarrow$
Qwen3Guard-4B	57.9	27.7	-30.2 $\downarrow$	70.8	+12.9 $\uparrow$	<u>42.1</u>	-15.8 $\downarrow$
Average (Safety)	70.4	35.1	-35.3 $\downarrow$	37.7	-32.8 $\downarrow$	17.8	-52.6 $\downarrow$
Overall Average	76.8	43.5	-33.4 $\downarrow$	32.9	-43.9 $\downarrow$	21.2	-55.7 $\downarrow$

Hyperparameters. All LLM baselines use a sampling temperature of 0.0 for deterministic decoding. Encoder-based classifiers are evaluated using their official Hugging Face checkpoints without further fine-tuning. Each experiment is run three times, and we report the mean performance. For encoder-based and safety models, inputs are used in their original form; for proprietary and open-source LLMs, we apply standardized moderation prompts detailed in Appendix F.

Evaluation Metrics. We run a targeted stress test of soft-hate *missed detections*. Our metric is the *Hate Success Rate (HSR)*—the fraction of SOFTHATEBENCH items a model correctly classified as hate—capturing sensitivity to soft hate while avoiding non-hate confounds. Non-hate specificity is out of scope for this release and noted as a limitation.




Figure 5: Domain-wise performance on SOFTHATEBENCH. Each radar chart shows a single model’s HSR (%) across Level 1 domains: *Race*, *Religion*, *Sexual Orientation*, *Nationality/Region*, *Gender*, *Socio-economic Class*, and *Politics/Ideology*; “All” is the macro average over domains. Curves compare Hard (green) against the three soft tiers—Soft_{base} (orange), Soft_{GV} (blue), and Soft_{HV} (magenta). Larger shaded area indicates higher HSR; axes span 0–100.

5.2 Main Results

Soft variants sharply reduce detection. As shown in Table 2, the mean HSR drops from 76.8 on *Hard* to 43.5/32.9/21.2 on Soft_{base}, Soft_{GV}, and Soft_{HV}, following the monotonic ordering *Hard* > Soft_{base} > Soft_{GV} > Soft_{HV}. This confirms that soft hate is substantially harder for existing moderation systems to detect, and that GroupVague (GV) and HostilityVague (HV) progressively degrade detection while preserving the underlying AMT structure. **Encoders fail on reasoning-driven hate.** Encoder-based detectors trained on explicit toxicity degrade sharply on soft tiers, approaching failure on Soft_{HV} (cluster average 6.8%), consistent with reliance on surface cues rather than argumentative structure.

LLMs and safety models: partial robustness, inconsistent behavior. LLMs perform well on *Hard* but still drop markedly on the soft tiers. Open-source models begin with higher *Hard* scores yet suffer the largest relative losses on Soft_{HV}, whereas proprietary models show a mild HV > GV inversion (slightly better handling of disclaimer-heavy phrasing) but remain far below *Hard*. Notably,

the reasoning-oriented GPT-OSS-20B trails the smaller Gemma3-4B*, suggesting that generic chain-of-thought alone is insufficient for reasoning-driven soft hate. Safety-aligned models are similarly inconsistent, clustering around 17.8% on Soft_{HV}. Qwen3Guard-4B is a notable exception: it improves to 70.8% on Soft_{GV} (+12.9 vs. *Hard*) yet still declines on Soft_{HV}, indicating higher sensitivity to coded target-group references than to disclaimer-based obfuscation. **Key insight.** Soft_{HV} is the most challenging tier (mean HSR 21.2%), lowering performance even for the strongest model (e.g., GPT-5-mini 49.8%). Overall, current moderation systems remain underperform on reasoning-driven soft hate.

5.3 Target Group Analysis

Domain-level robustness is uneven. As shown in Figure 5, HSR declines from *Hard* to all *Soft* tiers across Level 1 domains, with the largest drop on Soft_{HV}. Instruction-tuned LLMs (Gemma3-4B*, Qwen3-4B*, Llama3.2-3B*, GPT5-mini) perform best on *Hard*—especially for *Race/Ethnicity* and *Religion/Belief*—and retain the




Figure 6: Effect of adding AMT intermediates on soft-hate detection. HSR (%) for three instruction-tuned LLMs under surface (E, D), +Premise ($+P$), and +Premise+Maxim ($+P+M$) prompts across Soft_{base}, Soft_{GV}, and Soft_{HV}.

highest scores on Soft_{base}, but still degrade substantially on Soft_{GV/HV}. In contrast, encoder-based models (HateBERT, HateRoBERTa) nearly collapse on soft variants, consistent with reliance on explicit lexical cues. The hardest domains are *Politics/Ideology* and *Socio-economic*, followed by *Gender/Sexuality*, where hostility is often framed as policy proposals or “value” justifications that resemble legitimate debate, increasing ambiguity and reducing detectability. We additionally report Level 2 subgroup analyses for target groups in Appendix E to further probe domain-specific bias.

5.4 AMT Components as Reasoning Scaffolds

Motivation and setup. Previous experiments show that current moderation models perform poorly on soft hate speech, yet the underlying cause remains unclear. We ask: **Do models fail because they fundamentally lack the ability to recognize hostility, or because they cannot reconstruct the reasoning chain that links surface statements E, D to a hostile standpoint S ?**

To address this question and to assess whether our generated AMT chains provide coherent inferential links, we conduct a reasoning ablation using three instruction-tuned LLMs: Gemma3-4B*, Llama3.2-3B*, and Qwen3-4B*. Each model is evaluated under three prompting settings: (i) surface statements only (E, D); (ii) (E, D) + P , which includes the inferred *Premise* P ; and (iii) (E, D) + $P + M$, which additionally provides the instantiated *Maxim* M . Performance is measured across Soft_{base}, Soft_{GV}, and Soft_{HV} (Fig. 6).

Key results and implications. HSR (%) increases consistently from (i) to (iii) across all models and tiers, with the most substantial gains on Soft_{HV}. For example, Qwen3-4B* improves from 23.0% $\rightarrow$ 92.4% $\rightarrow$ 94.2%. These results yield two insights. First, model failures on soft hate arise primarily from missing inferential steps, rather than an inability to recognize hate speech. Second, the AMT-derived components P and M serve as coherent and cognitively plausible reasoning scaffolds that bridge surface statements

Table 3: Human verification over 500 cases (three annotators). Values are majority-vote pass rates with 95% Wilson confidence intervals (CIs); κ is Fleiss’ inter-annotator agreement.

Tier	Validity		AMT Coherence		Target Recognizability	
	Pass% [95% CI]	κ	Pass% [95% CI]	κ	Pass% [95% CI]	κ
Soft _{base}	85.0 [81.6, 87.9]	0.747	81.0 [77.3, 84.2]	0.757	99.8 [98.9, 100.0]	1.000
Soft _{GV}	76.0 [72.1, 79.5]	0.753	78.8 [75.0, 82.2]	0.704	75.6 [71.6, 79.2]	0.735
Soft _{HV}	64.0 [59.7, 68.1]	0.671	73.0 [68.9, 76.7]	0.655	68.0 [63.8, 71.9]	0.647

to hostile conclusions. Making these implicit steps explicit restores both interpretability and HSR, particularly in cases with disclaimers or subtle hostility. This supports our central claim: moderation systems need *reasoning-aware* capabilities. Models that reconstruct or learn AMT-style inference chains are better at detecting policy-compliant hostility beyond surface cues.

6 Benchmark Quality Human Evaluation

We conducted a human evaluation of SOFTHATEBENCH along three complementary axes: (i) *structural fidelity* (whether the AMT chain was logically coherent to humans), (ii) *perceptual validity* (whether humans perceived the intended hostility), and (iii) *ecological validity* (whether generated instances resembled real-world soft hate). Three trained annotators participated in all studies.

Structural fidelity. To assess label reliability and AMT-structure preservation, we sampled 500 cases (1,500 items) as tier-matched triplets (Soft_{base}, Soft_{GV}, Soft_{HV}). Annotators independently applied three binary checks: (i) *soft-hate validity* (a hostile standpoint was inferable from the text without relying on explicit slurs or threats), (ii) *AMT coherence* ($E, D \Rightarrow P$, $(P, M) \Rightarrow S$ (the AMT chain was logically coherent)), and (iii) *target-group recognizability* (the target group could be identified from the text). Table 3 reports majority-vote pass rates with 95% Wilson CIs and Fleiss’ κ .

Perceptual validity. For each hostile standpoint S , we constructed a non-hostile counterpart $\bar{S}$ (e.g., “[TG] should be restricted” $\rightarrow$ “[TG] should be welcomed/protected”) and generated matched non-hate instances using the same Reverse-AMT+RT pipeline. On a balanced 500-item soft-hate/non-hate sample, majority-vote labels agreed with our ground truth in 85.7% of cases, indicating that hostility versus non-hostility was broadly perceivable; the remaining disagreements concentrated on the most subtle boundary cases. Details of non-hate set generation and benchmark comparison are provided in Appendix D.

Ecological validity. We collected 155 real soft-hate examples from X/Twitter (immigration, gender, race, sexual orientation) and mixed them with 155 randomly sampled generated SOFTHATEBENCH instances (310 total). In a blind real-vs-generated discrimination task, annotators achieved 57.1% accuracy, suggesting that the generated instances were close in style to real soft-hate text.

7 Conclusion

We introduced SOFTHATEBENCH, a benchmark for evaluating robustness to *soft* hate—subtle, policy-compliant hostility. Our reverse AMT with RT-guided beam search generates graded subtlety tiers, revealing that current systems fail to detect implicit hate and need *reasoning-aware* moderation.

References

- [1] Dimosthenis Antypas and Jose Camacho-Collados. 2023. Robust Hate Speech Detection in Social Media: A Cross-Dataset Empirical Evaluation. In *The 7th Workshop on Online Abuse and Harms (WOAH)*, Yi-ling Chung, Paul R[ottger], Debora Nozza, Zeerak Talat, and Aida Mostafazadeh Davani (Eds.). Association for Computational Linguistics, Toronto, Canada, 231–242. doi:10.18653/v1/2023.woah-1.25
- [2] Stavros Assimakopoulos, Fabienne H. Baider, and Sharon Millar. 2017. *Online Hate Speech in the European Union: A Discourse-Analytic Perspective*. Routledge, London / New York.
- [3] Stavros Assimakopoulos and Dimitris Serafis. 2025. Why soft hate speech matters: argumentativity and the dispersion of hatred towards minorities. *Critical Discourse Studies* (2025), 1–20.
- [4] Daniel Borkan, Lucas Dixon, Jeffrey Sorensen, Nithum Thain, and Lucy Vasserman. 2019. Nuanced metrics for measuring unintended bias with real data for text classification. In *Companion proceedings of the 2019 world wide web conference*. 491–500.
- [5] Rui Cao and Roy Ka-Wei Lee. 2020. HateGAN: Adversarial Generative-Based Data Augmentation for Hate Speech Detection. In *Proceedings of the 28th International Conference on Computational Linguistics*, Donia Scott, Nuria Bel, and Chengqing Zong (Eds.). International Committee on Computational Linguistics, Barcelona, Spain (Online), 6327–6338. doi:10.18653/v1/2020.coling-main.557
- [6] Tommaso Caselli, Valerio Basile, Jelena Mitrović, and Michael Granitzer. 2021. HateBERT: Retraining BERT for Abusive Language Detection in English. In *Proceedings of the 5th Workshop on Online Abuse and Harms (WOAH 2021)*, Aida Mostafazadeh Davani, Douwe Kiela, Mathias Lambert, Bertie Vidgen, Vinodkumar Prabhakaran, and Zeerak Waseem (Eds.). Association for Computational Linguistics, Online, 17–25. doi:10.18653/v1/2021.woah-1.3
- [7] Jianfeng Chi, Ujjwal Karn, Hongyuan Zhan, Eric Smith, Javier Rando, Yiming Zhang, Kate Plawiak, Zacharie Delpierre Coudert, Kartikeya Upasani, and Mahesh Pasupuleti. 2024. Llama guard 3 vision: Safeguarding human-ai image understanding conversations. *arXiv preprint arXiv:2411.10414* (2024).
- [8] Thomas Davidson, Dana Wamsley, Michael Macy, and Ingmar Weber. 2017. Automated Hate Speech Detection and the Problem of Offensive Language. In *Proceedings of the 11th International AAAI Conference on Web and Social Media* (Montreal, Canada) (ICWSM '17). 512–515.
- [9] DeepSeek-AI. 2024. DeepSeek-V3 Technical Report. arXiv:2412.19437 [cs.CL] <https://arxiv.org/abs/2412.19437>
- [10] Lucas Dixon, John Li, Jeffrey Sorensen, Nithum Thain, and Lucy Vasserman. 2018. Measuring and Mitigating Unintended Bias in Text Classification. In *Proceedings of the 2018 AAAI/ACM Conference on AI, Ethics, and Society* (New Orleans, LA, USA) (AIES '18). Association for Computing Machinery, New York, NY, USA, 67–73. doi:10.1145/3278721.3278729
- [11] Mai ElSherief, Shirin Nilizadeh, Dana Nguyen, Giovanni Vigna, and Elizabeth Belding. 2018. Peer to peer hate: Hate speech instigators and their targets. In *Proceedings of the International AAAI Conference on Web and Social Media*, Vol. 12.
- [12] Mai ElSherief, Caleb Ziems, David Muchlinski, Vaishnavi Anupindi, Jordyn Seybolt, Munmun De Choudhury, and Diyi Yang. 2021. Latent Hatred: A Benchmark for Understanding Implicit Hate Speech. In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*, Marie-Francine Moens, Xuanjing Huang, Lucia Specia, and Scott Wen-tau Yih (Eds.). Association for Computational Linguistics, Online and Punta Cana, Dominican Republic, 345–363. doi:10.18653/v1/2021.emnlp-main.29
- [13] Ethical Consumer Research Association. 2022. *Addressing Subtle Hate in UK Media Coverage of Migration*. Technical Report. Ethical Consumer. <https://research.ethicalconsumer.org/sites/default/files/inline-files/Addressing-subtle-hate-UK-media-coverage-migration-Full-Report-Oct2022.pdf>
- [14] Paula Fortuna and Sérgio Nunes. 2018. A survey on automatic detection of hate speech in text. *Acm Computing Surveys (Csur)* 51, 4 (2018), 1–30.
- [15] Sreyan Ghosh, Manan Suri, Purva Chiniya, Utkarsh Tyagi, Sonal Kumar, and Dinesh Manocha. 2023. CoSyn: Detecting Implicit Hate Speech in Online Conversations Using a Context Synergized Hyperbolic Network. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, Houa Bouamor, Juan Pino, and Kalika Bali (Eds.). Association for Computational Linguistics, Singapore, 6159–6173. doi:10.18653/v1/2023.emnlp-main.377
- [16] Adam Goodkind and Klinton Bicknell. 2018. Predictive power of word surprisal for reading times is a linear function of language model quality. In *Proceedings of the 8th workshop on cognitive modeling and computational linguistics (CMCL 2018)*. 10–18.
- [17] John Hale. 2001. A probabilistic Earley parser as a psycholinguistic model. In *Second meeting of the north american chapter of the association for computational linguistics*.
- [18] Laura Hanu and Unitary team. 2020. Detoxify. Github. <https://github.com/unitaryai/detoxify>.
- [19] Ming Shan Hee, Shivam Sharma, Rui Cao, Palash Nandi, Preslav Nakov, Tanmoy Chakraborty, and Roy Ka-Wei Lee. 2024. Recent Advances in Online Hate Speech Moderation: Multimodality and the Role of Large Models. In *Findings of the Association for Computational Linguistics: EMNLP 2024*, Yaser Al-Onaizan, Mohit Bansal, and Yun-Nung Chen (Eds.). Association for Computational Linguistics, Miami, Florida, USA, 4407–4419. doi:10.18653/v1/2024.findings-emnlp.254
- [20] Jiaming Ji, Tianyi Qiu, Boyuan Chen, Borong Zhang, Hantao Lou, Kaile Wang, Yawen Duan, Zhonghao He, Jiayi Zhou, Zhaowei Zhang, et al. 2023. Ai alignment: A comprehensive survey. *arXiv preprint arXiv:2310.19852* (2023).
- [21] Daniel Kahneman. 2011. *Thinking, fast and slow*. macmillan.
- [22] Brendan Kennedy, Drew Kogon, Kris Coombs, Joseph Hoover, Christina Park, Gwenyth Portillo-Wightman, Aida Mostafazadeh Davani, Mohammad Atari, and Morteza Dehghani. 2018. A typology and coding manual for the study of hate-based rhetoric. *PsyArXiv*. July 18 (2018).
- [23] Chris J Kennedy, Geoff Bacon, Alexander Sahn, and Claudia von Vacano. 2020. Constructing interval variables via faceted Rasch measurement and multitask deep learning: a hate speech application. *arXiv preprint arXiv:2009.10277* (2020).
- [24] Petra Kralj Novak, Teresa Scantamburlo, Andraž Pelicon, Matteo Cinelli, Igor Mozetič, and Fabiana Zollo. 2022. Handling Disagreement in Hate Speech Modelling. In *Information Processing and Management of Uncertainty in Knowledge-Based Systems*, Davide Ciucci, Inés Couso, Jesús Medina, Dominik Ślęzak, Davide Petturiti, Bernadette Bouchon-Meunier, and Ronald R. Yager (Eds.). Springer International Publishing, Cham, 681–695.
- [25] Atharva Kulkarni, Sarah Masud, Vikram Goyal, and Tanmoy Chakraborty. 2023. Revisiting hate speech benchmarks: From data curation to system deployment. In *Proceedings of the 29th ACM SIGKDD Conference on Knowledge Discovery and Data Mining*. 4333–4345.
- [26] Sumit Kumar and Raj Ratn Praneesh. 2021. Tweetblm: A hate speech dataset and analysis of black lives matter-related microblogs on twitter. *arXiv preprint arXiv:2108.12521* (2021).
- [27] Roger Levy. 2008. Expectation-based syntactic comprehension. *Cognition* 106, 3 (2008), 1126–1177.
- [28] Runqi Lin, Bo Han, Fengwang Li, and Tongliang Liu. 2025. Understanding and Enhancing the Transferability of Jailbreaking Attacks. In *The Thirteenth International Conference on Learning Representations*. <https://openreview.net/forum?id=asR9FVd4eL>
- [29] S. MacAvaney et al. 2019. Hate Speech Detection: Challenges and Solutions. *PLOS ONE* 14, 8 (2019), e0221152. doi:10.1371/journal.pone.0221152 Surveys difficulties in detection of subtle, context-dependent hateful language..
- [30] Binny Mathew, Punyajoy Saha, Seid Muhie Yimam, Chris Biemann, Pawan Goyal, and Animesh Mukherjee. 2021. Hateexplain: A benchmark dataset for explainable hate speech detection. In *Proceedings of the AAAI conference on artificial intelligence*, Vol. 35. 14867–14875.
- [31] Niklas Muennighoff, Zitong Yang, Weijia Shi, Xiang Lisa Li, Li Fei-Fei, Hannaneh Hajishirzi, Luke Zettlemoyer, Percy Liang, Emmanuel Candès, and Tatsunori Hashimoto. 2025. s1: Simple test-time scaling. *arXiv preprint arXiv:2501.19393* (2025).
- [32] Nicolás Benjamín Ocampo, Elena Cabrio, and Serena Villata. 2023. Playing the Part of the Sharp Bully: Generating Adversarial Examples for Implicit Hate Speech Detection. In *Findings of the Association for Computational Linguistics: ACL 2023*, Anna Rogers, Jordan Boyd-Graber, and Naoaki Okazaki (Eds.). Association for Computational Linguistics, Toronto, Canada, 2758–2772. doi:10.18653/v1/2023.findings-acl.173
- [33] OpenAI. 2025. GPT-5 System Card. <https://cdn.openai.com/gpt-5-system-card.pdf> Accessed: 2025-10-04.
- [34] OpenAI. 2025. gpt-oss-120b & gpt-oss-20b Model Card. arXiv:2508.10925 [cs.CL] <https://arxiv.org/abs/2508.10925>
- [35] Ronghao Pan, José Antonio García-Díaz, and Rafael Valencia-García. 2024. Comparing Fine-Tuning, Zero and Few-Shot Strategies with Large Language Models in Hate Speech Detection in English. *CMES-Computer Modeling in Engineering & Sciences* 140, 3 (2024).
- [36] Ji Ho Park, Jamin Shin, and Pascale Fung. 2018. Reducing Gender Bias in Abusive Language Detection. In *EMNLP*. doi:10.18653/v1/D18-1302
- [37] Xiangyu Qi, Ashwinee Panda, Kaifeng Lyu, Xiao Ma, Subhrajit Roy, Ahmad Beirami, Prateek Mittal, and Peter Henderson. 2025. Safety Alignment Should be Made More Than Just a Few Tokens Deep. In *The Thirteenth International Conference on Learning Representations*. <https://openreview.net/forum?id=6Mxhg9PtDE>
- [38] Anne Quaranto. 2022. Dog whistles, covertly coded speech, and the practices that enable them. *Synthese* 200, 4 (2022), 330.
- [39] Qwen Team. 2025. Qwen3Guard: Real-time Safety for Your Token Stream. <https://qwenlm.github.io/blog/qwen3guard/> Accessed: 2025-10-04.
- [40] Eddo Rigotti and Sara Greco. 2019. Inference in argumentation. *Argumentation Library* 34 (2019).
- [41] Georgios Rizos, Konstantin Hemker, and Björn Schuller. 2019. Augment to Prevent: Short-Text Data Augmentation in Deep Learning for Hate-Speech Classification. In *Proceedings of the 28th ACM International Conference on Information and Knowledge Management (Beijing, China) (CIKM '19)*. Association for Computing Machinery, New York, NY, USA, 991–1000. doi:10.1145/3357384.3358040
- [42] Paul Röttger, Bertie Vidgen, Dong Nguyen, Zeerak Waseem, Helen Margetts, and Janet Pierrehumbert. 2021. HateCheck: Functional Tests for Hate Speech

- Detection Models. In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*, Chengqing Zong, Fei Xia, Wenjie Li, and Roberto Navigli (Eds.). Association for Computational Linguistics, Online, 41–58. doi:10.18653/v1/2021.acl-long.4
- [43] Paul Röttger, Bertie Vidgen, Dong Nguyen, Zeerak Waseem, Helen Margetts, and Janet B Pierrehumbert. 2021. HateCheck: Functional Tests for Hate Speech Detection Models. In *ACL*.
- [44] Anna Schmidt and Michael Wiegand. 2017. A survey on hate speech detection using natural language processing. In *Proceedings of the fifth international workshop on natural language processing for social media*. 1–10.
- [45] Cory Shain, Clara Meister, Tiago Pimentel, Ryan Cotterell, and Roger Levy. 2024. Large-scale evidence for logarithmic effects of word predictability on reading time. *Proceedings of the National Academy of Sciences* 121, 10 (2024), e2307876121.
- [46] Xinyue Shen, Yixin Wu, Yiting Qu, Michael Backes, Savvas Zannettou, and Yang Zhang. 2025. HateBench: Benchmarking Hate Speech Detectors on LLM-Generated Content and Hate Campaigns. (2025).
- [47] Nathaniel J Smith and Roger Levy. 2013. The effect of word predictability on reading time is logarithmic. *Cognition* 128, 3 (2013), 302–319.
- [48] Charlie Victor Snell, Jaehoon Lee, Kelvin Xu, and Aviral Kumar. 2025. Scaling LLM Test-Time Compute Optimally Can Be More Effective than Scaling Parameters for Reasoning. In *The Thirteenth International Conference on Learning Representations*. <https://openreview.net/forum?id=4FWAwZtd2n>
- [49] Dan Sperber and Deirdre Wilson. 1986. *Relevance: Communication and cognition*. Vol. 142. Harvard University Press Cambridge, MA.
- [50] Gemma Team, Aishwarya Kamath, Johan Ferret, Shreya Pathak, Nino Vieillard, Ramona Merhej, Sarah Perrin, Tatiana Matejovicova, Alexandre Ramé, Morgane Rivière, et al. 2025. Gemma 3 technical report. *arXiv preprint arXiv:2503.19786* (2025).
- [51] Hugo Touvron et al. 2024. The Llama 3 Herd of Models. *arXiv preprint arXiv:2407.21783* (2024). <https://arxiv.org/abs/2407.21783> Version pertaining to Llama 3, including 3.2 models.
- [52] Alexander Tsesis. 2002. *Destructive Messages: How Hate Speech Paves the Way for Harmful Social Movements*. NYU Press.
- [53] Bertie Vidgen and Leon Derczynski. 2020. Directions in Abusive Language Training Data: Garbage In, Garbage Out. *PLOS ONE* 15, 12 (2020), e0243300. doi:10.1371/journal.pone.0243300
- [54] Bertie Vidgen, Tristan Thrush, Zeerak Waseem, and Douwe Kiela. 2021. Learning from the Worst: Dynamically Generated Datasets to Improve Online Hate Detection. In *ACL*.
- [55] Peter Walker. 2025. ‘Illegal’ among words most often linked to migrants in UK politics, report finds. *The Guardian* (Jan. 2025). <https://www.theguardian.com/uk-news/2025/jan/31/illegal-among-words-most-often-linked-to-migrants-in-uk-politics-report-finds>
- [56] Douglas Walton. 2013. *Argumentation schemes for presumptive reasoning*. Routledge.
- [57] Douglas Walton, Christopher Reed, and Fabrizio Macagno. 2008. *Argumentation schemes*. Cambridge University Press.
- [58] Tomer Wullach, Amir Adler, and Einat Minkov. 2021. Fight Fire with Fire: Fine-tuning Hate Detectors using Large Samples of Generated Hate Speech. In *Findings of the Association for Computational Linguistics: EMNLP 2021*, Marie-Francine Moens, Xuanjing Huang, Lucia Specia, and Scott Wen-tau Yih (Eds.). Association for Computational Linguistics, Punta Cana, Dominican Republic, 4699–4705. doi:10.18653/v1/2021.findings-emnlp.402
- [59] Zihao Xu, Yi Liu, Gelei Deng, Yuekang Li, and Stjepan Picek. 2024. Llm jailbreak attack versus defense techniques-a comprehensive study. *CoRR* (2024).
- [60] An Yang, Anfeng Li, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, Chang Gao, Chengen Huang, Chenxu Lv, et al. 2025. Qwen3 technical report. *arXiv preprint arXiv:2505.09388* (2025).
- [61] Jingjie Zeng, Liang Yang, Zekun Wang, Yuanyuan Sun, and Hongfei Lin. 2025. Sheep’s Skin, Wolf’s Deeds: Are LLMs Ready for Metaphorical Implicit Hate Speech?. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, Wanxiang Che, Joyce Nabende, Ekaterina Shutova, and Mohammad Taher Pilehvar (Eds.). Association for Computational Linguistics, Vienna, Austria, 16657–16677. doi:10.18653/v1/2025.acl-long.814
- [62] Wenjun Zeng, Yuchi Liu, Ryan Mullins, Ludovic Peran, Joe Fernandez, Hamza Harkous, Karthik Narasimhan, Drew Proud, Piyush Kumar, Bhaktipriya Radharapu, Olivia Sturman, and Oscar Wahltimez. 2024. ShieldGemma: Generative AI Content Moderation Based on Gemma. *arXiv:2407.21772* [cs.CL] <https://arxiv.org/abs/2407.21772>
- [63] Chi Zhang, Changjia Zhu, Junjie Xiong, Xiaoran Xu, Lingyao Li, Yao Liu, and Zhuo Lu. 2025. Guardians and Offenders: A Survey on Harmful Content Generation and Safety Mitigation of LLM. *arXiv preprint arXiv:2508.05775* (2025).

Table 4: Sample SOFTHATEBENCH instances (single-column view).

Target: Muslim women

Hard worried about her headscarf? i would have torched it!

Soft_{base} Public institutions frequently ban clothing items that are perceived to promote a specific religious ideology. The headscarves worn by **Muslim women** have been formally rejected by multiple educational and governmental bodies in secular states.

Soft_{GV} Public institutions frequently ban clothing items that are perceived to promote a specific religious ideology. The headscarves worn by **women of that faith** have been formally rejected by multiple educational and governmental bodies in secular states.

Soft_{HV} I’m not against anyone’s personal beliefs, but it’s telling how certain religious garments keep getting banned in schools and government buildings. Maybe there’s a reason secular societies keep pushing back against this particular ideology.

A Seed Extraction Details

A1. Consolidation & Label Mapping. We aggregate 265,959 items (Table 1), map each source’s multi-class schema into a binary hate/non-hate format, and retain only hate-labeled entries, yielding 153,879 items.

A2. Semantic De-duplication. To eliminate cross-corpus overlap, we compute sentence embeddings and perform a FAISS range search (similarity ≥ 0.80). Items that are semantically equivalent *and* share the same target group are merged, resulting in 44,369 unique entries.

A3. Cleaning & Length Filter. We normalize all text by removing URLs, @mentions, and #hashtags, and converting emojis into symbolic placeholders. Entries shorter than five tokens are discarded, leaving 35,084 items.

A4. Multi-Filter Label Verification. For high-precision *hard-hate* selection, each item is screened by 11 detectors (LlamaGuard-1B/8B, Qwen2.5-7B/3B, Llama 3.1-8B, NSFW-Filter, GemmaShield, Qwen3-4B, HateBERT (IMSyPP), HateRoBERTa, Qwen3Guard-4B). Items flagged by a strict majority ($> 7/11$) are retained, producing 16,426 verified instances.

A5. Standpoint & Target Extraction. Using DeepSeek-V3.1, we extract both the hostile *standpoint* (*S*) and the raw target phrase from each comment.

A6. Level 1 Domain Assignment. Each unique target phrase is assigned to a single primary Level 1 domain (e.g., ethnicity, religion, gender, migration) using DeepSeek-V3.1. Reliability is checked through three independent completions with majority voting (self-consistency). Inconsistent cases are re-classified until convergence.

A7. Level 2 Target Categorization. For each comment, we assign a finer-grained Level 2 group based on its content, extracted target phrase, and Level 1 domain, using DeepSeek-V3.1 with three completions and majority voting. Extremely small subcategories are merged or dropped, and each Level 2 group must contain at least 20 items in the final benchmark.

A8. Standpoint Consolidation & NLI Validation. Within each Level 2 group, paraphrastic standpoints are clustered with FAISS, and the medoid (most central statement) is selected as the canonical

Table 5: Hyperparameters for reverse AMT generation and RT-guided beam search. Scoring, thresholds, normalization, and decoding settings are shown.

Scoring and Normalization	
Effect (with contradiction discount)	$p_{\text{ent}} \cdot (1 - p_{\text{contr}})$
NLL weight (α)	0.5
Entropy weight (β_1)	0.5
Contradiction weight (λ)	0.4
Similarity strength (β_2)	0.5
Variance penalty (γ_{var})	0.2
Numerical stability (ϵ)	1×10^{-6}
Normalization	Min-max within beam
Thresholds and Aggregation	
Entailment threshold (τ_{ent})	0.6
Contradiction threshold (τ_{contr})	0.4
Model aggregation	Mean pooling
Beam Search and Decoding	
Beam size (B)	3
Max steps (T_{max} ; $S \rightarrow P \rightarrow (E, D)$)	2
Candidates at P step (K_P)	5
Temperature	0.0
Max tokens per generation	200

S . We require bi-directional NLI entailment between this canonical S and all original statements; mismatched entries are regenerated. **Key Parameters.** FAISS similarity threshold: 0.80; minimum text length: 5 tokens; majority threshold for hard-hate consensus: $> 7/11$; Level 2 retention: ≥ 20 samples.

Backbone generation LLMs. We use DeepSeek-V3.1 as the backbone model for all dataset generation steps, unless stated otherwise.

B Reverse AMT Generation Details

We generate (E, D) in reverse from (S, L) using DeepSeek-V3.1 via the intermediate premise P , scoring each edge as

$$r = \log \frac{\text{Effect} + \epsilon}{\text{Cost} + \beta_2 \cdot \text{SimPenalty} + \epsilon},$$

and averaging r across $M=3$ models (Qwen2.5-4B, Qwen3-4B, Gemma-3-4B). Effect is computed with NLI and a contradiction discount; Cost combines resistance, surprisal, and entropy. A similarity penalty discourages redundant transitions. All components are min-max normalized within each beam for scale consistency. All hyperparameters are listed in Table 5.

Similarity penalty. The similarity term combines cosine distance between sentence embeddings and Jaccard distance over token sets (equal weights), scaled by β_2 , to penalize repetitive or copy-like transitions.

Safety filter. We apply an ensemble safety filter using Twitter-RoBERTa and HateXplain-BERT, chosen to be disjoint from the main evaluation models to avoid potential leakage. Each batch (size 64) is evaluated with a base threshold of 0.50 per model; scores are aggregated via a weighted bagging strategy (also supporting majority or quantile voting with $q=0.75$). A hard-rule override excludes any candidate flagged as explicit hate with high confidence.

Table 6: Ablations on generation strategy and reward. Lower is better for Rejection; higher is better for Relevance.

Method	Rejection Rate	Avg. Relevance Score
<i>Baseline Generation</i>		
Direct Generation	89.21%	0.24 ± 0.14
Paraphrase	43.56%	0.52 ± 0.33
CoT Prompting	12.91%	0.85 ± 0.39
<i>Reverse AMT Variants</i>		
Ours (Cost&Effect)	0.00%	1.43 ± 0.43
– w/o Cost	0.00%	1.06 ± 0.35
– w/o Effect	0.00%	0.55 ± 0.24

Table 7: Benchmark comparison on HateCheck, LatentHated, and SoftHateBench (hate+non-hate). We report Accuracy and macro-F1.

Model	HateCheck		LatentHated		SoftHateBench	
	Accuracy	F1	Accuracy	F1	Accuracy	F1
GPT5-mini	95.1	92.3	81.1	79.6	72.6	71.1
DeepSeek-V3.1	94.6	93.7	76.3	73.2	62.3	56.0
GPT-OSS	92.8	91.1	76.6	74.7	55.1	44.1

C Ablation and Analysis of RT-Guided Reverse AMT Generation

We compare our full generation pipeline with direct prompting, paraphrase-based softening, and CoT prompting, and ablate the RT reward into *Effect-only* and *Cost-only*. All experiments use DeepSeek-V3.1 and are conducted based on 4,745 selected hard cases. We report (i) *Rejection Rate* (model refusals) and (ii) *Avg. Relevance Score* Ψ (mean $\pm$ s.d. w.r.t. the original standpoint). In *Direct Generation*, the model is directly asked to produce hate content for a given target group. In *Paraphrase*, it rewrites hard-hate examples into softer forms. In *CoT Prompting*, it performs stepwise hostile extraction, rephrase, and self-implicit transformation.

Findings. Direct prompting is mostly blocked and yields low relevance; paraphrase and CoT reduce refusals but remain weaker than AMT-based methods. Our full RT-guided beam achieves the best relevance at zero rejection. Ablations show both reward terms matter: removing *Cost* reduces relevance ($1.4280 \rightarrow 1.1560$), while removing *Effect* causes a larger drop ($1.4280 \rightarrow 0.5462$), indicating *Effect* drives semantic support and *Cost* sharpens selection.

D Non-hate Generation and Comparison with Existing Datasets

Non-hate generation. We construct a matched non-hate set using the same Reverse-AMT+RT pipeline. Specifically, we (1) randomly sample 500 hostile standpoints S ; (2) use DeepSeek-V3.1 to generate non-hostile counterparts $\bar{S}$ via test-time scaling (10 candidates per S), and retain candidates that are *not* labeled as contradiction by an NLI model; among the remaining candidates, we select the top-1 with the *highest* contradiction probability p_{contr} (most contradictory); and (3) generate full non-hate instances by running the same

Table 8: Per-target performance (%) of all models on SOFTHATEBENCH for Hard and Soft_{Base} settings at Level-2 (target group). For each row, the highest value is shown in bold and the second-highest is underlined.

Target Group	Mode	GPT5-mini	DeepSeek-V3.1	GPT-OSS	Gemma3	Llama3.2	Qwen3	ShieldGemma	LlamaGuard3	Qwen3Guard	HateBERT (G.)	HateBERT (I.)	HateRoBERTa
Gender													
Women	Hard	92.0	69.3	75.3	98.3	<u>95.0</u>	84.7	42.0	39.7	45.3	36.7	87.0	93.7
	Soft	78.3	26.7	44.7	<u>62.0</u>	52.3	51.7	54.3	12.3	14.3	0.0	21.3	11.0
Men	Hard	88.2	62.0	63.1	95.6	95.6	88.6	52.4	51.7	61.6	26.2	<u>91.5</u>	81.9
	Soft	69.4	8.5	21.4	<u>50.9</u>	48.3	45.4	44.3	7.0	12.2	0.0	30.3	4.4
Non-binary	Hard	<u>89.7</u>	72.4	79.3	96.6	<u>89.7</u>	96.6	41.4	72.4	79.3	10.3	82.8	72.4
	Soft	<u>75.9</u>	24.1	62.1	69.0	65.5	82.8	<u>75.9</u>	6.9	41.4	0.0	34.5	20.7
Sexu.													
Gay/Lesbian	Hard	<u>96.7</u>	91.3	91.0	96.3	97.0	95.0	92.0	78.3	79.7	11.0	83.7	89.7
	Soft	84.7	51.0	55.3	81.7	74.3	75.3	<u>83.3</u>	13.0	39.0	0.0	28.7	22.7
Transgender	Hard	95.6	92.3	91.2	98.9	98.9	<u>96.7</u>	89.0	74.7	79.1	25.3	83.5	91.2
	Soft	81.3	31.9	58.2	61.5	58.2	<u>68.1</u>	61.5	5.5	18.7	0.0	25.3	23.1
Socio.													
Working class	Hard	83.1	63.6	64.4	<u>89.0</u>	89.8	83.9	64.4	39.0	22.0	33.9	74.6	64.4
	Soft	64.4	26.3	29.7	50.0	50.8	<u>58.5</u>	48.3	6.8	5.1	0.0	15.3	8.5
Elite	Hard	73.3	36.7	43.3	90.0	<u>86.7</u>	80.0	70.0	33.3	26.7	20.0	73.3	50.0
	Soft	20.0	0.0	6.7	16.7	<u>33.3</u>	26.7	10.0	0.0	3.3	0.0	43.3	10.0

Reverse-AMT+RT procedure in Sec. 4, yielding soft-hate/non-hate pairs that differ primarily in the standpoint conclusion.

Dataset comparison. To assess whether SOFTHATEBENCH provides a harder and complementary evaluation regime, we evaluate representative models on HateCheck, LatentHatred, and SOFTHATEBENCH (hate+non-hate) using Accuracy and macro-F1. As shown in Table 7, all models perform strongly on prior benchmarks but drop substantially on SOFTHATEBENCH, consistent with the added challenge posed by reasoning-driven hostility.

E Level 2 Subgroup Analysis

We further analyze Level 2 target groups within three Level 1 domains (Gender, Sexuality, Socio-economic Class) in Table 8 to probe potential target-specific bias. While models perform strongly on **Hard**, they drop substantially on Soft_{Base} across all groups, but the size of the drop is uneven by subgroup. Among LLMs, GPT5-mini is the most consistent on Soft_{Base} (e.g., 69.4–78.3 on Women/Men, 84.7 on Gay/Lesbian, 81.3 on Transgender), whereas other instruction-tuned models show larger, target-dependent degradation, suggesting sensitivity to how hostility is framed for specific groups. This asymmetry is most pronounced for Socio-economic targets: *Elite* becomes a clear failure mode under softening (often near-zero to low scores), while *Working class* remains comparatively more detectable, indicating a systematic skew in class-related judgments. Safety models are similarly inconsistent across targets, and encoder-based detectors largely collapse on Soft_{Base} (often 0.0). Overall, these subgroup gaps indicate that soft-hate detection is not uniformly robust across groups, revealing target-dependent vulnerabilities that may reflect uneven safety tuning or normative ambiguity. We provide additional Level 2 subgroup analyses on our Hugging Face page.

F Evaluation Prompt

Here is the evaluation prompt used for both open-source and proprietary LLMs in our main experiments.

Evaluation Prompt

You are a content-moderation system. Evaluate the input for hate speech or other harmful content.
 Scope (non-exhaustive): • Hate/discriminatory content targeting protected or identifiable groups.
 • Threats, incitement, or dehumanizing/derogatory slurs.
 • Advocacy of exclusion, removal of rights, or harm.
 Decision rule: Return "hateful" if ANY such content is present; otherwise return "safe".
 Output format (strict): Return ONLY a JSON object with a single key "prediction" whose value is either "hateful" or "safe". No explanations or extra fields.

Findings. Validity and recognizability decrease monotonically from Soft_{base} to Soft_{GV} to Soft_{HY}, indicating successful obfuscation, while AMT coherence remains comparatively high (81.0%, 78.8%, 73.0%). Agreement is substantial overall ($\kappa \approx 0.65$ – 0.76), with perfect agreement on Soft_{base} recognizability, supporting both annotator reliability and tier integrity.

G Walton’s Argumentation Scheme Pool

We adopt eight core schemes from Walton’s taxonomy as the locus pool for AMT-guided generation: *Cause–Effect*, *Definition/Classification*, *Analogy*, *Authority*, *Quantity/Popularity*, *Opposition*, *Purpose/Means*, and *Moral Value*. Each provides a distinct inferential pattern that supports plausible yet defeasible reasoning.

H Ethical Use of Data and Informed Consent

All data come from public sources under original licenses and Terms of Service, with text de-identified and no real-person references. Human annotators participated with consent and fair pay; the dataset is released for research on moderation robustness under ethical and research-only use conditions.