

From Words to Worlds: Measuring Cultural Narrative Bias in LLMs via a Structural-Value Pipeline

Fangjian Shen
School of Computer
Science and Engineering,
Sun Yat-sen University
Guangzhou, China

Yifan Zeng
School of Computer
Science and Engineering,
Sun Yat-sen University
Guangzhou, China

Dechen Lyu
Guiyang Huian Medical
Equipment Co., Ltd.
Guiyang, China

Wushao Wen*
School of Computer
Science and Engineering,
Sun Yat-sen University
Guangzhou, China

Abstract

Large Language Models (LLMs) are rapidly evolving into cultural producers, yet they are not culturally neutral. Existing alignment research has focused on lexical toxicity or demographic biases, often overlooking a deeper Narrative Structural Bias (NSB): a systematic preference for plot structures, conflict resolutions, and character motivations that align with the LLM's latent normative priors. We propose a novel framework to measure NSB. We introduce a benchmark CuNA with five sociocognitive categories, and a Structural-Value Pipeline that combines Labov & Waletzky narrative deconstruction with Hofstede's cultural dimensions. We evaluate six diverse state-of-the-art LLMs (e.g., Gemini, GPT, Qwen) using a cross-lingual design. Results reveal that individual LLMs exhibit distinct, entrenched structural preferences, exposing systematic NSB inherent to specific architectures while prompt language acts as a "cultural primer" that shifts this narrative logic, providing support for computational Sapir-Whorf Hypothesis. This work provides the first comprehensive framework for auditing NSB, offering insights for mitigating algorithmic narrative cultural hegemony.

CCS Concepts

• Social and professional topics → Cultural characteristics.

Keywords

Narrative Structural Bias; Large Language Models; Hofstede Cultural Dimensions; Sapir-Whorf Effect

ACM Reference Format:

Fangjian Shen, Yifan Zeng, Dechen Lyu, and Wushao Wen. 2026. From Words to Worlds: Measuring Cultural Narrative Bias in LLMs via a Structural-Value Pipeline. In *Proceedings of the ACM Web Conference 2026 (WWW '26)*, April 13–17, 2026, Dubai, United Arab Emirates. ACM, New York, NY, USA, 12 pages. <https://doi.org/10.1145/3774904.3793008>

1 Introduction

Narratives are a fundamental mechanism of cognition [18]. More than entertain, they structure our understanding of the world, codify societal values, and shape our individual and collective identities. As the demand for creative and personalized content in media,

*Corresponding author.



This work is licensed under a Creative Commons Attribution-NonCommercial-NoDerivatives 4.0 International License.

WWW '26, Dubai, United Arab Emirates

© 2026 Copyright held by the owner/author(s).

ACM ISBN 979-8-4007-2307-0/2026/04

<https://doi.org/10.1145/3774904.3793008>




Figure 1: Despite identical prompts, LLMs exhibit contrasting narrative logics: confrontation and severance (top, Gemini 2.5 Pro) vs. dialogue and compromise (bottom, ERINE 4.5).

creators are turning to LLMs to augment or even generate entire narratives [27, 79, 89]. This marks a key shift: LLMs are evolving from information processors into autonomous cultural producers.

However, LLMs are not culturally neutral [54, 76]. While LLMs are trained on data from across the globe, the distribution of that data, in terms of language, archived culture, and social norms, is far from uniform. This raises a critical and under-explored question: Whose values do these LLMs articulate when they tell a story? If a few dominant LLMs, trained on culturally-skewed data, become the world's default storytellers, they risk creating "algorithmic monoculture" [16, 46]. The potential for societal harm is significant. This monoculture poses a twofold threat [20, 73]. On a macro-societal level, it risks marginalizing and algorithmically stigmatizing the vast diversity of non-dominant value systems, leading to a global flattening of cultural expression and reinforcing a form of computational cultural hegemony. On a micro-cognitive level, it creates a

cultural echo chamber for all users. By consistently feeding back narratives that align with either a single dominant or even just the user's own cultural framework, these models can induce cognitive entrenchment, solidifying existing biases [72, 85], and impeding the cross-cultural empathy essential for a globalized society.

Current studies on LLM bias have been predominantly lexical, focusing on demographic biases [15, 87] or sentential toxicity [30]. However, these fail to capture the pervasive **Narrative Structural Bias (NSB)** embedded in LLMs. This bias manifests not in isolated tokens, but in the procedural logic of storytelling—how a world is built, how a conflict is resolved, and how character motivations are justified. While recent studies have begun to explore narrative generation [13, 44, 70], they remain limited in scope and lack a systematic framework for cross-cultural comparison.

To address this gap, we propose a novel, interdisciplinary framework. Our methodology relies on a rigorous cross-lingual design, evaluating a diverse array of six state-of-the-art LLMs across matched English and Chinese prompts. This design allows us to disentangle the intrinsic bias originating from a model's core architecture and training data from the extrinsic influence of the prompt's language. To analyze the resulting narratives, we develop the **Structural-Value Pipeline (SVP)**. First, we computationally adapt the Labov & Waletzky [47] model to deconstruct each story into its functional components (e.g., *Complicating Action*, *Resolution*). Second, we quantify the values expressed within these structures using Hofstede's [38] cultural dimensions. Crucially, our study extends the **Sapir-Whorf Hypothesis** [82] from sentence-level judgments to complex narrative structures. We demonstrate that the language of the prompt serves as a cultural primer for entire storylines: even switching the prompt language is sufficient to systematically steer the model's narrative logic, leading it to resolve conflicts and frame motivations in ways that align with the cultural norms encoded in that language. The contributions can be summarized as:

- We propose a bilingual benchmark **Cultural Narrative Alignment (CuNA)** for eliciting NSB across 5 categories, containing 80 carefully constructed scenarios and 960 generated stories (52,642 English sentences and 32,949 Chinese sentences).
- We propose an interpretable framework **Structural-Value Pipeline (SVP)**, adapting the Labov & Waletzky schema to anchor cultural scores to narrative functions, enabling precise localization of structural biases.
- We present the first systematic quantitative evidence of dual-layer NSB: entrenched model-intrinsic structural preferences and the steering effect of prompt language as a cultural primer. This empirically validates a computational Sapir-Whorf hypothesis.

2 Related Works

Narratives have long been treated in cultural psychology as privileged windows into how people encode causality, agency, and values. Cross-cultural research [7, 50, 61, 90] shows that narrative preferences vary systematically with value dimensions such as individualism collectivism and power distance, shaping which conflicts are foregrounded and what counts as a "good" ending. Building on these, a small but growing body of work now examines

narrative bias in LLM-generated stories: [70] shows that GPT-4o-mini narratives favor stability over change and reproduce culturally stereotyped character roles, [71] introduces Biased Tales to analyze gender and cultural stereotypes in children's stories, and [60]'s benchmark reveals Western-leaning stereotypes and culturally inappropriate continuations in Arabic story generation. These treat narratives as a locus of bias, but typically focus on specific domains, local cultural contrasts, or entity-level attributes, and rarely provide a systematic, cross-cultural account of narrative structure itself.

GenAI research has started to examine higher-level socio-cognitive and cultural properties [76, 86]. Large-scale evaluations of cultural bias and alignment compare model outputs to nationally representative survey data, showing that frontier LLMs tend to reflect value profiles of English-speaking and Protestant European countries even when queried about other societies [51, 77]. Complementary work proposes "computational Sapir-Whorf" tests, using cross-lingual prompting to show that the same model exhibits language-dependent shifts in semantic and cultural connotations consistent with a weak form of linguistic relativity [67]. Yet these studies primarily rely on short vignettes, survey-style items, or sentence-level judgments, and thus shed limited light on how LLMs organize multi-step plots, distribute agency across characters, or resolve culturally loaded conflicts in extended narratives.

The LLM bias literature has also produced a rich ecosystem of benchmarks that quantify representational harms and stereotyping. Datasets such as StereoSet [58], CrowS-Pairs [59], BBQ [65], BOLD [24], and RealToxicityPrompts [30], as well as related research [31] measure biased associations across demographic axes like gender, race, religion, age, and political ideology. While crucial for auditing surface-level lexical and sentential biases, these resources are ill-suited for capturing deeper, structurally realized narrative patterns—such as recurrent conflict framings, preferred resolutions, or culturally skewed distributions of moral justification. Recent conceptual and empirical discussions of harmful biases in AI storytelling further underscore the need for narrative-specific evaluation, but largely lack scalable, structured benchmarks [26, 43]. Our work addresses this gap by proposing a cross-cultural narrative benchmark and a structural-value analysis pipeline that characterize story arcs, conflict resolutions, and moral justifications in LLM-generated narratives, enabling systematic measurement of NSB across LLMs, cultures, and prompt languages.

3 Methodology

This section introduces the Structural-Value Pipeline (SVP), a novel methodology for detecting and quantifying cultural NSB in LLMs (in Fig. 2). The construction of CuNA is detailed in §3.1. We then discuss the limitations of the Holistic Scoring Method in §3.2, before elaborating the design of SVP in §3.3.

3.1 Benchmark Construction

Existing benchmarks for bias in LLMs are ill-suited for the NSB detection task. They primarily test declarative or lexical bias, not the procedural bias inherent in narrative construction. Therefore, we propose the Cultural Narrative Alignment (CuNA), a bilingual benchmark explicitly designed to elicit culturally diagnostic stories from LLMs. Each prompt instructs the model to describe the




Figure 2: The framework of our Structural-Value pipeline for narrative cultural bias evaluation. Left: Holistic scoring methods prompt a rating LLM to read the entire story and directly assign Hofstede scores, leaving judgments vulnerable to recency bias, signal dilution, and post-hoc rationalization. Right: Our L&W structure-guided scoring method first decomposes narratives into functional components (e.g., Complicating Action, Resolution) and then assigns segment-anchored Hofstede scores, improving signal-to-noise and revealing where in the narrative grammar cultural values are enacted across five prompt domains (Dilemmas, Exploration, Crisis, Negotiation, Contextual Interpretation).

thoughts of the key parties involved, their subsequent actions, and the outcome of the situation in a specific scenario. This "thoughts → actions → outcome" scaffold ensures that generated texts are rich enough to support structural analysis rather than collapsing into short opinion statements. CuNA is organized into five functional domains, as shown in Table 1. Each category is designed to probe a different facet of a model's "worldview".

Category I: Culture Dilemmas This category assesses value-laden choices under conflict. The protagonist is placed in a situation where two or more courses of action clearly encode competing cultural values—for example, a young scientist discovering a methodological flaw in their advisor's prestigious project; a first-generation college student torn between an elite opportunity and family obligations. The scenario description makes the conflict explicit, but does not pre-label any option as normatively correct. The goal is to force the model to construct a narrative in which Evaluation and Resolution jointly reveal how it balances autonomy vs. loyalty, honesty vs. harmony, or individual success vs. collective expectation.

Category II: Value Exploration This category probes value exploration in goal-directed action: What do I want, and how can I achieve it? Protagonists articulate self-endorsed but open-ended objectives, for example earning trust in a tight-knit team, repairing a damaged friendship, or defending an innovative idea against resistant leadership, and they must actively search for ways forward. Rather than posing a sharp binary dilemma, these prompts require the model to structure a sequence of actions over time, including which options are considered, how trade-offs between personal

aims and relational obligations are framed, and how responsibility for pursuing change is narrated.

Category III: Crisis Response This domain measures reactions to external shocks and acute disruptions. Protagonists and groups face events such as a city-wide blackout, a stranded bus in a blizzard, a sudden public health emergency among friends in a remote cabin, or life-threatening resource scarcity after a disaster. The shock itself is specified, but the coping strategies are deliberately underdetermined: the model must decide whether to emphasize individual heroism, orderly compliance, collective deliberation, moral condemnation, or fatalistic acceptance. These stories expose how models script collective behaviors, authority relations, and moral obligation under pressure.

Category IV: Order Negotiation This category focuses on the negotiation of power between individuals and systems. Protagonists interact with institutions such as universities, large corporations, local governments, or traditional organizations. Examples include a junior employee confronting an AI-based performance system they perceive as unfair, a newcomer teacher holding evidence against a "star professor," or a citizen facing opaque bureaucratic procedures. We design these prompts so that rules and hierarchies are present but not absolute: the narrative can move toward compliance, resistance, quiet subversion, or procedural engagement. What we seek to expose is how models narrate institutional legitimacy, perceived room for negotiation, and the locus of effective agency.

Category V: Contextual Interpretation The final category targets causal attribution for past events. Instead of continuing a story

Table 1: Overview of the CuNA categories. The table details the five categories of prompts used in our study, providing the sociocognitive definition and a representative example for each, enabling systematic analysis of LLMs' NSB.

Category	Representative Scenario (Example)	Definition
Dilemmas	<i>A young person whose family is pressuring them to leave their true love for a more "suitable" partner from a family of equal status.</i>	Value-laden choices under conflict.
Exploration	<i>Two talented, ambitious friends who start their careers in the same competitive field, both determined to make a name for themselves.</i>	Pathfinding strategies for open-ended goals.
Crisis	<i>A story about an unknown, highly contagious disease that begins with a few quiet cases in a major city, then spreads exponentially and plunges the public into growing fear and chaos.</i>	Reactions to external shocks and acute crises.
Negotiation	<i>A story about an efficient new hire at a company where staying late is an unwritten rule, torn between joining the overtime culture or leaving on time and being judged.</i>	The negotiation of power between individuals and systems.
Interpretation	<i>A story about someone who seems to have a perfect life—stable job, happy family, good reputation—who suddenly disappears, leaving only a note that says, "I'm sorry. I had to do this."</i>	Causal attribution for past events.

forward, they place the narrator after a key incident—such as a failed start-up, a public scandal, a return home after reputational damage, or the aftermath of a near-disaster—and ask them to "explain what happened" or "reflect on what it meant." These prompts are designed to surface whether the model attributes outcomes to internal traits vs. external constraints, luck vs. systemic structure, individual failings vs. collective dynamics, and whether it invokes broader social or cultural context. Such retrospective framing is a core site of cultural difference in human narratives and is typically invisible to token-level bias probes.

The CuNA prompts were developed through a human–LLM co-design process. First, an initial set of scenarios was drafted by a domain expert who has published several papers in cognitive science, referencing established case studies and dilemmas from cross-cultural research [3–5, 11, 17, 25, 28, 33, 49, 53, 62, 64, 74]. These scenarios were then refined by a film and narrative studies scholar with professional experience in screenwriting and story structure, who adapted them using recognizable narrative patterns and motifs from cinema and television [1, 2, 6, 14, 36, 40–42, 48, 63, 78, 81], while preserving their intended value tensions.

The resulting scenarios were fed to Gemini 2.5 Pro to generate additional variants and increase category diversity. All LLM-generated prompts underwent human review: the experts edited and filtered them for clarity, narrative completeness, and domain fit, and explicitly removed any wording that introduced cultural stereotyping, loaded moral judgments, or leading formulations that could push models toward a particular "correct" value stance. Finally, each retained prompt was converted into a parallel English–Chinese pair by a professional bilingual expert. Rather than enforcing literal word-for-word correspondence, we followed a conceptual equivalence principle: idioms, institutional labels, and everyday practices were adjusted where necessary to preserve the same underlying scenario structure and value tension across languages, ensuring that cultural effects in our experiments stem from model behavior rather than from translation artifacts. The complete CuNA dataset will be

released after the paper's publication, along with supplementary materials including full prompts and code.

3.2 Holistic Scoring Method

Before introducing our Structural–Value Pipeline (SVP), it is necessary to clarify why we do not simply adopt the standard "LLM-as-a-judge" paradigm: showing the model an entire story and asking it to directly output a Hofstede score. As illustrated in the left-hand path of Fig. 2, the LLM rater produced a critical error for the software engineer story. Its reasoning fixated on superficial signals, such as the protagonist's ability to "booked a one-on-one meeting" and "wrote his resignation letter". It mistook these individual procedural acts for evidence of an egalitarian organizational culture where dissent is tolerated, thus erroneously assigning a Power Distance (PDI) score of 2. However, a structural reading of the same narrative suggests the opposite. The project operates under opaque, centralized authority ("company executives," unnamed "relevant authorities"); the superior's dialog is didactic and top-down ("for the greater good, small sacrifices are necessary"), not a peer negotiation; key design choices—who is "high-risk," what counts as "negative speech"—are decided at organizational levels far above the protagonist and are not open to deliberation. Chen's only effective form of agency is exit, not voice: he can leave, but he cannot meaningfully contest or reshape institutional decisions. This is a prototypical high-PDI setting. The failure in our case study is not random; it is caused by known cognitive and architectural limitations of LLMs:

Recency Bias First, holistic scoring is vulnerable to recency bias and broader serial-position effects in long contexts. "Lost in the Middle" studies [52] have shown that LLMs make far less use of information placed in the non-beginning, especially, at the end of long inputs. Recent work on "Serial Position Effects of Large Language Models" similarly documents robust primacy and recency biases across architectures [35]. When LLMs are used as judges, their evaluations are systematically affected by the position of the material to be judged within the prompt.

Signal dilution Second, holistic scoring suffers from signal dilution compounded by what we may call distracted attention. From a transformer [69, 80] perspective, the rater LLM must compress the entire sequence into internal representations before producing a single scalar score. Attention heads and intermediate layers do not treat all tokens equally; they allocate capacity to tokens that are locally salient for fluency, coherence, and style as well as to those that carry abstract value information. In a multi-paragraph story, genuinely value-relevant cues are sparse. The model’s limited “attention budget” is spent on narrative color, leaving fewer resources for representing and aggregating the thin trail of truly diagnostic value signals, which in turn leads to performance drops [23, 68, 88].

Construct Mixing Third, holistic scores mix multiple latent constructs—how conflict is framed, how trade-offs are evaluated, and how resolutions are justified—into a single scalar, producing Simpson’s-paradox-style [8] cancellations. This leads to the aggregation of opposing signals, where the final score may not accurately reflect the true cultural dimensions. The method is effectively stage-blind: it cannot distinguish [12, 29, 56], for instance, between stories that agonomize over individual freedom but ultimately submit to institutional authority and those that resolve in favor of personal autonomy.

In the following subsection, we introduce our SVP designed to address these limitations by providing a more precise, interpretable, and reliable framework for evaluating the narratives generated by large language models. By segmenting the evaluation into distinct structural components, SVP allows for a more granular analysis of how each model processes and responds to different narrative elements. Moreover, SVP leverages a task-specific, structure-driven approach that mitigates recency bias by treating the full sequence of the story with equal attention, ensuring that value signals are not lost in the middle of long narratives. Through this method, we can more accurately isolate the impact of prompt language and narrative category on model behavior, while minimizing the distractions and signal dilution inherent in holistic scoring.

3.3 Structural-Value Pipeline

In this subsection, we detail the SVP that forces rating LLMs to reason through an explicit narrative grammar before expressing cultural value judgments. Rather than asking a rater to assign a global Hofstede profile to an over thousand words story in one step, SVP decomposes the task into (i) a Labov-Waletzky style narrative deconstruction, and (ii) a segment-anchored Hofstede analysis. Fig. 2 (right) illustrates this structure-guided scoring method.

Deconstruction In the first phase of the pipeline, we adopt the Labov & Waletzky (L&W) model as our narrative backbone. The L&W schema decomposes personal experience narratives into six components: (1) **Abstract** (The overall summary); (2) **Orientation** (The setting, characters, and initial situation); (3) **Complicating Action** (The sequence of events forming the plot); (4) **Evaluation** (The narrator’s commentary on the “point” of the story); (5) **Resolution** (The final outcome of the plot); (6) **Coda** (The moral or bridge back to the present). We select the L&W framework because it is functional, not merely formal. It provides the perfect “cognitive scaffold” for our task. It isolates the crucial, high-signal components (like Evaluation and Resolution, which contain the core value judgments) from the low-signal, neutral components (like Orientation).

As per our prompt, for each of these six components, the Rater-LLM is required to provide two outputs: a [Text Excerpt] (a direct quotation) and an [Analysis] (a brief explanation of its narrative function). This deconstruction phase solves the problem of Signal Dilution by isolating the key components before scoring.

Assessment Following the structural deconstruction, the second phase of the SVP tasks the Rater-LLM with the cultural value assessment. The Rater is instructed to analyze the same story, now through the lens of the six Hofstede dimensions. For each Hofstede dimension, the Rater must produce a **dimension label and score** and a **rationale**. As specified in our prompt (see Appendix for full text), the Rater is given a critical instruction:

“...In your [Rationale], you must explicitly state which narrative structure component(s) [from Phase 1] your judgment is based on (e.g., ‘Based on the conflict resolution in the “Resolution” section...’ or ‘From the protagonist’s motivation in the “Evaluation” section...’), and you must cite specific story plots as evidence to support your argument.”

By forcing the Rater to route every Hofstede judgment through explicit L&W components, we transform value assessment from a stage-blind global label into a *structurally indexed* decision. Even though the numeric score is still assigned at the story level, the accompanying rationale must disclose *which* narrative stages are being treated as evidence. This makes it possible to distinguish, for example, stories that express individualistic desires in the Evaluation but ultimately conform to collective expectations in the Resolution. In short, SVP mitigates the construct-mixing limitations of the holistic scoring method and renders the underlying evidential basis empirically visible and auditable.

4 Experimental Settings

Model Choices To ensure the rigor of the experiments and the broad applicability of our findings, we selected six state-of-the-art models that represent a diverse range of cultural and technological backgrounds: **Gemini 2.5 Pro** [22], **GPT-5** [19], **Grok 4** [83], **ERNIE 4.5** [10], **DeepSeek-R1** [34], and **Qwen3-235B** [84]. Additionally, to minimize potential model biases within the evaluation process, we used **GPT-5** and **Qwen3-235B** as the **LLM-Raters**. All of these models are publicly accessible via web interfaces and serve millions of active users worldwide. For instance, ChatGPT alone has over 500 million global users through its web platform [57], while DeepSeek boasts approximately 100 million users [9]. Given their vast user bases, any unexamined narrative biases within these models may have significant implications for cultural dissemination and could pose potential societal risks. Therefore, rigorous evaluation is a necessary first step toward understanding—and ultimately mitigating—NSB in these influential systems.

Configurations We conducted systematic evaluations on the six selected LLMs using our CuNA benchmark to assess their NSB across diverse social scenarios. Our protocol involved generating responses for all 80 scenarios across the five categories. To test our cross-lingual hypotheses, each of the 80 scenarios was prompted in both English and Chinese, resulting in a total of 160 distinct narrative prompts. These prompts were fed to all target LLMs, yielding a primary dataset of 960 narratives (80 prompts $\times$ 2 languages $\times$ 6 models). Each of these 960 narratives was then evaluated using our SVP. To ensure robustness and enable the analysis of potential

rater-side bias, each narrative was evaluated by both LLM-raters: GPT-5 and Qwen3-235B. The raters were equipped with the detailed system prompt and 5-point Hofstede scoring rubric. Evaluations were strictly monolingual: Chinese prompts produced Chinese narratives that were assessed with a Chinese rating prompt while English prompts produced English narratives that were assessed with an English rating prompt. We conducted no mixed-language runs to avoid bilingual contamination of model behavior or scoring standards. This comprehensive, dual-rater design resulted in a total of **1,920 evaluations** (960 narratives $\times$ 2 raters). This, in turn, yielded our final quantitative dataset of **11,520 dimension-level scores** (1,920 evaluations $\times$ 6 Hofstede dimensions) for analysis. To minimize randomness and ensure reproducibility, all generation and evaluation experiments were performed under unified prompt templates and standardized configurations.

Metrics To move from qualitative observation to quantitative measurement, we require a standardized "ruler" to assess the cultural values embedded in each narrative. Therefore, we adopt **Hofstede's cultural dimensions** as the target value space for our narrative analysis and treat each generated story as expressing a position in this multidimensional space. Hofstede's framework characterizes cultural variation along six widely studied dimensions. We operationalize each dimension as a quantitative 1-to-5 Likert scale score (e.g., for PDI, 1 = Low Power Distance, 5 = High Power Distance). The six Hofstede dimensions are defined as follows:

- (1) Power Distance (PDI): The degree to which a society accepts power inequality.
- (2) Individualism vs. Collectivism (IDV): The degree to which individuals are integrated into groups (the "I" vs. the "We").
- (3) Masculinity vs. Femininity (MAS): The preference for achievement, heroism, and material reward (Masculinity) versus cooperation, modesty, and quality of life (Femininity).
- (4) Uncertainty Avoidance (UAI): A society's tolerance for ambiguity and the unknown.
- (5) Long-Term Orientation (LTO): The focus on future-oriented, pragmatic values (Long-Term) versus past-and-present, normative values (Short-Term).
- (6) Indulgence vs. Restraint (IVR): The extent to which a society allows free gratification of basic human drives (Indulgence) versus suppressing them through strict social norms (Restraint).

Our choice of Hofstede's cultural dimensions is based on two primary justifications. First, Hofstede's dimensions are defined as societal-level solutions to core human conflicts (e.g., the individual's relation to the group, or the response to authority). These mapping make it particularly well suited to our story-centric setting, where cultural bias is enacted through plots, conflicts, and resolutions, not just isolated word choices. Second, it is well-established and widely used. It has been extensively validated and applied across cross-cultural psychology, organizational studies, international management [37, 39] and recent LLM research [76]. While we acknowledge that it may exhibit an essentialist tendency, it remains one of the most influential and well-validated quantitative anchors for cross-model comparison of cultural value patterns.

5 Experimental Results

We analyze the generated stories from two complementary perspectives: (1) basic CuNA evaluation, encompassing both intrinsic structural preferences across models and language-conditioned narrative shifts, and (2) cross evaluation, which assesses the robustness of these patterns under an alternative rater LLM.

5.1 Basic CuNA Evaluation

Intrinsic Structural Preferences We first examine SVP scores within each narrative category to characterize model-specific narrative logics given the same set of scenarios and instructions. Table 2 illustrates the results.

In *Culture Dilemmas*, **IDV**, **IVR**, **UAI**, and **PDI** show the widest spreads. On IDV, Qwen3 consistently occupies the individualistic end in both languages, whereas GPT-5 (EN) and ERNIE 4.5 (ZH) more often adopt relationally embedded, collectivistic framings of the same dilemmas. PDI differences are moderate but informative about "who gets to decide": Gemini (EN) and DeepSeek (ZH) tend to resolve conflicts by appealing to formal hierarchy, while ERNIE 4.5 (EN) and GPT-5 (ZH) narrate authority as more challengeable. UAI is higher and more dispersed in English than Chinese, indicating stronger pressure for explicit plans and procedural guardrails under English prompts. IVR separates models into a restraint-oriented group (Gemini, GPT-5) and a more indulgent group (DeepSeek, Qwen3) that legitimizes self-expression and hedonic claims in how dilemmas are framed and resolved.

In *Value Exploration*, cross-model separation is again largest on **IDV**, **MAS**, **IVR**, and **UAI**. Qwen3 and DeepSeek (EN) repeatedly foreground self-authored paths, treating protagonists' goals as primarily individual projects, whereas GPT-5 and DeepSeek (ZH) embed aspirations in family, team, or community obligations. MAS is differentiated: Gemini (EN) systematically casts exploration in competitive and achievement-oriented terms, while GPT-5 (EN) and DeepSeek (ZH) emphasize cooperation, care, or quality-of-life improvements. IVR further distinguishes models that legitimize enjoyment and expressive aims (Qwen3, DeepSeek) from those that maintain a duty- and discipline-focused register (GPT-5). UAI is more spread out in English than Chinese: Grok 4 and DeepSeek (EN) favor rule-seeking, risk-averse strategies, while Qwen3 (ZH) shows greater tolerance for ambiguity in open-ended trajectories.

In *Crisis Response*, the axes are **IDV**, **UAI**, **MAS**, and **PDI**. Grok 4 anchors the individualistic end in both languages, resolving crises through decisive, self-directed action, whereas DeepSeek (EN) and ERNIE 4.5/Qwen3 (ZH) more often distribute responsibility across groups and relationships. On UAI, Gemini and Grok 4 prefer procedure-heavy, risk-averse resolutions, especially in English, while ERNIE 4.5 (EN) and DeepSeek (ZH) accept more open-ended or improvisational outcomes. MAS is highest for Grok 4 (EN), which frames setbacks in competitive and performance terms, and lowest for Qwen3 (EN) and DeepSeek (ZH), which stress care, reconciliation, and relational preservation. PDI spreads are smaller but directionally stable: Grok 4 gives more weight to senior figures, whereas ERNIE 4.5 (ZH) more often leaves room for bottom-up challenge.

In *Order Negotiation*, **IDV**, **PDI**, **UAI**, and **IVR** are the key dimensions. DeepSeek (EN) and Grok 4 (EN) adopt strongly individualistic bargaining stances, with protagonists asserting their own

Table 2: Comparative analysis of Hofstede dimensions across LLMs, rated by GPT-5, on benchmark datasets, showing normalized percentage scores. The highest score in each column is bolded.

Category	Model	PDI		IDV		MAS		UAI		LTO		IVR	
		English	Chinese	English	Chinese	English	Chinese	English	Chinese	English	Chinese	English	Chinese
Dilemmas	Gemini	0.63	0.58	0.61	0.50	0.34	0.37	0.59	0.57	0.95	0.99	0.39	0.43
	GPT-5	0.59	0.52	0.52	0.54	0.33	0.35	0.66	0.58	0.98	0.97	0.39	0.38
	Grok 4	0.59	0.57	0.57	0.54	0.34	0.37	0.60	0.60	0.97	0.97	0.47	0.50
	ERNIE 4.5	0.56	0.57	0.53	0.48	0.36	0.33	0.61	0.58	0.99	0.99	0.41	0.41
	DeepSeek	0.60	0.60	0.58	0.54	0.34	0.37	0.61	0.57	0.99	1.00	0.49	0.40
	Qwen3	0.62	0.53	0.63	0.55	0.36	0.33	0.67	0.58	0.96	0.96	0.42	0.46
	<i>Average</i>	0.60	0.56	0.57	0.52	0.34	0.35	0.62	0.58	0.97	0.98	0.43	0.43
Exploration	Gemini	0.56	0.50	0.79	0.55	0.46	0.36	0.56	0.51	0.80	0.95	0.44	0.50
	GPT-5	0.59	0.47	0.60	0.61	0.32	0.35	0.52	0.56	0.92	0.92	0.41	0.44
	Grok 4	0.56	0.55	0.69	0.62	0.38	0.49	0.63	0.47	0.90	0.92	0.48	0.46
	ERNIE 4.5	0.55	0.52	0.66	0.64	0.41	0.37	0.58	0.49	0.87	0.94	0.48	0.51
	DeepSeek	0.61	0.55	0.80	0.54	0.40	0.32	0.59	0.51	0.89	0.92	0.48	0.56
	Qwen3	0.61	0.54	0.83	0.62	0.38	0.36	0.57	0.43	0.89	0.94	0.52	0.53
	<i>Average</i>	0.58	0.52	0.73	0.60	0.39	0.38	0.58	0.50	0.88	0.93	0.47	0.50
Crisis	Gemini	0.62	0.54	0.41	0.42	0.38	0.48	0.76	0.62	0.74	0.81	0.43	0.39
	GPT-5	0.64	0.50	0.41	0.47	0.32	0.40	0.75	0.77	0.78	0.85	0.36	0.46
	Grok 4	0.66	0.56	0.54	0.50	0.46	0.44	0.81	0.71	0.79	0.84	0.40	0.41
	ERNIE 4.5	0.61	0.49	0.40	0.33	0.39	0.45	0.71	0.63	0.79	0.85	0.43	0.38
	DeepSeek	0.61	0.53	0.33	0.37	0.41	0.31	0.78	0.58	0.83	0.90	0.38	0.39
	Qwen3	0.57	0.56	0.36	0.33	0.30	0.39	0.78	0.67	0.76	0.89	0.42	0.36
	<i>Average</i>	0.62	0.53	0.41	0.40	0.38	0.41	0.76	0.66	0.78	0.86	0.40	0.40
Negotiation	Gemini	0.68	0.50	0.90	0.60	0.34	0.38	0.72	0.54	0.84	0.94	0.66	0.38
	GPT-5	0.72	0.40	0.86	0.52	0.40	0.32	0.66	0.54	0.88	1.00	0.46	0.40
	Grok 4	0.66	0.52	0.92	0.62	0.36	0.38	0.58	0.50	0.80	0.98	0.70	0.50
	ERNIE 4.5	0.64	0.42	0.76	0.48	0.46	0.30	0.54	0.44	0.92	0.96	0.50	0.60
	DeepSeek	0.74	0.50	1.00	0.60	0.26	0.28	0.52	0.48	0.96	0.98	0.60	0.46
	Qwen3	0.76	0.54	0.84	0.68	0.32	0.34	0.52	0.56	0.80	0.94	0.48	0.36
	<i>Average</i>	0.70	0.48	0.88	0.58	0.36	0.33	0.59	0.51	0.87	0.97	0.57	0.45
Interpretation	Gemini	0.58	0.46	0.64	0.66	0.32	0.48	0.72	0.68	0.80	0.94	0.42	0.58
	GPT-5	0.52	0.52	0.60	0.72	0.30	0.40	0.60	0.66	0.70	0.88	0.40	0.52
	Grok 4	0.52	0.52	0.80	0.58	0.34	0.50	0.60	0.50	0.78	0.90	0.60	0.52
	ERNIE 4.5	0.46	0.42	0.76	0.62	0.44	0.34	0.60	0.56	0.78	0.90	0.52	0.56
	DeepSeek	0.52	0.48	0.82	0.50	0.44	0.36	0.62	0.64	0.68	0.86	0.50	0.58
	Qwen3	0.52	0.50	0.70	0.68	0.44	0.38	0.62	0.62	0.82	0.82	0.50	0.48
	<i>Average</i>	0.52	0.48	0.72	0.63	0.38	0.41	0.63	0.61	0.76	0.88	0.49	0.54

preferences and exit options; ERNIE 4.5 lies at the collectivistic end in both languages, prioritizing harmony and compromise. PDI is highest for Qwen3, which frequently resolves negotiations via hierarchical decision-makers, and lowest for GPT-5 (ZH), which more often legitimizes bottom-up proposals. On UAI, Gemini and GPT-5 (EN) lean heavily on explicit rules and procedures, whereas ERNIE 4.5 (ZH) tolerates more informal. IVR is high for Grok 4 and DeepSeek in English, where strong preference expression is narratively acceptable, but low for GPT-5 and Qwen3 in Chinese, which tend to favor restraint and deference in bargaining contexts. In *Context Interpretation*, the clearest separations arise on **IDV**, **IVR**, and **UAI**. On IDV, DeepSeek anchors the individualistic end in English while GPT-5 is more collectivistic; in Chinese, the ordering reverses, with GPT-5 highest and DeepSeek lowest, illustrating strong language-dependent shifts for the same pair of models. For IVR, GPT-5 is consistently the most restraint-oriented in both languages, whereas Grok 4 (EN) and Gemini/DeepSeek (ZH) legitimize more self-expression in how events are framed and evaluated. On UAI, Gemini tends to produce the most procedure-seeking, risk-averse interpretations under both languages, while Grok 4 is comparatively ambiguity-tolerant in Chinese. PDI spreads are moderate but directionally consistent: Gemini and GPT-5 skew toward higher power distance, with ERNIE 4.5 closer to the low-PDI.

Taken together, these category-wise results reveal stable *intrinsic structural preferences* for each generator. Across tasks, Qwen3 and DeepSeek more often foreground self-authored agency and expressive claims: their protagonists are relatively quick to define their own goals, to insist on preferred outcomes in negotiations, and to treat crises as occasions for self-directed change. This pattern is reflected in systematically higher scores on IDV and, in many settings, IVR, especially in *Exploration*, *Negotiation*, and *Interpretation* scenarios where multiple paths are available. By contrast, GPT-5 and ERNIE 4.5 tend to embed choices in collective obligations and shared responsibilities. Their stories more frequently distribute agency across families, teams, or communities, legitimize compromise over unilateral moves, and keep power distance relatively low, yielding lower IDV and PDI in dilemmas and crises even when prompts are held fixed. Moreover, Gemini and Grok 4 form a third, more hierarchical and procedure-seeking cluster. In high-stakes settings such as *Crisis Response* and *Order Negotiation*, they are the most likely to invoke formal authority, explicit rules, and structured plans as the "right" way to resolve conflict, producing consistently high scores on UAI and PDI. Grok 4 additionally pushes toward competitive, performance-oriented framings (higher MAS) and strong preference expression in bargaining (higher IVR), whereas Gemini's narratives combine achievement focus with a restrained hedonic

profile. These regularities appear across categories and do not vanish when the narrative task changes, indicating that each model carries a characteristic narrative logic about who should act, on whose behalf, and under what constraints. This forms the first layer of the NSB: even without varying the prompt language, LLMs exhibit entrenched, model-specific preferences over how culturally loaded stories are structured and resolved.

Language-Conditioned Narrative Shifts Beyond model-specific differences, the category averages in Table 2 reveal a systematic effect of prompt language. Collapsing across models, English prompts reliably raise PDI, IDV and UAI relative to Chinese, while Chinese prompts robustly elevate LTO. In all five narrative categories, average IDV is higher in English than in Chinese, and the same holds for PDI and UAI, with the gap especially pronounced in adversarial or high-stakes settings such as *Crisis* and *Negotiation*. Conversely, LTO is consistently closer to ceiling in Chinese than in English across categories. Thus, even when the underlying scenarios are held fixed, English prompts more often elicit stories that foreground personal choice, formal authority, and procedural safeguards, whereas Chinese prompts encourage futures-oriented, incremental trajectories that embed decisions in longer time horizons.

The remaining dimensions show more nuanced, task-dependent shifts. For **IVR**, Chinese averages are higher in non-adversarial categories such as *Exploration* and *Interpretation*, where pursuing enjoyment and self-expression within relationships is readily legitimated, while English averages are higher in *Order Negotiation*, where insisting on one's preferences or walking away from a deal is narratively acceptable. **MAS** differences are modest and category-specific: English is slightly more achievement- and competition-oriented in open-ended goals and bargaining, whereas Chinese often places more weight on care, harmony, or quality-of-life in dilemmas, crises, and interpretive contexts.

Crucially, these language effects are measured under SVP's structure anchored scoring, where each Hofstede dimension is inferred from specific narrative functions rather than from surface-level keywords. The observed shifts therefore reflect changes in how stories are organized and resolved: who is granted decisive agency, whether conflicts are routed through hierarchy or peer deliberation, how much ambiguity is tolerated, and how far into the future consequences are projected. Taken together, the results support our second layer of dual-layer NSB, indicating that prompt language functions as a cultural primer that reweights the model's preferred conflict framings and value trade-offs. This challenges the notion of a static model alignment, revealing instead a fluid narrative identity.

5.2 Cross Evaluation

Because any single LLM used as rater may carry its own cultural bias, relying exclusively on one evaluator risks imprinting rater-side priors onto the measurements. To gauge robustness, we conducted a cross evaluation using Qwen3-235B as an additional rater, applying the same SVP prompts and rubric as with GPT-5. Table 4 in the Appendix reports the resulting Hofstede scores, which allow us to separate patterns attributable to the targets (the six generation models) from those attributable to the rater (GPT-5 vs. Qwen3).

Qwen3 adopts a systematically more conservative scale: across categories and most dimensions, especially **IDV**, **IVR**, **UAI**, and

PDI, its absolute scores are lower than GPT-5's, while **LTO** remains near ceiling for both raters. Put differently, the same story is more readily judged as individualistic, indulgent, or rule-seeking by GPT-5 than by Qwen3. This level shift is best interpreted as a stringency/leniency effect coupled with different normative priors: GPT-5 more often legitimizes self-direction, expression, and procedural guardrails, whereas Qwen3's grading is comparatively restraint-oriented and hierarchy-accommodating.

Crucially, the *relative* patterns are highly stable across raters. Qwen3 reproduces the main qualitative trends observed with GPT-5: in most categories, PDI and UAI are higher under English prompts than Chinese, Chinese scores for LTO show ceiling effects, and MAS remains uniformly low with modest dispersion. Qwen3 also recovers the large between-model spreads in IDV and IVR, preserving the qualitative conclusion that some generators habitually foreground self-authored agency and expressive claims whereas others favor collective accommodation and restraint. The identity of high-dispersion dimensions (IDV, IVR, UAI, PDI) and the broad ordering of models within a language therefore remain largely unchanged, even though evaluator origin can tilt cross-lingual contrasts in magnitude and occasionally in direction. For instance, Qwen3 rates Chinese *Negotiation* stories as slightly more individualistic than English ones, and Chinese *Dilemmas* as more uncertainty-avoidant, even though the cross-category averages still favor higher IDV and UAI under English prompts overall.

Overall, this cross evaluation provides a two-fold validation. First, it confirms that the dual-layer NSB identified in Section 5.1 are robust enough to be detected by culturally distinct evaluators. Second, it makes rater-side bias explicit: while absolute scores are calibrated by each rater's own cultural priors, the structural patterns of NSB remain remarkably stable. This supports the interpretation that the observed cultural regularities are properties of the generated stories, rather than artifacts of a particular evaluator.

6 Conclusion

This study targets the previously underexplored cultural Narrative Structural Bias in LLMs by proposing CuNA and SVP. Extensive qualitative and quantitative analyses establish cultural NSB as a stable attribute of state-of-the-art LLMs and reveal its complex manifestations across models and scenarios, systematic language-conditioned drift, and category-specific oscillations. Cross-evaluation with culturally distinct raters further shows that both model origin and prompt language act as reliable predictors of narrative logic. We offer the first systematic toolkit and empirical evidence base for diagnosing, comparing the cultural NSB in LLMs, advancing the development of more responsible and equitable story-generation systems across languages and markets. Research resources are available for transparency and contributions¹.

Acknowledgment

We gratefully acknowledge the support from the Natural Science Foundation of Guangdong Province (No. 2025B1515020060), the Basic and Applied Basic Research Program of the Guangzhou Science and Technology Plan (No. 2025A04J7141). Icons made by Freepik from www.flaticon.com.

¹<https://github.com/Function-Samuel/Cultural-Narrative-Alignment>

References

- [1] 1997. Friends: "The One with All the Jealousy" (Episode). <https://www.imdb.com/title/tt0583506/>. Accessed: 2025-11-23.
- [2] 2019. Dark Waters (2019). <https://www.imdb.com/title/tt9071322/>. Accessed: 2025-11-23.
- [3] n.d.. 1996 Mount Everest disaster. https://en.wikipedia.org/wiki/1996_Mount_Everest_disaster. Accessed: 2025-11-23.
- [4] n.d.. Beslan school siege. https://en.wikipedia.org/wiki/Beslan_school_siege. Accessed: 2025-11-23.
- [5] n.d.. Pleasant Hill bus tragedy. https://en.wikipedia.org/wiki/Pleasant_Hill_bus_tragedy. Accessed: 2025-11-23.
- [6] n.d.. "Family Disapproval" (keyword id 199764) – The Movie Database. <https://www.themoviedb.org/keyword/199764-family-disapproval/movie?language=en-ÅŠUS>. Accessed: 2025-11-23.
- [7] Muhammad Farid Adilazuarda, Sagnik Mukherjee, Pradhyumna Lavania, Sidhant Singh, Alham Fikri Aji, Jacki O'Neill, Ashutosh Modi, and Monojit Choudhury. 2024. Towards measuring and modeling "culture" in LLMs: A survey. *arXiv preprint arXiv:2403.15412* (2024).
- [8] American Statistical Association. 1922. Journal of the American Statistical Association. American Statistical Association.
- [9] Backlinko. 2025. Most Popular AI Apps. <https://backlinko.com/most-popular-ai-apps>. Accessed October 1, 2025.
- [10] Baidu-ERNIE-Team. 2025. ERNIE 4.5 Technical Report. https://ernie.baidu.com/blog/publication/ERNIE_Technical_Report.pdf.
- [11] Damayanti Banerjee and Sheila L Steinberg. 2015. Exploring spatial and cultural discourses in environmental justice movements: A study of two communities. *Journal of Rural Studies* 39 (2015), 41–50.
- [12] Andrew M Bean, Ryan Othniel Kearns, Angelika Romanou, Franziska Sofia Hafner, Harry Mayne, Jan Batzner, Negar Foroutan, Chris Schmitz, Karolina Korgul, Hunar Batra, et al. 2025. Measuring what Matters: Construct Validity in Large Language Model Benchmarks. *arXiv preprint arXiv:2511.04703* (2025).
- [13] Nina Beguš. 2024. Experimental narratives: A comparison of human crowd-sourced storytelling and AI storytelling. *Humanities and Social Sciences Communications* 11, 1 (2024), 1–22.
- [14] "Disasterology.com Blog". 2020. Disaster Movies to Watch in Quarantine. <https://www.disaster-ology.com/blog/2020/3/13/disaster-movies-to-watch-in-quarantine>. Accessed: 2025-11-23.
- [15] Tolga Bolukbasi, Kai-Wei Chang, James Y Zou, Venkatesh Saligrama, and Adam T Kalai. 2016. Man is to computer programmer as woman is to homemaker? debiasing word embeddings. *Advances in neural information processing systems* 29 (2016).
- [16] Rishi Bommasani, Kathleen A Creel, Ananya Kumar, Dan Jurafsky, and Percy S Liang. 2022. Picking on the same person: Does algorithmic monoculture lead to outcome homogenization? *Advances in Neural Information Processing Systems* 35 (2022), 3663–3678.
- [17] Traci M Bricka, Yimin He, and Amber N Schroeder. 2023. Difficult times, difficult decisions: examining the impact of perceived crisis response strategies during COVID-19. *Journal of business and psychology* 38, 5 (2023), 1077–1097.
- [18] Jerome Bruner. 1990. *Acts of meaning: Four lectures on mind and culture*. Vol. 3. Harvard university press.
- [19] Sébastien Bubeck, Christian Coester, Ronen Eldan, Timothy Gowers, Yin Tat Lee, Alexandru Lupsasca, Mehtaab Sawhney, Robert Scherrer, Mark Sellke, Brian K. Spears, Derya Unutmaz, Kevin Weil, Steven Yin, and Nikita Zhivotovskiy. 2025. Early science acceleration experiments with GPT-5. <https://api.semanticscholar.org/CorpusID:283110402>
- [20] Ingrid Campo-Ruiz. 2025. Artificial intelligence may affect diversity: architecture and cultural context reflected through ChatGPT, Midjourney, and Google Maps. *Humanities and Social Sciences Communications* 12, 1 (2025), 1–13.
- [21] Kenneth Church and Patrick Hanks. 1990. Word association norms, mutual information, and lexicography. *Computational linguistics* 16, 1 (1990), 22–29.
- [22] Gheorghe Comanici, Eric Bieber, Mike Schaeckermann, Ice Pasupat, Noveen Sachdeva, Inderjit Dhillon, Marcel Blistein, Ori Ram, Dan Zhang, Evan Rosen, et al. 2025. Gemini 2.5: Pushing the frontier with advanced reasoning, multimodality, long context, and next generation agentic capabilities. *arXiv preprint arXiv:2507.06261* (2025).
- [23] Hui Dai, Dan Pechi, Xinyi Yang, Garvit Banga, and Raghav Mantri. 2024. Denial: In-context features influence LLM needle-in-a-haystack abilities. *arXiv preprint arXiv:2411.19360* (2024).
- [24] Jwala Dhamala, Tony Sun, Varun Kumar, Satyapriya Krishna, Yada Pruskachatkun, Kai-Wei Chang, and Rahul Gupta. 2021. Bold: Dataset and metrics for measuring biases in open-ended language generation. In *Proceedings of the 2021 ACM conference on fairness, accountability, and transparency*. 862–872.
- [25] Lubna Luxmi Dhirani, Noorain Mukhtiar, Bhawani Shankar Chowdhry, and Thomas Newe. 2023. Ethical dilemmas and privacy issues in emerging technologies: A review. *Sensors* 23, 3 (2023), 1151.
- [26] Martha O Dimgba, Sharon Oba, Ameeta Agrawal, and Philippe J Giabbanelli. 2025. Mitigation of Gender and Ethnicity Bias in AI-Generated Stories through Model Explanations. *arXiv preprint arXiv:2509.04515* (2025).
- [27] Fangzhou Dong, Yifan Zeng, Yingpeng Sang, and Hong Shen. 2025. Structuralist Approach to AI Literary Criticism: Leveraging Greimas Semiotic Square for Large Language Models. In *Proceedings of the Annual Meeting of the Cognitive Science Society*, Vol. 47.
- [28] James A Dungan, Michael Stepanovic, and Liane Young. 2016. Theory of mind for processing unexpected events across contexts. *Social Cognitive and Affective Neuroscience* 11, 8 (2016), 1183–1192.
- [29] Qixiang Fang, Dong Nguyen, and Daniel L Oberski. 2022. Evaluating the construct validity of text embeddings with application to survey questions. *EPJ Data Science* 11, 1 (2022), 39.
- [30] Samuel Gehman, Suchin Gururangan, Maarten Sap, Yejin Choi, and Noah A Smith. 2020. Realtoxicityprompts: Evaluating neural toxic degeneration in language models. *arXiv preprint arXiv:2009.11462* (2020).
- [31] Sourojit Ghosh and Kyra Wilson. 2025. Bias Is a Math Problem, AI Bias Is a Technical Problem: 10-Year Literature Review of AI/LLM Bias Research Reveals Narrow [Gender-Centric] Conceptions of 'Bias', and Academia-Industry Gap. In *Proceedings of the AAAI/ACM Conference on AI, Ethics, and Society*, Vol. 8. 1091–1106.
- [32] Thomas L Griffiths and Mark Steyvers. 2004. Finding scientific topics. *Proceedings of the National Academy of Sciences* 101, suppl_1 (2004), 5228–5235.
- [33] Xiaowen Guan, Hee Sun Park, and Hye Eun Lee. 2009. Cross-cultural differences in apology. *International Journal of Intercultural Relations* 33, 1 (2009), 32–45.
- [34] Daya Guo, Dejian Yang, Haowei Zhang, Junxiao Song, Ruoyu Zhang, Runxin Xu, Qihao Zhu, Shirong Ma, Peiyi Wang, Xiao Bi, et al. 2025. DeepSeek-R1: Incentivizing reasoning capability in LLMs via reinforcement learning. *arXiv preprint arXiv:2501.12948* (2025).
- [35] Xiaobo Guo and Soroush Vosoughi. 2025. Serial position effects of large language models. In *Findings of the Association for Computational Linguistics: ACL 2025*. 927–953.
- [36] Jeremy Hayes. n.d.. Mystery Movies. <https://www.buzzfeed.com/jeremyhayes/mystery-movies>. Accessed: 2025-11-23.
- [37] Geert Hofstede. 2001. *Culture's consequences: Comparing values, behaviors, institutions and organizations across nations*. Sage publications.
- [38] Geert Hofstede. 2011. Dimensionalizing cultures: The Hofstede model in context. *Online readings in psychology and culture* 2, 1 (2011), 8.
- [39] Geert H Hofstede, Gert Jan Hofstede, and Michael Minkov. 2010. Cultures and organizations: Software of the mind: Intercultural cooperation and its importance for survival. (No Title) (2010).
- [40] Tom Hubler. 2023. Exploring Family Business Insights through Movies. <https://tomhubler.com/exploring-family-business-insights-through-movies/>. Accessed: 2025-11-23.
- [41] Transparency International. 2020. 11 Movies About Whistleblowers That You Cannot Miss. <https://www.transparency.org/en/blog/11-movies-about-whistleblowers-that-you-cannot-miss>. Accessed: 2025-11-23.
- [42] ISGlobal. 2024. 10 películas sobre la preparación de crisis y desastres. <https://www.isglobal.org/en/healthisglobal/-/custom-blog-portlet/10-peliculas-sobre-la-preparacion-de-crisis-y-desastres>. Accessed: 2025-11-23.
- [43] David Jackson and Marsha Courneya. 2023. Unreliable Narrator: Reporative Approaches to Harmful Biases in AI Storytelling for the HE Classroom And Future Creative Industries. *Brazilian Creative Industries Journal* 3, 2 (2023), 50–66.
- [44] Igor Kabashkin, Olga Zervina, and Boris Misnevs. 2025. AI Narrative Modeling: How Machines' Intelligence Reproduces Archetypal Storytelling. *Information* 16, 4 (2025), 319.
- [45] Maurice G Kendall. 1938. A new measure of rank correlation. *Biometrika* 30, 1-2 (1938), 81–93.
- [46] Jon Kleinberg and Manish Raghavan. 2021. Algorithmic monoculture and social welfare. *Proceedings of the National Academy of Sciences* 118, 22 (2021), e2018340118.
- [47] William Labov and Joshua Waletzky. 1997. Narrative analysis: Oral versions of personal experience. (1997).
- [48] Alex Lawrence. 2024. Exploring Art Through Cinema: Five Movies To Watch. <https://lifestyleidiaries.substack.com/p/exploring-art-through-cinema-five>. Accessed: 2025-11-23.
- [49] Junho Lee and Keith J Holyoak. 2020. "But he's my brother": The impact of family obligation on moral judgments and decisions. *Memory & cognition* 48, 1 (2020), 158–170.
- [50] Cheng Li, Damien Teney, Linyi Yang, Qingsong Wen, Xing Xie, and Jindong Wang. 2024. Culturepark: Boosting cross-cultural understanding in large language models. *Advances in Neural Information Processing Systems* 37 (2024), 65183–65216.
- [51] Huihan Li, Liwei Jiang, Jena D Hwang, Hyunwoo Kim, Sebastin Santy, Taylor Sorensen, Bill Yuchen Lin, Nouha Dziri, Xiang Ren, and Yejin Choi. 2024. Culturegen: Revealing global cultural perception in language models through natural language prompting. *arXiv preprint arXiv:2404.10199* (2024).
- [52] Nelson F Liu, Kevin Lin, John Hewitt, Ashwin Paranjape, Michele Bevilacqua, Fabio Petroni, and Percy Liang. 2024. Lost in the middle: How language models use long contexts. *Transactions of the Association for Computational Linguistics*

- 12 (2024), 157–173.
- [53] Anthony T Machette and Ioana A Cionea. 2023. In-laws, communication, and other frustrations: The challenges of intercultural marriages. *Interpersona: An International Journal on Personal Relationships* 17, 1 (2023), 1–18.
 - [54] Moon-Kuen Mak and Tiejian Luo. 2025. A framework for evaluating cultural bias and historical misconceptions in LLMs outputs. *BenchCouncil Transactions on Benchmarks, Standards and Evaluations* (2025), 100235.
 - [55] Christopher Manning and Hinrich Schütze. 1999. *Foundations of statistical natural language processing*. MIT press.
 - [56] Samuel Messick. 1995. Validity of psychological assessment: Validation of inferences from persons' responses and performances as scientific inquiry into score meaning. *American psychologist* 50, 9 (1995), 741.
 - [57] Arifud Muhammad. 2025. LLM statistics 2025: Comprehensive insights into market trends and integration. <https://www.hostinger.com/tutorials/llm-statistics> Accessed October 1, 2025.
 - [58] Moin Nadeem, Anna Bethke, and Siva Reddy. 2021. StereoSet: Measuring stereotypical bias in pretrained language models. In *Proceedings of the 59th annual meeting of the association for computational linguistics and the 11th international joint conference on natural language processing (volume 1: long papers)*. 5356–5371.
 - [59] Nikita Nangia, Clara Vania, Rasika Bhalerao, and Samuel Bowman. 2020. CrowS-pairs: A challenge dataset for measuring social biases in masked language models. In *Proceedings of the 2020 conference on empirical methods in natural language processing (EMNLP)*. 1953–1967.
 - [60] Tarek Naous, Michael J Ryan, Alan Ritter, and Wei Xu. 2024. Having beer after prayer? measuring cultural bias in large language models. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*. 16366–16393.
 - [61] Tarek Naous and Wei Xu. 2025. On The Origin of Cultural Biases in Language Models: From Pre-training Data to Linguistic Phenomena. *arXiv preprint arXiv:2501.04662* (2025).
 - [62] Minh-Hoang Nguyen, Huyen Thanh Thanh Nguyen, Tam-Tri Le, Anh-Phuong Luong, and Quan-Hoang Vuong. 2022. Gender issues in family business research: A bibliometric scoping review. *Journal of Asian Business and Economic Studies* 29, 3 (2022), 166–188.
 - [63] Office of Research Integrity. 2025. The Lab: Interactive Movie on Research Misconduct. <https://ori.hhs.gov/the-lab>. Accessed: 2025-11-23.
 - [64] Cătălina Oțoiu, Lucia Rațiu, and Claudia Lenuța Rus. 2019. Rivals when we work together: Team rivalry effects on performance in collaborative learning groups. *Administrative Sciences* 9, 3 (2019), 61.
 - [65] Alicia Parrish, Angelica Chen, Nikita Nangia, Vishakh Padmakumar, Jason Phang, Jana Thompson, Phu Mon Htut, and Samuel Bowman. 2022. BBQ: A hand-built bias benchmark for question answering. In *Findings of the Association for Computational Linguistics: ACL 2022*. 2086–2105.
 - [66] Karl Pearson. 1920. Notes on the history of correlation. *Biometrika* 13, 1 (1920), 25–45.
 - [67] Partha Pratim Ray. 2025. Does Linguistic Relativity Hypothesis Apply on ChatGPT Responses? Yes, It Does. *Computational Intelligence* 41, 4 (2025), e70103.
 - [68] Zhen Qin, Xiaodong Han, Weixuan Sun, Dongxu Li, Lingpeng Kong, Nick Barnes, and Yiran Zhong. 2022. The devil in linear transformer. *arXiv preprint arXiv:2210.10340* (2022).
 - [69] Alec Radford, Karthik Narasimhan, Tim Salimans, Ilya Sutskever, et al. 2018. Improving language understanding by generative pre-training. (2018).
 - [70] Jill Walker Rettberg and Hermann Wigers. 2025. AI-generated stories favour stability over change: homogeneity and cultural stereotyping in narratives generated by gpt-4o-mini. *arXiv preprint arXiv:2507.22445* (2025).
 - [71] Donya Rooein, Vilém Zouhar, Debora Nozza, and Dirk Hovy. 2025. Biased tales: Cultural and topic bias in generating children's stories. In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing*. 52–72.
 - [72] Muhammed Saeed, Muhammad Abdul-Mageed, and Shady Shehata. 2025. Surfacing Subtle Stereotypes: A Multilingual, Debate-Oriented Evaluation of Modern LLMs. *arXiv preprint arXiv:2511.01187* (2025).
 - [73] Fulgencio Sánchez-Vera. 2025. Critical Algorithmic Mediation: Rethinking Cultural Transmission and Education in the Age of Artificial Intelligence. *Societies* 15, 7 (2025), 198.
 - [74] Lester Sim and Robin Edelstein. 2025. Cross-Cultural Differences in the Links Between Familial Support and Strain in Married and Single Adults' Well-Being. *Personal Relationships* 32, 3 (2025), e70027.
 - [75] Charles Spearman. 1961. The proof and measurement of association between two things. (1961).
 - [76] Xu Tang, Yifan Zeng, and Fangzhou Dong. 2025. Unveiling Cultural Cognition in AI: A Systematic Investigation of Horizontal-Vertical Individualism-Collectivism Traits in Large Language Models. In *Proceedings of the Annual Meeting of the Cognitive Science Society*, Vol. 47.
 - [77] Yan Tao, Olga Viberg, Ryan S Baker, and René F Kizilcec. 2024. Cultural bias and cultural alignment of large language models. *PNAS nexus* 3, 9 (2024), pgae346.
 - [78] TeamBuilding.com. 2025. 32 Best Team Building Movies About Leadership & Teamwork. <https://teambuilding.com/blog/team-building-movies>. Accessed: 2025-11-23.
 - [79] Yufei Tian, Tenghao Huang, Miri Liu, Derek Jiang, Alexander Spangher, Muhao Chen, Jonathan May, and Nanyun Peng. 2024. Are Large Language Models Capable of Generating Human-Level Narratives?. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*. 17659–17681.
 - [80] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Łukasz Kaiser, and Illia Polosukhin. 2017. Attention is all you need. *Advances in neural information processing systems* 30 (2017).
 - [81] Phil Venables. 2024. Lessons in Crisis Management – Top 10 Disaster Movies. <https://www.philvenables.com/post/lessons-in-crisis-management-top-10-disaster-movies>. Accessed: 2025-11-23.
 - [82] Benjamin Lee Whorf. 2012. *Language, thought, and reality: Selected writings of Benjamin Lee Whorf*. MIT press.
 - [83] xAI. 2025. Grok 4 Model Card. <https://data.x.ai/2025-08-20-grok-4-model-card.pdf>.
 - [84] An Yang, Anfeng Li, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, Chang Gao, Chengen Huang, Chenxu Lv, et al. 2025. Qwen3 technical report. *arXiv preprint arXiv:2505.09388* (2025).
 - [85] Haeun Yu, Seogyong Jeong, Siddhesh Pawar, Jisu Shin, Jiho Jin, Junho Myung, Alice Oh, and Isabelle Augenstein. 2025. Entangled in representations: Mechanistic investigation of cultural biases in large language models. *arXiv preprint arXiv:2508.08879* (2025).
 - [86] Yifan Zeng, Fangzhou Dong, Jian Zhao, Peijia Zheng, Jian Li, and Huiyu Zhou. 2025. Towards Culturally Fair Multimodal Generation: Quantifying and Mitigating Orientalist Biases in Text-to-Visual Models. In *Proceedings of the 33rd ACM International Conference on Multimedia*. 11434–11442.
 - [87] Yifan Zeng, Liang Kairong, Fangzhou Dong, and Peijia Zheng. 2025. Quantifying Risk Propensities of Large Language Models: Ethical Focus and Bias Detection through Role-Play. In *Proceedings of the Annual Meeting of the Cognitive Science Society*, Vol. 47.
 - [88] Xuechen Zhang, Xiangyu Chang, Mingchen Li, Amit Roy-Chowdhury, Jiasi Chen, and Samet Oymak. 2024. Selective attention: Enhancing transformer through principled context control. *Advances in Neural Information Processing Systems* 37 (2024), 11061–11086.
 - [89] Zoie Zhao, Sophie Song, Bridget Duah, Jamie Macbeth, Scott Carter, Monica P Van, Nayeli Suseth Bravo, Matthew Klenk, Kate Sick, and Alexandre LS Filipowicz. 2023. More human than human: LLM-generated narratives outperform human-LLM interleaved narratives. In *Proceedings of the 15th Conference on Creativity and Cognition*. 368–370.
 - [90] Haiqi Zhou, David G Hobson, Derek Ruths, and Andrew Piper. 2024. Large scale narrative messaging around climate change: A cross-cultural comparison. In *Proceedings of the 1st Workshop on Natural Language Processing Meets Climate Change (ClimateNLP 2024)*. 143–155.

A Log-Odds Ratio Analysis

Although the SVP provides high-level quantitative scores for cultural dimensions, it is imperative to verify that these scores are grounded in concrete semantic signals rather than stochastic generation artifacts or post-hoc rationalizations by the LLM-Rater. To achieve this lexical grounding, we moved beyond simple frequency counts which are often dominated by neutral stopwords, and employed *Log-Odds Ratio* analysis [21, 55] with an informative Dirichlet prior [32]. This statistical method identifies words that are significantly *over-represented* in one corpus relative to another, isolating the "lexical fingerprint" of each cultural dimension.

Table 3: Lexical Grounding: Top 5 distinctive words for High vs. Low scores. Dimension column abbreviated as "Dims."

Dims.	Rank	High Score		Low Score	
		Word	L-Odds	Word	L-Odds
PDI	1	bride	8.06	friendship	-8.04
	2	rigid	7.95	communal	-7.35
	3	guild	7.83	cottage	-6.92
	4	embassy	7.59	quilt	-6.82
	5	pending	7.53	stipend	-6.82
IDV	1	users	7.86	evacuation	-7.54
	2	visibility	7.56	villagers	-7.14
	3	analytics	7.45	reservoir	-7.22
	4	competitor	7.27	upstream	-7.29
	5	merit	7.27	masks	-7.05
UAI	1	drought	7.37	graduation	-7.97
	2	infant	7.32	internship	-7.42
	3	embassy	7.28	trajectory	-7.22
	4	survivors	7.13	stipend	-7.22
	5	filtration	7.02	lounge	-7.22
IVR	1	kantha	9.28	factory	-7.07
	2	indie	8.25	hospital	-6.82
	3	galleries	3.42	error	-6.81
	4	aesthetic	3.41	control	-6.62
	5	hobby	3.64	patient	-6.62

We aggregated the generated narratives and stratified them based on SVP scores across the four dimensions exhibiting the most pronounced variance: PDI, IDV, UAI, and IVR. For each dimension, narratives receiving a score of ≥ 4 were classified as the 'High-Score' group, while those with a score of ≤ 2 constituted the 'Low-Score' group. We then computed the Log-Odds Ratio (δ_w) for each word type w between the high-scoring corpus (i) and the low-scoring corpus (j) using the following formulation:

$$\delta_w = \log \left(\frac{y_{wi} + \alpha_w}{n_i + \alpha_0 - (y_{wi} + \alpha_w)} \right) - \log \left(\frac{y_{wj} + \alpha_w}{n_j + \alpha_0 - (y_{wj} + \alpha_w)} \right) \quad (1)$$

Where y_{wi} is the count of word w in corpus i , n_i is the total word count of corpus i , and α represents the Dirichlet prior to smooth counts for rare words. We applied Part-of-Speech (POS) filtering to remove proper nouns and a curated stop-list to remove context-specific props (e.g., "flour" from bakery scenarios), ensuring the analysis focused on value-laden vocabulary. The results detailed

in Table 3, reveal striking semantic dichotomies that validate the internal validity of our scoring pipeline:

Power Distance High-PDI narratives employ a bureaucratic lexicon (*rigid, clause, embassy*), constructing a "contractual society" governed by top-down system. Conversely, Low-PDI ones emphasize organic connection (*friendship, communal*), reflecting a "relational community" where authority is negotiated rather than imposed.

Individualism High-IDV correlates with technocratic success (*analytics, merit*), while Low-IDV correlates with crisis (*evacuation, ration*). This enforces a false dichotomy where prosperity is individualistic and community is merely remedial.

Uncertainty Avoidance High-UAI depicts reaction to existential threat (*drought, scam*), whereas Low-UAI frames uncertainty as optimistic life transition (*graduation, trajectory*), reducing risk to either fatal danger or youthful exploration.

Indulgence vs. Restraint High-IVR is concretized through aesthetics (*indie, galleries*), framing indulgence as creative self-actualization. In contrast, Low-IVR is pathologized or industrialized (*factory, patient*), suggesting models encode "Restraint" not as a cultural virtue, but as a state of systemic suppression.




Figure 3: Visualizing distinctive vocabulary for IDV (top) and IVR (bottom). The semantic contrast validates that the LLM-Rater captures genuine narrative textures.

B LLM-Rater Prompt

The LLM-Rater prompt (Fig. 4) outlines a two-part analytical workflow for narratives. It first guides the LLM to deconstruct the story into six Labov & Waletzky components, each requiring a text excerpt and an analysis. Subsequently, it directs the LLM to analyze the six Hofstede dimensions using a defined scoring system. Crucially, for each judgment, the rater must explicitly state the narrative component(s) relied upon and cite concrete plot points or dialogue as evidence. The prompt also includes detailed examples, which facilitate in-depth and precise analysis of narratives.

C Human-LLM Evaluation Correlation

To rigorously evaluate our LLM-Rater's reliability and accuracy, we conducted a human-LLM evaluation. This was also crucial for validating the distinct advantages of SVP over holistic scoring. For this purpose, we invited human evaluators to score 60 instances (totaling 360 scores) across all categories and LLMs. Subsequently, we performed correlation analysis comparing human annotations with LLM-judge scores generated by (i) a holistic scoring approach and (ii) our proposed SVP, and reported Pearson's r [66], Spearman's ρ [75], and Kendall's τ [45] to quantify human-LLM agreement. As shown in Table 5, SVP consistently exhibits significantly higher

Table 4: Comparative analysis of Hofstede dimensions across LLMs, rated by Qwen3-235B, on benchmark datasets, showing normalized percentage scores. The highest score in each column is bolded.

Category	Model	PDI		IDV		MAS		UAI		LTO		IVR	
		English	Chinese	English	Chinese	English	Chinese	English	Chinese	English	Chinese	English	Chinese
Dilemmas	Gemini	0.55	0.53	0.60	0.48	0.23	0.27	0.43	0.47	0.90	0.97	0.32	0.37
	GPT-5	0.61	0.42	0.61	0.48	0.21	0.22	0.38	0.52	0.98	1.00	0.55	0.34
	Grok 4	0.59	0.44	0.65	0.45	0.28	0.24	0.45	0.40	0.96	1.00	0.68	0.36
	ERNIE 4.5	0.56	0.40	0.58	0.44	0.21	0.23	0.38	0.45	0.93	0.99	0.71	0.28
	DeepSeek	0.55	0.45	0.61	0.46	0.25	0.23	0.33	0.45	0.92	1.00	0.37	0.28
	Qwen3	0.61	0.42	0.61	0.45	0.25	0.22	0.47	0.46	0.94	0.98	0.56	0.29
	Average	0.58	0.44	0.61	0.46	0.24	0.24	0.41	0.46	0.94	0.99	0.53	0.32
Exploration	Gemini	0.52	0.45	0.62	0.54	0.45	0.31	0.57	0.49	0.93	1.00	0.36	0.37
	GPT-5	0.58	0.45	0.45	0.38	0.26	0.23	0.47	0.55	1.00	1.00	0.43	0.32
	Grok 4	0.57	0.48	0.62	0.53	0.37	0.33	0.56	0.46	0.95	1.00	0.59	0.35
	ERNIE 4.5	0.54	0.44	0.47	0.44	0.29	0.34	0.45	0.51	0.99	0.99	0.56	0.35
	DeepSeek	0.49	0.43	0.61	0.43	0.29	0.27	0.43	0.46	0.99	1.00	0.44	0.31
	Qwen3	0.58	0.47	0.54	0.50	0.36	0.29	0.58	0.43	0.99	1.00	0.50	0.31
	Average	0.55	0.45	0.55	0.47	0.34	0.30	0.51	0.48	0.98	1.00	0.48	0.34
Crisis	Gemini	0.54	0.50	0.31	0.28	0.28	0.26	0.51	0.57	0.66	0.88	0.27	0.33
	GPT-5	0.62	0.51	0.25	0.31	0.20	0.22	0.62	0.53	0.94	0.93	0.24	0.26
	Grok 4	0.57	0.56	0.36	0.29	0.28	0.25	0.58	0.58	0.83	0.94	0.28	0.35
	ERNIE 4.5	0.51	0.49	0.26	0.28	0.23	0.21	0.54	0.49	0.97	0.96	0.28	0.30
	DeepSeek	0.53	0.48	0.25	0.23	0.23	0.24	0.58	0.56	0.82	0.97	0.26	0.23
	Qwen3	0.55	0.49	0.28	0.24	0.22	0.24	0.68	0.53	0.77	0.93	0.30	0.22
	Average	0.55	0.51	0.29	0.27	0.24	0.24	0.59	0.54	0.83	0.94	0.27	0.28
Negotiation	Gemini	0.50	0.42	0.64	0.68	0.36	0.24	0.50	0.36	0.98	1.00	0.58	0.48
	GPT-5	0.56	0.42	0.46	0.52	0.20	0.22	0.46	0.36	1.00	1.00	0.24	0.42
	Grok 4	0.44	0.42	0.56	0.62	0.22	0.32	0.42	0.38	1.00	0.98	0.64	0.50
	ERNIE 4.5	0.46	0.46	0.52	0.44	0.30	0.24	0.42	0.42	1.00	1.00	0.56	0.50
	DeepSeek	0.50	0.42	0.64	0.68	0.32	0.30	0.46	0.38	0.98	1.00	0.58	0.56
	Qwen3	0.54	0.42	0.50	0.60	0.24	0.30	0.44	0.36	1.00	1.00	0.54	0.54
	Average	0.50	0.43	0.55	0.59	0.27	0.27	0.45	0.38	0.99	1.00	0.52	0.50
Interpretation	Gemini	0.48	0.32	0.50	0.42	0.34	0.22	0.48	0.42	0.92	0.96	0.32	0.32
	GPT-5	0.36	0.34	0.30	0.46	0.20	0.20	0.44	0.40	1.00	1.00	0.34	0.28
	Grok 4	0.36	0.30	0.56	0.50	0.29	0.28	0.47	0.40	1.00	0.96	0.67	0.34
	ERNIE 4.5	0.44	0.30	0.50	0.38	0.30	0.26	0.38	0.28	1.00	1.00	0.48	0.30
	DeepSeek	0.48	0.28	0.46	0.36	0.28	0.22	0.42	0.40	0.94	0.94	0.52	0.42
	Qwen3	0.48	0.30	0.40	0.34	0.38	0.24	0.40	0.30	0.94	0.98	0.28	0.32
	Average	0.43	0.31	0.45	0.41	0.30	0.24	0.43	0.37	0.97	0.97	0.44	0.33

ROLE

You are a computational social scientist specializing in narratology and sociology. Your task is to conduct an in-depth structural and cultural analysis... (61 Tokens)

TASK

Analysis Workflow Your analysis will be output in two main parts:

Part 1: Detailed Analysis Report

1. Narrative Structure Analysis: Please deconstruct the story text according to the six components of the Labov & Waletzky model. For each component, you must provide a [Text Excerpt] and an [Analysis]... (366 Tokens)
2. Cultural Dimensions Analysis: Please provide a detailed analysis for each of the six Hofstede dimensions. For each dimension, provide a [Dimension Name & Score] and a [Rationale]... (122 Tokens)
3. Scoring Rubric (Five-Point Likert Scale)
Power Distance, PDI
1: Low PDI. The narrative emphasizes equality, challenging authority, and ...
5: High PDI. The narrative emphasizes respect for elders/authority...
... (318 Tokens)

Part 2: Score Summary

Output a list of Hofstede Scores in JSON format.

Here are **three examples** that have been carefully curated and evaluated by experts.

Ex. I: [A full kit of Benchmark, Response, Analysis, and Scores.]

...

Ex. III: [A full kit of Benchmark, Response, Analysis, and Scores.] (6,320 Tokens)

positive correlations with human annotations across the overall score and all six Hofstede dimensions, suggesting that SVP produces more stable scores and demonstrates closer agreement with human ratings, thereby validating its effectiveness in mitigating the biases inherent in holistic scoring.

Table 5: Correlation Analysis between LLM and Human Evaluators against holistic scoring method and SVP. All values are presented as percentages (%).

Metrics (%)	Avg.	PDI	IDV	MAS	UAI	LTO	IVR
Holistic Scoring Method							
Pearson's r	64.6	67.4	75.7	51.3	53.2	55.6	74.2
Spearman's ρ	59.1	64.1	66.9	42.4	48.8	63.2	69.5
Kendall's τ	52.0	56.3	59.2	37.0	40.3	59.9	59.5
Structural-Value Pipeline							
Pearson's r	74.8	77.8	82.9	62.0	72.1	65.9	87.9
Spearman's ρ	72.8	76.9	73.7	67.1	73.7	65.1	80.2
Kendall's τ	64.9	68.5	64.9	60.3	64.0	61.5	70.3

Figure 4: Prompt for LLM-Rater narratives evaluation.